

ML LAB

ASSIGNMENT 3

NAME: SWAPNIL GHOSH
ROLL: 001911001061
DEPT: IT
YEAR: 4TH

PART 1



HELPER FUNCTIONS 1 -

```
helper_functions
```

```
helper_functions > plot_macro_avg_roc_curve

1 from sklearn.metrics import confusion_matrix
2 import numpy as np
3 import matplotlib.pyplot as plt
4 from sklearn.base import clone
5 from sklearn.model_selection import train_test_split
6 from sklearn.metrics import roc_curve, auc
7 from sklearn.preprocessing import label_binarize
8
9 def index_of_max(arr, not_allowed):
10     best = float('-inf')
11     best_index = -1
12     for i in range(len(arr)):
13         if arr[i] <= best or i in not_allowed:
14             continue
15         best = arr[i]
16         best_index = i
17     return best_index
18
19 #returns the mapping of hidden state to true class
20 def get_mapping(classifier, X_train, y_train):
21
22     y_pred = classifier.predict(X_train)
23     matrix = confusion_matrix(y_train, y_pred)
24     mapping = dict()
25     used = set() #set of all the hidden states that have already been mapped
26
27     for i in range(len(matrix)):
28         best_guess = index_of_max(matrix[i], used)
29         used.add(best_guess)
30         mapping[best_guess] = i
31     return mapping
32
33 def custom_scoring_func(y_true, y_pred):
34     matrix = confusion_matrix(y_true, y_pred)
35     mapping = dict()
36     used = set() #set of all the hidden states that have already been mapped
37
38     for i in range(len(matrix)):
39         best_guess = index_of_max(matrix[i], used)
40         used.add(best_guess)
41         mapping[best_guess] = i
42
43     y_pred = list(map(lambda hs : mapping[hs], y_pred))
44
45     return np.count_nonzero(y_true==y_pred)/len(y_true)
46
47 def custom_learning_curve(base_classifier, X, y, train_ratios):
48
```

```
helper_functions
```

```
helper_functions > plot_macro_avg_roc_curve

47 def custom_learning_curve(base_classifier, X, y, train_ratios):
48
49     def helper(X_train, y_train, X_test, y_test):
50         classifier = clone(base_classifier)
51         classifier.fit(X_train)
52         mapper = get_mapping(classifier, X_train, y_train)
53         y_pred = list(map(lambda hs : mapper[hs], classifier.predict(X_test)))
54         y_train_pred = list(map(lambda hs : mapper[hs],
55                                classifier.predict(X_train)))
56         score_test = np.count_nonzero(y_test==y_pred)/len(y_test)
57         score_train = np.count_nonzero(y_train==y_train_pred)/len(y_train)
58
59         return score_train, score_test
60
61     train_sizes = np.int_(train_ratios*len(X))
62     train_scores = []
63     test_scores = []
64
65     for ratio in train_ratios:
66         train = []
67         test = []
68         for random_state in range(5):
69             X_train, X_test, y_train, y_test = train_test_split(X, y,
70                                                               train_size=ratio, random_state=random_state)
71             score_train, score_test = helper(X_train, y_train, X_test, y_test)
72             train.append(score_train)
73             test.append(score_test)
74         train_scores.append(train)
75         test_scores.append(test)
76
77     return train_sizes, train_scores, test_scores
78
79 def save_learning_curve(base_classifier, title, X, y, ylim=(0.4, 1.01),
80                        train_ratios=np.linspace(0.18, 0.8, 4), filepath="lc.png"):
81     plt.figure()
82     ax = plt.gca()
83     ax.set_title(title)
84     if ylim is not None:
85         ax.set_ylim(*ylim)
86     ax.set_xlabel("Training examples")
87     ax.set_ylabel("Score")
88
89     train_sizes, train_scores, test_scores = custom_learning_curve(
90         base_classifier,
91         X,
92         y,
93         train_ratios=train_ratios
94     )
```

HELPER FUNCTIONS 2 -

```
helper_functions / plot_macro_avg_roc_curve
77 def save_learning_curve(base_classifier, title, X, y, ylim=(0.4, 1.01),
78     train_ratios=np.linspace(0.18, 0.8, 4), filepath="lc.png"):
79     plt.figure()
80     ax = plt.gca()
81     ax.set_title(title)
82     if ylim is not None:
83         ax.set_ylim(*ylim)
84     ax.set_xlabel("Training examples")
85     ax.set_ylabel("Score")
86
87     train_sizes, train_scores, test_scores = custom_learning_curve(
88         base_classifier,
89         X,
90         y,
91         train_ratios=train_ratios
92     )
93
94     train_scores_mean = np.mean(train_scores, axis=1)
95     train_scores_std = np.std(train_scores, axis=1)
96     test_scores_mean = np.mean(test_scores, axis=1)
97     test_scores_std = np.std(test_scores, axis=1)
98
99     # Plot learning curve
100    ax.grid()
101    ax.fill_between(
102        train_sizes,
103        train_scores_mean - train_scores_std,
104        train_scores_mean + train_scores_std,
105        alpha=0.1,
106        color="r",
107    )
108    ax.fill_between(
109        train_sizes,
110        test_scores_mean - test_scores_std,
111        test_scores_mean + test_scores_std,
112        alpha=0.1,
113        color="g",
114    )
115    ax.plot(
116        train_sizes, train_scores_mean, "o-", color="r", label="Training score"
117    )
118    ax.plot(
119        train_sizes, test_scores_mean, "o-", color="g", label="Validation score"
120    )
121    ax.legend(loc="best")
122    plt.savefig(filepath)
123
```

```
helper_functions / plot_macro_avg_roc_curve
124 def plot_macro_avg_roc_curve(classifier, n_classes, name, X_train, X_test, y_train, y_test):
125     classifier.fit(X_train)
126     y_pred = classifier.predict(X_test)
127     mapping = get_mapping(classifier, X_train, y_train)
128     y_pred_changed = list(map(lambda hs : mapping[hs], y_pred))
129     y_test_binarized = label_binarize(y_test, classes=[0, 1, 2])
130     y_score = classifier.predict_proba(X_test)
131
132     #Rearranging the columns of y_score matrix
133     perm_list = [0 for _ in range(n_classes)]
134
135     for key in mapping:
136         perm_list[mapping[key]] = key
137
138     y_score= y_score[:, perm_list]
139
140     # Compute ROC curve and ROC area for each class
141     fpr = dict()
142     tpr = dict()
143     roc_auc = dict()
144     for i in range(n_classes):
145         fpr[i], tpr[i], _ = roc_curve(y_test_binarized[:, i], y_score[:, i])
146         roc_auc[i] = auc(fpr[i], tpr[i])
147
148     # First aggregate all false positive rates
149     all_fpr = np.unique(np.concatenate([fpr[i] for i in range(n_classes)]))
150
151     # Then interpolate all ROC curves at this points
152     mean_tpr = np.zeros_like(all_fpr)
153     for i in range(n_classes):
154         mean_tpr += np.interp(all_fpr, fpr[i], tpr[i])
155
156     # Finally average it and compute AUC
157     mean_tpr /= n_classes
158
159     fpr["macro"] = all_fpr
160     tpr["macro"] = mean_tpr
161     roc_auc["macro"] = auc(fpr["macro"], tpr["macro"])
162
163
164     plt.plot(
165         fpr["macro"],
166         tpr["macro"],
167         label="{0} (AUC = {1:0.2f})".format(name, roc_auc["macro"]),
168         alpha=0.7,
169         linewidth=1,
170     )
171
172
173 return plt
```

WINE - CLASSIFICATION

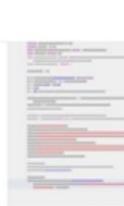


GAUSSIAN HMM



GAUSSIAN HMM (NO PARAMETER TUNING)- CODE & EVALUATION

```
gaussian_no_param.py No Selection
1 import matplotlib.pyplot as plt
2 import pandas as pd
3 from sklearn.model_selection import train_test_split
4 from hmmlearn import hmm
5 from sklearn.metrics import classification_report, confusion_matrix,
   accuracy_score, ConfusionMatrixDisplay
6 from helper_functions import *
7
8 NUM_CLASSES = 3
9
10 df = pd.read_csv('./wine.data', header=None)
11 df = df.sample(frac = 1, random_state=0)
12 X = df.drop([0], axis=1)
13 y = df[0]
14 y = y-1 #so that classes are 0,1,2 instead of 1,2,3
15
16 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.6,
   random_state=0)
17 classifier = hmm.GaussianHMM(n_components=NUM_CLASSES)
18 classifier.fit(X_train)
19
20 mapping = get_mapping(classifier, X_train, y_train)
21 y_pred = list(map(lambda hs : mapping[hs], classifier.predict(X_test)))
22
23 print("Confusion Matrix:")
24 print(confusion_matrix(y_test, y_pred))
25 print("-----")
26 print("-----")
27 print("Performance Evaluation:")
28 print(classification_report(y_test, y_pred))
29 print("-----")
30 print("-----")
31 print("Accuracy Score:")
32 print(accuracy_score(y_test, y_pred))
33
34 plt.figure()
35 ConfusionMatrixDisplay.from_predictions(y_test, y_pred)
36 plt.savefig('./no_param.png')
37
38 plt.figure()
39 save_learning_curve(filepath='./no_param_lc.png',
   base_classifier=classifier, title="Learning curve for default
   Gaussian HMM", X=X, y=y)
40
```



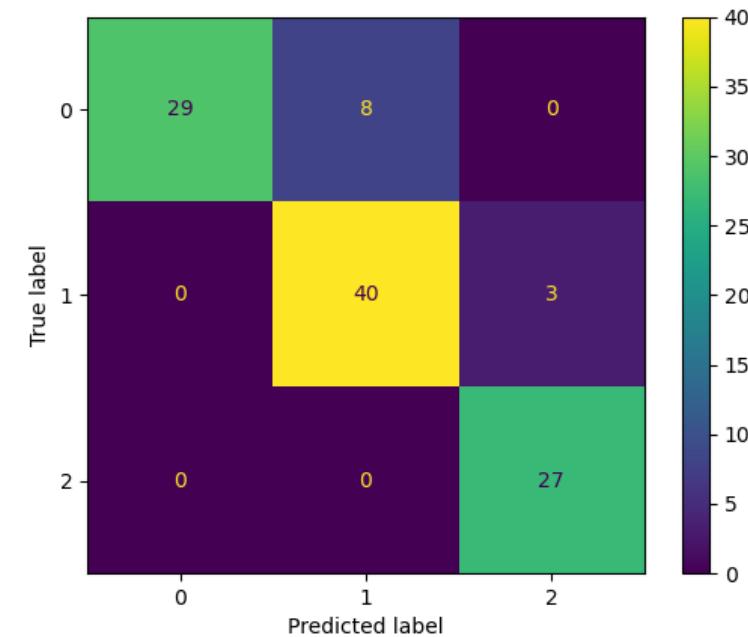
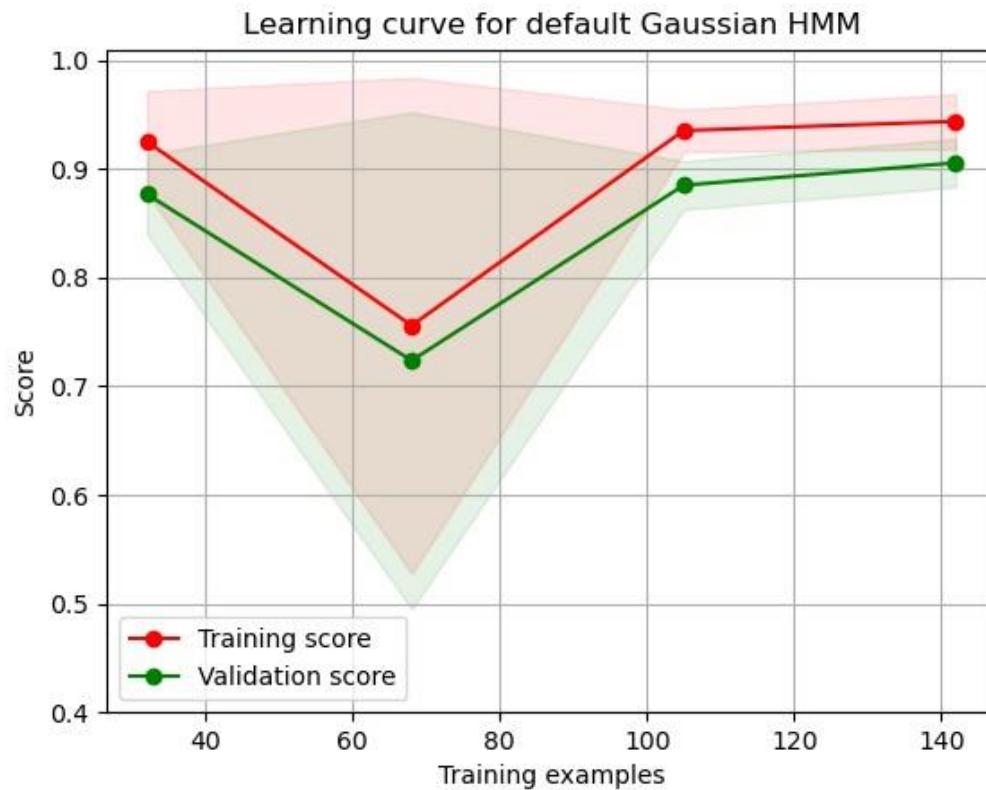
```
no_param.py
Confusion Matrix:
[[9 8 0]
 [0 40 3]
 [0 0 27]]
```

Performance Evaluation:

	precision	recall	f1-score	support
0	1.00	0.78	0.88	37
1	0.83	0.93	0.88	43
2	0.90	1.00	0.95	27
accuracy			0.90	107
macro avg	0.91	0.90	0.90	107
weighted avg	0.91	0.90	0.90	107

Accuracy Score:
.897196261682243

GAUSSIAN HMM (NO PARAMETER TUNING)- CONFUSION MATRIX AND LEARNING CURVE

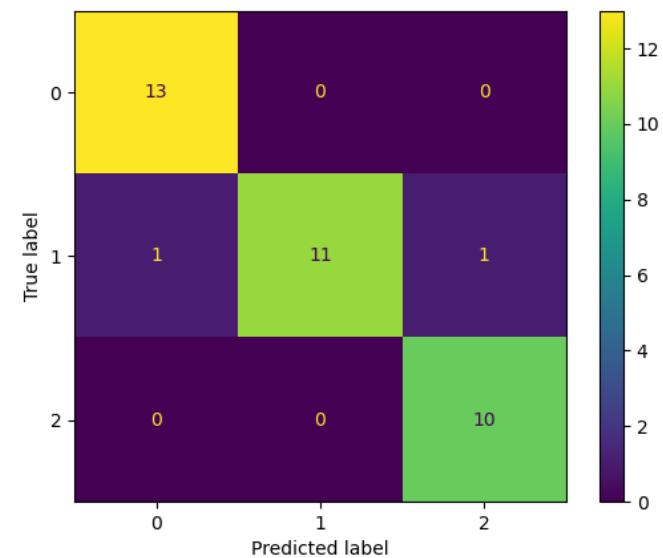
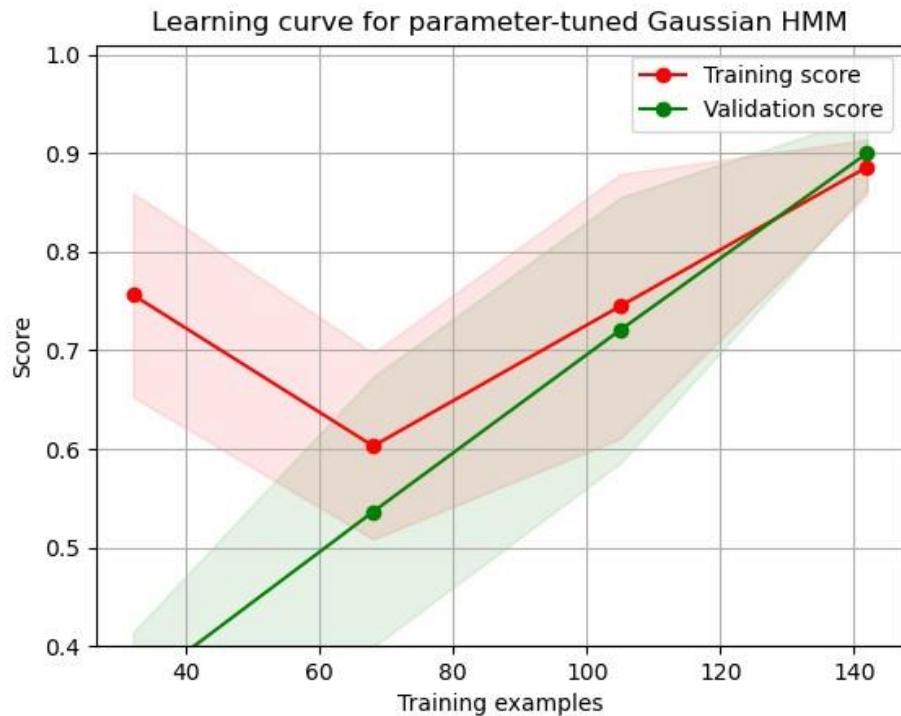


GAUSSIAN HMM (WITH PARAMETER TUNING)- CODE & EVALUATION

```
# gaussian_param > No Selection
1 import matplotlib.pyplot as plt
2 import pandas as pd
3 from sklearn.model_selection import train_test_split
4 from hmmlearn import hmm
5 from sklearn.metrics import classification_report, confusion_matrix,
   accuracy_score, ConfusionMatrixDisplay
6 from helper_functions import *
7
8 NUM_CLASSES = 3
9
10 df = pd.read_csv('./wine.data', header=None)
11 df = df.sample(frac = 1, random_state=0)
12 X = df.drop([0], axis=1)
13 y = df[0]
14 y = y-1 #so that classes are 0,1,2 instead of 1,2,3
15
16 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
   random_state=0)
17 classifier = hmm.GaussianHMM(n_components=NUM_CLASSES,
   covariance_type="full", min_covar=0.001, tol=0.01,
   implementation="log", random_state=0)
18 classifier.fit(X_train)
19
20 mapping = get_mapping(classifier, X_train, y_train)
21 y_pred = list(map(lambda hs : mapping[hs], classifier.predict(X_test)))
22
23 print("Confusion Matrix:")
24 print(confusion_matrix(y_test, y_pred))
25 print("-----")
26 print("-----")
27 print("Performance Evaluation:")
28 print(classification_report(y_test, y_pred))
29 print("-----")
30 print("-----")
31 print("Accuracy Score:")
32 print(accuracy_score(y_test, y_pred))
33
34 plt.figure()
35 ConfusionMatrixDisplay.from_predictions(y_test, y_pred)
36 plt.savefig('./param.png')
37
38 plt.figure()
39 save_learning_curve(filepath='./param_lc.png',
   base_classifier=classifier, title="Learning curve for parameter-tuned
   Gaussian HMM", X=X, y=y)
40
```

```
confusion Matrix:
[[3  0  0]
 [1 11  1]
 [0  0 10]]
-----
performance Evaluation:
precision      recall    f1-score   support
          0       0.93      1.00      0.96      13
          1       1.00      0.85      0.92      13
          2       0.91      1.00      0.95      10
                                               accuracy           0.94      36
                                               macro avg       0.95      0.94      36
                                               weighted avg    0.95      0.94      36
-----
accuracy Score:
.9444444444444444
```

GAUSSIAN HMM (WITH PARAMETER TUNING)- CONFUSION MATRIX AND LEARNING CURVE



HMM WITH GAUSSIAN MIXTURE MODEL EMISSIONS



GMM HMM (NO PARAMETER TUNING)- CODE & EVALUATION

```
gmm_hmm_no_param > No Selection
1 import matplotlib.pyplot as plt
2 import pandas as pd
3 from sklearn.model_selection import train_test_split
4 from hmmlearn import hmm
5 from sklearn.metrics import classification_report, confusion_matrix,
   accuracy_score, ConfusionMatrixDisplay
6 from helper_functions import *
7
8 NUM_CLASSES = 3
9
10 df = pd.read_csv('./wine.data', header=None)
11 df = df.sample(frac = 1, random_state=0)
12 X = df.drop([0], axis=1)
13 y = df[0]
14 y = y-1 #so that classes are 0,1,2 instead of 1,2,3
15
16 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.6,
   random_state=0)
17 classifier = hmm.GMMHMM(n_components=NUM_CLASSES)
18 classifier.fit(X_train)
19
20 mapping = get_mapping(classifier, X_train, y_train)
21 y_pred = list(map(lambda hs : mapping[hs], classifier.predict(X_test)))
22
23 print("Confusion Matrix:")
24 print(confusion_matrix(y_test, y_pred))
25 print("-----")
26 print("-----")
27 print("Performance Evaluation:")
28 print(classification_report(y_test, y_pred))
29 print("-----")
30 print("-----")
31 print("Accuracy Score:")
32 print(accuracy_score(y_test, y_pred))
33
34 plt.figure()
35 ConfusionMatrixDisplay.from_predictions(y_test, y_pred)
36 plt.savefig('./no_param.png')
37
38 plt.figure()
39 save_learning_curve(filepath='./no_param_lc.png',
   base_classifier=classifier, title="Learning curve for default GMM
   HMM", X=X, y=y)
40
```

Confusion Matrix:

```
[[29  8  0]
 [ 0 40  3]
 [ 0  0 27]]
```

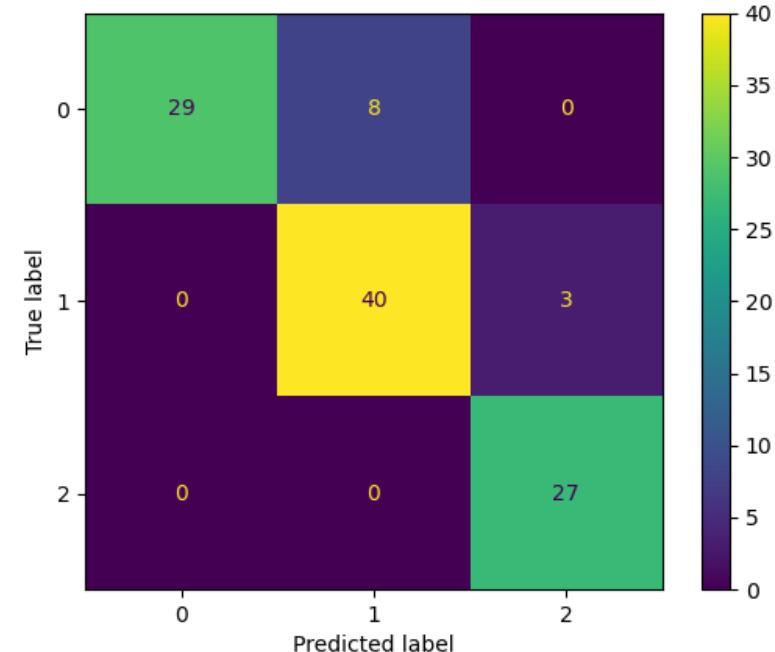
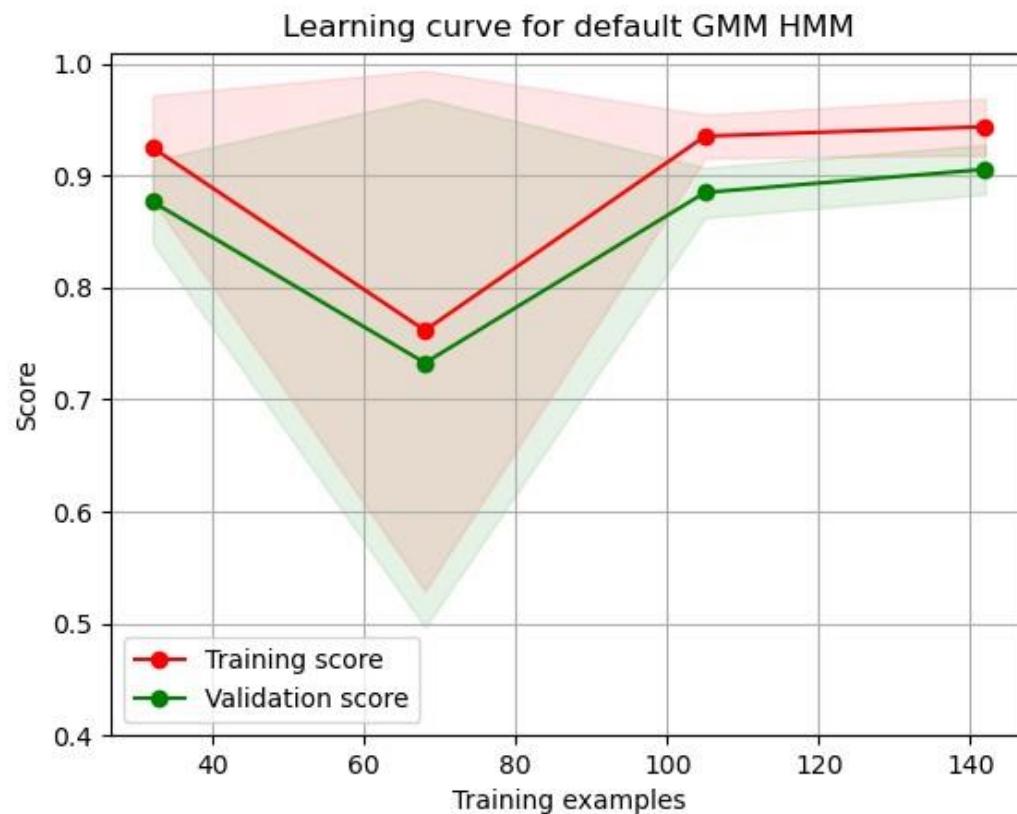
Performance Evaluation:

	precision	recall	f1-score	support
0	1.00	0.78	0.88	37
1	0.83	0.93	0.88	43
2	0.90	1.00	0.95	27
accuracy			0.90	107
macro avg	0.91	0.90	0.90	107
weighted avg	0.91	0.90	0.90	107

Accuracy Score:

```
0.897196261682243
```

GMM HMM (NO PARAMETER TUNING)- CONFUSION MATRIX AND LEARNING CURVE



GMM HMM (WITH PARAMETER TUNING)- CODE & EVALUATION

```
# gmm_hmm_param / no selection
1 import matplotlib.pyplot as plt
2 import pandas as pd
3 from sklearn.model_selection import train_test_split
4 from hmmlearn import hmm
5 from sklearn.metrics import classification_report, confusion_matrix,
   accuracy_score, ConfusionMatrixDisplay
6 from helper_functions import *
7
8 NUM_CLASSES = 3
9
10 df = pd.read_csv('../wine.data', header=None)
11 df = df.sample(frac = 1, random_state=0)
12 X = df.drop([0], axis=1)
13 y = df[0]
14 y = y-1 #so that classes are 0,1,2 instead of 1,2,3
15
16 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
   random_state=0)
17 classifier = hmm.GMMHMM(n_components=3, covariance_type="diag", n_mix=2,
   min_covar=0.01, tol=0.01, implementation="log", random_state=1)
18 classifier.fit(X_train)
19
20 mapping = get_mapping(classifier, X_train, y_train)
21 y_pred = list(map(lambda hs : mapping[hs], classifier.predict(X_test)))
22
23 print("Confusion Matrix:")
24 print(confusion_matrix(y_test, y_pred))
25 print("-----")
26 print("Performance Evaluation:")
27 print(classification_report(y_test, y_pred))
28 print("-----")
29 print("Accuracy Score:")
30 print(accuracy_score(y_test, y_pred))
31
32 plt.figure()
33 ConfusionMatrixDisplay.from_predictions(y_test, y_pred)
34 plt.savefig('../param.png')
35
36 plt.figure()
37 save_learning_curve(filepath='../param_lc.png',
   base_classifier=classifier, title="Learning curve for parameter-tuned
   GMM HMM", X=X, y=y)
38
39
```

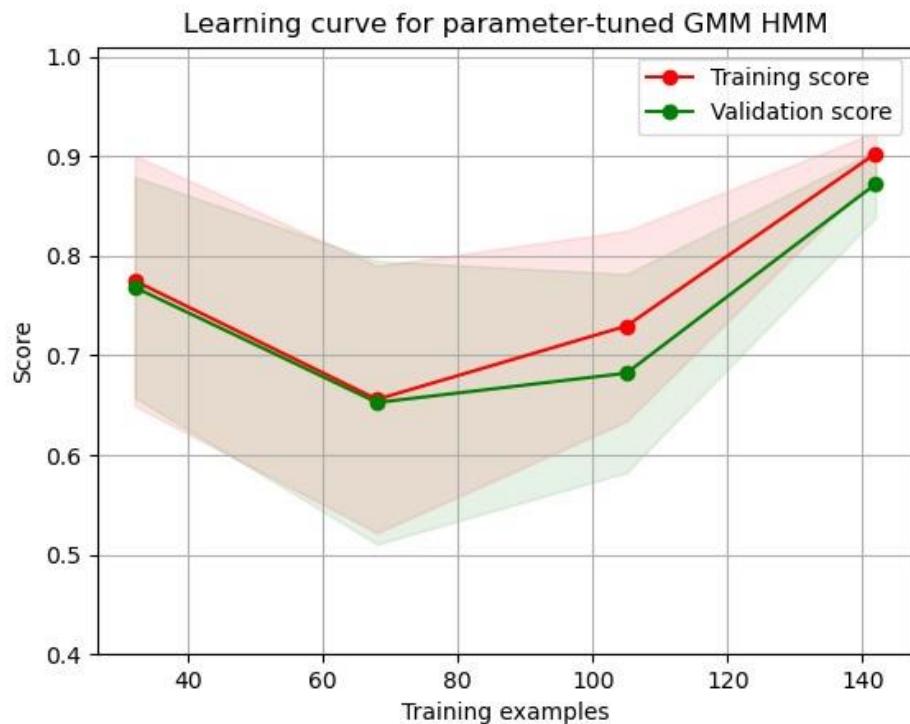
Confusion Matrix:
[[13 0 0]
 [1 10 2]
 [0 0 10]]

Performance Evaluation:

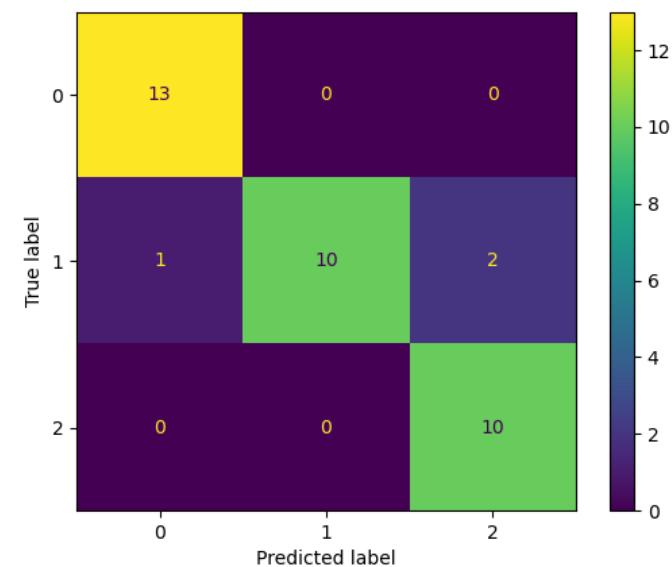
	precision	recall	f1-score	support
0	0.93	1.00	0.96	13
1	1.00	0.77	0.87	13
2	0.83	1.00	0.91	10
accuracy			0.92	36
macro avg	0.92	0.92	0.91	36
weighted avg	0.93	0.92	0.91	36

accuracy Score:
0.9166666666666666

GMM HMM (WITH PARAMETER TUNING)- CONFUSION MATRIX AND LEARNING CURVE

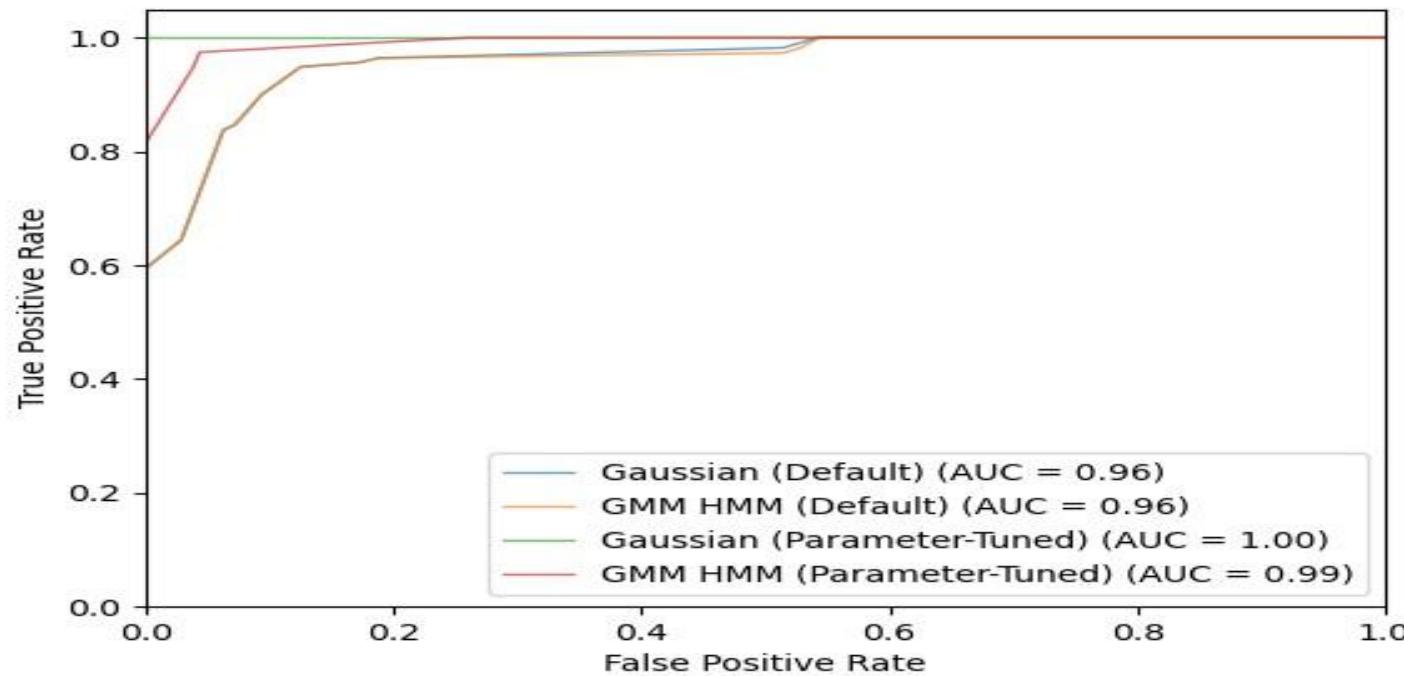


CONFUSION MATRIX
HeatMap



COMPARISON BETWEEN GAUSSIAN HMM AND GMM HMM.

ROC CURVES AND AUC COMPARISON



ROC CURVES AND AUC COMPARISON - CODE

```
make_roc_auc_curves > No Selection
1 import matplotlib.pyplot as plt
2 import pandas as pd
3 from sklearn.model_selection import train_test_split
4 from hmmlearn import hmm
5 from sklearn.metrics import classification_report, confusion_matrix, accuracy_score, ConfusionMatrixDisplay
6 from helper_functions import *
7
8 NUM_CLASSES = 3
9
10 df = pd.read_csv('../wine.data', header=None)
11 df = df.sample(frac = 1, random_state=0)
12 X = df.drop([0], axis=1)
13 y = df[0]
14 y = y-1 #so that classes are 0,1,2 instead of 1,2,3
15
16
17 classifier = dict()
18 classifier['gaussian_no_param'] = hmm.GaussianHMM(n_components=NUM_CLASSES)
19 classifier['gaussian_param'] = hmm.GaussianHMM(n_components=NUM_CLASSES, covariance_type="full",
19     min_covar=0.001, tol=0.01, implementation="log", random_state=0)
20
21 classifier['gmm_hmm_no_param'] = hmm.GMMHMM(n_components=NUM_CLASSES)
22 classifier['gmm_hmm_param'] = hmm.GMMHMM(n_components=3, covariance_type="diag", n_mix=2, min_covar=0.01,
22     tol=0.01, implementation="log", random_state=1)
23
24 #division for default Gaussian and GMM HMM models, since they perform best in this partition
25 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.6, random_state=0)
26
27
28 plot_macro_avg_roc_curve(classifier['gaussian_no_param'], n_classes=NUM_CLASSES, name="Gaussian
28     (Default)", X_train=X_train, X_test=X_test, y_train=y_train, y_test=y_test)
29 plot_macro_avg_roc_curve(classifier['gmm_hmm_no_param'], n_classes=NUM_CLASSES, name="GMM HMM (Default)",
29     X_train=X_train, X_test=X_test, y_train=y_train, y_test=y_test)
30
31 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=0)
32
33 plot_macro_avg_roc_curve(classifier['gaussian_param'], n_classes=NUM_CLASSES, name="Gaussian
33     (Parameter-Tuned)", X_train=X_train, X_test=X_test, y_train=y_train, y_test=y_test)
34 plot_macro_avg_roc_curve(classifier['gmm_hmm_param'], n_classes=NUM_CLASSES, name="GMM HMM
34     (Parameter-Tuned)", X_train=X_train, X_test=X_test, y_train=y_train, y_test=y_test)
35
36 plt.xlim([0.0, 1.0])
37 plt.ylim([0.0, 1.05])
38 plt.xlabel("False Positive Rate")
39 plt.ylabel("True Positive Rate")
40 plt.legend(loc="lower right")
41
42 plt.savefig('roc_auc.png')
```

Line: 1 Col: 1 |

COMPARISON OF PERFORMANCE

CLASSIFIER		Accuracy	P ec	Recall	F -Sco	AUC
GAUSSI AN HMM	DEFAULT	0.93	0.93	0.93	0.93	0.96
	TUNED	0.94	0.95	0.94	0.94	.00
GMM HMM	DEFAULT	0.90	0.9	0.90	0.90	0.96
	TUNED	0.92	0.93	0.92	0.9	0.99

IONOSPHERE - CLASSIFICATION



GAUSSIAN HMM



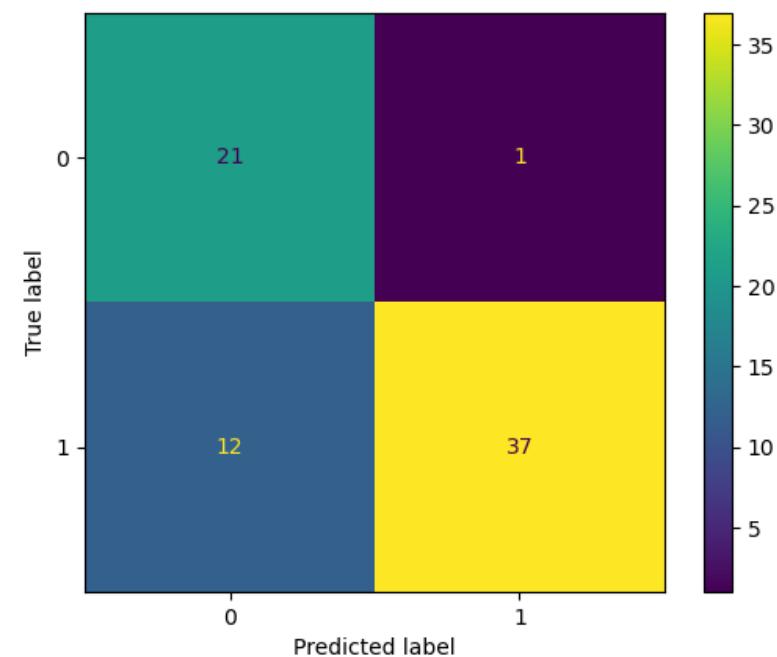
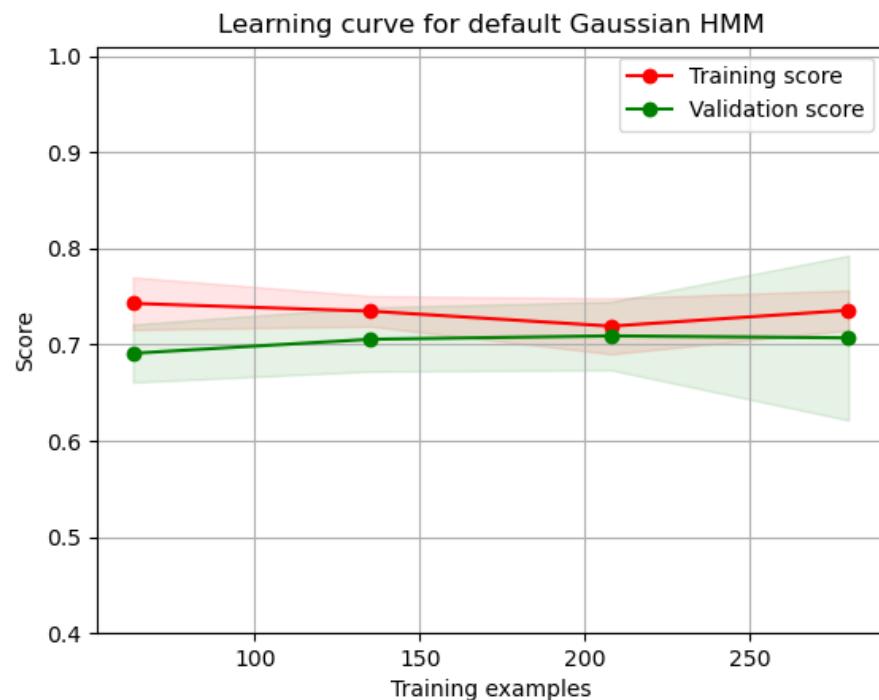
GAUSSIAN HMM (NO PARAMETER TUNING)- CODE & EVALUATION

```
gaussian_no_param() No Selection
1 import matplotlib.pyplot as plt
2 import pandas as pd
3 from sklearn.model_selection import train_test_split
4 from hmmlearn import hmm
5 from sklearn.metrics import classification_report, confusion_matrix,
   accuracy_score, ConfusionMatrixDisplay
6 from helper_functions import *
7
8 NUM_CLASSES = 2
9
10 df = pd.read_csv('./ionosphere.data', header=None)
11 df = df.sample(frac = 1, random_state=0)
12 X = df.drop([34], axis=1)
13 y = df[34]
14 y = y.map({'g':1, 'b':0})
15
16 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
   random_state=3)
17 classifier = hmm.GaussianHMM(n_components=NUM_CLASSES)
18 classifier.fit(X_train)
19
20 mapping = get_mapping(classifier, X_train, y_train)
21 y_pred = list(map(lambda hs : mapping[hs], classifier.predict(X_test)))
22
23 print("Confusion Matrix:")
24 print(confusion_matrix(y_test, y_pred))
25 print("-----")
26 print("-----")
27 print("Performance Evaluation:")
28 print(classification_report(y_test, y_pred))
29 print("-----")
30 print("-----")
31 print("Accuracy Score:")
32 print(accuracy_score(y_test, y_pred))
33
34 plt.figure()
35 ConfusionMatrixDisplay.from_predictions(y_test, y_pred)
36 plt.savefig('./no_param.png')
37
38 plt.figure()
39 save_learning_curve(filepath='./no_param_lc.png',
   base_classifier=classifier, title="Learning curve for default
   Gaussian HMM", X=X, y=y)
40
```

```
confusion matrix:
[[21  1]
 [12 37]]
-----
performance evaluation:
      precision    recall  f1-score   support
0          0.64     0.95     0.76      22
1          0.97     0.76     0.85      49

accuracy                           0.82      71
macro avg       0.81     0.85     0.81      71
weighted avg    0.87     0.82     0.82      71
-----
accuracy score:
0.8169014084507842
```

GAUSSIAN HMM (NO PARAMETER TUNING)- CONFUSION MATRIX AND

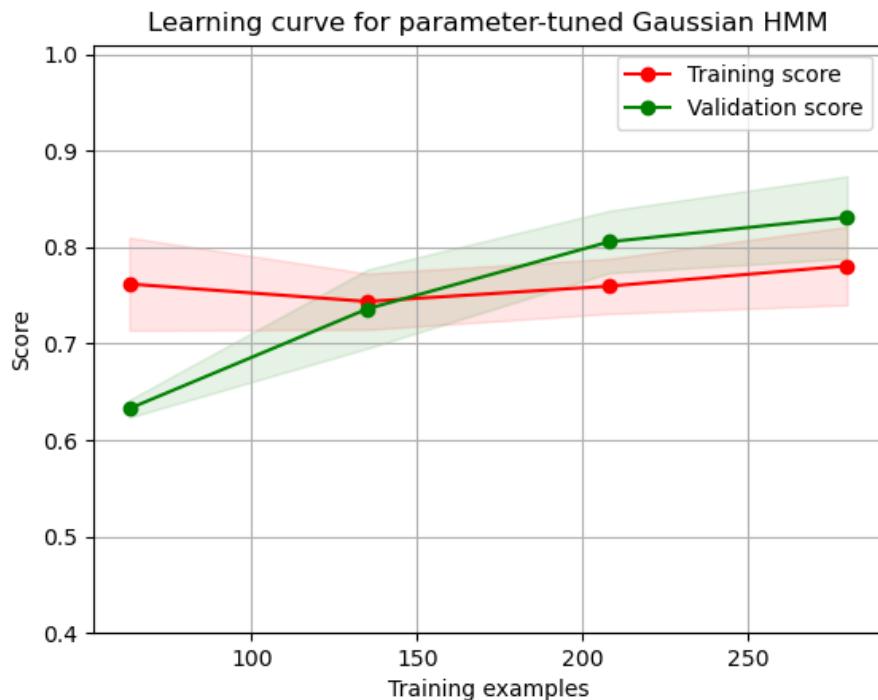


GAUSSIAN HMM (WITH PARAMETER TUNING)- CODE & EVALUATION

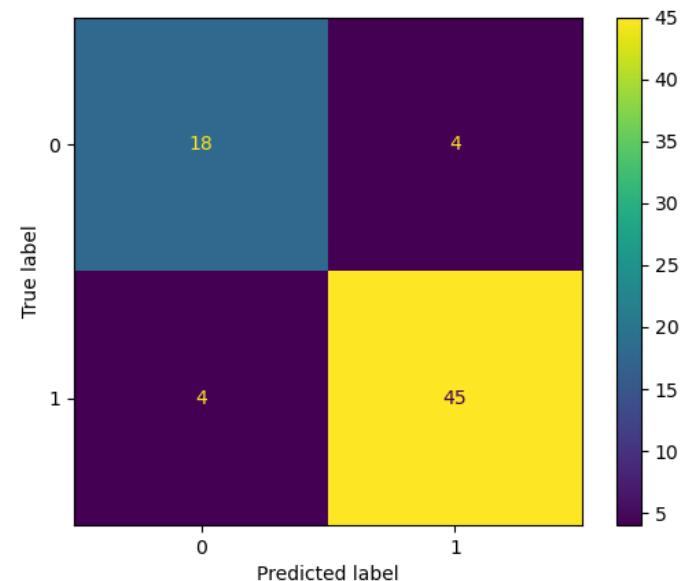
```
gaussian_param / test.ipynb
1 import matplotlib.pyplot as plt
2 import pandas as pd
3 from sklearn.model_selection import train_test_split
4 from hmmlearn import hmm
5 from sklearn.metrics import classification_report, confusion_matrix,
   accuracy_score, ConfusionMatrixDisplay
6 from helper_functions import *
7
8 NUM_CLASSES = 2
9
10 df = pd.read_csv('./ionosphere.data', header=None)
11 df = df.sample(frac = 1, random_state=0)
12 X = df.drop([34], axis=1)
13 y = df[34]
14 y = y.map({'g':1, 'b':0})
15
16 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
   random_state=3)
17 classifier = hmm.GaussianHMM(n_components=NUM_CLASSES,
   covariance_type="full", min_covar=0.01, tol=0.01,
   implementation="log", random_state=0)
18 classifier.fit(X_train)
19
20 mapping = get_mapping(classifier, X_train, y_train)
21 y_pred = list(map(lambda hs : mapping[hs], classifier.predict(X_test)))
22
23 print("Confusion Matrix:")
24 print(confusion_matrix(y_test, y_pred))
25 print("-----")
26 print("-----")
27 print("Performance Evaluation:")
28 print(classification_report(y_test, y_pred))
29 print("-----")
30 print("-----")
31 print("Accuracy Score:")
32 print(accuracy_score(y_test, y_pred))
33
34 plt.figure()
35 ConfusionMatrixDisplay.from_predictions(y_test, y_pred)
36 plt.savefig('./param.png')
37
38 plt.figure()
39 save_learning_curve(filepath='./param_lc.png',
   base_classifier=classifier, title="Learning curve for parameter-tuned
   Gaussian HMM", X=X, y=y)
40
```

```
-----,
Confusion Matrix:
[[18  4]
 [ 4 46]]
-----
Performance Evaluation:
precision    recall  f1-score   support
          0       0.82      0.82      0.82      22
          1       0.92      0.92      0.92      49
   accuracy                           0.89      71
  macro avg       0.87      0.87      0.87      71
weighted avg       0.89      0.89      0.89      71
-----
Accuracy Score:
0.8873239436619719
```

GAUSSIAN HMM (WITH PARAMETER TUNING)- CONFUSION MATRIX AND LEARNING CURVE



CONFUSION MATRIX
HeatMap



HMM WITH GAUSSIAN MIXTURE MODEL EMISSIONS



GMM HMM (NO PARAMETER TUNING)- CODE & EVALUATION

```
 1 import matplotlib.pyplot as plt
 2 import pandas as pd
 3 from sklearn.model_selection import train_test_split
 4 from hmmlearn import hmm
 5 from sklearn.metrics import classification_report, confusion_matrix,
 6     accuracy_score, ConfusionMatrixDisplay
 7 from helper_functions import *
 8
 9 NUM_CLASSES = 2
10
11 df = pd.read_csv('./ionosphere.data', header=None)
12 df = df.sample(frac = 1, random_state=0)
13 X = df.drop([34], axis=1)
14 y = df[34]
15 y = y.map({'g':1, 'b':0})
16
17 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.8,
18     random_state=3)
19 classifier = hmm.GMMHMM(n_components=NUM_CLASSES)
20 classifier.fit(X_train)
21
22 mapping = get_mapping(classifier, X_train, y_train)
23 y_pred = list(map(lambda hs : mapping[hs], classifier.predict(X_test)))
24
25 print("Confusion Matrix:")
26 print(confusion_matrix(y_test, y_pred))
27 print("-----")
28 print("Performance Evaluation:")
29 print(classification_report(y_test, y_pred))
30 print("-----")
31 print("Accuracy Score:")
32 print(accuracy_score(y_test, y_pred))
33
34 plt.figure()
35 ConfusionMatrixDisplay.from_predictions(y_test, y_pred)
36 plt.savefig('./no_param.png')
37
38 plt.figure()
39 save_learning_curve(filepath='./no_param_lc.png',
40     base_classifier=classifier, title="Learning curve for default GMM
41     HMM", X=X, y=y)
```

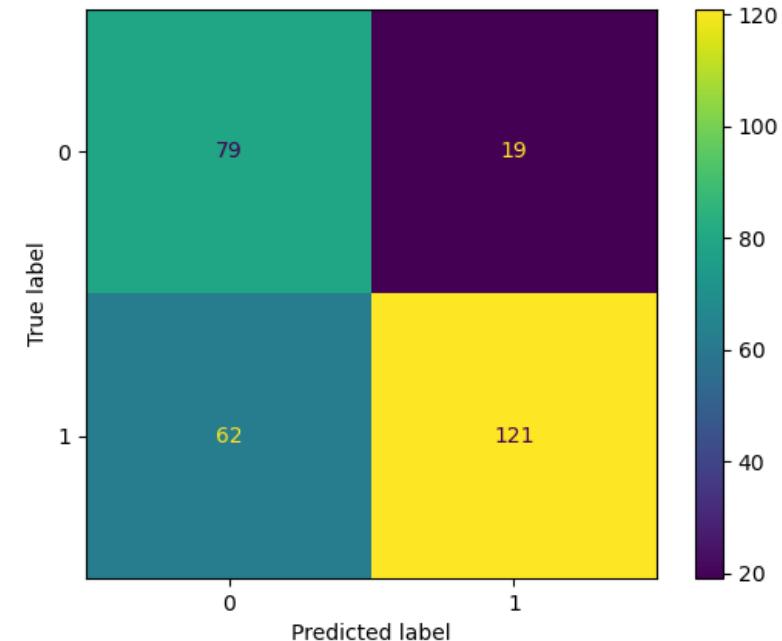
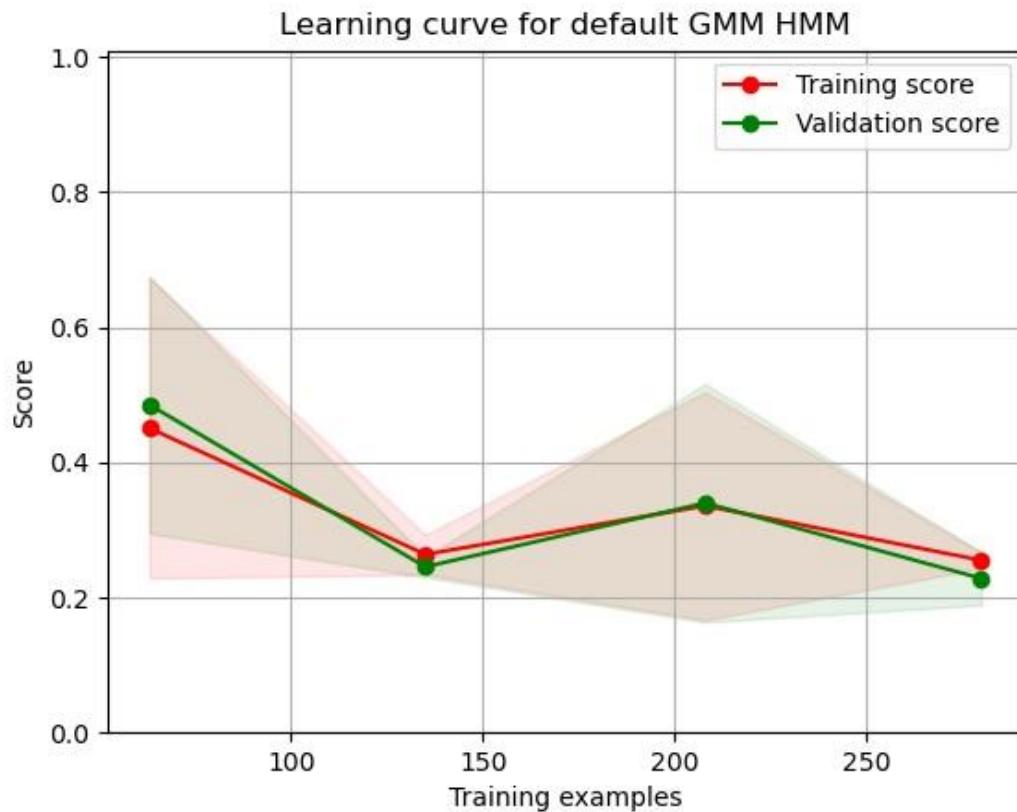
Confusion Matrix:
[[79 19]
 [62 121]]

Performance Evaluation:

	precision	recall	f1-score	support
0	0.56	0.81	0.66	98
1	0.86	0.66	0.75	183
accuracy			0.71	281
macro avg	0.71	0.73	0.71	281
weighted avg	0.76	0.71	0.72	281

Accuracy Score:
0.7117437722419929

GMM HMM (NO PARAMETER TUNING)- CONFUSION MATRIX AND LEARNING CURVE



GMM HMM (WITH PARAMETER TUNING)- CODE & EVALUATION

```
 1 import matplotlib.pyplot as plt
 2 import pandas as pd
 3 from sklearn.model_selection import train_test_split
 4 from hmmlearn import hmm
 5 from sklearn.metrics import classification_report, confusion_matrix,
 6     accuracy_score, ConfusionMatrixDisplay
 7 from helper_functions import *
 8
 9 NUM_CLASSES = 2
10
11 df = pd.read_csv('./ionosphere.data', header=None)
12 df = df.sample(frac = 1, random_state=0)
13 X = df.drop([34], axis=1)
14 y = df[34]
15 y = y.map({'g':1, 'b':0})
16
17 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
18     random_state=3)
19 classifier = hmm.GMMHMM(n_components=NUM_CLASSES,
20     covariance_type="spherical", n_mix=3, min_covar=0.001, tol=0.01,
21     implementation="log", random_state=1)
22 classifier.fit(X_train)
23
24 mapping = get_mapping(classifier, X_train, y_train)
25 y_pred = list(map(lambda hs : mapping[hs], classifier.predict(X_test)))
26
27 print("Confusion Matrix:")
28 print(confusion_matrix(y_test, y_pred))
29 print("-----")
30 print("Performance Evaluation:")
31 print(classification_report(y_test, y_pred))
32 print("-----")
33 print("Accuracy Score:")
34 print(accuracy_score(y_test, y_pred))
35
36 plt.figure()
37 ConfusionMatrixDisplay.from_predictions(y_test, y_pred)
38 plt.savefig('./param.png')
39
40 plt.figure()
41 save_learning_curve(filepath='./param_lc.png',
42     base_classifier=classifier, title="Learning curve for parameter-tuned
43     GMM HMM", X=X, y=y)
```

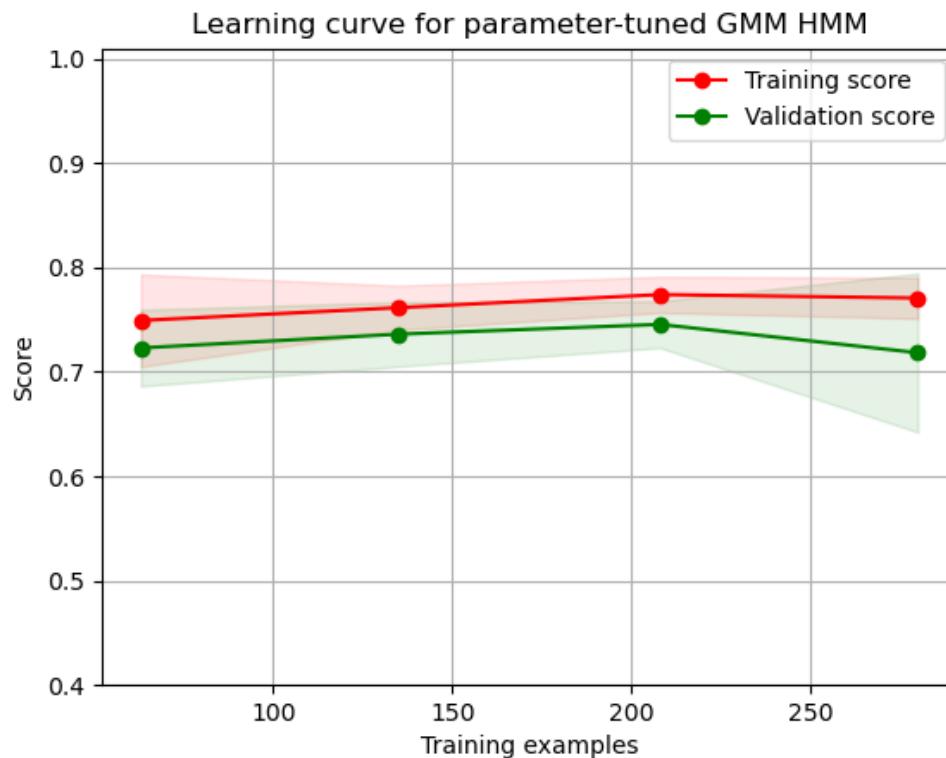
Confusion Matrix:
[[20 2]
 [11 38]]

Performance Evaluation:

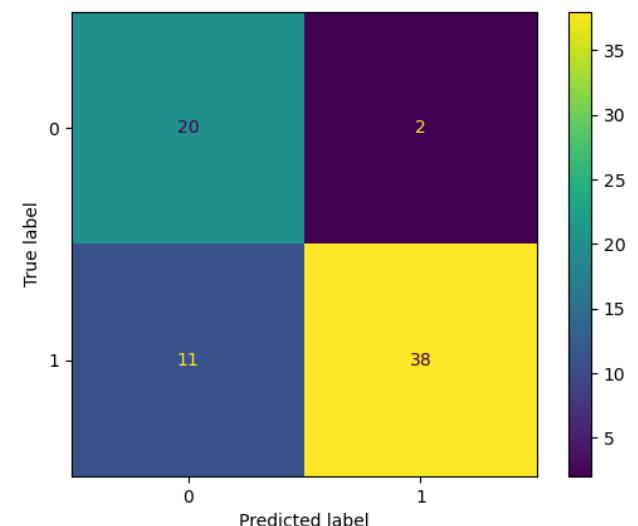
	precision	recall	f1-score	support
0	0.65	0.91	0.75	22
1	0.95	0.78	0.85	49
accuracy			0.82	71
macro avg	0.80	0.84	0.80	71
weighted avg	0.86	0.82	0.82	71

Accuracy Score:
0.8169014084507042

GMM HMM (WITH PARAMETER TUNING)- CONFUSION MATRIX AND LEARNING CURVE

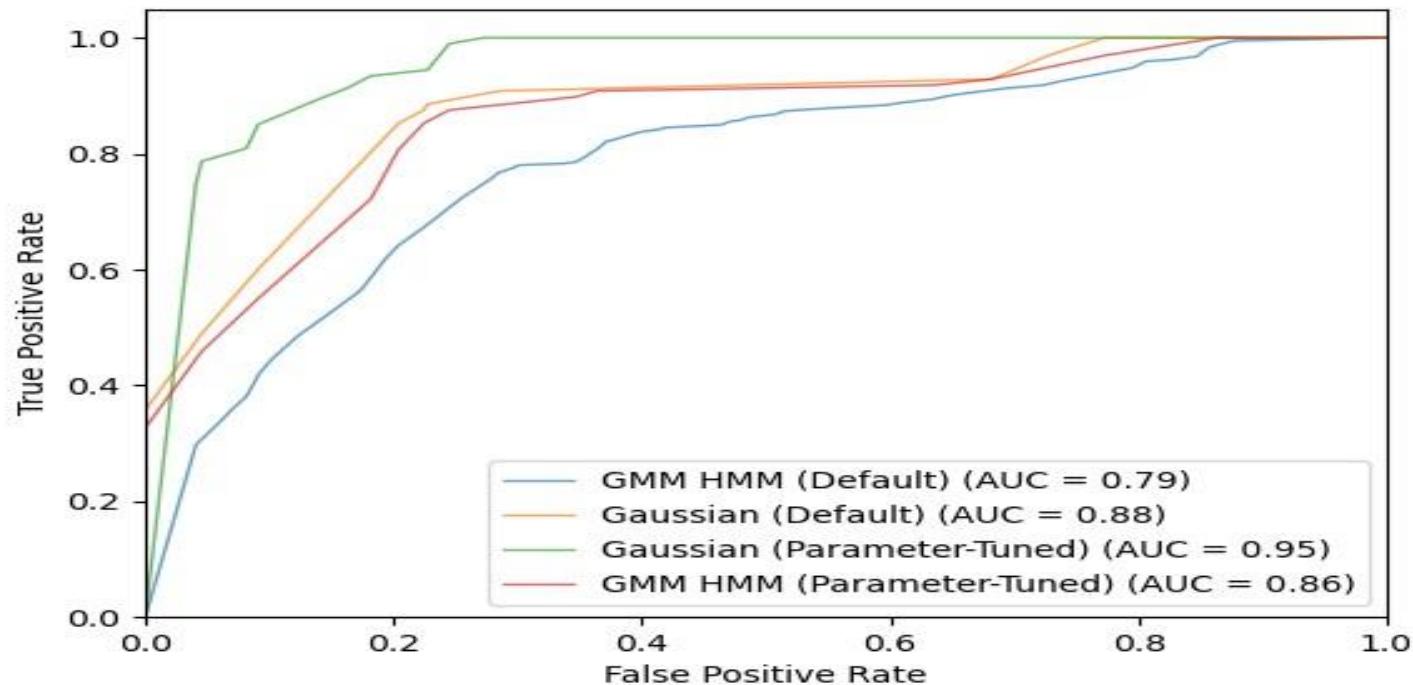


CONFUSION MATRIX
HeatMap

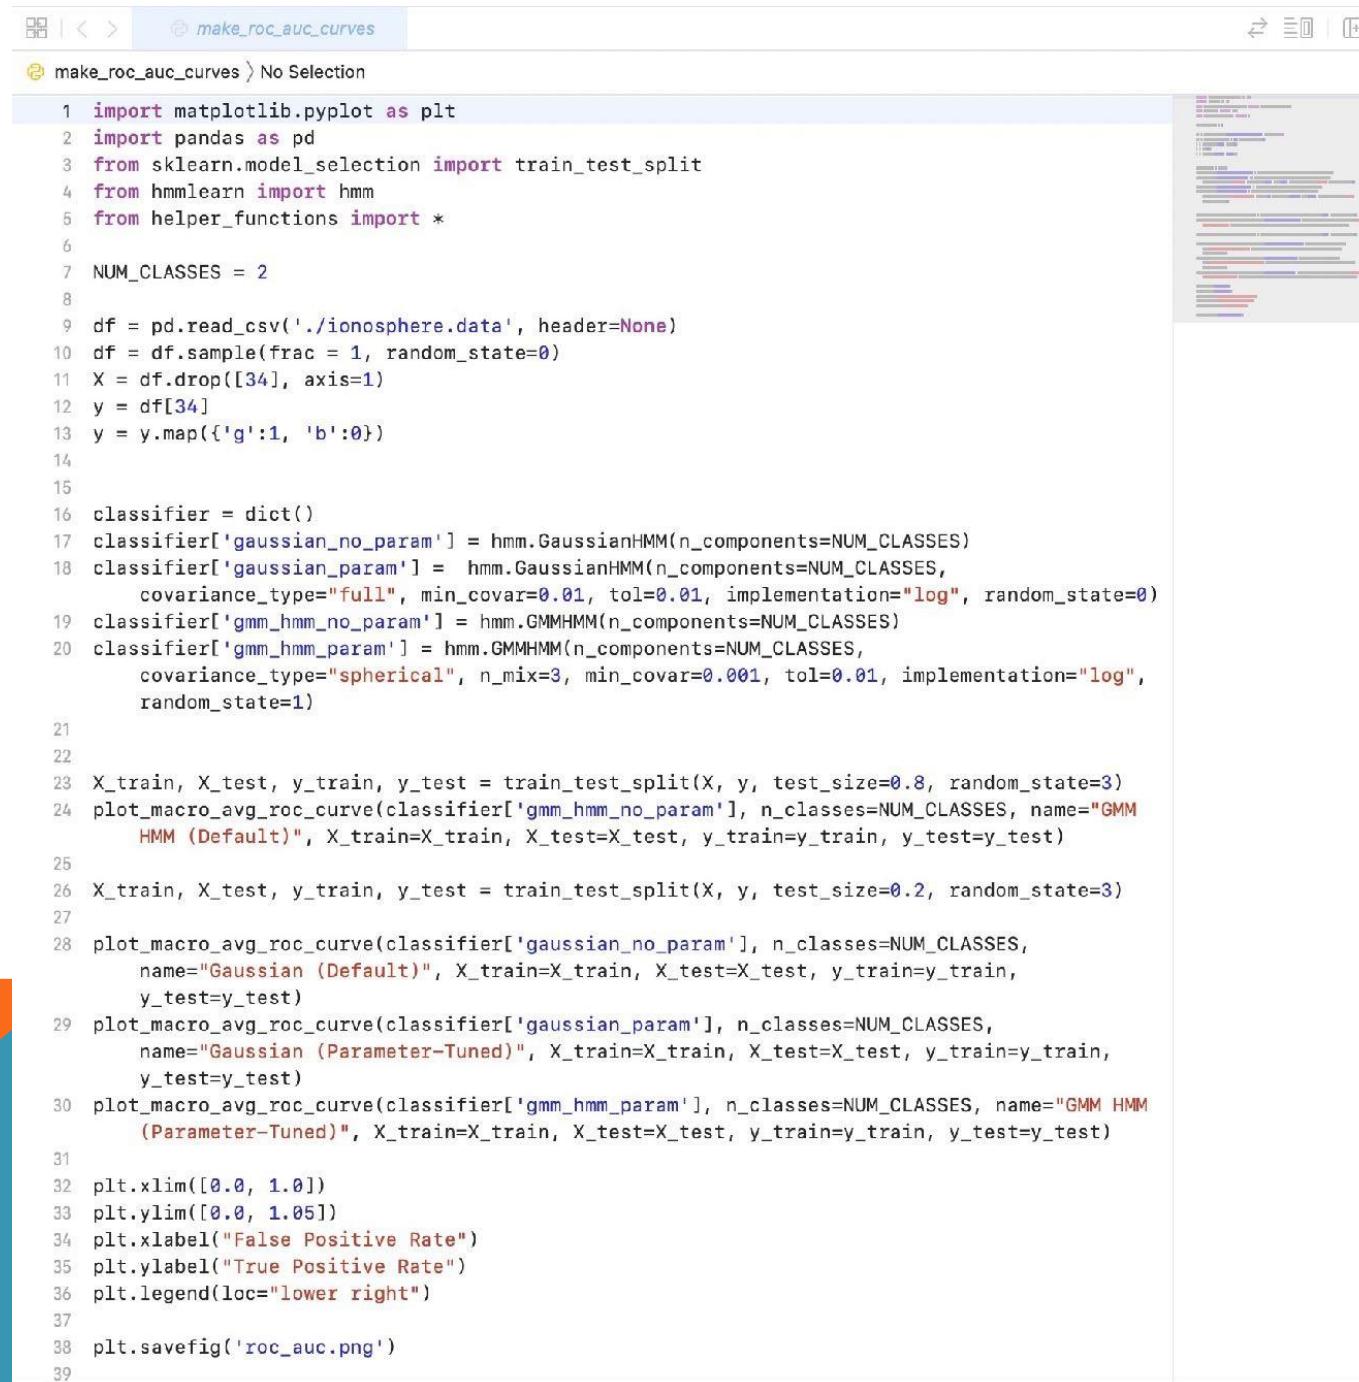


COMPARISON BETWEEN GAUSSIAN HMM AND GMM HMM.

ROC CURVES AND AUC COMPARISON



ROC CURVES AND AUC COMPARISON - CODE



```
make_roc_auc_curves
```

```
make_roc_auc_curves > No Selection
```

```
1 import matplotlib.pyplot as plt
2 import pandas as pd
3 from sklearn.model_selection import train_test_split
4 from hmmlearn import hmm
5 from helper_functions import *
6
7 NUM_CLASSES = 2
8
9 df = pd.read_csv('./ionosphere.data', header=None)
10 df = df.sample(frac = 1, random_state=0)
11 X = df.drop([34], axis=1)
12 y = df[34]
13 y = y.map({'g':1, 'b':0})
14
15
16 classifier = dict()
17 classifier['gaussian_no_param'] = hmm.GaussianHMM(n_components=NUM_CLASSES)
18 classifier['gaussian_param'] = hmm.GaussianHMM(n_components=NUM_CLASSES,
19     covariance_type="full", min_covar=0.01, tol=0.01, implementation="log", random_state=0)
20 classifier['gmm_hmm_no_param'] = hmm.GMMHMM(n_components=NUM_CLASSES)
21 classifier['gmm_hmm_param'] = hmm.GMMHMM(n_components=NUM_CLASSES,
22     covariance_type="spherical", n_mix=3, min_covar=0.001, tol=0.01, implementation="log",
23     random_state=1)
24
25
26 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.8, random_state=3)
27 plot_macro_avg_roc_curve(classifier['gmm_hmm_no_param'], n_classes=NUM_CLASSES, name="GMM
28 HMM (Default)", X_train=X_train, X_test=X_test, y_train=y_train, y_test=y_test)
29
30 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=3)
31
32 plot_macro_avg_roc_curve(classifier['gaussian_no_param'], n_classes=NUM_CLASSES,
33     name="Gaussian (Default)", X_train=X_train, X_test=X_test, y_train=y_train,
34     y_test=y_test)
35 plot_macro_avg_roc_curve(classifier['gaussian_param'], n_classes=NUM_CLASSES,
36     name="Gaussian (Parameter-Tuned)", X_train=X_train, X_test=X_test, y_train=y_train,
37     y_test=y_test)
38 plot_macro_avg_roc_curve(classifier['gmm_hmm_param'], n_classes=NUM_CLASSES, name="GMM HMM
39 (Parameter-Tuned)", X_train=X_train, X_test=X_test, y_train=y_train, y_test=y_test)
40
41 plt.xlim([0.0, 1.0])
42 plt.ylim([0.0, 1.05])
43 plt.xlabel("False Positive Rate")
44 plt.ylabel("True Positive Rate")
45 plt.legend(loc="lower right")
46
47
48 plt.savefig('roc_auc.png')
```

COMPARISON OF PERFORMANCE

CLASSIFIER		Accuracy	P ec	Recall	F -Sco	AUC
GAUSSIAN HMM	DEFAULT	0.82	0.87	0.82	0.82	0.88
	TUNED	0.89	0.89	0.89	0.89	0.95
GMM HMM	DEFAULT	0.7	0.76	0.7	0.72	0.79
	TUNED	0.82	0.86	0.82	0.82	0.86

WISCONSIN BREAST CANCER - CLASSIFICATION



GAUSSIAN HMM



GAUSSIAN HMM (NO PARAMETER TUNING)- CODE & EVALUATION

```
gaussian_no_param / no_selection
1 import matplotlib.pyplot as plt
2 from sklearn.datasets import load_breast_cancer
3 from sklearn.model_selection import train_test_split
4 from hmmlearn import hmm
5 from sklearn.metrics import classification_report, confusion_matrix,
   accuracy_score, ConfusionMatrixDisplay
6 from helper_functions import *
7
8 NUM_CLASSES = 2
9
10 X, y = load_breast_cancer(return_X_y=True)
11
12 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.6,
   random_state=0)
13 classifier = hmm.GaussianHMM(n_components=NUM_CLASSES)
14 classifier.fit(X_train)
15
16 mapping = get_mapping(classifier, X_train, y_train)
17 y_pred = list(map(lambda hs : mapping[hs], classifier.predict(X_test)))
18
19 print("Confusion Matrix:")
20 print(confusion_matrix(y_test, y_pred))
21 print("-----")
22 print("-----")
23 print("Performance Evaluation:")
24 print(classification_report(y_test, y_pred))
25 print("-----")
26 print("-----")
27 print("Accuracy Score:")
28 print(accuracy_score(y_test, y_pred))
29
30 plt.figure()
31 ConfusionMatrixDisplay.from_predictions(y_test, y_pred)
32 plt.savefig('./no_param.png')
33
34 plt.figure()
35 save_learning_curve(filepath='./no_param_lc.png',
   base_classifier=classifier, title="Learning curve for default
   Gaussian HMM", X=X, y=y)
36
```

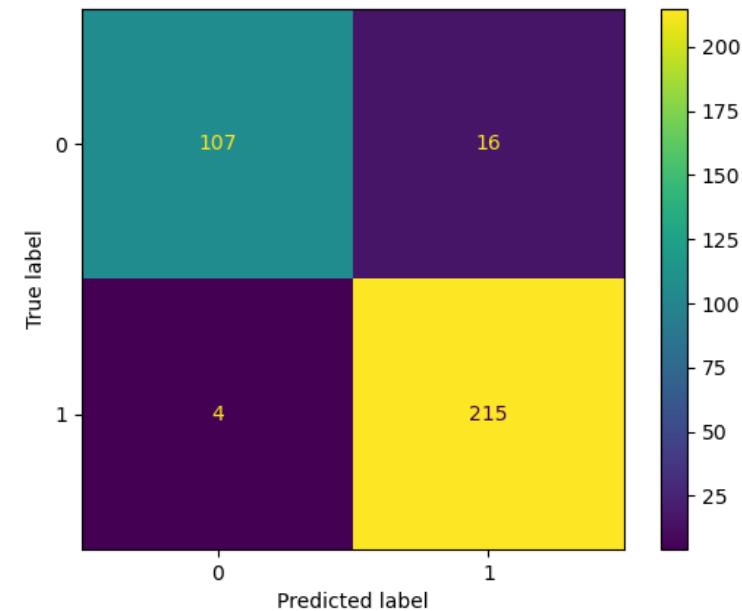
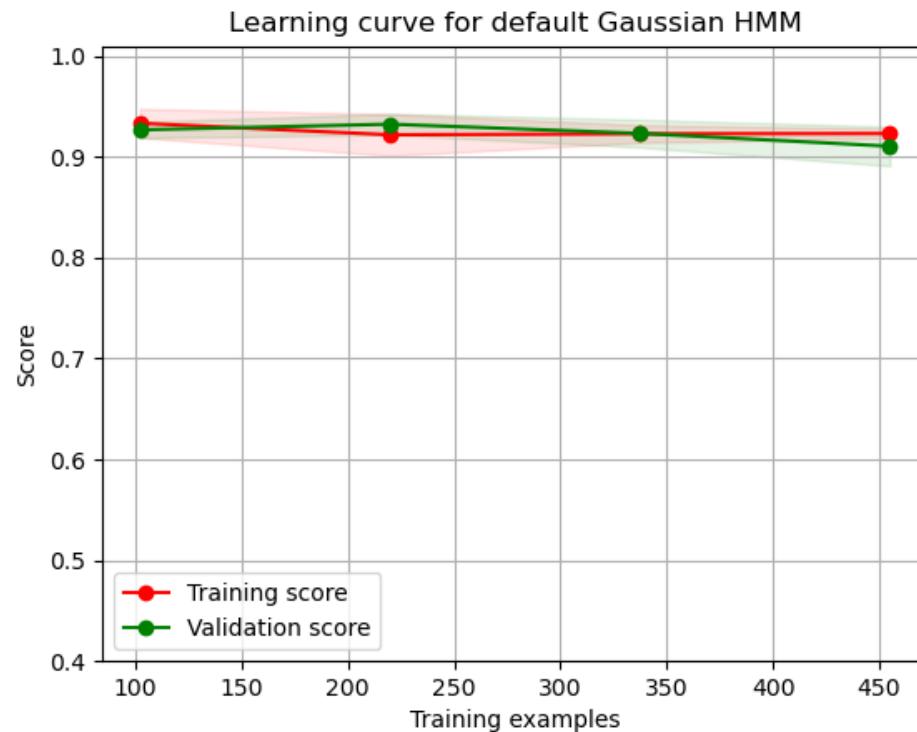
```
:confusion Matrix:
[[107 16]
 [ 4 215]]
```

```
:Performance Evaluation:
```

	precision	recall	f1-score	support
0	0.96	0.87	0.91	123
1	0.93	0.98	0.96	219
accuracy			0.94	342
macro avg	0.95	0.93	0.94	342
weighted avg	0.94	0.94	0.94	342

```
:accuracy Score:
0.9415204678362573
```

GAUSSIAN HMM (NO PARAMETER TUNING)- CONFUSION MATRIX AND LEARNING CURVE



GAUSSIAN HMM (WITH PARAMETER TUNING)- CODE & EVALUATION

```
gaussian_param > No Selection
1 import matplotlib.pyplot as plt
2 from sklearn.datasets import load_breast_cancer
3 from sklearn.model_selection import train_test_split
4 from hmmlearn import hmm
5 from sklearn.metrics import classification_report, confusion_matrix,
accuracy_score, ConfusionMatrixDisplay
6 from helper_functions import *
7
8 NUM_CLASSES = 2
9
10 X, y = load_breast_cancer(return_X_y=True)
11
12 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.4,
random_state=0)
13 classifier = hmm.GaussianHMM(n_components=NUM_CLASSES,
covariance_type="full", min_covar=0.001, tol=0.01,
implementation="log", random_state=0)
14 classifier.fit(X_train)
15
16 mapping = get_mapping(classifier, X_train, y_train)
17 y_pred = list(map(lambda hs : mapping[hs], classifier.predict(X_test)))
18
19 print("Confusion Matrix:")
20 print(confusion_matrix(y_test, y_pred))
21 print("-----")
22 print("-----")
23 print("Performance Evaluation:")
24 print(classification_report(y_test, y_pred))
25 print("-----")
26 print("-----")
27 print("Accuracy Score:")
28 print(accuracy_score(y_test, y_pred))
29
30 plt.figure()
31 ConfusionMatrixDisplay.from_predictions(y_test, y_pred)
32 plt.savefig('./param.png')
33
34 plt.figure()
35 save_learning_curve(filepath='./param_lc.png',
base_classifier=classifier, title="Learning curve for parameter-tuned
Gaussian HMM", X=X, y=y)
36
```

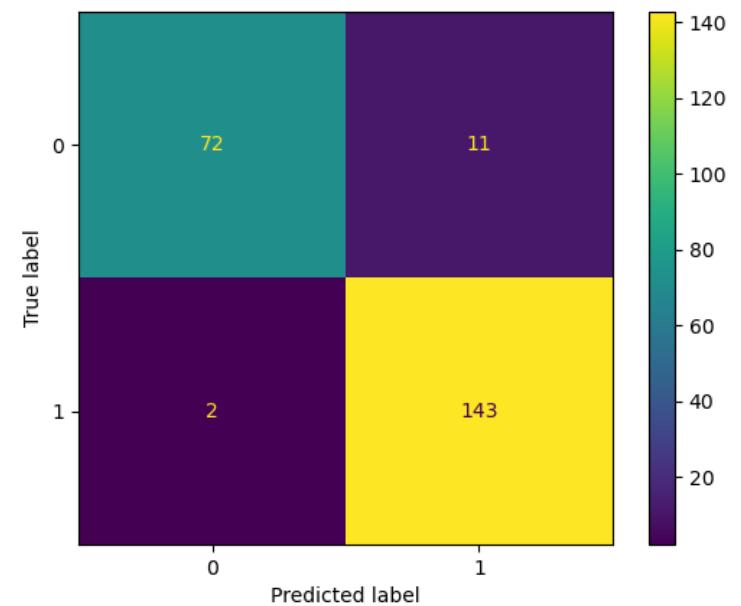
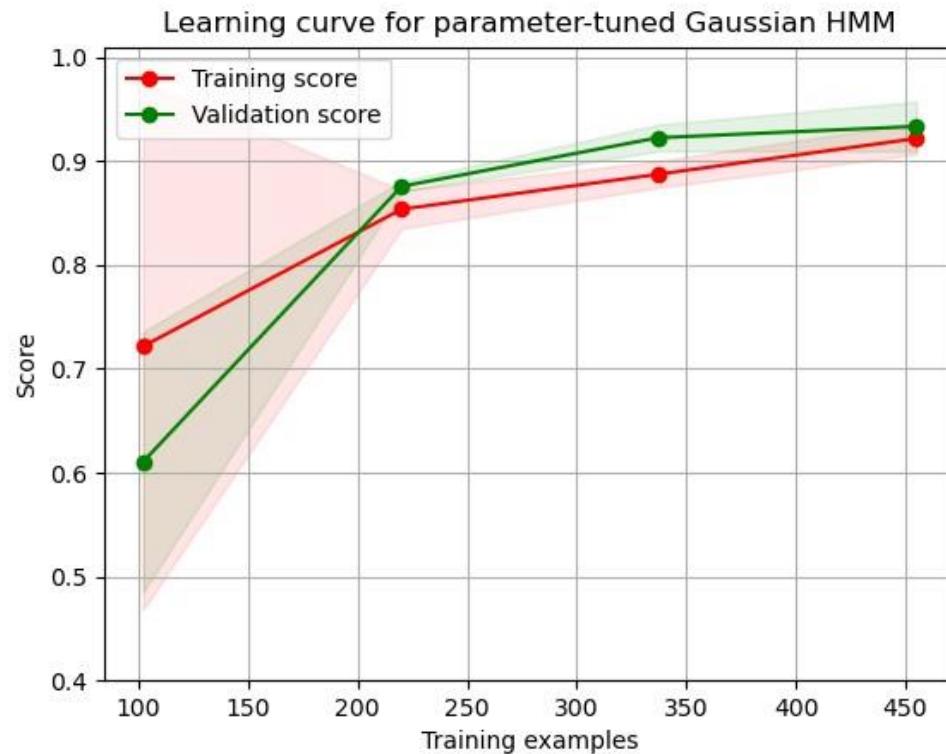
Confusion Matrix:
[[72 11]
 [2 143]]

Performance Evaluation:

	precision	recall	f1-score	support
0	0.97	0.87	0.92	83
1	0.93	0.99	0.96	145
accuracy			0.94	228
macro avg	0.95	0.93	0.94	228
weighted avg	0.94	0.94	0.94	228

Accuracy Score:
0.9409090909090909

GAUSSIAN HMM (WITH PARAMETER TUNING)- CONFUSION MATRIX AND LEARNING CURVE



HMM WITH GAUSSIAN MIXTURE MODEL EMISSIONS



GMM HMM (NO PARAMETER TUNING)- CODE & EVALUATION

```
1 import matplotlib.pyplot as plt
2 from sklearn.datasets import load_breast_cancer
3 from sklearn.model_selection import train_test_split
4 from hmmlearn import hmm
5 from sklearn.metrics import classification_report, confusion_matrix,
   accuracy_score, ConfusionMatrixDisplay
6 from helper_functions import *
7
8 NUM_CLASSES = 2
9
10 X, y = load_breast_cancer(return_X_y=True)
11
12 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.6,
   random_state=0)
13 classifier = hmm.GMMHMM(n_components=NUM_CLASSES)
14 classifier.fit(X_train)
15
16 mapping = get_mapping(classifier, X_train, y_train)
17 y_pred = list(map(lambda hs : mapping[hs], classifier.predict(X_test)))
18
19 print("Confusion Matrix:")
20 print(confusion_matrix(y_test, y_pred))
21 print("-----")
22 print("-----")
23 print("Performance Evaluation:")
24 print(classification_report(y_test, y_pred))
25 print("-----")
26 print("-----")
27 print("Accuracy Score:")
28 print(accuracy_score(y_test, y_pred))
29
30 plt.figure()
31 ConfusionMatrixDisplay.from_predictions(y_test, y_pred)
32 plt.savefig('./no_param.png')
33
34 plt.figure()
35 save_learning_curve(filepath='./no_param_lc.png',
   base_classifier=classifier, title="Learning curve for default GMM
   HMM", X=X, y=y, ylim=(0.4, 1.01))
36
```



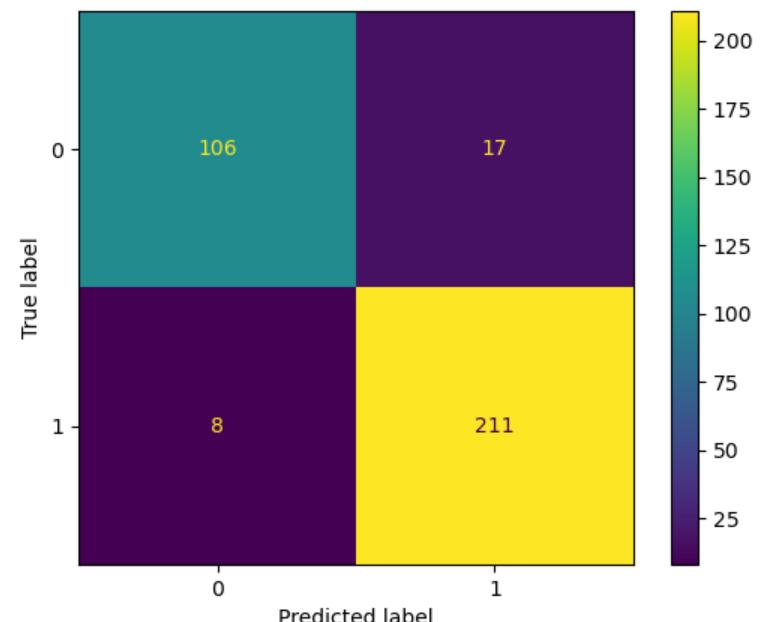
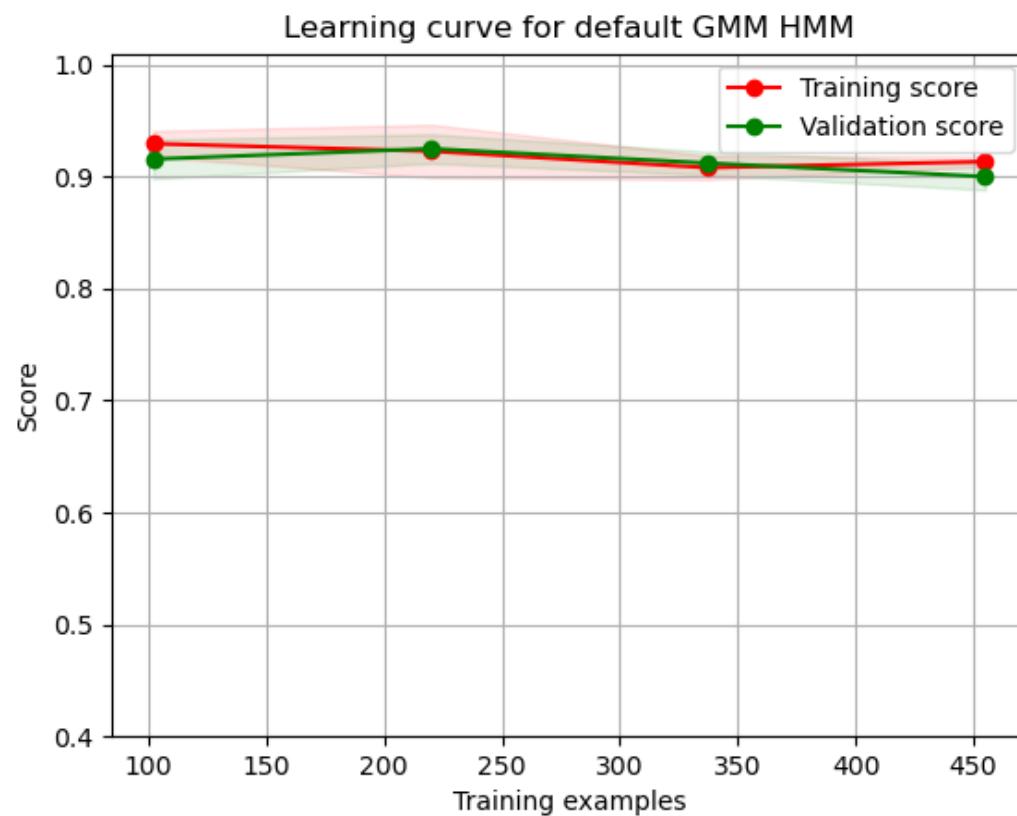
```
Confusion Matrix:
[[106 17]
 [ 8 211]]
```

```
Performance Evaluation:
          precision    recall  f1-score   support
0           0.93     0.86     0.89     123
1           0.93     0.96     0.94     219

   accuracy                           0.93    342
   macro avg       0.93     0.91     0.92    342
weighted avg       0.93     0.93     0.93    342
```

```
accuracy Score:
0.9269005847953217
```

GMM HMM (NO PARAMETER TUNING)- CONFUSION MATRIX AND LEARNING CURVE



GMM HMM (WITH PARAMETER TUNING)- CODE & EVALUATION

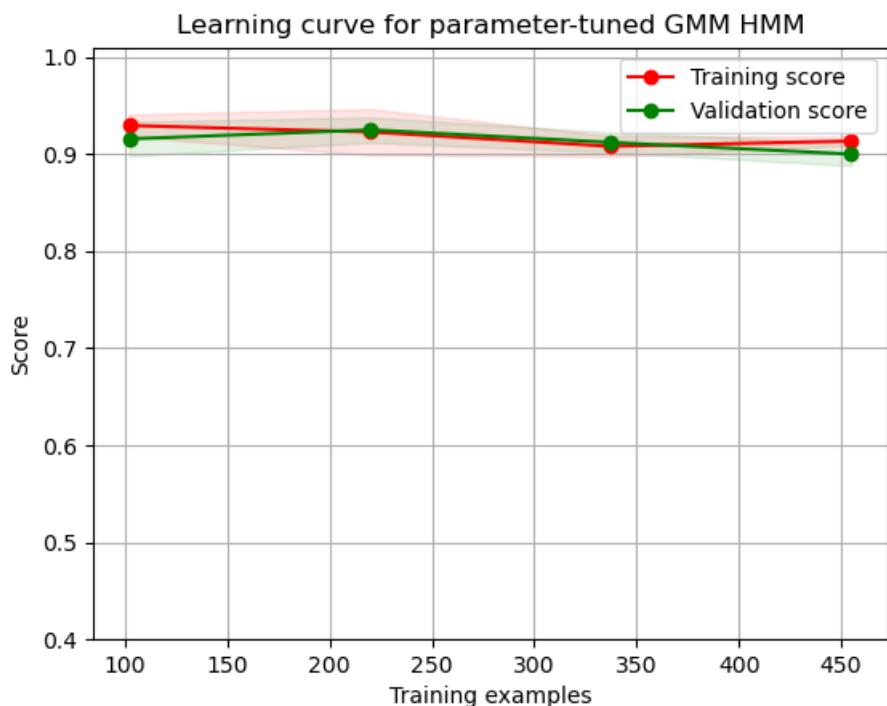
```
1 import matplotlib.pyplot as plt
2 from sklearn.datasets import load_breast_cancer
3 from sklearn.model_selection import train_test_split
4 from hmmlearn import hmm
5 from sklearn.metrics import classification_report, confusion_matrix,
6     accuracy_score, ConfusionMatrixDisplay
7 from helper_functions import *
8
9 NUM_CLASSES = 2
10 X, y = load_breast_cancer(return_X_y=True)
11
12 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.6,
13     random_state=0)
14 classifier = hmm.GMMHMM(n_components=NUM_CLASSES, n_mix=1,
15     covariance_type="diag", min_covar=0.001, tol=0.01,
16     implementation="log")
17 classifier.fit(X_train)
18
19 mapping = get_mapping(classifier, X_train, y_train)
20 y_pred = list(map(lambda hs : mapping[hs], classifier.predict(X_test)))
21 print("Confusion Matrix:")
22 print(confusion_matrix(y_test, y_pred))
23 print("-----")
24 print("Performance Evaluation:")
25 print(classification_report(y_test, y_pred))
26 print("-----")
27 print("Accuracy Score:")
28 print(accuracy_score(y_test, y_pred))
29
30 plt.figure()
31 ConfusionMatrixDisplay.from_predictions(y_test, y_pred)
32 plt.savefig('./param.png')
33
34 plt.figure()
35 save_learning_curve(filepath='./param_lc.png',
36     base_classifier=classifier, title="Learning curve for parameter-tuned
37     GMM HMM", X=X, y=y)
```

Confusion Matrix:
[[106 17]
 [8 211]]

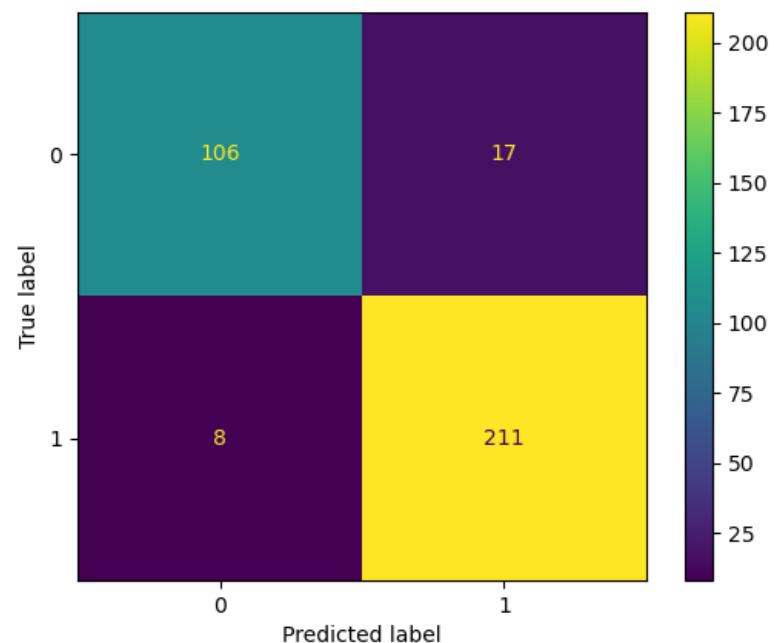
	precision	recall	f1-score	support
0	0.93	0.86	0.89	123
1	0.93	0.96	0.94	219
accuracy			0.93	342
macro avg	0.93	0.91	0.92	342
weighted avg	0.93	0.93	0.93	342

Accuracy Score:
0.9269005847953217

GMM HMM (WITH PARAMETER TUNING)- CONFUSION MATRIX AND LEARNING CURVE

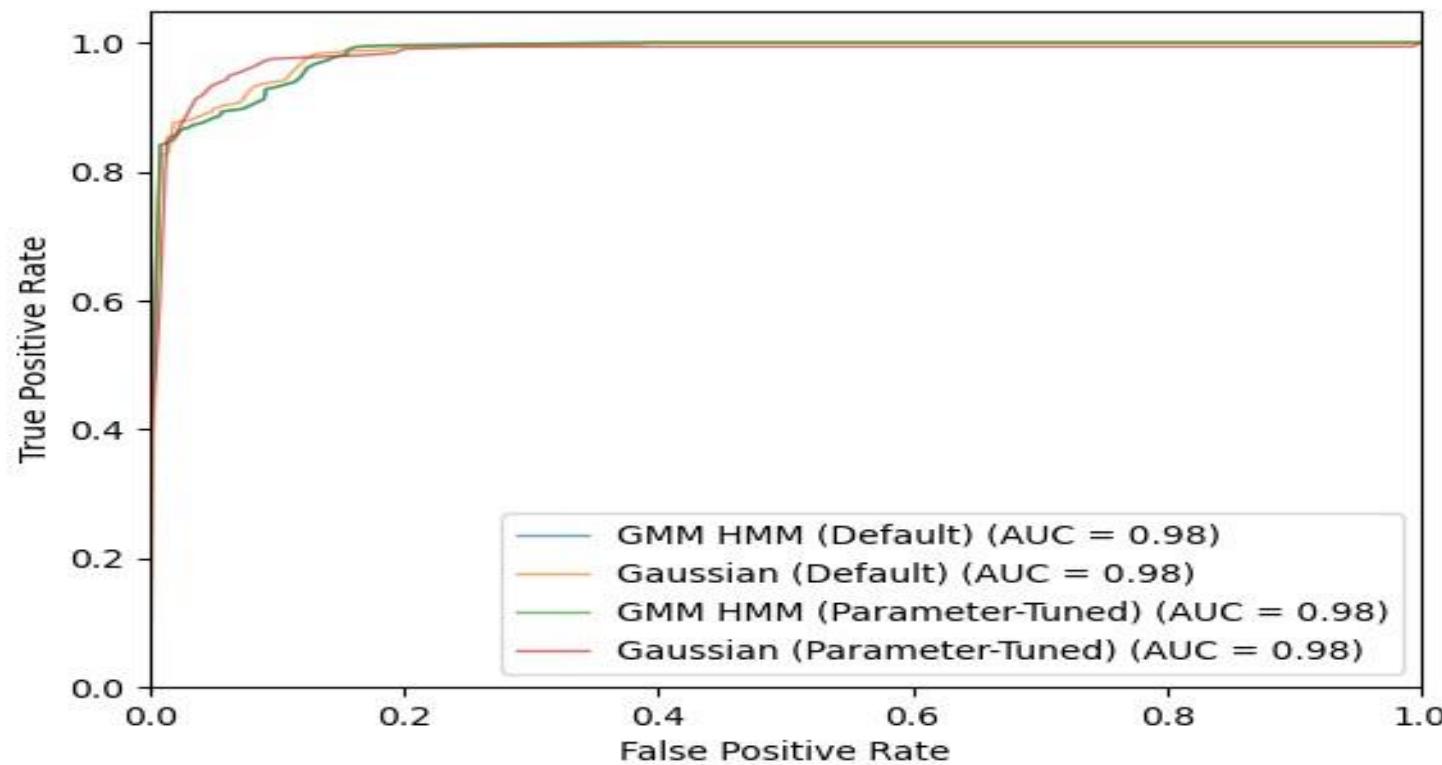


CONFUSION MATRIX

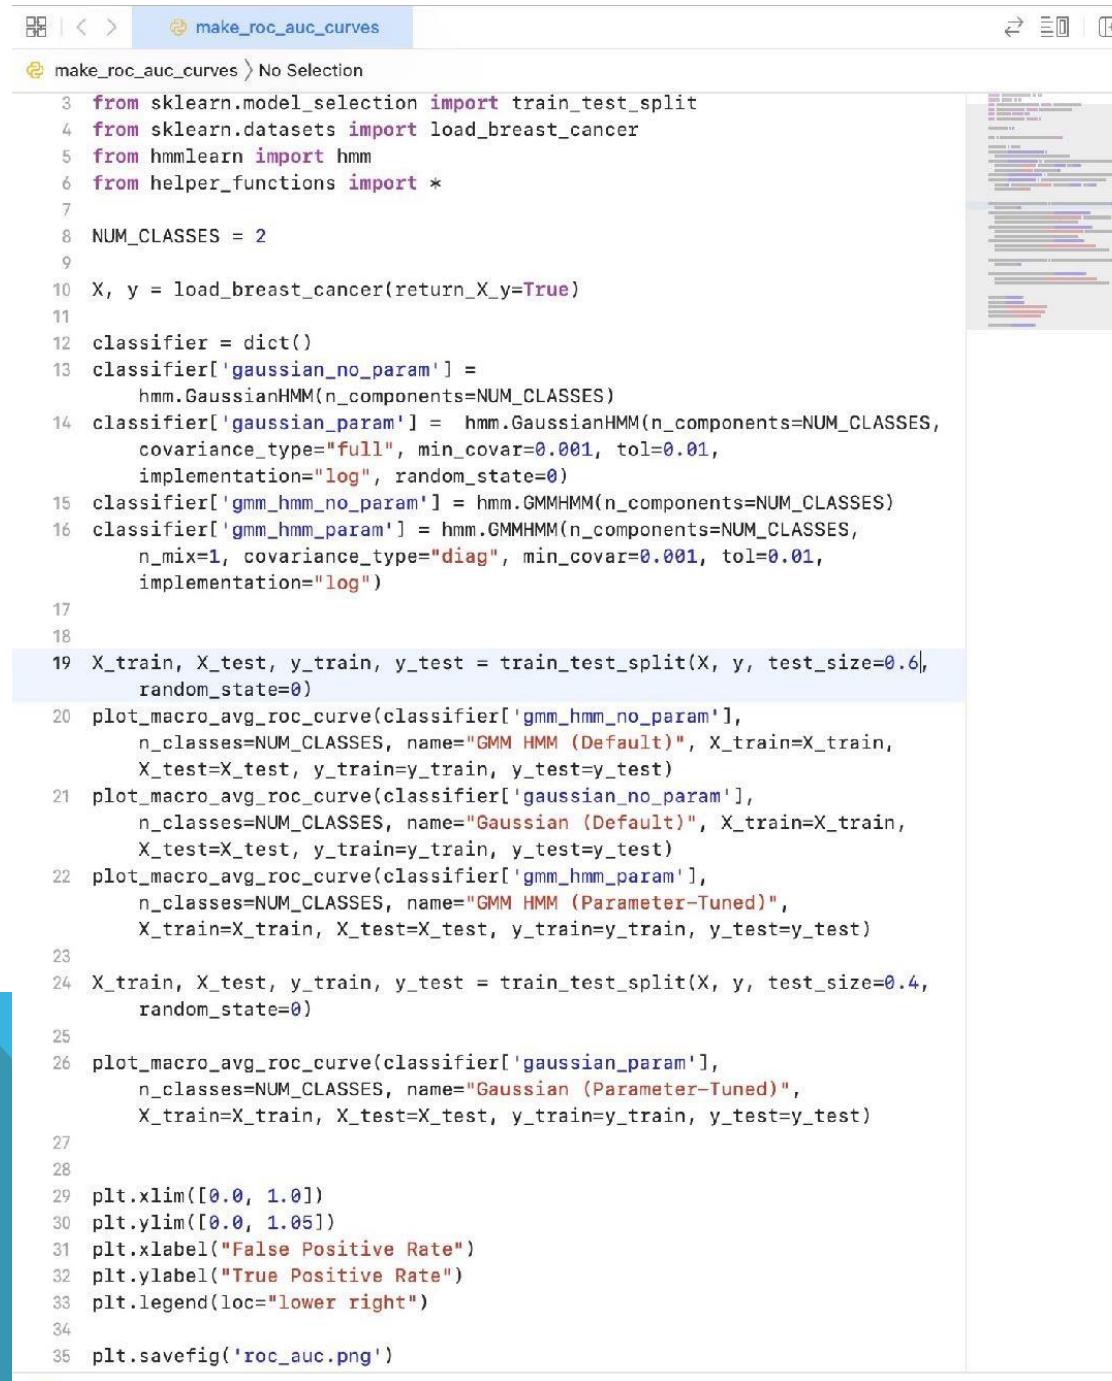


COMPARISON BETWEEN GAUSSIAN HMM AND GMM HMM.

ROC CURVES AND AUC COMPARISON



ROC CURVES AND AUC COMPARISON - CODE



```
from sklearn.model_selection import train_test_split
from sklearn.datasets import load_breast_cancer
from hmmlearn import hmm
from helper_functions import *

NUM_CLASSES = 2

X, y = load_breast_cancer(return_X_y=True)

classifier = dict()
classifier['gaussian_no_param'] = hmm.GaussianHMM(n_components=NUM_CLASSES)
classifier['gaussian_param'] = hmm.GaussianHMM(n_components=NUM_CLASSES,
                                               covariance_type="full", min_covar=0.001, tol=0.01,
                                               implementation="log", random_state=0)
classifier['gmm_hmm_no_param'] = hmm.GMMHMM(n_components=NUM_CLASSES)
classifier['gmm_hmm_param'] = hmm.GMMHMM(n_components=NUM_CLASSES,
                                         n_mix=1, covariance_type="diag", min_covar=0.001, tol=0.01,
                                         implementation="log")

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.6,
                                                    random_state=0)
plot_macro_avg_roc_curve(classifier['gmm_hmm_no_param'],
                         n_classes=NUM_CLASSES, name="GMM HMM (Default)", X_train=X_train,
                         X_test=X_test, y_train=y_train, y_test=y_test)
plot_macro_avg_roc_curve(classifier['gaussian_no_param'],
                         n_classes=NUM_CLASSES, name="Gaussian (Default)", X_train=X_train,
                         X_test=X_test, y_train=y_train, y_test=y_test)
plot_macro_avg_roc_curve(classifier['gmm_hmm_param'],
                         n_classes=NUM_CLASSES, name="GMM HMM (Parameter-Tuned)",
                         X_train=X_train, X_test=X_test, y_train=y_train, y_test=y_test)

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.4,
                                                    random_state=0)
plot_macro_avg_roc_curve(classifier['gaussian_param'],
                         n_classes=NUM_CLASSES, name="Gaussian (Parameter-Tuned)",
                         X_train=X_train, X_test=X_test, y_train=y_train, y_test=y_test)

plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.legend(loc="lower right")

plt.savefig('roc_auc.png')
```

COMPARISON OF PERFORMANCE

CLASSIFIER		Accuracy	Prec	Recall	F -Sco	AUC
GAUSSIAN HMM	DEFAULT	0.94	0.94	0.94	0.94	0.98
	TUNED	0.94	0.94	0.94	0.94	0.98
GMM HMM	DEFAULT	0.93	0.93	0.93	0.93	0.98
	TUNED	0.93	0.93	0.93	0.93	0.98

PART 2 , 3



CIFAR - CLASSIFICATION



CNN



IMPORTS AND DATA PREPROCESSING

Convolutional Neural Network

Importing the libraries

```
In [4]: import numpy as np
import tensorflow as tf

from tensorflow import keras
from tensorflow.keras import layers, datasets
from tensorflow.keras.models import Sequential
import matplotlib.pyplot as plt
```

Part 1 - Data Preprocessing

Normalization and augmentation will be done in the CNN itself

```
In [5]: img_height = 32
img_width = 32
num_classes = 10
```

```
In [6]: (train_images, train_labels), (test_images, test_labels) = datasets.cifar10.load_data()
```

MODEL ARCHITECTURE

Building the CNN

```
In [38]: cnn = Sequential([
    layers.Rescaling(1./255, input_shape=(img_height, img_width, 3)),
    layers.Conv2D(filters=32, kernel_size=3, padding='same', activation='relu'),
    layers.Conv2D(filters=32, kernel_size=3, padding='same', activation='relu'),
    layers.MaxPooling2D(pool_size=2, strides=2),
    layers.Conv2D(32, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(pool_size=2, strides=2),
    layers.Dropout(0.2),
    layers.Flatten(),
    layers.Dense(128, activation='relu'),
    layers.Dense(num_classes, activation="softmax")
])
```

COMPILATION AND SUMMARY

Compiling the CNN

```
In [39]: cnn.compile(optimizer=keras.optimizers.Adam(learning_rate=0.0005, epsilon=0.01),
                  loss=tf.keras.losses.SparseCategoricalCrossentropy(),
                  metrics=['accuracy'])
```

```
In [40]: cnn.summary()
```

Model: "sequential_4"

Layer (type)	Output Shape	Param #
rescaling_1 (Rescaling)	(None, 32, 32, 3)	0
conv2d_10 (Conv2D)	(None, 32, 32, 32)	896
conv2d_11 (Conv2D)	(None, 32, 32, 32)	9248
max_pooling2d_8 (MaxPooling 2D)	(None, 16, 16, 32)	0
conv2d_12 (Conv2D)	(None, 16, 16, 32)	9248
max_pooling2d_9 (MaxPooling 2D)	(None, 8, 8, 32)	0
dropout_7 (Dropout)	(None, 8, 8, 32)	0
flatten_2 (Flatten)	(None, 2048)	0
dense_8 (Dense)	(None, 128)	262272
dense_9 (Dense)	(None, 10)	1290
<hr/>		
Total params: 282,954		
Trainable params: 282,954		
Non-trainable params: 0		

TRAINING AND SAVING

Training the CNN on the Training set and evaluating it on the Test set

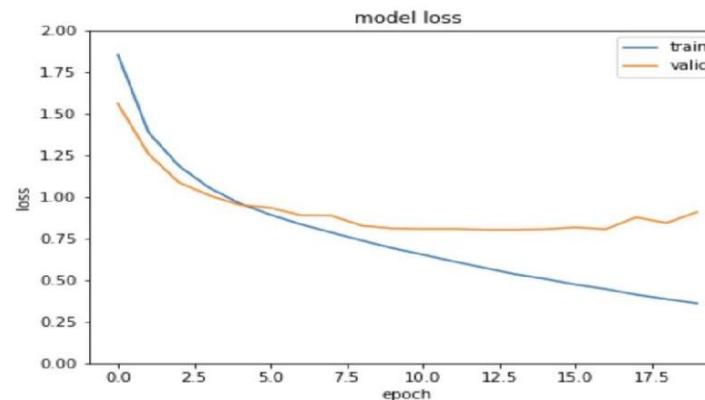
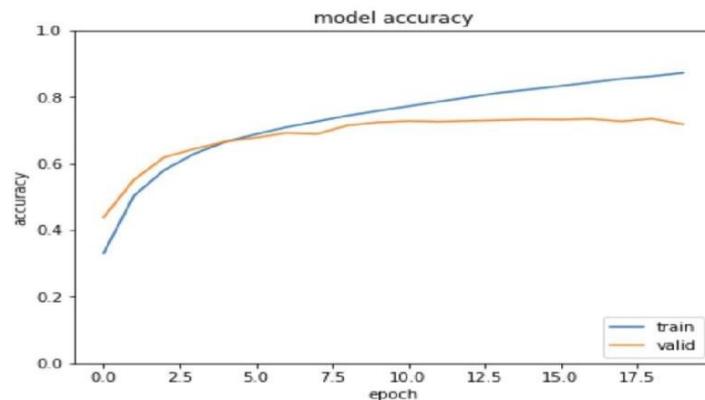
```
In [41]: history = cnn.fit(train_images, train_labels, epochs=20, validation_data=(test_images, test_labels))
1563/1563 [=====] - 52s 33ms/step - loss: 0.5073 - accuracy: 0.8229 - val_loss: 0.8052 - val_accuracy: 0.7304
Epoch 15/20
1563/1563 [=====] - 52s 33ms/step - loss: 0.5073 - accuracy: 0.8229 - val_loss: 0.8052 - val_accuracy: 0.7329
Epoch 16/20
1563/1563 [=====] - 50s 32ms/step - loss: 0.4732 - accuracy: 0.8331 - val_loss: 0.8167 - val_accuracy: 0.7322
Epoch 17/20
1563/1563 [=====] - 54s 34ms/step - loss: 0.4459 - accuracy: 0.8441 - val_loss: 0.8052 - val_accuracy: 0.7342
Epoch 18/20
1563/1563 [=====] - 53s 34ms/step - loss: 0.4118 - accuracy: 0.8551 - val_loss: 0.8775 - val_accuracy: 0.7267
Epoch 19/20
1563/1563 [=====] - 54s 35ms/step - loss: 0.3850 - accuracy: 0.8624 - val_loss: 0.8425 - val_accuracy: 0.7348
Epoch 20/20
1563/1563 [=====] - 54s 34ms/step - loss: 0.3596 - accuracy: 0.8725 - val_loss: 0.9090 - val_accuracy: 0.7184
```

```
In [42]: cnn.save('cnn.h5')
```

TRAINING AND LOSS CURVES

```
In [47]: def save_train_loss_curves(history, filename):
    plt.figure(figsize=(15,5))
    plt.subplot(1,2,1)
    plt.plot(history.history['accuracy'])
    plt.plot(history.history['val_accuracy'])
    plt.ylim(0, 1)
    plt.title('model accuracy')
    plt.ylabel('accuracy')
    plt.xlabel('epoch')
    plt.legend(['train', 'valid'], loc='lower right')
    plt.subplot(1,2,2)
    plt.plot(history.history['loss'])
    plt.plot(history.history['val_loss'])
    plt.title('model loss')
    plt.ylabel('loss')
    plt.xlabel('epoch')
    plt.legend(['train', 'valid'], loc='upper right')
    plt.ylim([0,2])
    plt.savefig(filename)
```

```
In [48]: save_train_loss_curves(history, 'cnn_loss_acc.png')
```



PERFORMANCE - CODE

Performance

```
In [44]: from sklearn.metrics import classification_report, confusion_matrix, accuracy_score, ConfusionMatrixDisplay
```

```
In [45]: def get_predictions(model, inputs):

    def index_max_arr_element(arr):
        best_index = -1
        best = float('-inf')
        for i, num in enumerate(arr):
            if num > best:
                best = num
                best_index = i
        return best_index

    pred_arr = model.predict(inputs)
    pred_labels = [index_max_arr_element(row) for row in pred_arr]
    return pred_labels
```

```
In [46]: pred_labels = get_predictions(cnn, test_images)

print("Confusion Matrix:")
print(confusion_matrix(test_labels, pred_labels))
print("-----")
print("-----")
print("Performance Evaluation:")
print(classification_report(test_labels, pred_labels))
print("-----")
print("-----")
print("Accuracy Score:")
print(accuracy_score(test_labels, pred_labels))

plt.figure()
ConfusionMatrixDisplay.from_predictions(test_labels, pred_labels)
plt.savefig('./cnn_matrix.png')
```

9/313 [.....] - ETA: 4s

2022-09-16 18:28:02.208514: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device type GPU is enabled.

PERFORMANCE

```
2022-09-16 18:28:02.208514: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
```

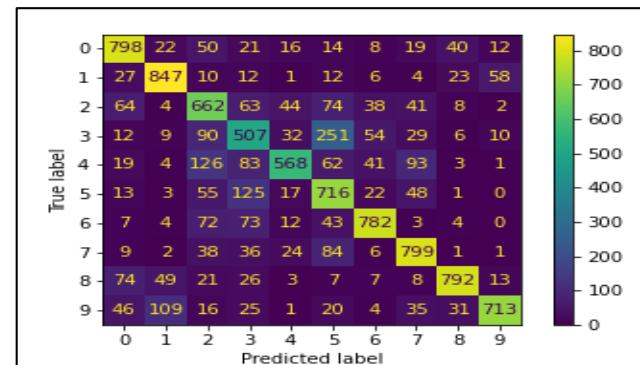
313/313 [=====] - 3s 8ms/step

Confusion Matrix:

```
[[798 22 50 21 16 14 8 19 40 12]
 [ 27 847 10 12 1 12 6 4 23 58]
 [ 64 4 662 63 44 74 38 41 8 2]
 [ 12 9 90 507 32 251 54 29 6 10]
 [ 19 4 126 83 568 62 41 93 3 1]
 [ 13 3 55 125 17 716 22 48 1 0]
 [ 7 4 72 73 12 43 782 3 4 0]
 [ 9 2 38 36 24 84 6 799 1 1]
 [ 74 49 21 26 3 7 7 8 792 13]
 [ 46 109 16 25 1 20 4 35 31 713]]
```

Performance Evaluation:

	precision	recall	f1-score	support
0	0.75	0.80	0.77	1000
1	0.80	0.85	0.83	1000
2	0.58	0.66	0.62	1000
3	0.52	0.51	0.51	1000
4	0.79	0.57	0.66	1000
5	0.56	0.72	0.63	1000
6	0.81	0.78	0.79	1000
7	0.74	0.80	0.77	1000
8	0.87	0.79	0.83	1000
9	0.88	0.71	0.79	1000
accuracy			0.72	10000
macro avg	0.73	0.72	0.72	10000
weighted avg	0.73	0.72	0.72	10000



Accuracy Score:
0.7184

VGG-16

IMPORTS, DATA PREPROCESSING & ARCHITECTURE

VGG-16

```
In [49]: from tensorflow.keras.applications import vgg16
```

```
In [50]: train_images_p, test_images_p = vgg16.preprocess_input(train_images), vgg16.preprocess_input(test_images)
```

```
In [51]: vgg_based_model = Sequential([
    layers.Rescaling(1./255, input_shape=(img_height, img_width, 3)),
    vgg16.VGG16(include_top=False, weights="imagenet", input_shape=(img_height, img_width, 3), pooling="max"),
    layers.Flatten(),
    layers.Dense(512, activation='relu'),
    layers.Dropout(0.2),
    layers.Dense(128, activation='relu'),
    layers.Dropout(0.2),
    layers.Dense(num_classes, activation="softmax")
])
```

COMPIRATION AND SUMMARY

```
In [52]: vgg_based_model.compile(optimizer=keras.optimizers.Adam(learning_rate=0.0002, epsilon=0.01),  
                               loss=tf.keras.losses.SparseCategoricalCrossentropy(),  
                               metrics=['accuracy'])
```

```
In [53]: vgg_based_model.summary()
```

Model: "sequential_5"

Layer (type)	Output Shape	Param #
<hr/>		
rescaling_2 (Rescaling)	(None, 32, 32, 3)	0
vgg16 (Functional)	(None, 512)	14714688
flatten_3 (Flatten)	(None, 512)	0
dense_10 (Dense)	(None, 512)	262656
dropout_8 (Dropout)	(None, 512)	0
dense_11 (Dense)	(None, 128)	65664
dropout_9 (Dropout)	(None, 128)	0
dense_12 (Dense)	(None, 10)	1290
<hr/>		
Total params: 15,044,298		
Trainable params: 15,044,298		
Non-trainable params: 0		

TRAINING AND SAVING

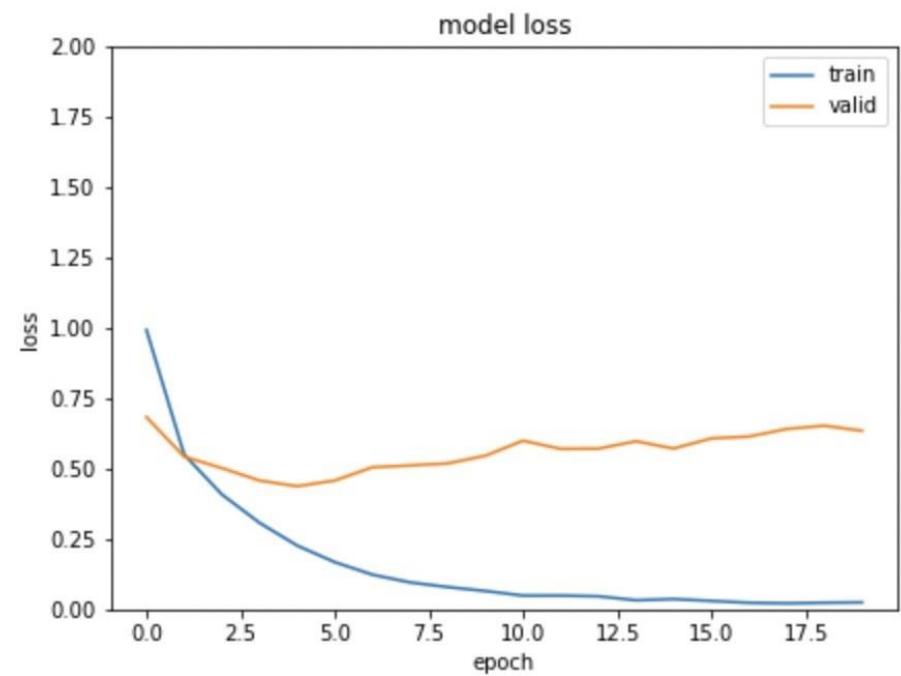
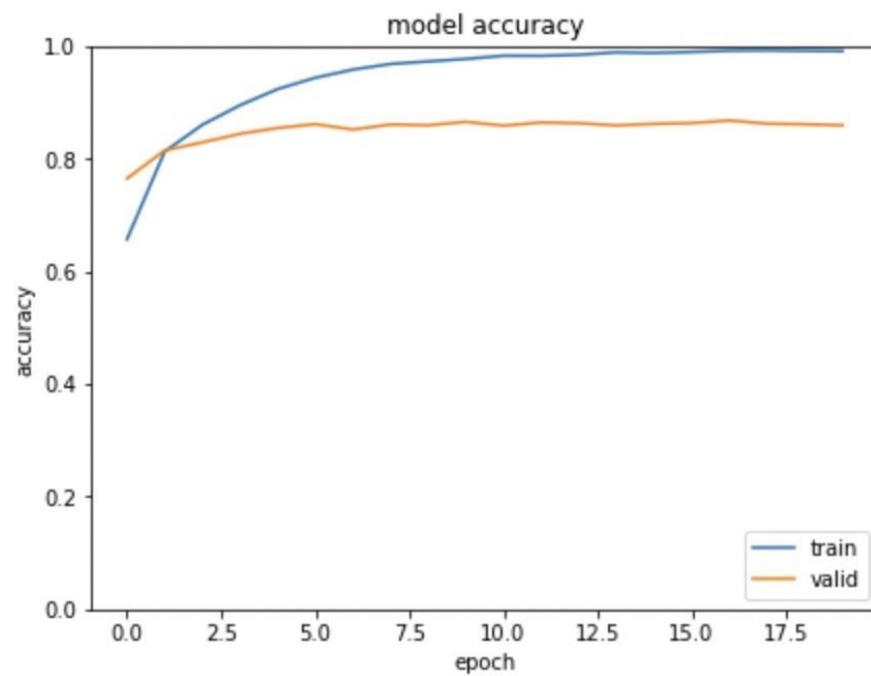
```
In [54]: history = vgg_based_model.fit(train_images_p, train_labels, epochs=20, validation_data=(test_images_p, test_labels))
```

1563/1563 [=====] - 2023 101ms/step - loss: 0.0319 - accuracy: 0.9885 - val_loss: 0.5713 -
val_accuracy: 0.8596
Epoch 15/20
1563/1563 [=====] - 288s 184ms/step - loss: 0.0361 - accuracy: 0.9885 - val_loss: 0.5713 -
val_accuracy: 0.8625
Epoch 16/20
1563/1563 [=====] - 282s 181ms/step - loss: 0.0296 - accuracy: 0.9902 - val_loss: 0.6073 -
val_accuracy: 0.8638
Epoch 17/20
1563/1563 [=====] - 297s 190ms/step - loss: 0.0230 - accuracy: 0.9925 - val_loss: 0.6136 -
val_accuracy: 0.8684
Epoch 18/20
1563/1563 [=====] - 292s 187ms/step - loss: 0.0211 - accuracy: 0.9929 - val_loss: 0.6407 -
val_accuracy: 0.8628
Epoch 19/20
1563/1563 [=====] - 295s 188ms/step - loss: 0.0232 - accuracy: 0.9922 - val_loss: 0.6523 -
val_accuracy: 0.8617
Epoch 20/20
1563/1563 [=====] - 293s 187ms/step - loss: 0.0244 - accuracy: 0.9918 - val_loss: 0.6342 -
val_accuracy: 0.8599

```
In [56]: vgg_based_model.save('vgg16.h5')
```

TRAINING AND LOSS CURVES

```
In [55]: save_train_loss_curves(history, 'vgg_loss_acc.png')
```



PERFORMANCE

```
2022-09-16 23:40:57.867625: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
```

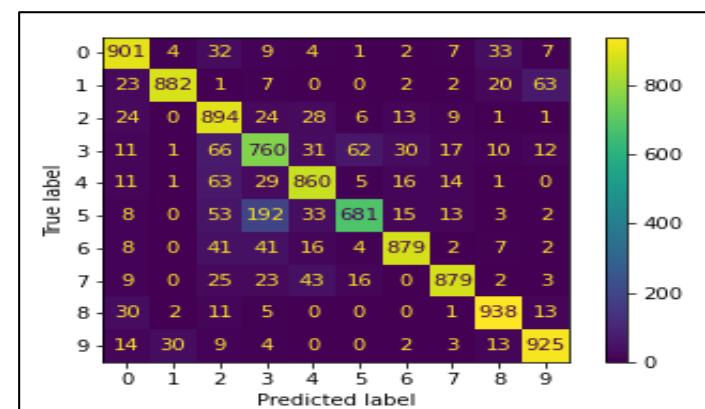
313/313 [=====] - 5s 15ms/step

Confusion Matrix:

```
[[901  4  32  9  4  1  2  7  33  7]
 [23 882  1  7  0  0  2  2  20  63]
 [24  0 894  24  28  6 13  9  1  1]
 [11  1 66 760  31  62  30  17  10  12]
 [11  1 63 29 860  5 16 14  1  0]
 [ 8  0 53 192  33  681  15  13  3  2]
 [ 8  0 41 41 16  4 879  2  7  2]
 [ 9  0 25 23 43  16  0 879  2  3]
 [30  2 11  5  0  0  1 938  13]
 [14 30  9  4  0  0  2  3 13 925]]
```

Performance Evaluation:

	precision	recall	f1-score	support
0	0.87	0.90	0.88	1000
1	0.96	0.88	0.92	1000
2	0.75	0.89	0.81	1000
3	0.69	0.76	0.73	1000
4	0.85	0.86	0.85	1000
5	0.88	0.68	0.77	1000
6	0.92	0.88	0.90	1000
7	0.93	0.88	0.90	1000
8	0.91	0.94	0.93	1000
9	0.90	0.93	0.91	1000
accuracy			0.86	10000
macro avg	0.87	0.86	0.86	10000
weighted avg	0.87	0.86	0.86	10000



Accuracy Score:

0.8599

RES NET- 50

IMPORTS, DATA PREPROCESSING & ARCHITECTURE

RESNET-50

```
In [57]: from tensorflow.keras.applications import resnet50
```

```
In [58]: train_images_pr, test_images_pr = resnet50.preprocess_input(train_images), resnet50.preprocess_input(test_images)
```

```
In [59]: resnet_based_model = Sequential([
    layers.Rescaling(1./255, input_shape=(img_height, img_width, 3)),
    resnet50.ResNet50(include_top=False, weights="imagenet", input_shape=(img_height, img_width, 3), pooling="max"),
    layers.Flatten(),
    layers.Dense(512, activation='relu'),
    layers.Dropout(0.2),
    layers.Dense(128, activation='relu'),
    layers.Dropout(0.2),
    layers.Dense(num_classes, activation="softmax")
])
```

COMPILATION AND SUMMARY

```
In [60]: resnet_based_model.compile(optimizer=keras.optimizers.Adam(learning_rate=0.0002, epsilon=0.01),  
                                loss=tf.keras.losses.SparseCategoricalCrossentropy(),  
                                metrics=['accuracy'])
```

```
In [61]: resnet_based_model.summary()
```

Model: "sequential_6"

Layer (type)	Output Shape	Param #
rescaling_3 (Rescaling)	(None, 32, 32, 3)	0
resnet50 (Functional)	(None, 2048)	23587712
flatten_4 (Flatten)	(None, 2048)	0
dense_13 (Dense)	(None, 512)	1049088
dropout_10 (Dropout)	(None, 512)	0
dense_14 (Dense)	(None, 128)	65664
dropout_11 (Dropout)	(None, 128)	0
dense_15 (Dense)	(None, 10)	1290
<hr/>		
Total params: 24,703,754		
Trainable params: 24,650,634		
Non-trainable params: 53,120		

TRAINING AND SAVING

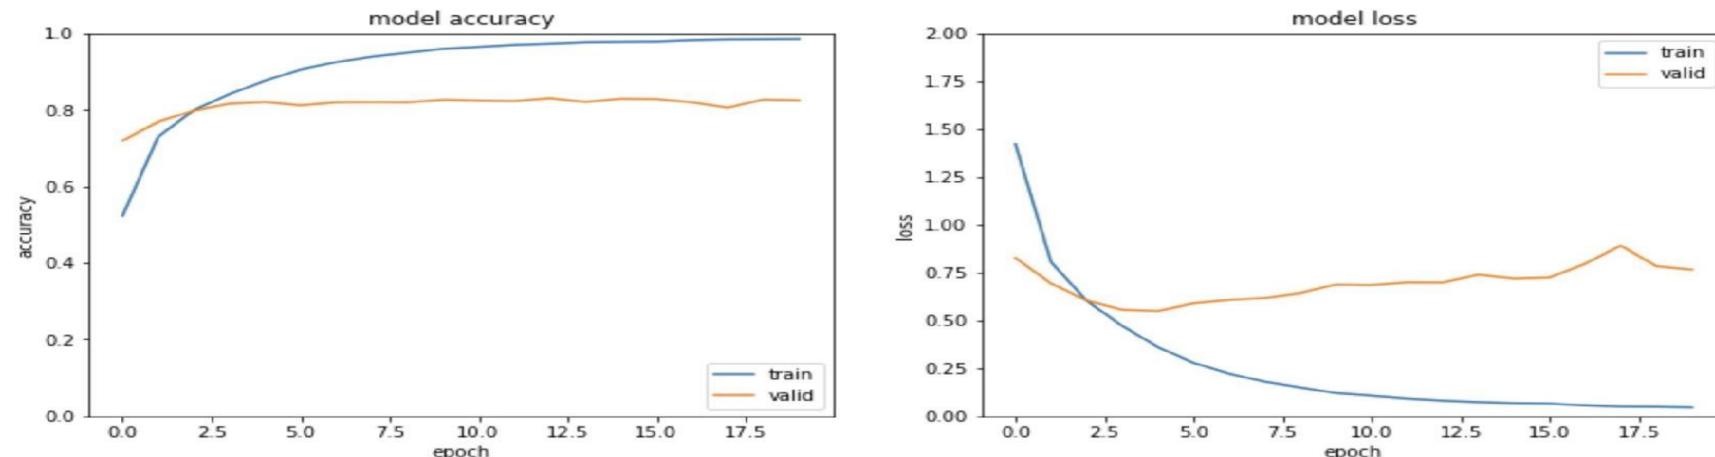
```
In [62]: history = resnet_based_model.fit(train_images_pr, train_labels, epochs=20, validation_data=(test_images_pr, test_labels))

1563/1563 [=====] - 212s 136ms/step - loss: 0.0728 - accuracy: 0.9768 - val_loss: 0.7386 - val_accuracy: 0.8207
Epoch 15/20
1563/1563 [=====] - 211s 135ms/step - loss: 0.0672 - accuracy: 0.9777 - val_loss: 0.7188 - val_accuracy: 0.8294
Epoch 16/20
1563/1563 [=====] - 213s 136ms/step - loss: 0.0651 - accuracy: 0.9784 - val_loss: 0.7238 - val_accuracy: 0.8285
Epoch 17/20
1563/1563 [=====] - 215s 137ms/step - loss: 0.0551 - accuracy: 0.9820 - val_loss: 0.7959 - val_accuracy: 0.8196
Epoch 18/20
1563/1563 [=====] - 215s 138ms/step - loss: 0.0494 - accuracy: 0.9839 - val_loss: 0.8906 - val_accuracy: 0.8058
Epoch 19/20
1563/1563 [=====] - 231s 148ms/step - loss: 0.0482 - accuracy: 0.9843 - val_loss: 0.7835 - val_accuracy: 0.8271
Epoch 20/20
1563/1563 [=====] - 279s 179ms/step - loss: 0.0450 - accuracy: 0.9849 - val_loss: 0.7657 - val_accuracy: 0.8252
```

```
In [64]: resnet_based_model.save('resnet50.h5')
```

TRAINING AND LOSS CURVES

```
In [63]: save_train_loss_curves(history, 'resnet_loss_acc.png')
```



```
In [64]: resnet_based_model.save('resnet50.h5')
```

PERFORMANCE

```
2022-09-16 23:47:09.495012: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
```

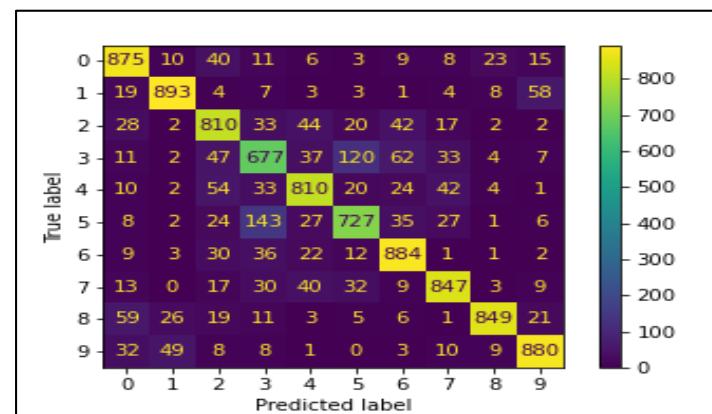
313/313 [=====] - 9s 25ms/step

Confusion Matrix:

```
[[875 10 40 11 6 3 9 8 23 15]
 [ 19 893 4 7 3 3 1 4 8 58]
 [ 28 2 810 33 44 20 42 17 2 2]
 [ 11 2 47 677 37 120 62 33 4 7]
 [ 10 2 54 33 810 20 24 42 4 11]
 [ 8 2 24 143 27 727 35 27 1 6]
 [ 9 3 30 36 22 12 884 1 1 2]
 [ 13 0 17 30 40 32 9 847 3 9]
 [ 59 26 19 11 3 5 6 1 849 21]
 [ 32 49 8 8 1 0 3 10 9 880]]
```

Performance Evaluation:

	precision	recall	f1-score	support
0	0.82	0.88	0.85	1000
1	0.90	0.89	0.90	1000
2	0.77	0.81	0.79	1000
3	0.68	0.68	0.68	1000
4	0.82	0.81	0.81	1000
5	0.77	0.73	0.75	1000
6	0.82	0.88	0.85	1000
7	0.86	0.85	0.85	1000
8	0.94	0.85	0.89	1000
9	0.88	0.88	0.88	1000
accuracy			0.83	10000
macro avg	0.83	0.83	0.83	10000
weighted avg	0.83	0.83	0.83	10000



Accuracy Score:

0.8252

RNN (LSTM)



IMAGE PROCESSING AND MODEL ARCHITECTURE

RNN (LSTM based)

```
In [10]: train_images_res, test_images_res = train_images.reshape(train_images.shape[0], 1024, 3), test_images.reshape(test_i  
In [11]: train_images_res.shape  
Out[11]: (50000, 1024, 3)  
  
In [12]: lstm_based_model = Sequential([  
    layers.Rescaling(1./255, input_shape=(1024, 3)),  
    layers.LSTM(units=50, return_sequences=True),  
    layers.Dropout(0.2),  
    layers.LSTM(units=50),  
    layers.Dropout(0.2),  
    layers.Dense(128, activation='relu'),  
    layers.Dropout(0.2),  
    layers.Dense(num_classes, activation="softmax")  
])  
  
Metal device set to: Apple M1  
  
2022-09-16 11:44:55.396775: I tensorflow/core/common_runtime/pluggable_device/pluggable_device_factory.cc:305] Could not identify NUMA node of platform GPU ID 0, defaulting to 0. Your kernel may not have been built with NUMA support.  
2022-09-16 11:44:55.397008: I tensorflow/core/common_runtime/pluggable_device/pluggable_device_factory.cc:271] Created TensorFlow device (/job:localhost/replica:0/task:0/device:GPU:0 with 0 MB memory) -> physical PluggableDevice (device: 0, name: METAL, pci bus id: <undefined>)
```

COMPILATION AND SUMMARY

```
In [13]: lstm_based_model.compile(optimizer=keras.optimizers.Adam(learning_rate=0.001),
                             loss=tf.keras.losses.SparseCategoricalCrossentropy(),
                             metrics=['accuracy'])
```

```
In [15]: lstm_based_model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
<hr/>		
rescaling (Rescaling)	(None, 1024, 3)	0
lstm (LSTM)	(None, 1024, 50)	10800
dropout (Dropout)	(None, 1024, 50)	0
lstm_1 (LSTM)	(None, 50)	20200
dropout_1 (Dropout)	(None, 50)	0
dense (Dense)	(None, 128)	6528
dropout_2 (Dropout)	(None, 128)	0
dense_1 (Dense)	(None, 10)	1290
<hr/>		
Total params: 38,818		
Trainable params: 38,818		
Non-trainable params: 0		

TRAINING AND SAVING

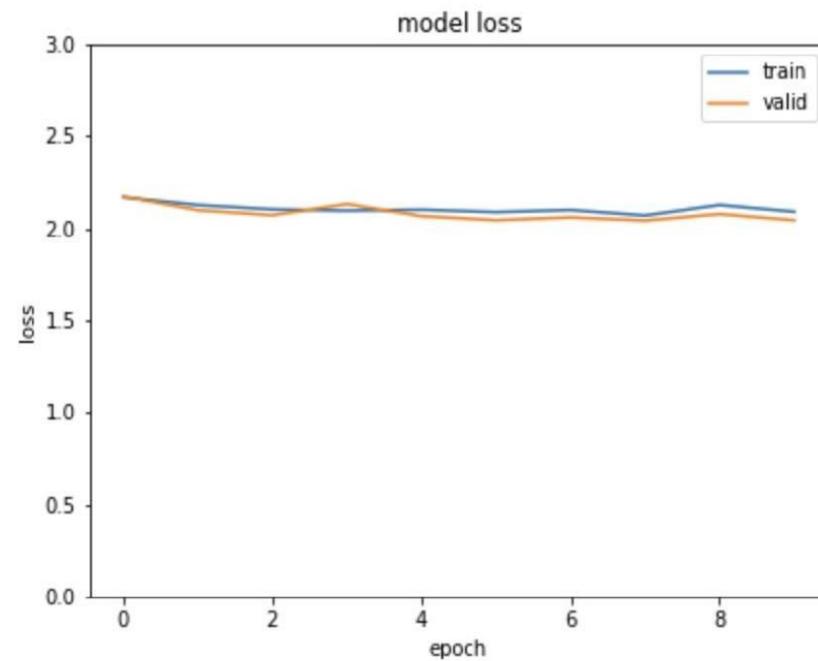
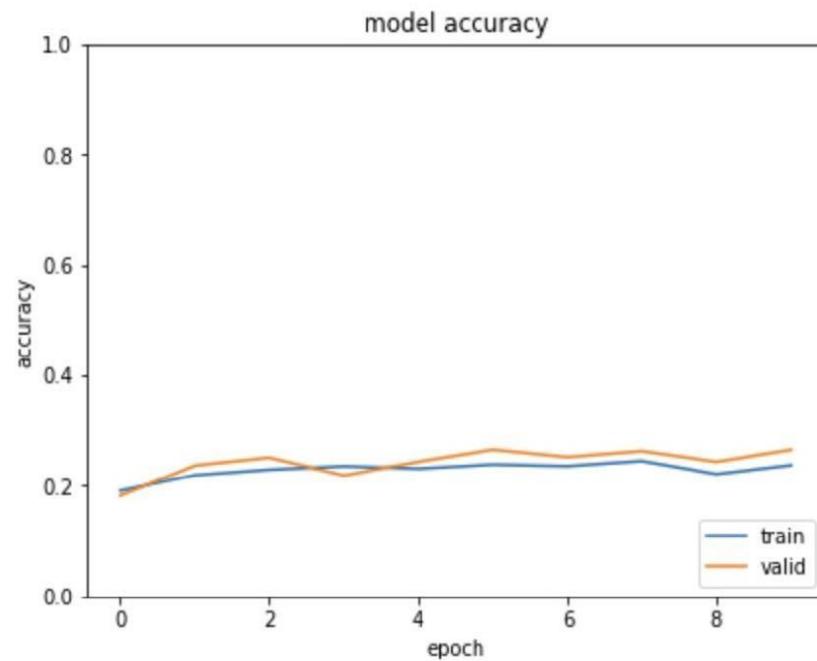
```
In [14]: history = lstm_based_model.fit(train_images_res, train_labels, epochs=10, validation_data=(test_images_res, test_labels))

val_accuracy: 0.2188
Epoch 5/10
1563/1563 [=====] - 795s 509ms/step - loss: 2.1007 - accuracy: 0.2312 - val_loss: 2.0660 -
val_accuracy: 0.2436
Epoch 6/10
1563/1563 [=====] - 807s 516ms/step - loss: 2.0879 - accuracy: 0.2386 - val_loss: 2.0442 -
val_accuracy: 0.2653
Epoch 7/10
1563/1563 [=====] - 877s 561ms/step - loss: 2.0990 - accuracy: 0.2357 - val_loss: 2.0592 -
val_accuracy: 0.2520
Epoch 8/10
1563/1563 [=====] - 927s 593ms/step - loss: 2.0696 - accuracy: 0.2450 - val_loss: 2.0422 -
val_accuracy: 0.2628
Epoch 9/10
1563/1563 [=====] - 917s 587ms/step - loss: 2.1258 - accuracy: 0.2214 - val_loss: 2.0771 -
val_accuracy: 0.2435
Epoch 10/10
1563/1563 [=====] - 815s 522ms/step - loss: 2.0896 - accuracy: 0.2373 - val_loss: 2.0438 -
val_accuracy: 0.2651
```

```
In [19]: lstm_based_model.save('lstm.h5')
```

TRAINING AND LOSS CURVES

```
In [18]: save_train_loss_curves(history, 'lstm_loss_acc.png')
```



PERFORMANCE

```
2022-09-16 23:56:19.986646: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.  
2022-09-16 23:56:23.320163: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
```

313/313 [=====] - 107s 318ms/step

Confusion Matrix:

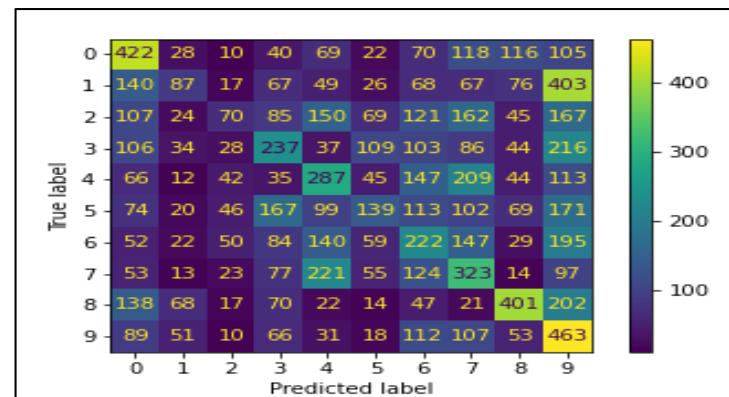
```
[[422 28 10 40 69 22 70 118 116 105]
 [140 87 17 67 49 26 68 67 76 403]
 [107 24 70 85 150 69 121 162 45 167]
 [106 34 28 237 37 109 103 86 44 216]
 [ 66 12 42 35 287 45 147 209 44 113]
 [ 74 20 46 167 99 139 113 102 69 171]
 [ 52 22 50 84 140 59 222 147 29 195]
 [ 53 13 23 77 221 55 124 323 14 97]
 [138 68 17 70 22 14 47 21 401 202]
 [ 89 51 10 66 31 18 112 107 53 463]]
```

Performance Evaluation:

	precision	recall	f1-score	support
0	0.34	0.42	0.38	1000
1	0.24	0.09	0.13	1000
2	0.22	0.07	0.11	1000
3	0.26	0.24	0.25	1000
4	0.26	0.29	0.27	1000
5	0.25	0.14	0.18	1000
6	0.20	0.22	0.21	1000
7	0.24	0.32	0.28	1000
8	0.45	0.40	0.42	1000
9	0.22	0.46	0.30	1000
accuracy			0.27	10000
macro avg	0.27	0.27	0.25	10000
weighted avg	0.27	0.27	0.25	10000

Accuracy Score:

0.2651



ALEXNET



MODEL ARCHITECTURE - I

```
[30]: def alexnet_model(img_shape=(32, 32, 3), n_classes=10, l2_reg=0.,
weights=None):

    # Initialize model
    alexnet = Sequential()

    # Layer 1
    alexnet.add(Conv2D(96, (11, 11), input_shape=img_shape, padding='same', kernel_regularizer=l2(l2_reg)))
    alexnet.add(BatchNormalization())
    alexnet.add(Activation('relu'))
    alexnet.add(MaxPooling2D(pool_size=(2, 2)))

    # Layer 2
    alexnet.add(Conv2D(256, (5, 5), padding='same'))
    alexnet.add(BatchNormalization())
    alexnet.add(Activation('relu'))
    alexnet.add(MaxPooling2D(pool_size=(2, 2)))

    # Layer 3
    alexnet.add(ZeroPadding2D((1, 1)))
    alexnet.add(Conv2D(512, (3, 3), padding='same'))
    alexnet.add(BatchNormalization())
    alexnet.add(Activation('relu'))
    alexnet.add(MaxPooling2D(pool_size=(2, 2)))

    # Layer 4
    alexnet.add(ZeroPadding2D((1, 1)))
    alexnet.add(Conv2D(1024, (3, 3), padding='same'))
    alexnet.add(BatchNormalization())
    alexnet.add(Activation('relu'))

    # Layer 5
    alexnet.add(ZeroPadding2D((1, 1)))
    alexnet.add(Conv2D(1024, (3, 3), padding='same'))
    alexnet.add(BatchNormalization())
    alexnet.add(Activation('relu'))
    alexnet.add(MaxPooling2D(pool_size=(2, 2)))
```

MODEL ARCHITECTURE - II

```
# Layer 6
alexnet.add(Flatten())
alexnet.add(Dense(3072))
alexnet.add(BatchNormalization())
alexnet.add(Activation('relu'))
alexnet.add(Dropout(0.5))

# Layer 7
alexnet.add(Dense(4096))
alexnet.add(BatchNormalization())
alexnet.add(Activation('relu'))
alexnet.add(Dropout(0.5))

# Layer 8
alexnet.add(Dense(n_classes))
alexnet.add(BatchNormalization())
alexnet.add(Activation('softmax'))

if weights is not None:
    alexnet.load_weights(weights)

return alexnet
```

```
[31]: alexnet_model = alexnet_model(img_shape=(32, 32, 3), n_classes=10)
```

```
[32]: alexnet_model.compile(optimizer=keras.optimizers.Adam(learning_rate=0.0002, epsilon=0.05),
                           loss=tf.keras.losses.SparseCategoricalCrossentropy(),
                           metrics=['accuracy'])
```

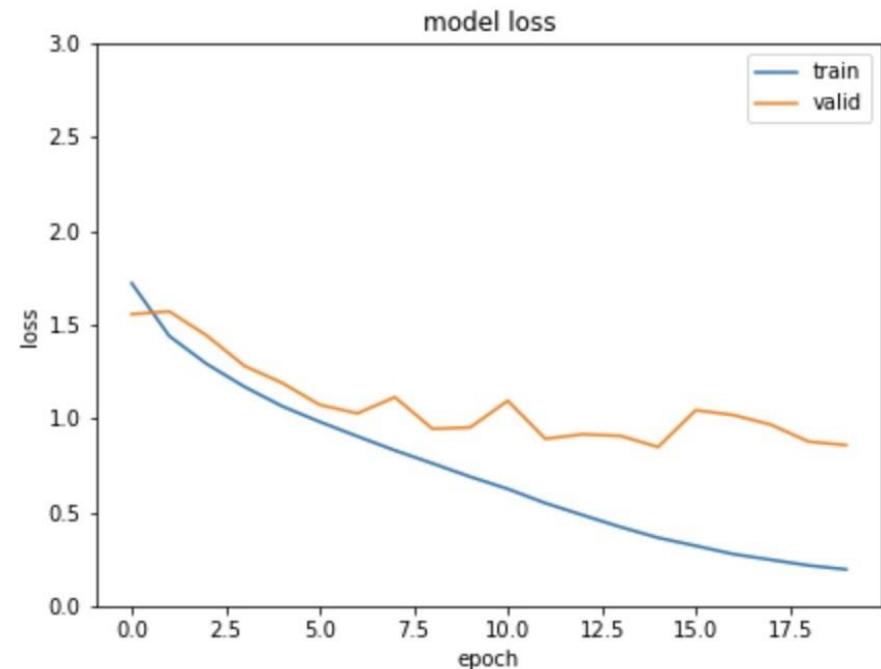
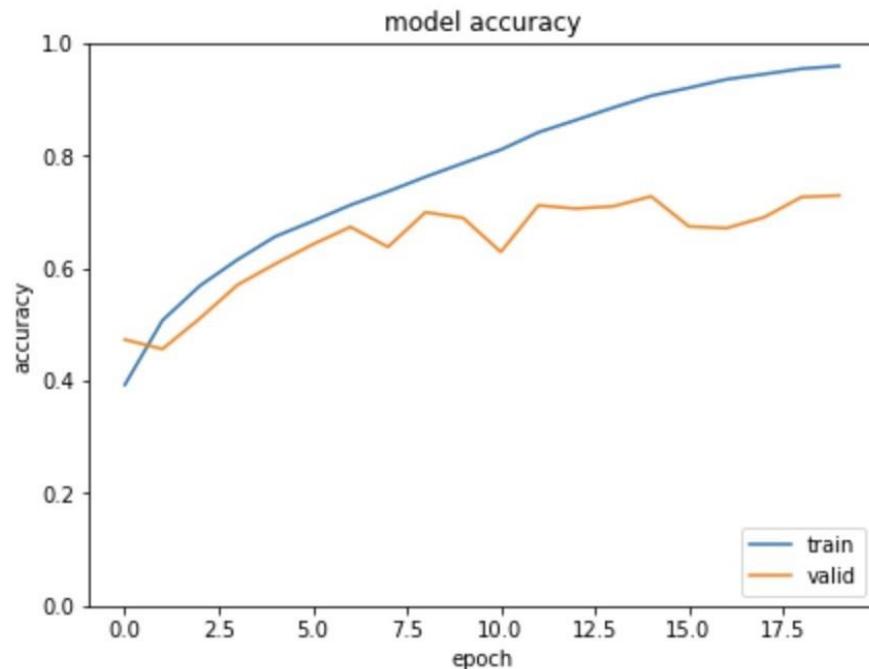
TRAINING AND SAVING

```
In [34]: history = alexnet_model.fit(train_images, train_labels, epochs=20, validation_data=(test_images, test_labels))
          val_accuracy: 0.7104
          Epoch 15/20
          1563/1563 [=====] - 639s 409ms/step - loss: 0.3634 - accuracy: 0.9068 - val_loss: 0.8472 -
          val_accuracy: 0.7281
          Epoch 16/20
          1563/1563 [=====] - 576s 369ms/step - loss: 0.3203 - accuracy: 0.9209 - val_loss: 1.0437 -
          val_accuracy: 0.6746
          Epoch 17/20
          1563/1563 [=====] - 569s 364ms/step - loss: 0.2760 - accuracy: 0.9361 - val_loss: 1.0177 -
          val_accuracy: 0.6715
          Epoch 18/20
          1563/1563 [=====] - 597s 382ms/step - loss: 0.2458 - accuracy: 0.9453 - val_loss: 0.9665 -
          val_accuracy: 0.6907
          Epoch 19/20
          1563/1563 [=====] - 656s 419ms/step - loss: 0.2156 - accuracy: 0.9549 - val_loss: 0.8762 -
          val_accuracy: 0.7269
          Epoch 20/20
          1563/1563 [=====] - 634s 406ms/step - loss: 0.1944 - accuracy: 0.9598 - val_loss: 0.8573 -
          val_accuracy: 0.7295
```

```
In [37]: alexnet_model.save('alexnet.h5')
```

TRAINING AND LOSS CURVES

```
In [35]: save_train_loss_curves(history, 'alexnet_loss_acc.png')
```



PERFORMANCE

```
2022-09-17 00:04:14.825904: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
```

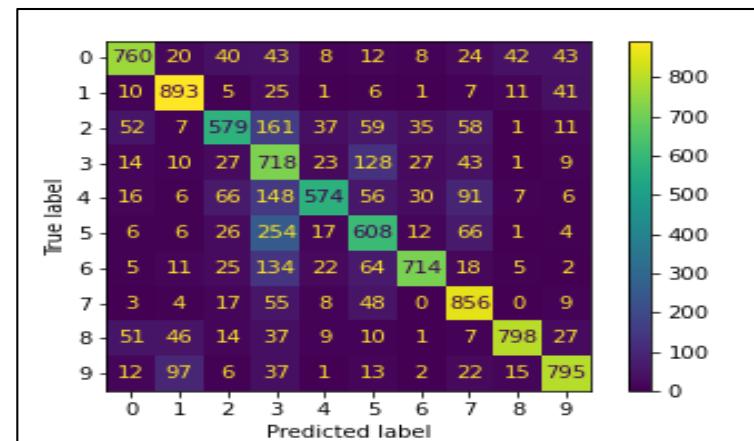
313/313 [=====] - 23s 72ms/step

Confusion Matrix:

```
[[760 20 40 43 8 12 8 24 42 43]
 [ 10 893 5 25 1 6 1 7 11 41]
 [ 52 7 579 161 37 59 35 58 1 11]
 [ 14 10 27 718 23 128 27 43 1 9]
 [ 16 6 66 148 574 56 30 91 7 6]
 [ 6 6 26 254 17 608 12 66 1 4]
 [ 5 11 25 134 22 64 714 18 5 2]
 [ 3 4 17 55 8 48 0 856 0 9]
 [ 51 46 14 37 9 10 1 7 798 27]
 [ 12 97 6 37 1 13 2 22 15 795]]
```

Performance Evaluation:

	precision	recall	f1-score	support
0	0.82	0.76	0.79	1000
1	0.81	0.89	0.85	1000
2	0.72	0.58	0.64	1000
3	0.45	0.72	0.55	1000
4	0.82	0.57	0.68	1000
5	0.61	0.61	0.61	1000
6	0.86	0.71	0.78	1000
7	0.72	0.86	0.78	1000
8	0.91	0.80	0.85	1000
9	0.84	0.80	0.82	1000
accuracy			0.73	10000
macro avg	0.75	0.73	0.73	10000
weighted avg	0.75	0.73	0.73	10000



Accuracy Score:
0.7295

GOOGLENET



IMPORTS, DATA PREPROCESSING AND INCEPTION MODULE

GOOGLENET

```
In [9]: from tensorflow.keras.layers import Flatten, Input, Conv2D, Dense, Dropout, MaxPooling2D, AveragePooling2D, GlobalAvgPool2D
from tensorflow.keras.layers import concatenate
from tensorflow.keras.models import Model
from tensorflow.keras.regularizers import l2

In [11]: train_images_normalized = train_images/255.0
test_images_normalized = test_images/255.0

In [12]: def inception(x,
                  filters_1x1,
                  filters_3x3_reduce,
                  filters_3x3,
                  filters_5x5_reduce,
                  filters_5x5,
                  filters_pool):
    path1 = layers.Conv2D(filters_1x1, (1, 1), padding='same', activation='relu')(x)
    path2 = layers.Conv2D(filters_3x3_reduce, (1, 1), padding='same', activation='relu')(x)
    path2 = layers.Conv2D(filters_3x3, (1, 1), padding='same', activation='relu')(path2)
    path3 = layers.Conv2D(filters_5x5_reduce, (1, 1), padding='same', activation='relu')(x)
    path3 = layers.Conv2D(filters_5x5, (1, 1), padding='same', activation='relu')(path3)
    path4 = layers.MaxPool2D((3, 3), strides=(1, 1), padding='same')(x)
    path4 = layers.Conv2D(filters_pool, (1, 1), padding='same', activation='relu')(path4)
    return tf.concat([path1, path2, path3, path4], axis=3)
```

MODEL ARCHITECTURE

```
In [14]: inp = layers.Input(shape=(32, 32, 3))
input_tensor = layers.experimental.preprocessing.Resizing(224, 224, interpolation="bilinear", input_shape=train_image_size)(inp)
x = layers.Conv2D(64, 7, strides=2, padding='same', activation='relu')(input_tensor)
x = layers.MaxPooling2D(3, strides=2)(x)
x = layers.Conv2D(64, 1, strides=1, padding='same', activation='relu')(x)
x = layers.Conv2D(192, 3, strides=1, padding='same', activation='relu')(x)
x = layers.MaxPooling2D(3, strides=2)(x)
x = inception(x, filters_1x1=64, filters_3x3_reduce=96, filters_3x3=128, filters_5x5_reduce=16, filters_5x5=32, filters_5x5_2=64, filters_5x5_3=96)
x = inception(x, filters_1x1=128, filters_3x3_reduce=128, filters_3x3=192, filters_5x5_reduce=32, filters_5x5=96, filters_5x5_2=64, filters_5x5_3=128)
x = layers.MaxPooling2D(3, strides=2)(x)
x = inception(x, filters_1x1=192, filters_3x3_reduce=96, filters_3x3=208, filters_5x5_reduce=16, filters_5x5=48, filters_5x5_2=64, filters_5x5_3=96)
aux1 = layers.AveragePooling2D((5, 5), strides=3)(x)
aux1 = layers.Conv2D(128, 1, padding='same', activation='relu')(aux1)
aux1 = layers.Flatten()(aux1)
aux1 = layers.Dense(1024, activation='relu')(aux1)
aux1 = layers.Dropout(0.7)(aux1)
aux1 = layers.Dense(10, activation='softmax')(aux1)
x = inception(x, filters_1x1=160, filters_3x3_reduce=112, filters_3x3=224, filters_5x5_reduce=24, filters_5x5=64, filters_5x5_2=64, filters_5x5_3=96)
x = inception(x, filters_1x1=128, filters_3x3_reduce=128, filters_3x3=256, filters_5x5_reduce=24, filters_5x5=64, filters_5x5_2=64, filters_5x5_3=128)
x = inception(x, filters_1x1=112, filters_3x3_reduce=144, filters_3x3=288, filters_5x5_reduce=32, filters_5x5=64, filters_5x5_2=64, filters_5x5_3=128)
aux2 = layers.AveragePooling2D((5, 5), strides=3)(x)
aux2 = layers.Conv2D(128, 1, padding='same', activation='relu')(aux2)
aux2 = layers.Flatten()(aux2)
aux2 = layers.Dense(1024, activation='relu')(aux2)
aux2 = layers.Dropout(0.7)(aux2)
aux2 = layers.Dense(10, activation='softmax')(aux2)
x = inception(x, filters_1x1=256, filters_3x3_reduce=160, filters_3x3=320, filters_5x5_reduce=32, filters_5x5=128, filters_5x5_2=128, filters_5x5_3=192)
x = layers.MaxPooling2D(3, strides=2)(x)
x = inception(x, filters_1x1=256, filters_3x3_reduce=160, filters_3x3=320, filters_5x5_reduce=32, filters_5x5=128, filters_5x5_2=128, filters_5x5_3=192)
x = inception(x, filters_1x1=384, filters_3x3_reduce=192, filters_3x3=384, filters_5x5_reduce=48, filters_5x5=128, filters_5x5_2=128, filters_5x5_3=192)
x = layers.GlobalAveragePooling2D()(x)
x = layers.Dropout(0.4)(x)
out = layers.Dense(10, activation='softmax')(x)

Metal device set to: Apple M1

2022-09-21 21:16:54.156905: I tensorflow/core/common_runtime/pluggable_device/pluggable_device_factory.cc:305] Could not identify NUMA node of platform GPU ID 0, defaulting to 0. Your kernel may not have been built with NUMA support.
```

COMPILATION

```
In [15]: googlenet_model = Model(inputs = inp, outputs = [out, aux1, aux2])
```

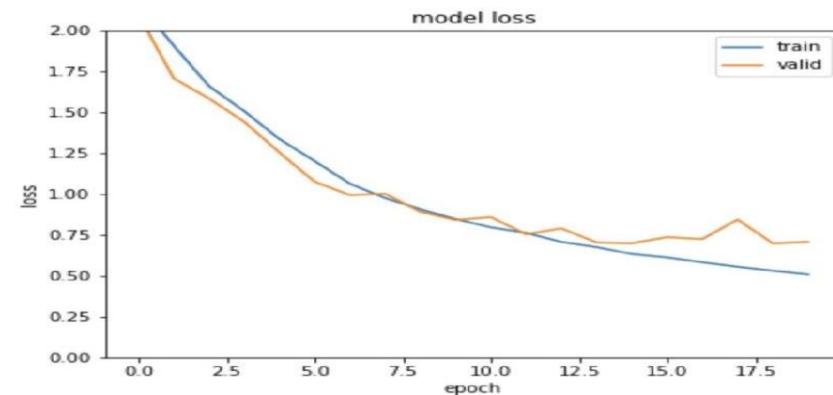
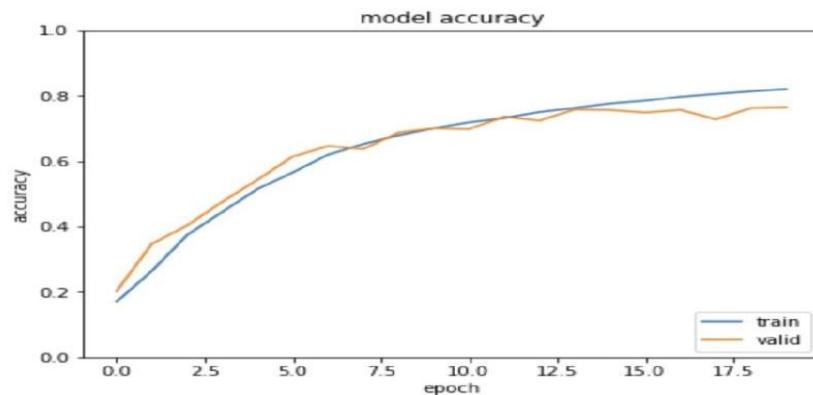
```
In [16]: googlenet_model.compile(optimizer=keras.optimizers.Adam(learning_rate=0.01, epsilon=0.05),
                           loss=[keras.losses.sparse_categorical_crossentropy,
                                 keras.losses.sparse_categorical_crossentropy,
                                 keras.losses.sparse_categorical_crossentropy],
                           loss_weights=[1, 0.3, 0.3],
                           metrics=['accuracy'])
```

TRAINING

```
In [18]: history = googlenet_model.fit(train_images_normalized, [train_labels, train_labels, train_labels], \
                                     validation_data=(test_images_normalized, [test_labels, test_labels, test_labels]), \
                                     batch_size=64, epochs=20)
82/782 [=====] - 1583s 2s/step - loss: 0.9042 - dense_4_loss: 0.5818 - dense_1_loss: 0.63
83 - dense_3_loss: 0.6166 - dense_4_accuracy: 0.7971 - dense_1_accuracy: 0.7675 - dense_3_accuracy: 0.7840 - val_lo
ss: 1.1571 - val_dense_4_loss: 0.7238 - val_dense_1_loss: 0.7220 - val_dense_3_loss: 0.7225 - val_dense_4_accuracy:
0.7569 - val_dense_1_accuracy: 0.7496 - val_dense_3_accuracy: 0.7564
Epoch 18/20
782/782 [=====] - 957s 1s/step - loss: 0.9217 - dense_4_loss: 0.5546 - dense_1_loss: 0.632
4 - dense_3_loss: 0.5915 - dense_4_accuracy: 0.8060 - dense_1_accuracy: 0.7782 - dense_3_accuracy: 0.7909 - val_lo
ss: 1.3123 - val_dense_4_loss: 0.8436 - val_dense_1_loss: 0.7811 - val_dense_3_loss: 0.7814 - val_dense_4_accuracy:
0.7280 - val_dense_1_accuracy: 0.7348 - val_dense_3_accuracy: 0.7318
Epoch 19/20
782/782 [=====] - 960s 1s/step - loss: 0.8858 - dense_4_loss: 0.5302 - dense_1_loss: 0.614
2 - dense_3_loss: 0.5714 - dense_4_accuracy: 0.8140 - dense_1_accuracy: 0.7841 - dense_3_accuracy: 0.8000 - val_lo
ss: 1.1291 - val_dense_4_loss: 0.6975 - val_dense_1_loss: 0.7393 - val_dense_3_loss: 0.6996 - val_dense_4_accuracy:
0.7623 - val_dense_1_accuracy: 0.7431 - val_dense_3_accuracy: 0.7600
Epoch 20/20
782/782 [=====] - 930s 1s/step - loss: 0.8481 - dense_4_loss: 0.5082 - dense_1_loss: 0.586
0 - dense_3_loss: 0.5471 - dense_4_accuracy: 0.8207 - dense_1_accuracy: 0.7938 - dense_3_accuracy: 0.8087 - val_lo
ss: 1.1265 - val_dense_4_loss: 0.7073 - val_dense_1_loss: 0.6846 - val_dense_3_loss: 0.7131 - val_dense_4_accuracy:
0.7648 - val_dense_1_accuracy: 0.7649 - val_dense_3_accuracy: 0.7601
```

TRAINING AND LOSS CURVES

```
In [26]: plt.figure(figsize=(15,5))
plt.subplot(1,2,1)
plt.plot(history.history['dense_4_accuracy'])
plt.plot(history.history['val_dense_4_accuracy'])
plt.ylim(0, 1)
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'valid'], loc='lower right')
plt.subplot(1,2,2)
plt.plot(history.history['dense_4_loss'])
plt.plot(history.history['val_dense_4_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'valid'], loc='upper right')
plt.ylim([0,2])
plt.savefig("googlenet_loss_acc.png")
```



PERFORMANCE

313/313 [=====] - 32s 104ms/step

Confusion Matrix:

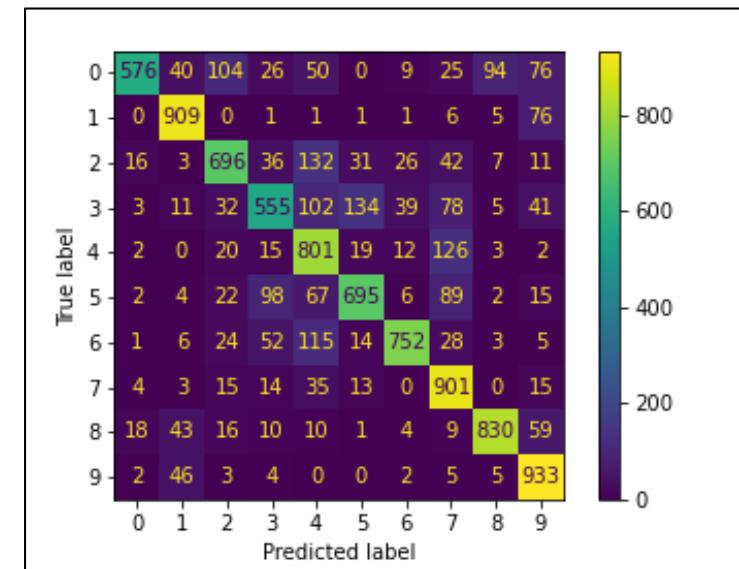
```
[[576 40 104 26 50 0 9 25 94 76]
 [ 0 909 0 1 1 1 6 5 76]
 [ 16 3 696 36 132 31 26 42 7 11]
 [ 3 11 32 555 102 134 39 78 5 41]
 [ 2 0 20 15 801 19 12 126 3 2]
 [ 2 4 22 98 67 695 6 89 2 15]
 [ 1 6 24 52 115 14 752 28 3 5]
 [ 4 3 15 14 35 13 0 901 0 15]
 [ 18 43 16 10 10 1 4 9 830 59]
 [ 2 46 3 4 0 0 2 5 5 933]]
```

Performance Evaluation:

	precision	recall	f1-score	support
0	0.92	0.58	0.71	1000
1	0.85	0.91	0.88	1000
2	0.75	0.70	0.72	1000
3	0.68	0.56	0.61	1000
4	0.61	0.80	0.69	1000
5	0.77	0.69	0.73	1000
6	0.88	0.75	0.81	1000
7	0.69	0.90	0.78	1000
8	0.87	0.83	0.85	1000
9	0.76	0.93	0.84	1000
accuracy			0.76	10000
macro avg	0.78	0.76	0.76	10000
weighted avg	0.78	0.76	0.76	10000

Accuracy Score:

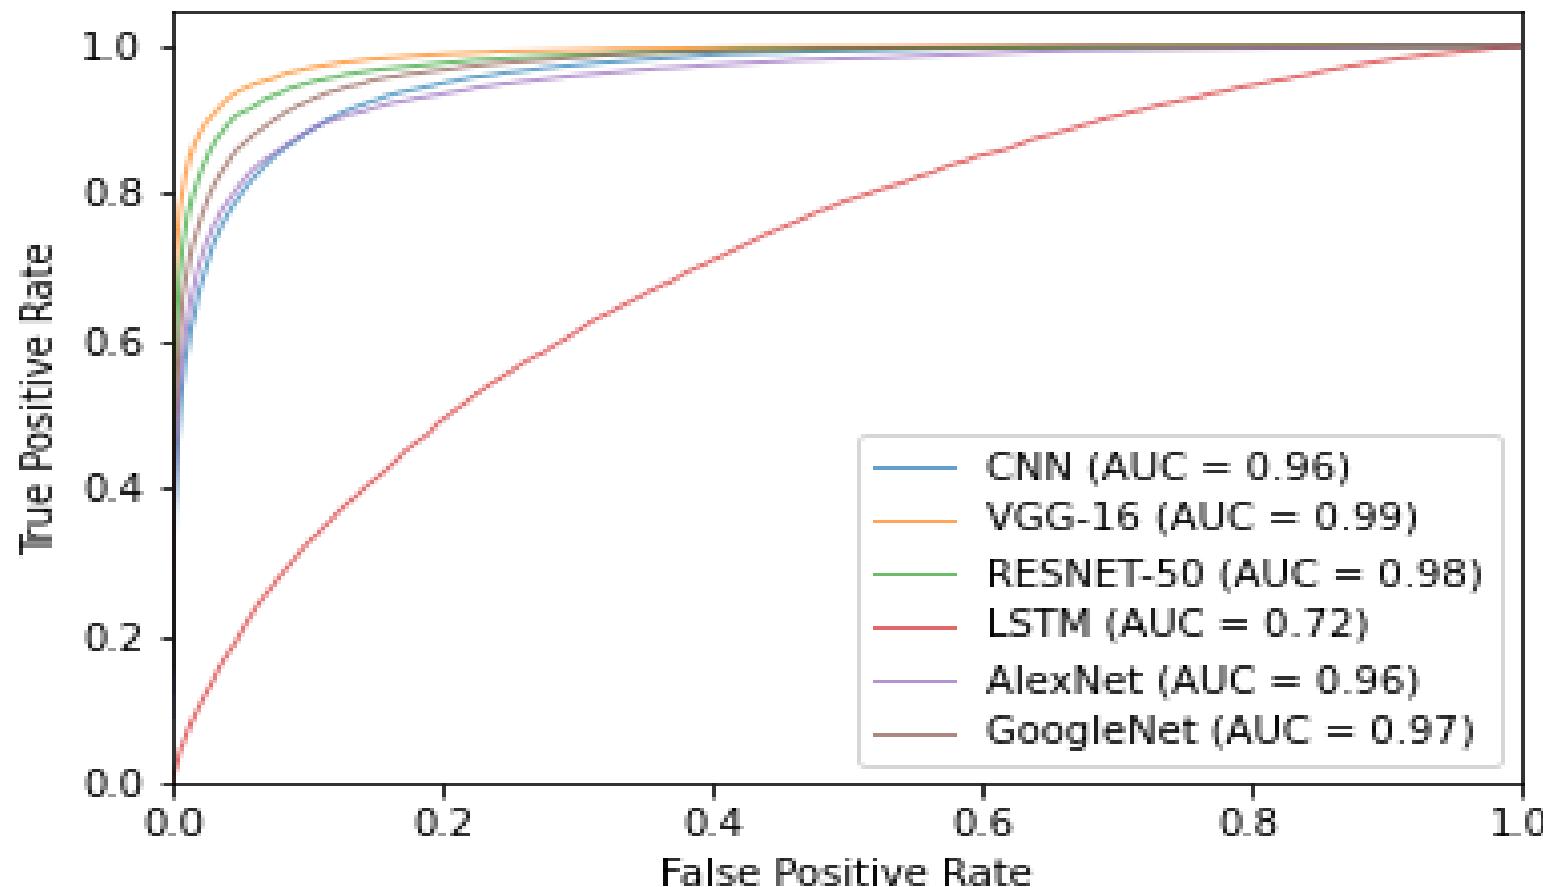
0.7648



COMPARISON BETWEEN CNN, VGG-16, RESNET-50, LSTM(RNN), ALEXNET AND GOOGLENET.



ROC CURVES AND AUC COMPARISON



ROC CURVES AND AUC COMPARISON - CODE (I)

```
In [51]: def plot_macro_avg_roc_curve(classifier, n_classes, name, X_test, y_test, single_output=True):
    y_test_binarized = label_binarize(y_test, classes=list(range(0,n_classes)))
    y_score = classifier.predict(X_test)

    #for cases where the model (like googlenet) has extra (auxillary) outputs
    if not single_output:
        y_score = y_score[0]

    # Compute ROC curve and ROC area for each class
    fpr = dict()
    tpr = dict()
    roc_auc = dict()
    for i in range(n_classes):
        fpr[i], tpr[i], _ = roc_curve(y_test_binarized[:, i], y_score[:, i])
        roc_auc[i] = auc(fpr[i], tpr[i])

    # First aggregate all false positive rates
    all_fpr = np.unique(np.concatenate([fpr[i] for i in range(n_classes)]))

    # Then interpolate all ROC curves at this points
    mean_tpr = np.zeros_like(all_fpr)
    for i in range(n_classes):
        mean_tpr += np.interp(all_fpr, fpr[i], tpr[i])

    # Finally average it and compute AUC
    mean_tpr /= n_classes

    fpr["macro"] = all_fpr
    tpr["macro"] = mean_tpr
    roc_auc["macro"] = auc(fpr["macro"], tpr["macro"])

    plt.plot(
        fpr["macro"],
        tpr["macro"],
        label="{0} (AUC = {1:0.2f})".format(name, roc_auc["macro"]),
        alpha=0.7,
        linewidth=1,
    )

    return plt
```

ROC CURVES AND AUC COMPARISON - CODE (II)

```
In [43]: cnn = load_model('cnn.h5')
vgg_based_model = load_model('vgg16.h5')
resnet_based_model = load_model('resnet50.h5')
lstm_based_model = load_model('lstm.h5')
alexnet_model = load_model('alexnet.h5')
googlenet_model = load_model('googlenet.h5')
```

```
In [58]: plot_macro_avg_roc_curve(cnn, 10, "CNN", test_images, test_labels)
plot_macro_avg_roc_curve(vgg_based_model, 10, "VGG-16", test_images_p, test_labels)
plot_macro_avg_roc_curve(resnet_based_model, 10, "RESNET-50", test_images_pr, test_labels)
plot_macro_avg_roc_curve(lstm_based_model, 10, "LSTM", test_images_res, test_labels)
plot_macro_avg_roc_curve(alexnet_model, 10, "AlexNet", test_images, test_labels)
plot_macro_avg_roc_curve(googlenet_model, 10, "GoogleNet", test_images_normalized, test_labels, single_output=False)

plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.legend(loc="lower right")

plt.savefig('roc_auc.png')
```

```
313/313 [=====] - 2s 6ms/step
313/313 [=====] - 5s 17ms/step
313/313 [=====] - 8s 25ms/step
313/313 [=====] - 92s 294ms/step
313/313 [=====] - 22s 72ms/step
```

```
2022-09-22 16:22:55.443709: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
```

```
313/313 [=====] - 35s 108ms/step
```

COMPARISON OF PERFORMANCE

CLASSIFIER	Accuracy	Precision	Recall	F1-Score	AUC
CNN	0.72	0.73	0.72	0.72	0.96
VGG-16	0.86	0.87	0.86	0.86	0.99
RESNET-50	0.83	0.83	0.83	0.83	0.98
LSTM	0.27	0.27	0.27	0.25	0.72
ALEXNET	0.73	0.75	0.73	0.73	0.96
GOOGLENET	0.76	0.78	0.76	0.76	0.97

CONCLUSION:

GoogleNet performed better than AlexNet which, in turn, performed better than a simpler CNN. This is expected. However, VGG-16 and RESNET-50 performed much better than all the other models. This is because they were directly imported from the Keras library, retaining the weights that they learned whilst being trained on the ‘ImageNet’. This probably enhanced their feature detection ability.

Also, LSTM performed poorly but this was expected. RNNs learn patterns within sequences and the CIFAR-10 dataset simply does not allow this. For example, two pictures of the same dog from different angles would have scarcely any common sequences of pixel intensities.

MNIST - CLASSIFICATION



CNN



IMPORTS AND DATA PREPROCESSING

Convolutional Neural Network

Importing the libraries

```
In [1]: import numpy as np
import tensorflow as tf

from tensorflow import keras
from tensorflow.keras import layers, datasets
from tensorflow.keras.models import Sequential
import matplotlib.pyplot as plt
```

Part 1 - Data Preprocessing

Normalization and augmentation will be done in the CNN itself

```
In [2]: img_height = 28
img_width = 28
num_classes = 10
```

```
In [34]: (train_images, train_labels), (test_images, test_labels) = datasets.mnist.load_data()
```

MODEL ARCHITECTURE

Part 2 - Building the CNN

Building the CNN

```
In [5]: cnn = Sequential([
    layers.Rescaling(1./255, input_shape=(img_height, img_width, 1)),
    layers.Conv2D(filters=32, kernel_size=3, padding='same', activation='relu'),
    layers.Conv2D(filters=32, kernel_size=3, padding='same', activation='relu'),
    layers.MaxPooling2D(pool_size=2, strides=2),
    layers.Conv2D(32, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(pool_size=2, strides=2),
    layers.Dropout(0.2),
    layers.Flatten(),
    layers.Dense(128, activation='relu'),
    layers.Dense(num_classes, activation="softmax")
])
```

COMPILATION AND SUMMARY

Compiling the CNN

```
In [6]: cnn.compile(optimizer=keras.optimizers.Adam(learning_rate=0.0005, epsilon=0.01),
                  loss=tf.keras.losses.SparseCategoricalCrossentropy(),
                  metrics=['accuracy'])
```

```
In [7]: cnn.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
<hr/>		
rescaling_1 (Rescaling)	(None, 28, 28, 1)	0
conv2d_3 (Conv2D)	(None, 28, 28, 32)	320
conv2d_4 (Conv2D)	(None, 28, 28, 32)	9248
max_pooling2d_2 (MaxPooling 2D)	(None, 14, 14, 32)	0
conv2d_5 (Conv2D)	(None, 14, 14, 32)	9248
max_pooling2d_3 (MaxPooling 2D)	(None, 7, 7, 32)	0
dropout_1 (Dropout)	(None, 7, 7, 32)	0
flatten_1 (Flatten)	(None, 1568)	0
dense_2 (Dense)	(None, 128)	200832
dense_3 (Dense)	(None, 10)	1290
<hr/>		
Total params: 220,938		
Trainable params: 220,938		
Non-trainable params: 0		

TRAINING AND SAVING

Training the CNN on the Training set and evaluating it on the Test set

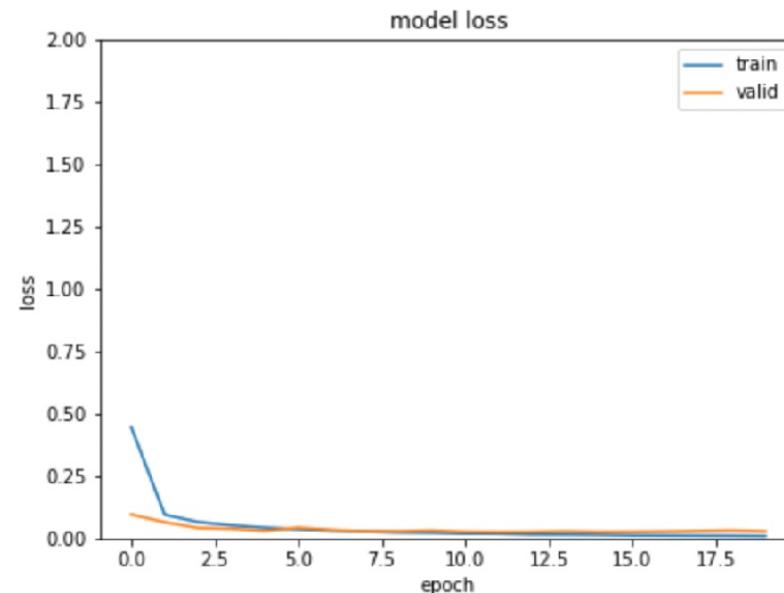
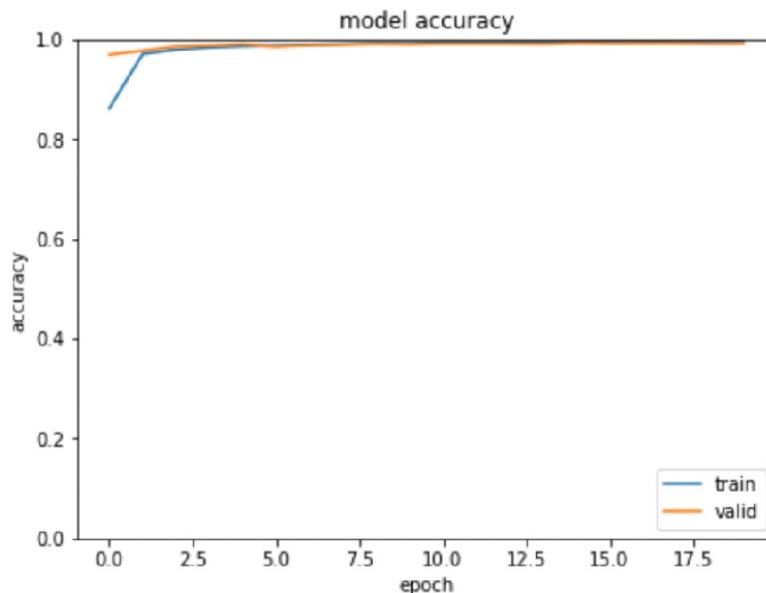
```
In [8]: history = cnn.fit(train_images, train_labels, epochs=20, validation_data=(test_images, test_labels))  
1875/1875 [=====] - 20s 11ms/step - loss: 0.0138 - accuracy: 0.9956 - val_loss: 0.0252 - v  
al_accuracy: 0.9908  
Epoch 15/20  
1875/1875 [=====] - 20s 11ms/step - loss: 0.0118 - accuracy: 0.9965 - val_loss: 0.0258 - v  
al_accuracy: 0.9927  
Epoch 16/20  
1875/1875 [=====] - 20s 11ms/step - loss: 0.0111 - accuracy: 0.9965 - val_loss: 0.0273 - v  
al_accuracy: 0.9918  
Epoch 17/20  
1875/1875 [=====] - 20s 11ms/step - loss: 0.0104 - accuracy: 0.9965 - val_loss: 0.0295 - v  
al_accuracy: 0.9919  
Epoch 18/20  
1875/1875 [=====] - 20s 11ms/step - loss: 0.0100 - accuracy: 0.9965 - val_loss: 0.0314 - v  
al_accuracy: 0.9920  
Epoch 19/20  
1875/1875 [=====] - 20s 11ms/step - loss: 0.0090 - accuracy: 0.9973 - val_loss: 0.0283 - v  
al_accuracy: 0.9916  
Epoch 20/20
```

```
In [9]: cnn.save('cnn.h5')
```

TRAINING AND LOSS CURVES

```
In [14]: def save_train_loss_curves(history, filename):
    plt.figure(figsize=(15,5))
    plt.subplot(1,2,1)
    plt.plot(history.history['accuracy'])
    plt.plot(history.history['val_accuracy'])
    plt.ylim(0, 1)
    plt.title('model accuracy')
    plt.ylabel('accuracy')
    plt.xlabel('epoch')
    plt.legend(['train', 'valid'], loc='lower right')
    plt.subplot(1,2,2)
    plt.plot(history.history['loss'])
    plt.plot(history.history['val_loss'])
    plt.title('model loss')
    plt.ylabel('loss')
    plt.xlabel('epoch')
    plt.legend(['train', 'valid'], loc='upper right')
    plt.ylim([0,2])
    plt.savefig(filename)
```

```
In [15]: save_train_loss_curves(history, 'cnn_loss_acc.png')
```



PERFORMANCE - CODE

Performance

```
In [11]: from sklearn.metrics import classification_report, confusion_matrix, accuracy_score, ConfusionMatrixDisplay
```

```
In [12]: def get_predictions(model, inputs):

    def index_max_arr_element(arr):
        best_index = -1
        best = float('-inf')
        for i, num in enumerate(arr):
            if num > best:
                best = num
                best_index = i
        return best_index

    pred_arr = model.predict(inputs)
    pred_labels = [index_max_arr_element(row) for row in pred_arr]
    return pred_labels
```

```
In [13]: pred_labels = get_predictions(cnn, test_images)

print("Confusion Matrix:")
print(confusion_matrix(test_labels, pred_labels))
print("-----")
print("-----")
print("Performance Evaluation:")
print(classification_report(test_labels, pred_labels))
print("-----")
print("-----")
print("Accuracy Score:")
print(accuracy_score(test_labels, pred_labels))

plt.figure()
ConfusionMatrixDisplay.from_predictions(test_labels, pred_labels)
plt.savefig('./cnn_matrix.png')
```

```
37/313 [==>.....] - ETA: 1s
```

```
2022-09-15 19:38:44.820728: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin op
```

PERFORMANCE

37/313 [==>.....] - ETA: 1s

2022-09-15 19:38:44.820728: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.

313/313 [=====] - 1s 4ms/step

Confusion Matrix:

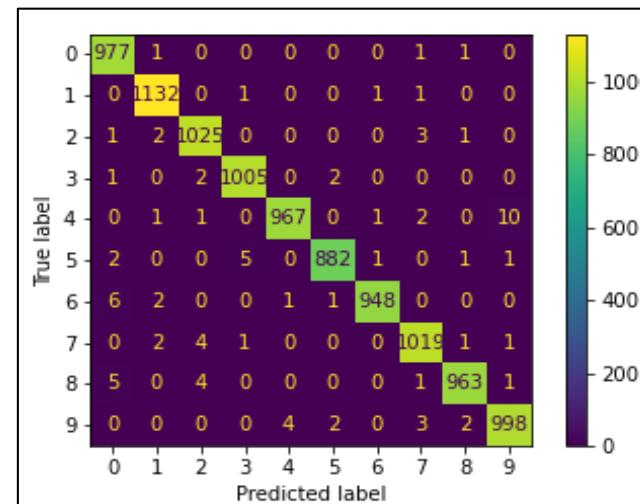
```
[[ 977   1   0   0   0   0   0   1   1   0]
 [  0 1132   0   1   0   0   1   1   0   0]
 [  1   2 1025   0   0   0   0   3   1   0]
 [  1   0   2 1005   0   2   0   0   0   0]
 [  0   1   1   0 967   0   1   2   0   10]
 [  2   0   0   5   0 882   1   0   1   1]
 [  6   2   0   0   1   1 948   0   0   0]
 [  0   2   4   1   0   0   0 1019   1   1]
 [  5   0   4   0   0   0   0   1 963   1]
 [  0   0   0   0   4   2   0   3   2 998]]
```

Performance Evaluation:

	precision	recall	f1-score	support
0	0.98	1.00	0.99	980
1	0.99	1.00	1.00	1135
2	0.99	0.99	0.99	1032
3	0.99	1.00	0.99	1010
4	0.99	0.98	0.99	982
5	0.99	0.99	0.99	892
6	1.00	0.99	0.99	958
7	0.99	0.99	0.99	1028
8	0.99	0.99	0.99	974
9	0.99	0.99	0.99	1009
accuracy			0.99	10000
macro avg	0.99	0.99	0.99	10000
weighted avg	0.99	0.99	0.99	10000

Accuracy Score:

0.9916



VGG-16

IMPORTS, DATA PREPROCESSING & ARCHITECTURE

VGG-16

```
In [16]: from tensorflow.keras.applications import vgg16
```

Preprocessing data to make it compatible to VGG and RESNET models

```
In [41]: train_images_stacked = np.dstack([train_images] * 3)
test_images_stacked = np.dstack([test_images]*3)

test_images_stacked = test_images_stacked.reshape(-1,28,28,3)
train_images_stacked = train_images_stacked.reshape(-1,28,28,3)

train_images_resized= tf.image.resize(train_images_stacked, tf.constant([32, 32]))
test_images_resized= tf.image.resize(test_images_stacked, tf.constant([32, 32]))
```

```
In [42]: train_images_p, test_images_p = vgg16.preprocess_input(train_images_resized), vgg16.preprocess_input(test_images_res
```

```
In [43]: vgg_based_model = Sequential([
    layers.Rescaling(1./255, input_shape=(32, 32, 3)),
    vgg16.VGG16(include_top=False, weights="imagenet", pooling="max", input_shape=(32, 32, 3)),
    layers.Flatten(),
    layers.Dense(512, activation='relu'),
    layers.Dropout(0.2),
    layers.Dense(128, activation='relu'),
    layers.Dropout(0.2),
    layers.Dense(num_classes, activation="softmax")
])
```

COMPIRATION AND SUMMARY

```
In [44]: vgg_based_model.compile(optimizer=keras.optimizers.Adam(learning_rate=0.0002, epsilon=0.01),  
                           loss=tf.keras.losses.SparseCategoricalCrossentropy(),  
                           metrics=['accuracy'])
```

```
In [45]: vgg_based_model.summary()
```

Model: "sequential_7"

Layer (type)	Output Shape	Param #
rescaling_13 (Rescaling)	(None, 32, 32, 3)	0
vgg16 (Functional)	(None, 512)	14714688
flatten_5 (Flatten)	(None, 512)	0
dense_17 (Dense)	(None, 512)	262656
dropout_14 (Dropout)	(None, 512)	0
dense_18 (Dense)	(None, 128)	65664
dropout_15 (Dropout)	(None, 128)	0
dense_19 (Dense)	(None, 10)	1290
<hr/>		
Total params: 15,044,298		
Trainable params: 15,044,298		
Non-trainable params: 0		

TRAINING AND SAVING

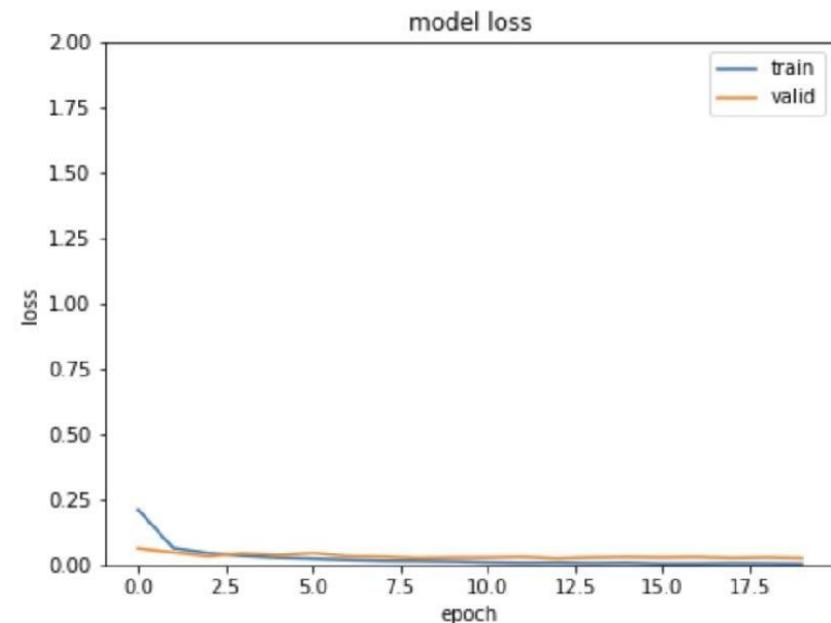
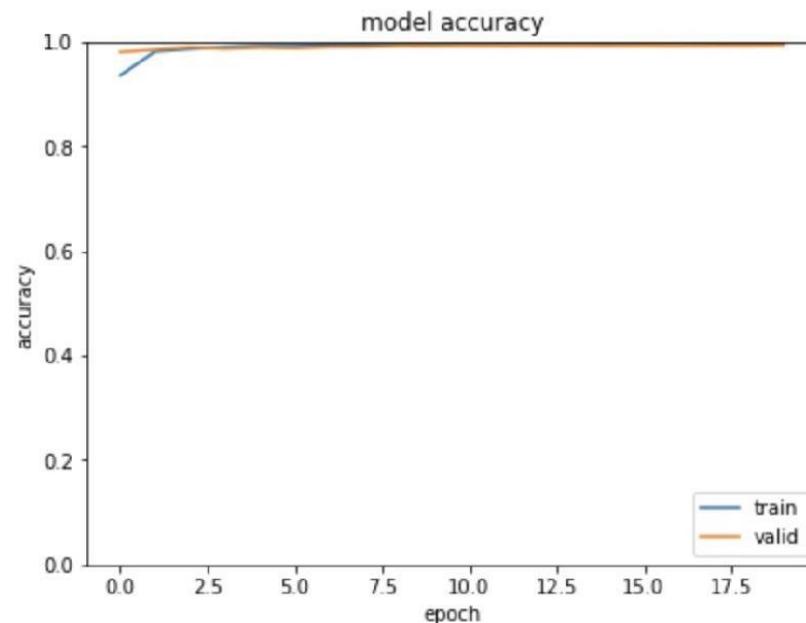
```
In [46]: history = vgg_based_model.fit(train_images_p, train_labels, epochs=20, validation_data=(test_images_p, test_labels))
```

```
1875/1875 [=====] - 1793 95ms/step - loss: 0.0005 - accuracy: 0.9995 - val_loss: 0.0274 - val_accuracy: 0.9927
Epoch 15/20
1875/1875 [=====] - 188s 100ms/step - loss: 0.0077 - accuracy: 0.9976 - val_loss: 0.0290 - val_accuracy: 0.9928
Epoch 16/20
1875/1875 [=====] - 180s 96ms/step - loss: 0.0040 - accuracy: 0.9989 - val_loss: 0.0276 - val_accuracy: 0.9933
Epoch 17/20
1875/1875 [=====] - 177s 95ms/step - loss: 0.0043 - accuracy: 0.9987 - val_loss: 0.0291 - val_accuracy: 0.9927
Epoch 18/20
1875/1875 [=====] - 204s 109ms/step - loss: 0.0057 - accuracy: 0.9985 - val_loss: 0.0254 - val_accuracy: 0.9934
Epoch 19/20
1875/1875 [=====] - 247s 132ms/step - loss: 0.0051 - accuracy: 0.9984 - val_loss: 0.0275 - val_accuracy: 0.9935
Epoch 20/20
1875/1875 [=====] - 277s 148ms/step - loss: 0.0033 - accuracy: 0.9990 - val_loss: 0.0247 - val_accuracy: 0.9942
```

```
In [48]: vgg_based_model.save('vgg16.h5')
```

TRAINING AND LOSS CURVES

```
In [47]: save_train_loss_curves(history, 'vgg_loss_acc.png')
```



PERFORMANCE - CODE

```
In [63]: pred_labels = get_predictions(vgg_based_model, test_images_p)

print("Confusion Matrix:")
print(confusion_matrix(test_labels, pred_labels))
print("-----")
print("-----")
print("Performance Evaluation:")
print(classification_report(test_labels, pred_labels))
print("-----")
print("-----")
print("Accuracy Score:")
print(accuracy_score(test_labels, pred_labels))

plt.figure()
ConfusionMatrixDisplay.from_predictions(test_labels, pred_labels)
plt.savefig('./vgg_matrix.png')
```

PERFORMANCE

5/313 [.....] - ETA: 4s

2022-09-16 11:40:18.080302: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.

313/313 [=====] - 5s 14ms/step

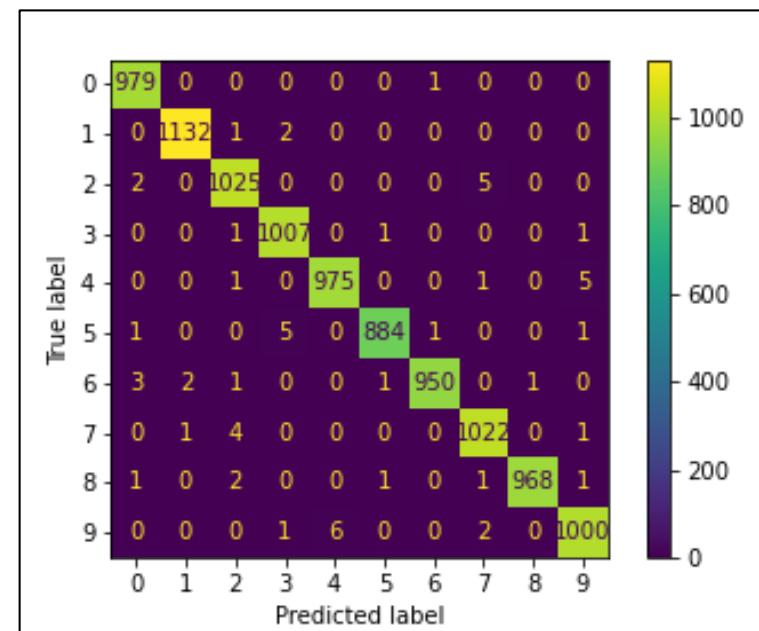
Confusion Matrix:

```
[[ 979  0  0  0  0  0  1  0  0  0]
 [ 0 1132  1  2  0  0  0  0  0  0]
 [ 2  0 1025  0  0  0  0  5  0  0]
 [ 0  0  1 1007  0  1  0  0  0  1]
 [ 0  0  1  0  975  0  0  1  0  5]
 [ 1  0  0  5  0  884  1  0  0  1]
 [ 3  2  1  0  0  1  950  0  1  0]
 [ 0  1  4  0  0  0  0  1022  0  1]
 [ 1  0  2  0  0  1  0  1  968  1]
 [ 0  0  0  1  6  0  0  2  0  1000]]
```

Performance Evaluation:

	precision	recall	f1-score	support
0	0.99	1.00	1.00	980
1	1.00	1.00	1.00	1135
2	0.99	0.99	0.99	1032
3	0.99	1.00	0.99	1010
4	0.99	0.99	0.99	982
5	1.00	0.99	0.99	892
6	1.00	0.99	0.99	958
7	0.99	0.99	0.99	1028
8	1.00	0.99	1.00	974
9	0.99	0.99	0.99	1009
accuracy			0.99	10000
macro avg	0.99	0.99	0.99	10000
weighted avg	0.99	0.99	0.99	10000

Accuracy Score:
0.9942



RES NET- 50

IMPORTS, DATA PREPROCESSING & ARCHITECTURE

RESNET-50

```
In [49]: from tensorflow.keras.applications import resnet50
```

```
In [50]: train_images_pr, test_images_pr = resnet50.preprocess_input(train_images_resized), resnet50.preprocess_input(test_im
```

```
In [51]: resnet_based_model = Sequential([
    layers.Rescaling(1./255, input_shape=(32, 32, 3)),
    resnet50.ResNet50(include_top=False, weights="imagenet", pooling="max", input_shape=(32, 32, 3)),
    layers.Flatten(),
    layers.Dense(512, activation='relu'),
    layers.Dropout(0.2),
    layers.Dense(128, activation='relu'),
    layers.Dropout(0.2),
    layers.Dense(num_classes, activation="softmax")
])
```

COMPILATION AND SUMMARY

```
In [52]: resnet_based_model.compile(optimizer=keras.optimizers.Adam(learning_rate=0.0002, epsilon=0.01),
                                loss=tf.keras.losses.SparseCategoricalCrossentropy(),
                                metrics=['accuracy'])
```

```
In [53]: resnet_based_model.summary()
```

Model: "sequential_8"

Layer (type)	Output Shape	Param #
rescaling_14 (Rescaling)	(None, 32, 32, 3)	0
resnet50 (Functional)	(None, 2048)	23587712
flatten_6 (Flatten)	(None, 2048)	0
dense_20 (Dense)	(None, 512)	1049088
dropout_16 (Dropout)	(None, 512)	0
dense_21 (Dense)	(None, 128)	65664
dropout_17 (Dropout)	(None, 128)	0
dense_22 (Dense)	(None, 10)	1290
<hr/>		
Total params: 24,703,754		
Trainable params: 24,650,634		
Non-trainable params: 53,120		

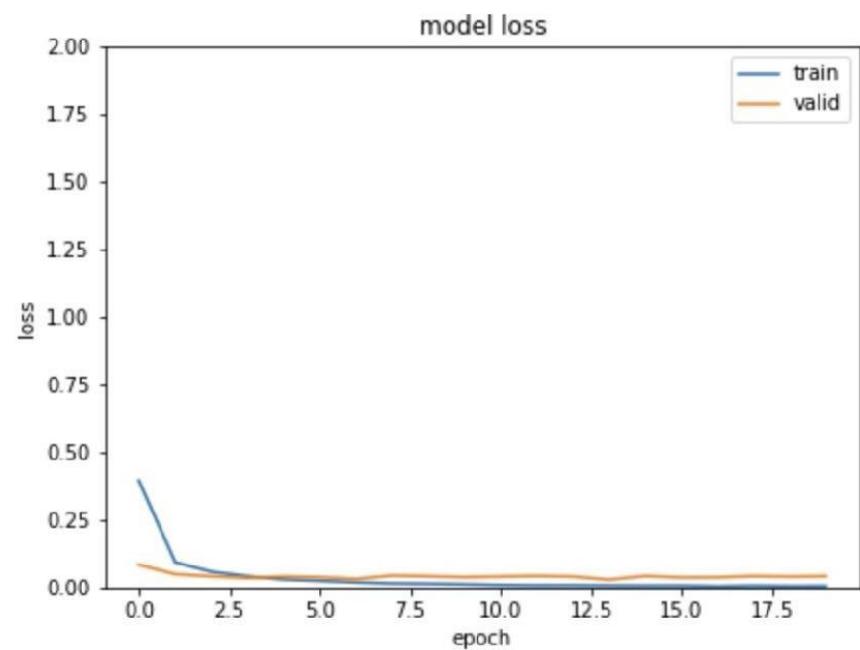
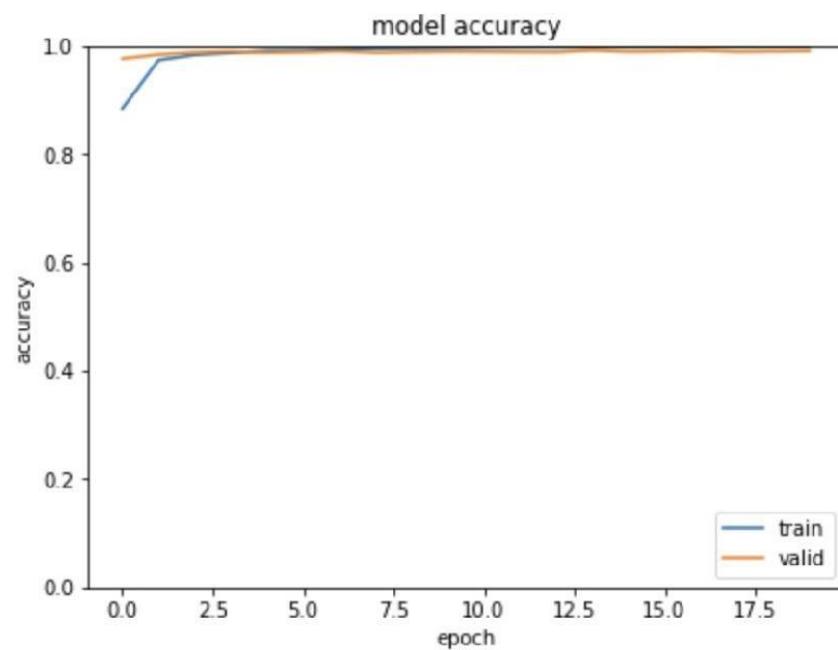
TRAINING AND SAVING

```
In [54]: history = resnet_based_model.fit(train_images_pr, train_labels, epochs=20, validation_data=(test_images_pr, test_labels))
Epoch 1/20
1875/1875 [=====] - 178s 92ms/step - loss: 0.0053 - accuracy: 0.9982 - val_loss: 0.0420 - val_accuracy: 0.9928
Epoch 15/20
1875/1875 [=====] - 196s 105ms/step - loss: 0.0053 - accuracy: 0.9985 - val_loss: 0.0410 - val_accuracy: 0.9903
Epoch 16/20
1875/1875 [=====] - 211s 112ms/step - loss: 0.0054 - accuracy: 0.9984 - val_loss: 0.0357 - val_accuracy: 0.9910
Epoch 17/20
1875/1875 [=====] - 235s 125ms/step - loss: 0.0030 - accuracy: 0.9991 - val_loss: 0.0363 - val_accuracy: 0.9926
Epoch 18/20
1875/1875 [=====] - 256s 136ms/step - loss: 0.0051 - accuracy: 0.9986 - val_loss: 0.0408 - val_accuracy: 0.9902
Epoch 19/20
1875/1875 [=====] - 314s 168ms/step - loss: 0.0035 - accuracy: 0.9988 - val_loss: 0.0389 - val_accuracy: 0.9908
Epoch 20/20
1875/1875 [=====] - 344s 184ms/step - loss: 0.0040 - accuracy: 0.9988 - val_loss: 0.0405 - val_accuracy: 0.9915
```

```
In [56]: resnet_based_model.save('resnet50.h5')
```

TRAINING AND LOSS CURVES

```
In [55]: save_train_loss_curves(history, 'resnet_loss_acc.png')
```



PERFORMANCE - CODE

```
In [64]: pred_labels = get_predictions(resnet_based_model, test_images_pr)

print("Confusion Matrix:")
print(confusion_matrix(test_labels, pred_labels))
print("-----")
print("-----")
print("Performance Evaluation:")
print(classification_report(test_labels, pred_labels))
print("-----")
print("-----")
print("Accuracy Score:")
print(accuracy_score(test_labels, pred_labels))

plt.figure()
ConfusionMatrixDisplay.from_predictions(test_labels, pred_labels)
plt.savefig('./resnet_matrix.png')
```

2022-09-16 11:41:07.726305: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.

PERFORMANCE

```
2022-09-16 11:41:07.726305: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
```

313/313 [=====] - 8s 25ms/step

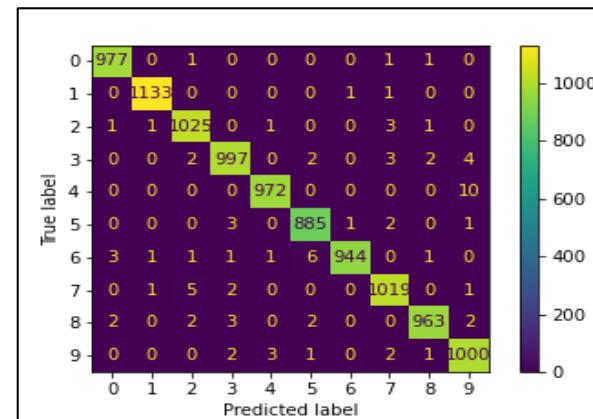
Confusion Matrix:

[977 0 1 0 0 0 0 1 1 0]
[0 1133 0 0 0 0 1 1 0 0]
[1 1 1025 0 1 0 0 3 1 0]
[0 0 2 997 0 2 0 3 2 4]
[0 0 0 0 972 0 0 0 0 10]
[0 0 0 3 0 885 1 2 0 1]
[3 1 1 1 1 6 944 0 1 0]
[0 1 5 2 0 0 0 1019 0 1]
[2 0 2 3 0 2 0 0 963 2]
[0 0 0 2 3 1 0 2 1 1000]]

Performance Evaluation:

	precision	recall	f1-score	support
0	0.99	1.00	1.00	980
1	1.00	1.00	1.00	1135
2	0.99	0.99	0.99	1032
3	0.99	0.99	0.99	1010
4	0.99	0.99	0.99	982
5	0.99	0.99	0.99	892
6	1.00	0.99	0.99	958
7	0.99	0.99	0.99	1028
8	0.99	0.99	0.99	974
9	0.98	0.99	0.99	1009
accuracy			0.99	10000
macro avg	0.99	0.99	0.99	10000
weighted avg	0.99	0.99	0.99	10000

Accuracy Score:
0.9915



RNN (LSTM)



ARCHITECTURE AND COMPIILATION

```
In [57]: lstm_based_model = Sequential([
    layers.Rescaling(1./255, input_shape=(28, 28)),
    layers.LSTM(units=50, return_sequences=True),
    layers.Dropout(0.2),
    layers.LSTM(units=50),
    layers.Dropout(0.2),
    layers.Dense(128, activation='relu'),
    layers.Dropout(0.2),
    layers.Dense(num_classes, activation="softmax")
])

In [58]: lstm_based_model.compile(optimizer=keras.optimizers.Adam(learning_rate=0.001),
                               loss=tf.keras.losses.SparseCategoricalCrossentropy(),
                               metrics=['accuracy'])

In [60]: lstm_based_model.summary()

Model: "sequential_9"


```

Layer (type) Output Shape Param #
=====
rescaling_15 (Rescaling) (None, 28, 28) 0
lstm_4 (LSTM) (None, 28, 50) 15800
dropout_18 (Dropout) (None, 28, 50) 0
lstm_5 (LSTM) (None, 50) 20200
dropout_19 (Dropout) (None, 50) 0
dense_23 (Dense) (None, 128) 6528
dropout_20 (Dropout) (None, 128) 0
dense_24 (Dense) (None, 10) 1290
=====
Total params: 43,818
Trainable params: 43,818

```


```

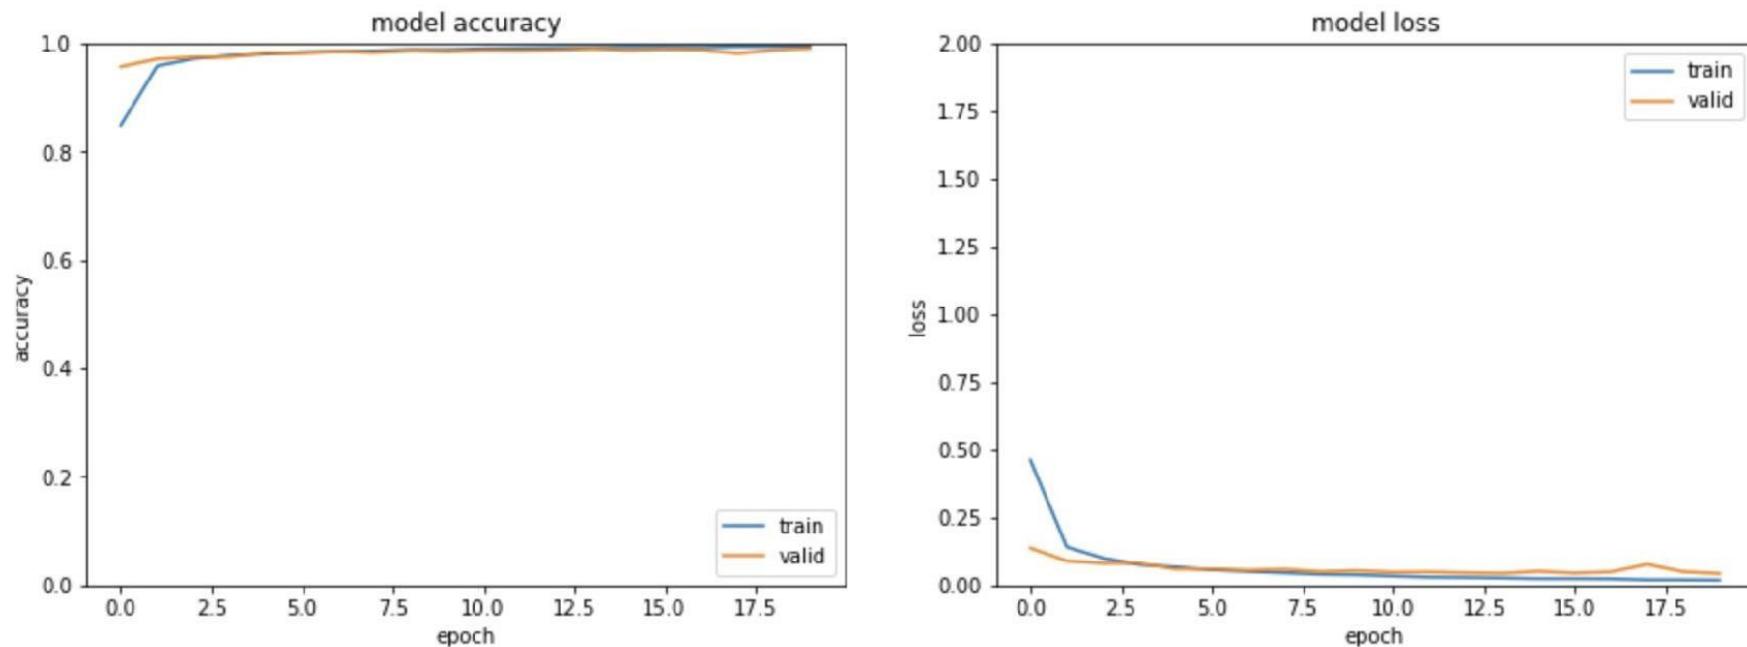
TRAINING AND SAVING

```
In [59]: history = lstm_based_model.fit(train_images, train_labels, epochs=20, validation_data=(test_images, test_labels))
1875/1875 [=====] - 46s 24ms/step - loss: 0.0236 - accuracy: 0.9929 - val_loss: 0.0506 - val_accuracy: 0.9895
Epoch 15/20
1875/1875 [=====] - 46s 24ms/step - loss: 0.0236 - accuracy: 0.9929 - val_loss: 0.0444 - val_accuracy: 0.9870
Epoch 16/20
1875/1875 [=====] - 45s 24ms/step - loss: 0.0236 - accuracy: 0.9929 - val_loss: 0.0444 - val_accuracy: 0.9884
Epoch 17/20
1875/1875 [=====] - 45s 24ms/step - loss: 0.0229 - accuracy: 0.9927 - val_loss: 0.0485 - val_accuracy: 0.9879
Epoch 18/20
1875/1875 [=====] - 46s 24ms/step - loss: 0.0192 - accuracy: 0.9941 - val_loss: 0.0766 - val_accuracy: 0.9816
Epoch 19/20
1875/1875 [=====] - 45s 24ms/step - loss: 0.0190 - accuracy: 0.9940 - val_loss: 0.0491 - val_accuracy: 0.9875
Epoch 20/20
1875/1875 [=====] - 45s 24ms/step - loss: 0.0180 - accuracy: 0.9945 - val_loss: 0.0426 - val_accuracy: 0.9905
```

```
In [76]: lstm_based_model.save('lstm.h5')
```

TRAINING AND LOSS CURVES

```
In [75]: save_train_loss_curves(history, 'lstm_loss_acc.png')
```



PERFORMANCE - CODE

```
In [65]: pred_labels = get_predictions(lstm_based_model, test_images)

print("Confusion Matrix:")
print(confusion_matrix(test_labels, pred_labels))
print("-----")
print("-----")
print("Performance Evaluation:")
print(classification_report(test_labels, pred_labels))
print("-----")
print("-----")
print("Accuracy Score:")
print(accuracy_score(test_labels, pred_labels))

plt.figure()
ConfusionMatrixDisplay.from_predictions(test_labels, pred_labels)
plt.savefig('./lstm_matrix.png')
```

PERFORMANCE

```
2022-09-16 11:41:58.096347: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin op
timer for device_type GPU is enabled.
2022-09-16 11:41:58.157487: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin op
timer for device_type GPU is enabled.
```

```
313/313 [=====] - 3s 9ms/step
```

Confusion Matrix:

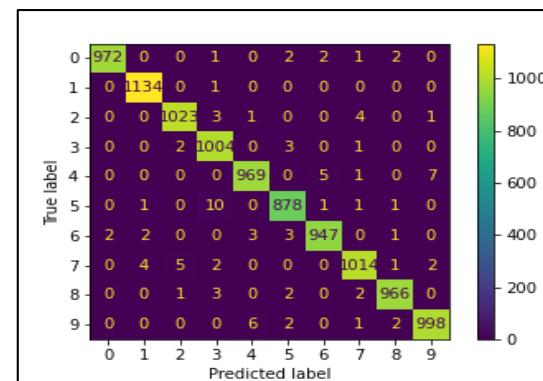
```
[[ 972   0   0   1   0   2   2   1   2   0]
 [  0 1134   0   1   0   0   0   0   0   0]
 [  0   0 1023   3   1   0   0   4   0   1]
 [  0   0   2 1004   0   3   0   1   0   0]
 [  0   0   0   0 969   0   5   1   0   7]
 [  0   1   0   10   0 878   1   1   1   0]
 [  2   2   0   0   3   3 947   0   1   0]
 [  0   4   5   2   0   0   0 1014   1   2]
 [  0   0   1   3   0   2   0   2 966   0]
 [  0   0   0   0   6   2   0   1   2 998]]
```

Performance Evaluation:

	precision	recall	f1-score	support
0	1.00	0.99	0.99	980
1	0.99	1.00	1.00	1135
2	0.99	0.99	0.99	1032
3	0.98	0.99	0.99	1010
4	0.99	0.99	0.99	982
5	0.99	0.98	0.99	892
6	0.99	0.99	0.99	958
7	0.99	0.99	0.99	1028
8	0.99	0.99	0.99	974
9	0.99	0.99	0.99	1009
accuracy			0.99	10000
macro avg	0.99	0.99	0.99	10000
weighted avg	0.99	0.99	0.99	10000

Accuracy Score:

0.9905



ALEXNET



IMPORTS & DATA PREPROCESSING

ALEXNET

```
In [66]: from tensorflow.keras.layers import Dense, Dropout, Activation, Flatten, Conv2D, MaxPooling2D, ZeroPadding2D, BatchN  
from tensorflow.keras.regularizers import l2
```

```
In [68]: train_images_stacked = np.dstack([train_images] * 3)  
test_images_stacked = np.dstack([test_images]*3)  
  
test_images_stacked = test_images_stacked.reshape(-1,28,28,3)  
train_images_stacked = train_images_stacked.reshape(-1,28,28,3)
```

MODEL ARCHITECTURE - I

```
In [69]: def alexnet_model(img_shape=(28, 28, 3), n_classes=10, l2_reg=0.,
weights=None):

    # Initialize model
    alexnet = Sequential()

    # Layer 1
    alexnet.add(Conv2D(96, (11, 11), input_shape=img_shape, padding='same', kernel_regularizer=l2(l2_reg)))
    alexnet.add(BatchNormalization())
    alexnet.add(Activation('relu'))
    alexnet.add(MaxPooling2D(pool_size=(2, 2)))

    # Layer 2
    alexnet.add(Conv2D(256, (5, 5), padding='same'))
    alexnet.add(BatchNormalization())
    alexnet.add(Activation('relu'))
    alexnet.add(MaxPooling2D(pool_size=(2, 2)))

    # Layer 3
    alexnet.add(ZeroPadding2D((1, 1)))
    alexnet.add(Conv2D(512, (3, 3), padding='same'))
    alexnet.add(BatchNormalization())
    alexnet.add(Activation('relu'))
    alexnet.add(MaxPooling2D(pool_size=(2, 2)))

    # Layer 4
    alexnet.add(ZeroPadding2D((1, 1)))
    alexnet.add(Conv2D(1024, (3, 3), padding='same'))
    alexnet.add(BatchNormalization())
    alexnet.add(Activation('relu'))

    # Layer 5
    alexnet.add(ZeroPadding2D((1, 1)))
    alexnet.add(Conv2D(1024, (3, 3), padding='same'))
    alexnet.add(BatchNormalization())
    alexnet.add(Activation('relu'))
    alexnet.add(MaxPooling2D(pool_size=(2, 2)))

    # Layer 6
    alexnet.add(Flatten())
    alexnet.add(Dense(3072))
```

MODEL ARCHITECTURE - II

```
# Layer 6
alexnet.add(Flatten())
alexnet.add(Dense(3072))
alexnet.add(BatchNormalization())
alexnet.add(Activation('relu'))
alexnet.add(Dropout(0.5))

# Layer 7
alexnet.add(Dense(4096))
alexnet.add(BatchNormalization())
alexnet.add(Activation('relu'))
alexnet.add(Dropout(0.5))

# Layer 8
alexnet.add(Dense(n_classes))
alexnet.add(BatchNormalization())
alexnet.add(Activation('softmax'))

if weights is not None:
    alexnet.load_weights(weights)

return alexnet
```

```
In [70]: alexnet_model = alexnet_model(img_shape=(28, 28, 3), n_classes=10)
```

```
In [71]: alexnet_model.compile(optimizer=keras.optimizers.Adam(learning_rate=0.0002, epsilon=0.05),
                           loss=tf.keras.losses.SparseCategoricalCrossentropy(),
                           metrics=['accuracy'])
```

TRAINING AND SAVING

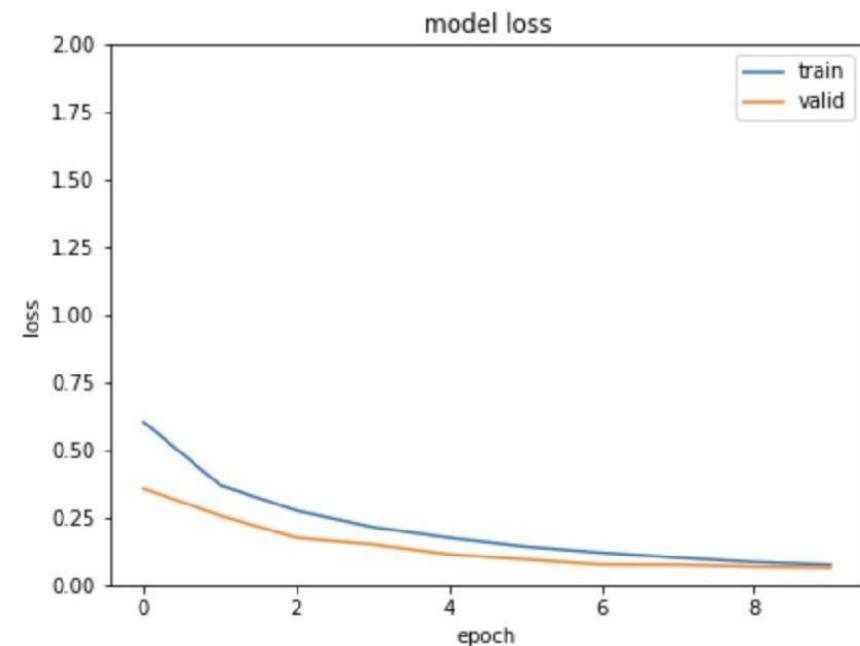
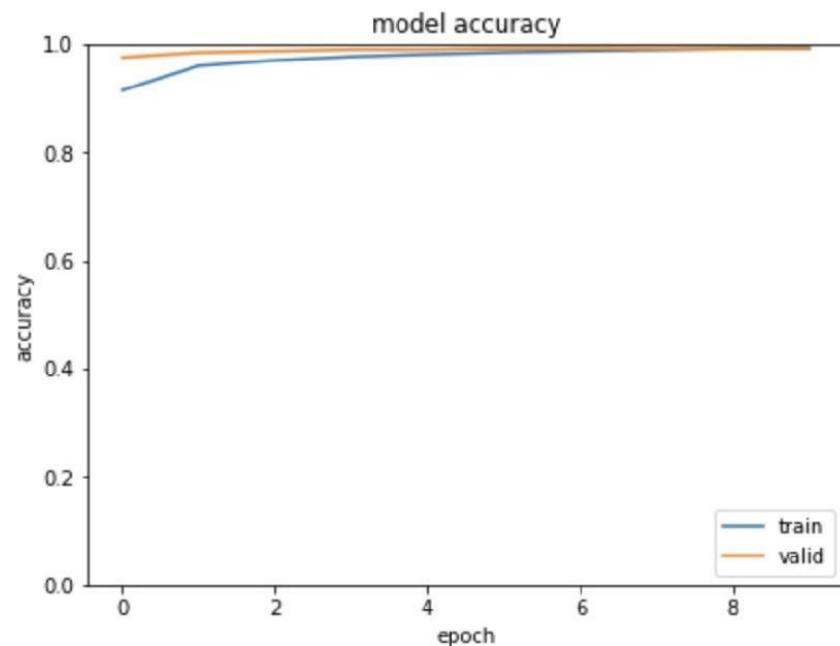
```
In [77]: history = alexnet_model.fit(train_images_stacked, train_labels, epochs=10, validation_data=(test_images_stacked, tes
Epoch 1/10
1875/1875 [=====] - ETA: 0s - loss: 0.6029 - accuracy: 0.9153
2022-09-16 18:14:56.536461: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin op
timizer for device_type GPU is enabled.

1875/1875 [=====] - 771s 411ms/step - loss: 0.6029 - accuracy: 0.9153 - val_loss: 0.3567 -
val_accuracy: 0.9757
Epoch 2/10
1875/1875 [=====] - 747s 398ms/step - loss: 0.3695 - accuracy: 0.9604 - val_loss: 0.2572 -
val_accuracy: 0.9848
Epoch 3/10
1875/1875 [=====] - 1021s 545ms/step - loss: 0.2755 - accuracy: 0.9704 - val_loss: 0.1749
- val_accuracy: 0.9873
Epoch 4/10
1875/1875 [=====] - 1054s 562ms/step - loss: 0.2144 - accuracy: 0.9772 - val_loss: 0.1506
- val_accuracy: 0.9902
Epoch 5/10
1875/1875 [=====] - 1223s 652ms/step - loss: 0.1748 - accuracy: 0.9814 - val_loss: 0.1136
- val_accuracy: 0.9902
Epoch 6/10
1875/1875 [=====] - 948s 505ms/step - loss: 0.1422 - accuracy: 0.9850 - val_loss: 0.0946 -
val_accuracy: 0.9920
Epoch 7/10
1875/1875 [=====] - 923s 492ms/step - loss: 0.1194 - accuracy: 0.9876 - val_loss: 0.0755 -
val_accuracy: 0.9923
Epoch 8/10
1875/1875 [=====] - 942s 503ms/step - loss: 0.1012 - accuracy: 0.9897 - val_loss: 0.0747 -
val_accuracy: 0.9920
Epoch 9/10
1875/1875 [=====] - 618s 330ms/step - loss: 0.0848 - accuracy: 0.9918 - val_loss: 0.0679 -
val_accuracy: 0.9917
Epoch 10/10
1875/1875 [=====] - 548s 292ms/step - loss: 0.0747 - accuracy: 0.9925 - val_loss: 0.0642 -
val_accuracy: 0.9920

In [78]: alexnet_model.save('alexnet.h5')
```

TRAINING AND LOSS CURVES

```
In [79]: save_train_loss_curves(history, 'alexnet_loss_acc.png')
```



PERFORMANCE

1/313 [.....] - ETA: 1:06

2022-09-16 20:44:23.437447: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.

313/313 [=====] - 17s 54ms/step

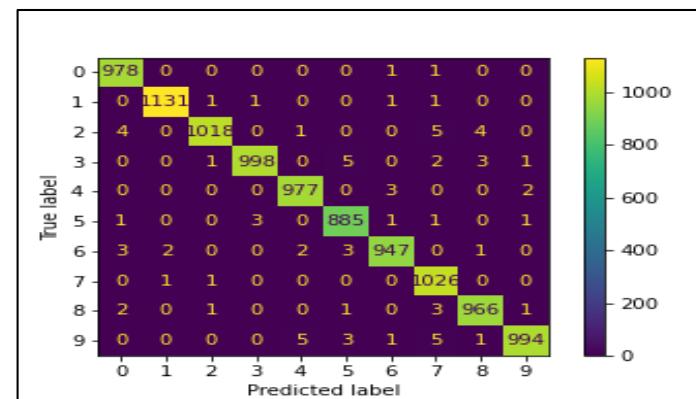
Confusion Matrix:

[978	0	0	0	0	0	1	1	0	0]
[0	1131	1	1	0	0	1	1	0	0]
[4	0	1018	0	1	0	0	5	4	0]
[0	0	1	998	0	5	0	2	3	1]
[0	0	0	0	977	0	3	0	0	2]
[1	0	0	3	0	885	1	1	0	1]
[3	2	0	0	2	3	947	0	1	0]
[0	1	1	0	0	0	0	1026	0	0]
[2	0	1	0	0	1	0	3	966	1]
[0	0	0	0	5	3	1	5	1	994]

Performance Evaluation:

	precision	recall	f1-score	support
0	0.99	1.00	0.99	980
1	1.00	1.00	1.00	1135
2	1.00	0.99	0.99	1032
3	1.00	0.99	0.99	1010
4	0.99	0.99	0.99	982
5	0.99	0.99	0.99	892
6	0.99	0.99	0.99	958
7	0.98	1.00	0.99	1028
8	0.99	0.99	0.99	974
9	0.99	0.99	0.99	1009
accuracy			0.99	10000
macro avg	0.99	0.99	0.99	10000
weighted avg	0.99	0.99	0.99	10000

Accuracy Score:
0.992



GOOGLENET



IMPORTS, DATA PREPROCESSING AND INCEPTION MODULE

GOOGLENET

```
In [7]: from tensorflow.keras.layers import Flatten, Input, Conv2D, Dense, Dropout, MaxPooling2D, AveragePooling2D, GlobalAv
from tensorflow.keras.layers import concatenate
from tensorflow.keras.models import Model
from tensorflow.keras.regularizers import l2

In [8]: train_images_normalized = train_images/255
test_images_normalized = test_images/255

In [9]: train_images_stacked = np.dstack([train_images] * 3)
test_images_stacked = np.dstack([test_images]*3)

test_images_stacked = test_images_stacked.reshape(-1,28,28,3)
train_images_stacked = train_images_stacked.reshape(-1,28,28,3)

In [10]: def inception(x,
                  filters_1x1,
                  filters_3x3_reduce,
                  filters_3x3,
                  filters_5x5_reduce,
                  filters_5x5,
                  filters_pool):
    path1 = layers.Conv2D(filters_1x1, (1, 1), padding='same', activation='relu')(x)
    path2 = layers.Conv2D(filters_3x3_reduce, (1, 1), padding='same', activation='relu')(x)
    path2 = layers.Conv2D(filters_3x3, (1, 1), padding='same', activation='relu')(path2)
    path3 = layers.Conv2D(filters_5x5_reduce, (1, 1), padding='same', activation='relu')(x)
    path3 = layers.Conv2D(filters_5x5, (1, 1), padding='same', activation='relu')(path3)
    path4 = layers.MaxPool2D((3, 3), strides=(1, 1), padding='same')(x)
    path4 = layers.Conv2D(filters_pool, (1, 1), padding='same', activation='relu')(path4)
    return tf.concat([path1, path2, path3, path4], axis=3)
```

MODEL ARCHITECTURE

```
In [11]: inp = layers.Input(shape=(28, 28, 3))
input_tensor = layers.experimental.preprocessing.Resizing(224, 224, interpolation="bilinear", input_shape=train_image)
x = layers.Conv2D(64, 7, strides=2, padding='same', activation='relu')(input_tensor)
x = layers.MaxPooling2D(3, strides=2)(x)
x = layers.Conv2D(64, 1, strides=1, padding='same', activation='relu')(x)
x = layers.Conv2D(192, 3, strides=1, padding='same', activation='relu')(x)
x = layers.MaxPooling2D(3, strides=2)(x)
x = inception(x, filters_1x1=64, filters_3x3_reduce=96, filters_3x3=128, filters_5x5_reduce=16, filters_5x5=32, filters_5x5=32, filters_5x5=32, filters_5x5=32)
x = inception(x, filters_1x1=128, filters_3x3_reduce=128, filters_3x3=192, filters_5x5_reduce=32, filters_5x5=96, filters_5x5=96, filters_5x5=96, filters_5x5=96)
x = layers.MaxPooling2D(3, strides=2)(x)
x = inception(x, filters_1x1=192, filters_3x3_reduce=96, filters_3x3=208, filters_5x5_reduce=16, filters_5x5=48, filters_5x5=48, filters_5x5=48, filters_5x5=48)
aux1 = layers.AveragePooling2D((5, 5), strides=3)(x)
aux1 = layers.Conv2D(128, 1, padding='same', activation='relu')(aux1)
aux1 = layers.Flatten()(aux1)
aux1 = layers.Dense(1024, activation='relu')(aux1)
aux1 = layers.Dropout(0.7)(aux1)
aux1 = layers.Dense(10, activation='softmax')(aux1)
x = inception(x, filters_1x1=160, filters_3x3_reduce=112, filters_3x3=224, filters_5x5_reduce=24, filters_5x5=64, filters_5x5=64, filters_5x5=64, filters_5x5=64)
x = inception(x, filters_1x1=128, filters_3x3_reduce=128, filters_3x3=256, filters_5x5_reduce=24, filters_5x5=64, filters_5x5=64, filters_5x5=64, filters_5x5=64)
x = inception(x, filters_1x1=112, filters_3x3_reduce=144, filters_3x3=288, filters_5x5_reduce=32, filters_5x5=64, filters_5x5=64, filters_5x5=64, filters_5x5=64)
aux2 = layers.AveragePooling2D((5, 5), strides=3)(x)
aux2 = layers.Conv2D(128, 1, padding='same', activation='relu')(aux2)
aux2 = layers.Flatten()(aux2)
aux2 = layers.Dense(1024, activation='relu')(aux2)
aux2 = layers.Dropout(0.7)(aux2)
aux2 = layers.Dense(10, activation='softmax')(aux2)
x = inception(x, filters_1x1=256, filters_3x3_reduce=160, filters_3x3=320, filters_5x5_reduce=32, filters_5x5=128, filters_5x5=128, filters_5x5=128, filters_5x5=128)
x = layers.MaxPooling2D(3, strides=2)(x)
x = inception(x, filters_1x1=256, filters_3x3_reduce=160, filters_3x3=320, filters_5x5_reduce=32, filters_5x5=128, filters_5x5=128, filters_5x5=128, filters_5x5=128)
x = inception(x, filters_1x1=384, filters_3x3_reduce=192, filters_3x3=384, filters_5x5_reduce=48, filters_5x5=128, filters_5x5=128, filters_5x5=128, filters_5x5=128)
x = layers.GlobalAveragePooling2D()(x)
x = layers.Dropout(0.4)(x)
out = layers.Dense(10, activation='softmax')(x)
```

Metal device set to: Apple M1

```
2022-09-21 22:36:22.919386: I tensorflow/core/common_runtime/pluggable_device/pluggable_device_factory.cc:305] Could not identify NUMA node of platform GPU ID 0, defaulting to 0. Your kernel may not have been built with NUMA support.
```

```
2022-09-21 22:36:22.919552: I tensorflow/core/common_runtime/pluggable_device/pluggable_device_factory.cc:271] Created TensorFlow device (/job:localhost/replica:0/task:0/device:GPU:0 with 0 MB memory) -> physical PluggableDevice (
```

COMPILATION

```
In [16]: googlenet_model = Model(inputs = inp, outputs = [out, aux1, aux2])
```

```
In [17]: googlenet_model.compile(optimizer=keras.optimizers.Adam(learning_rate=0.01, epsilon=0.05),  
                           loss=[keras.losses.sparse_categorical_crossentropy,  
                                 keras.losses.sparse_categorical_crossentropy,  
                                 keras.losses.sparse_categorical_crossentropy],  
                           loss_weights=[1, 0.3, 0.3],  
                           metrics=['accuracy'])
```

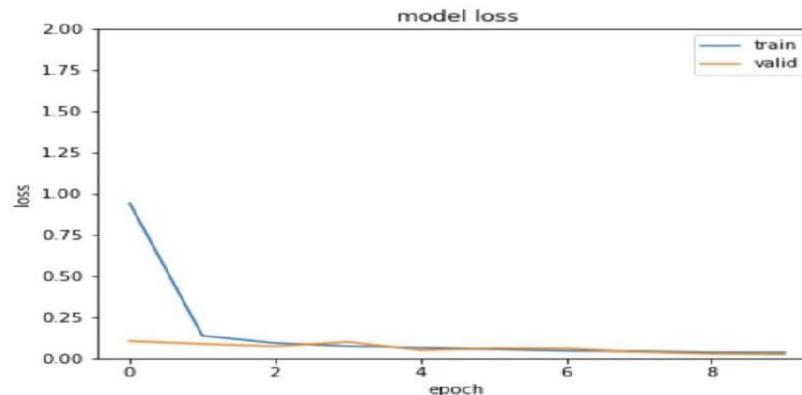
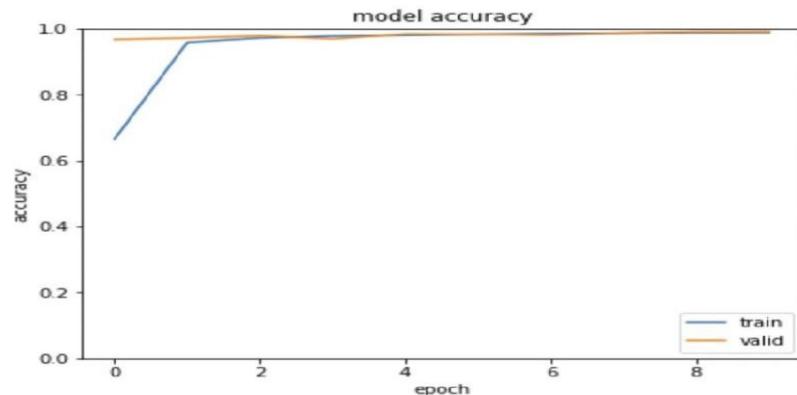
TRAINING

```
In [19]: history = googlenet_model.fit(train_images_stacked, [train_labels, train_labels, train_labels], \
                                     validation_data=(test_images_stacked, [test_labels, test_labels, test_labels]), \
                                     batch_size=64, epochs=10)

Epoch 1/10
938/938 [=====] - 2443s 3s/step - loss: 0.0771 - dense_4_loss: 0.0465 - dense_1_loss: 0.05
34 - dense_3_loss: 0.0486 - dense_4_accuracy: 0.9857 - dense_1_accuracy: 0.9841 - dense_3_accuracy: 0.9853 - val_lo
ss: 0.0955 - val_dense_4_loss: 0.0598 - val_dense_1_loss: 0.0592 - val_dense_3_loss: 0.0600 - val_dense_4_accuracy:
0.9810 - val_dense_1_accuracy: 0.9801 - val_dense_3_accuracy: 0.9804
Epoch 8/10
938/938 [=====] - 2489s 3s/step - loss: 0.0738 - dense_4_loss: 0.0441 - dense_1_loss: 0.05
23 - dense_3_loss: 0.0467 - dense_4_accuracy: 0.9867 - dense_1_accuracy: 0.9841 - dense_3_accuracy: 0.9855 - val_lo
ss: 0.0597 - val_dense_4_loss: 0.0401 - val_dense_1_loss: 0.0329 - val_dense_3_loss: 0.0324 - val_dense_4_accuracy:
0.9872 - val_dense_1_accuracy: 0.9882 - val_dense_3_accuracy: 0.9891
Epoch 9/10
938/938 [=====] - 2519s 3s/step - loss: 0.0626 - dense_4_loss: 0.0367 - dense_1_loss: 0.04
57 - dense_3_loss: 0.0406 - dense_4_accuracy: 0.9886 - dense_1_accuracy: 0.9856 - dense_3_accuracy: 0.9875 - val_lo
ss: 0.0508 - val_dense_4_loss: 0.0289 - val_dense_1_loss: 0.0427 - val_dense_3_loss: 0.0304 - val_dense_4_accuracy:
0.9909 - val_dense_1_accuracy: 0.9864 - val_dense_3_accuracy: 0.9904
Epoch 10/10
938/938 [=====] - 2530s 3s/step - loss: 0.0621 - dense_4_loss: 0.0366 - dense_1_loss: 0.04
54 - dense_3_loss: 0.0394 - dense_4_accuracy: 0.9892 - dense_1_accuracy: 0.9865 - dense_3_accuracy: 0.9880 - val_lo
ss: 0.0423 - val_dense_4_loss: 0.0263 - val_dense_1_loss: 0.0280 - val_dense_3_loss: 0.0252 - val_dense_4_accuracy:
0.9914 - val_dense_1_accuracy: 0.9899 - val_dense_3_accuracy: 0.9924
```

TRAINING AND LOSS CURVES

```
In [25]: plt.figure(figsize=(15,5))
plt.subplot(1,2,1)
plt.plot(history.history['dense_4_accuracy'])
plt.plot(history.history['val_dense_4_accuracy'])
plt.ylim(0, 1)
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'valid'], loc='lower right')
plt.subplot(1,2,2)
plt.plot(history.history['dense_4_loss'])
plt.plot(history.history['val_dense_4_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'valid'], loc='upper right')
plt.ylim([0,2])
plt.savefig("googlenet_loss_acc.png")
```



PERFORMANCE

313/313 [=====] - 33s 104ms/step

Confusion Matrix:

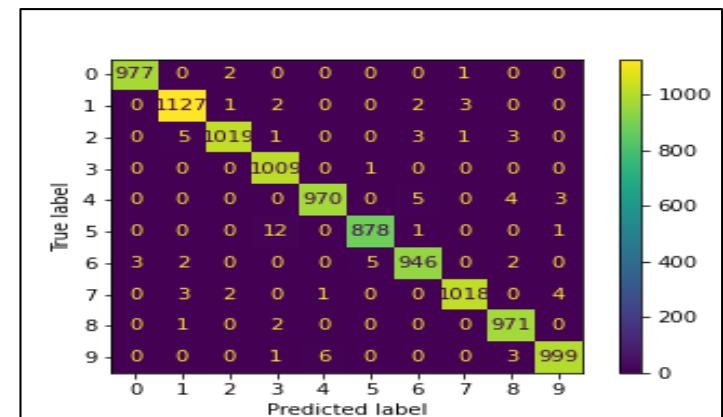
```
[[ 977   0   2   0   0   0   1   0   0   0]
 [ 0 1127   1   2   0   0   2   3   0   0]
 [ 0   5 1019   1   0   0   3   1   3   0]
 [ 0   0   0 1009   0   1   0   0   0   0]
 [ 0   0   0   0  970   0   5   0   4   3]
 [ 0   0   0   12   0  878   1   0   0   1]
 [ 0   0   0   0   5  946   0   2   0   0]
 [ 3   2   0   0   0   1   0   0 1018   0   4]
 [ 0   3   2   0   0   0   0   0  971   0   0]
 [ 0   0   0   1   6   0   0   0   3  999]]
```

Performance Evaluation:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	980
1	0.99	0.99	0.99	1135
2	1.00	0.99	0.99	1032
3	0.98	1.00	0.99	1010
4	0.99	0.99	0.99	982
5	0.99	0.98	0.99	892
6	0.99	0.99	0.99	958
7	1.00	0.99	0.99	1028
8	0.99	1.00	0.99	974
9	0.99	0.99	0.99	1009
accuracy			0.99	10000
macro avg	0.99	0.99	0.99	10000
weighted avg	0.99	0.99	0.99	10000

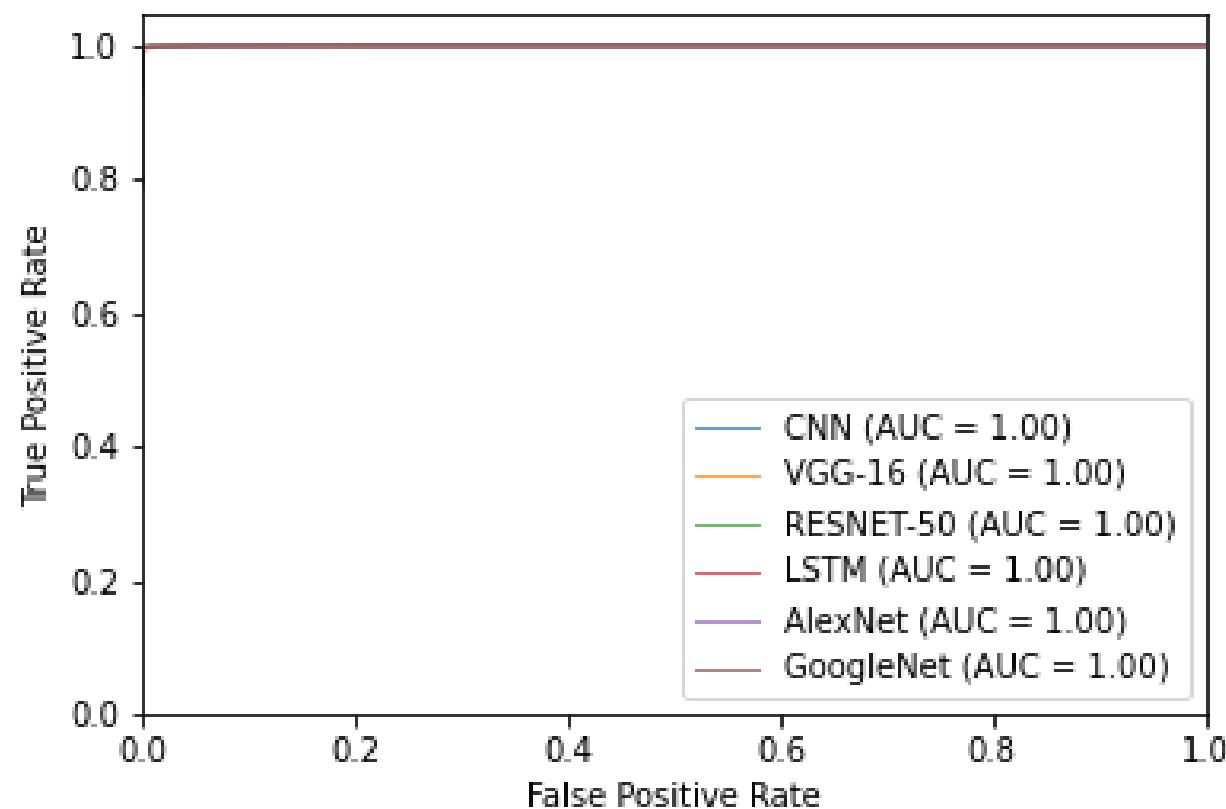
Accuracy Score:

0.9914



COMPARISON BETWEEN CNN, VGG-16, RESNET-50, LSTM(RNN), ALEXNET AND GOOGLENET.

ROC CURVES AND AUC COMPARISON



ROC CURVES AND AUC COMPARISON - CODE (I)

```
In [40]: def plot_macro_avg_roc_curve(classifier, n_classes, name, X_test, y_test, single_output=True):

    y_test_binarized = label_binarize(y_test, classes=list(range(0,n_classes)))
    y_score = classifier.predict(X_test)

    #for cases where the model (like googlenet) has extra (auxillary) outputs
    if not single_output:
        y_score = y_score[0]

    # Compute ROC curve and ROC area for each class
    fpr = dict()
    tpr = dict()
    roc_auc = dict()
    for i in range(n_classes):
        fpr[i], tpr[i], _ = roc_curve(y_test_binarized[:, i], y_score[:, i])
        roc_auc[i] = auc(fpr[i], tpr[i])

    # First aggregate all false positive rates
    all_fpr = np.unique(np.concatenate([fpr[i] for i in range(n_classes)]))

    # Then interpolate all ROC curves at this points
    mean_tpr = np.zeros_like(all_fpr)
    for i in range(n_classes):
        mean_tpr += np.interp(all_fpr, fpr[i], tpr[i])

    # Finally average it and compute AUC
    mean_tpr /= n_classes

    fpr["macro"] = all_fpr
    tpr["macro"] = mean_tpr
    roc_auc["macro"] = auc(fpr["macro"], tpr["macro"])

    plt.plot(
        fpr["macro"],
        tpr["macro"],
        label="{0} (AUC = {1:0.2f})".format(name, roc_auc["macro"]),
        alpha=0.7,
        linewidth=1,
    )

    return plt
```

ROC CURVES AND AUC COMPARISON - CODE (II)

```
In [41]: cnn = load_model('cnn.h5')
vgg_based_model = load_model('vgg16.h5')
resnet_based_model = load_model('resnet50.h5')
lstm_based_model = load_model('lstm.h5')
alexnet_model = load_model('alexnet.h5')
googlenet_model = load_model('googlenet.h5')
```

```
In [44]: plot_macro_avg_roc_curve(cnn, 10, "CNN", test_images, test_labels)
plot_macro_avg_roc_curve(vgg_based_model, 10, "VGG-16", test_images_p, test_labels)
plot_macro_avg_roc_curve(resnet_based_model, 10, "RESNET-50", test_images_pr, test_labels)
plot_macro_avg_roc_curve(lstm_based_model, 10, "LSTM", test_images, test_labels)
plot_macro_avg_roc_curve(alexnet_model, 10, "AlexNet", test_images_stacked, test_labels)
plot_macro_avg_roc_curve(googlenet_model, 10, "GoogleNet", test_images_n_stacked, test_labels, single_output=False)

plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.legend(loc="lower right")

plt.savefig('roc_auc.png')
```

COMPARISON OF PERFORMANCE

CLASSIFIER	Accuracy	Precision	Recall	F1-Score	AUC
CNN	0.99	0.99	0.99	0.99	1.00
VGG-16	0.99	0.99	0.99	0.99	1.00
RESNET-50	0.99	0.99	0.99	0.99	1.00
LSTM	0.99	0.99	0.99	0.99	1.00
ALEXNET	0.99	0.99	0.99	0.99	1.00
GOOGLENET	0.99	0.99	0.99	0.99	1.00

CONCLUSION:

All the networks performed very well on the MNIST dataset. The images present patterns in all dimensions for different models to exploit. Even the LSTM gave a good performance.

For such simple classifications, (less) time spent in training and predicting is more important as a distinguishing factor. Hence, it can be said that a simple CNN is ideal for this problem.

SAVEE - CLASSIFICATION



CONVERSION TO WAVEFORM (IMAGE)



IMPORTS AND CLASS NAMES

```
In [1]: import matplotlib.pyplot as plt
import tensorflow as tf
from matplotlib.backend_bases import RendererBase
from scipy import signal
from scipy.io import wavfile
# import soundfile as sf
import os
import numpy as np
from PIL import Image
from scipy.fftpack import fft

%matplotlib inline
```

Some mappings

```
In [2]: class_names = ['anger', 'disgust', 'fear', 'happiness', 'neutral', 'sadness', 'surprise']
lsymbols = ['a', 'd', 'f', 'h', 'n', 'sa', 'su']
full = {
    'a': 'anger',
    'd': 'disgust',
    'f': 'fear',
    'h': 'happiness',
    'n': 'neutral',
    'sa': 'sadness',
    'su': 'surprise'
}
```

AUDIO PREPROCESSING

Audio Preprocessing (Conversion to images of waveform)

```
In [3]: audio_path = './AudioData/'
pict_Path = './ImageData/'
samples = []

In [4]: if not os.path.exists(pict_Path):
    os.makedirs(pict_Path)

subFolderList = []
for x in os.listdir(audio_path):
    if os.path.isdir(audio_path + '/' + x):
        subFolderList.append(x)

In [5]: for x in class_names:
    os.makedirs(pict_Path + '/' + x)

In [6]: def wav2img_waveform(wav_path, targetdir='', figsize=(4,4), num=0):
    samplerate,test_sound = wavfile.read(wav_path)
    fig = plt.figure(figsize=figsize)
    plt.plot(test_sound)
    plt.axis('off')
    output_file = wav_path.split('/')[-1].split('.wav')[0]
    output_file = targetdir + '/' + output_file
    plt.savefig(f'{output_file}_{num}.png')
    plt.close()

In [7]: num = 0
for i, x in enumerate(subFolderList):
    # get all the wave files
    all_files = [y for y in os.listdir(audio_path + x) if '.wav' in y]
    for file in all_files:
        if file[1].isdigit():
            directory = full[file[0]]
        else:
            directory = full[file[0:2]]
        wav2img_waveform(audio_path + x + '/' + file, pict_Path + '/' + directory, num=num)
    num += 1
```

LOADING THE IMAGES INTO TRAIN AND TEST DATASETS

```
In [4]: class_names = ['anger', 'disgust', 'fear', 'happiness', 'neutral', 'sadness', 'surprise']
img_height = 224
img_width = 224
num_train_images = 384
num_test_images = 96
num_classes = len(class_names)
```

```
In [5]: train_ds = tf.keras.utils.image_dataset_from_directory(
    pict_Path,
    validation_split=0.2,
    subset="training",
    seed=123,
    image_size=(img_height, img_width),
    batch_size=num_train_images,
    label_mode='int')
```

Found 480 files belonging to 7 classes.
Using 384 files for training.
Metal device set to: Apple M1

```
2022-09-23 17:51:24.731527: I tensorflow/core/common_runtime/pluggable_device/pluggable_device_factory.cc:305] Could not identify NUMA node of platform GPU ID 0, defaulting to 0. Your kernel may not have been built with NUMA support.
2022-09-23 17:51:24.731713: I tensorflow/core/common_runtime/pluggable_device/pluggable_device_factory.cc:271] Created TensorFlow device (/job:localhost/replica:0/task:0/device:GPU:0 with 0 MB memory) -> physical PluggableDevice (device: 0, name: METAL, pci bus id: <undefined>)
```

```
In [6]: test_ds = tf.keras.utils.image_dataset_from_directory(
    pict_Path,
    validation_split=0.2,
    subset="validation",
    seed=123,
    image_size=(img_height, img_width),
    batch_size=num_test_images,
    label_mode='int')
```

Found 480 files belonging to 7 classes.
Using 96 files for validation.

BREAKING THE DATASETS INTO IMAGES AND LABELS

```
In [7]: train_images = train_labels = test_images = test_labels = None
```

```
In [60]: train_images, train_labels = next(train_ds.as_numpy_iterator())
```

```
In [62]: test_images, test_labels = next(test_ds.as_numpy_iterator())
```

CNN



IMPORTS AND MODEL ARCHITECTURE

Convolutional Neural Network

Importing the libraries

```
In [10]: from tensorflow import keras
from tensorflow.keras import layers, datasets
from tensorflow.keras.models import Sequential
import matplotlib.pyplot as plt
```

Building the CNN

```
In [37]: cnn = Sequential([
    layers.Rescaling(1./255, input_shape=(img_height, img_width, 3)),
    layers.Conv2D(filters=32, kernel_size=3, padding='same', activation='relu'),
    layers.Conv2D(filters=32, kernel_size=3, padding='same', activation='relu'),
    layers.MaxPooling2D(pool_size=2, strides=2),
    layers.Conv2D(32, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(pool_size=2, strides=2),
    layers.Dropout(0.2),
    layers.Flatten(),
    layers.Dense(128, activation='relu'),
    layers.Dense(num_classes, activation="softmax")
])
```

COMPILATION AND SUMMARY

Compiling the CNN

```
In [38]: cnn.compile(optimizer=keras.optimizers.Adam(learning_rate=0.0005, epsilon=0.01),  
                  loss=tf.keras.losses.SparseCategoricalCrossentropy(),  
                  metrics=['accuracy'])
```

```
In [39]: cnn.summary()
```

Model: "sequential_3"

Layer (type)	Output Shape	Param #
<hr/>		
rescaling_3 (Rescaling)	(None, 224, 224, 3)	0
conv2d_9 (Conv2D)	(None, 224, 224, 32)	896
conv2d_10 (Conv2D)	(None, 224, 224, 32)	9248
max_pooling2d_6 (MaxPooling 2D)	(None, 112, 112, 32)	0
conv2d_11 (Conv2D)	(None, 112, 112, 32)	9248
max_pooling2d_7 (MaxPooling 2D)	(None, 56, 56, 32)	0
dropout_3 (Dropout)	(None, 56, 56, 32)	0
flatten_3 (Flatten)	(None, 100352)	0
dense_6 (Dense)	(None, 128)	12845184
dense_7 (Dense)	(None, 7)	903
<hr/>		
Total params: 12,865,479		
Trainable params: 12,865,479		
Non-trainable params: 0		

TRAINING AND SAVING

Training the CNN on the Training set and evaluating it on the Test set

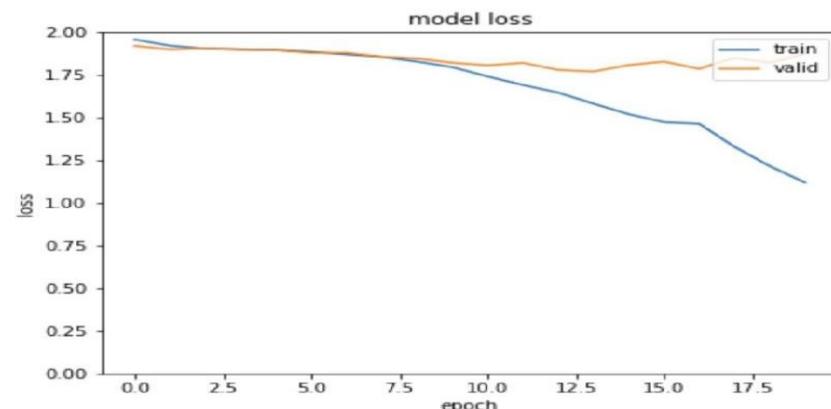
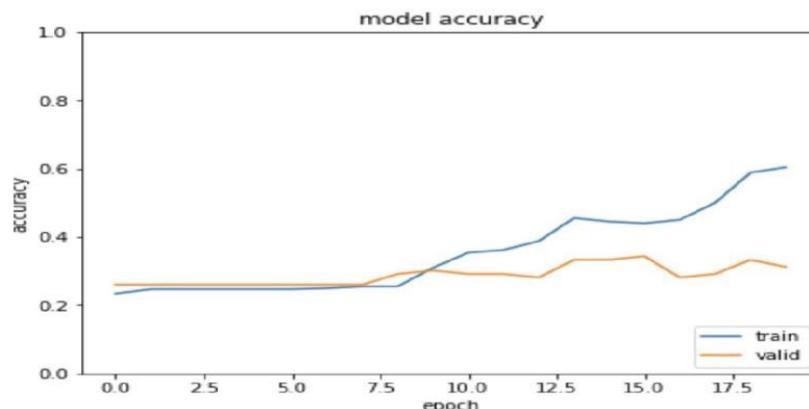
```
In [40]: history = cnn.fit(train_images, train_labels, epochs=20, validation_data=(test_images, test_labels))
          accuracy: 0.3333
Epoch 15/20
12/12 [=====] - 3s 254ms/step - loss: 1.5198 - accuracy: 0.4453 - val_loss: 1.8071 - val_a
ccuracy: 0.3333
Epoch 16/20
12/12 [=====] - 3s 254ms/step - loss: 1.4734 - accuracy: 0.4401 - val_loss: 1.8280 - val_a
ccuracy: 0.3438
Epoch 17/20
12/12 [=====] - 3s 255ms/step - loss: 1.4640 - accuracy: 0.4505 - val_loss: 1.7854 - val_a
ccuracy: 0.2812
Epoch 18/20
12/12 [=====] - 3s 258ms/step - loss: 1.3295 - accuracy: 0.5000 - val_loss: 1.8504 - val_a
ccuracy: 0.2917
Epoch 19/20
12/12 [=====] - 3s 255ms/step - loss: 1.2158 - accuracy: 0.5885 - val_loss: 1.8208 - val_a
ccuracy: 0.3333
Epoch 20/20
12/12 [=====] - 3s 258ms/step - loss: 1.1194 - accuracy: 0.6042 - val_loss: 1.8733 - val_a
ccuracy: 0.3125
```

```
In [41]: cnn.save('cnn.h5')
```

TRAINING AND LOSS CURVES

```
In [30]: def save_train_loss_curves(history, filename):
    plt.figure(figsize=(15,5))
    plt.subplot(1,2,1)
    plt.plot(history.history['accuracy'])
    plt.plot(history.history['val_accuracy'])
    plt.ylim(0, 1)
    plt.title('model accuracy')
    plt.ylabel('accuracy')
    plt.xlabel('epoch')
    plt.legend(['train', 'valid'], loc='lower right')
    plt.subplot(1,2,2)
    plt.plot(history.history['loss'])
    plt.plot(history.history['val_loss'])
    plt.title('model loss')
    plt.ylabel('loss')
    plt.xlabel('epoch')
    plt.legend(['train', 'valid'], loc='upper right')
    plt.ylim([0,3])
    plt.savefig(filename)
```

```
In [47]: save_train_loss_curves(history, 'cnn_loss_acc.png')
```



PERFORMANCE - CODE

Performance

```
In [11]: from sklearn.metrics import classification_report, confusion_matrix, accuracy_score, ConfusionMatrixDisplay
```

```
In [12]: def get_predictions(model, inputs):

    def index_max_arr_element(arr):
        best_index = -1
        best = float('-inf')
        for i, num in enumerate(arr):
            if num > best:
                best = num
                best_index = i
        return best_index

    pred_arr = model.predict(inputs)
    pred_labels = [index_max_arr_element(row) for row in pred_arr]
    return pred_labels
```

```
In [45]: pred_labels = get_predictions(cnn, test_images)

print("Confusion Matrix:")
print(confusion_matrix(test_labels, pred_labels))
print("-----")
print("-----")
print("Performance Evaluation:")
print(classification_report(test_labels, pred_labels))
print("-----")
print("-----")
print("Accuracy Score:")
print(accuracy_score(test_labels, pred_labels))

plt.figure()
ConfusionMatrixDisplay.from_predictions(test_labels, pred_labels)
plt.savefig('./cnn_matrix.png')
```

PERFORMANCE

Confusion Matrix:

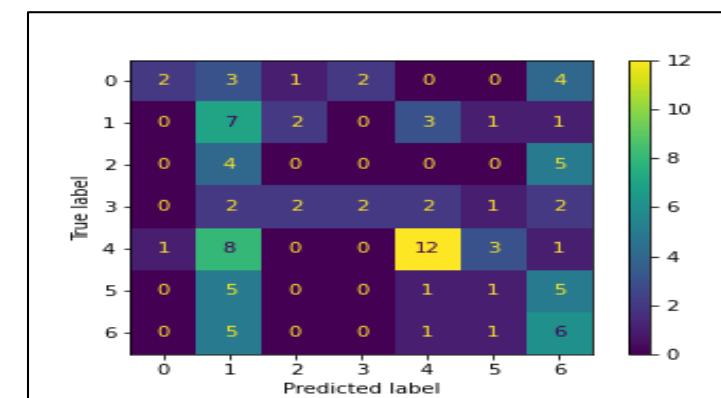
```
[[ 2  3  1  2  0  0  4]
 [ 0  7  2  0  3  1  1]
 [ 0  4  0  0  0  0  5]
 [ 0  2  2  2  2  1  2]
 [ 1  8  0  0  12  3  1]
 [ 0  5  0  0  1  1  5]
 [ 0  5  0  0  1  1  6]]
```

Performance Evaluation:

	precision	recall	f1-score	support
0	0.67	0.17	0.27	12
1	0.21	0.50	0.29	14
2	0.00	0.00	0.00	9
3	0.50	0.18	0.27	11
4	0.63	0.48	0.55	25
5	0.14	0.08	0.11	12
6	0.25	0.46	0.32	13
accuracy			0.31	96
macro avg	0.34	0.27	0.26	96
weighted avg	0.39	0.31	0.31	96

Accuracy Score:

0.3125



VGG-16

IMPORTS, DATA PREPROCESSING & ARCHITECTURE

```
In [55]: from tensorflow.keras.applications import vgg16
```

```
In [56]: train_images_p, test_images_p = vgg16.preprocess_input(np.copy(train_images)), vgg16.preprocess_input(np.copy(test_i
```

```
In [50]: vgg_based_model = Sequential([
    layers.Rescaling(1./255, input_shape=(img_height, img_width, 3)),
    vgg16.VGG16(include_top=False, weights="imagenet", input_shape=(img_height, img_width, 3), pooling="max"),
    layers.Flatten(),
    layers.Dense(512, activation='relu'),
    layers.Dropout(0.2),
    layers.Dense(128, activation='relu'),
    layers.Dropout(0.2),
    layers.Dense(num_classes, activation="softmax")
])
```

COMPIRATION AND SUMMARY

```
In [51]: vgg_based_model.compile(optimizer=keras.optimizers.Adam(learning_rate=0.0005, epsilon=0.01),  
                           loss=tf.keras.losses.SparseCategoricalCrossentropy(),  
                           metrics=['accuracy'])
```

```
In [52]: vgg_based_model.summary()
```

Model: "sequential_4"

Layer (type)	Output Shape	Param #
<hr/>		
rescaling_4 (Rescaling)	(None, 224, 224, 3)	0
vgg16 (Functional)	(None, 512)	14714688
flatten_4 (Flatten)	(None, 512)	0
dense_8 (Dense)	(None, 512)	262656
dropout_4 (Dropout)	(None, 512)	0
dense_9 (Dense)	(None, 128)	65664
dropout_5 (Dropout)	(None, 128)	0
dense_10 (Dense)	(None, 7)	903
<hr/>		
Total params: 15,043,911		
Trainable params: 15,043,911		
Non-trainable params: 0		

TRAINING AND SAVING

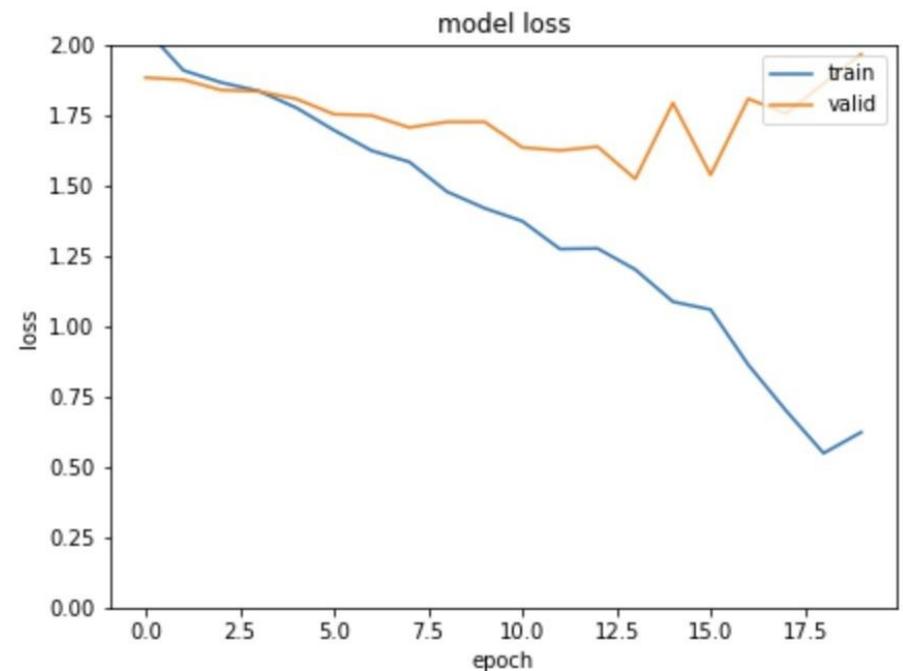
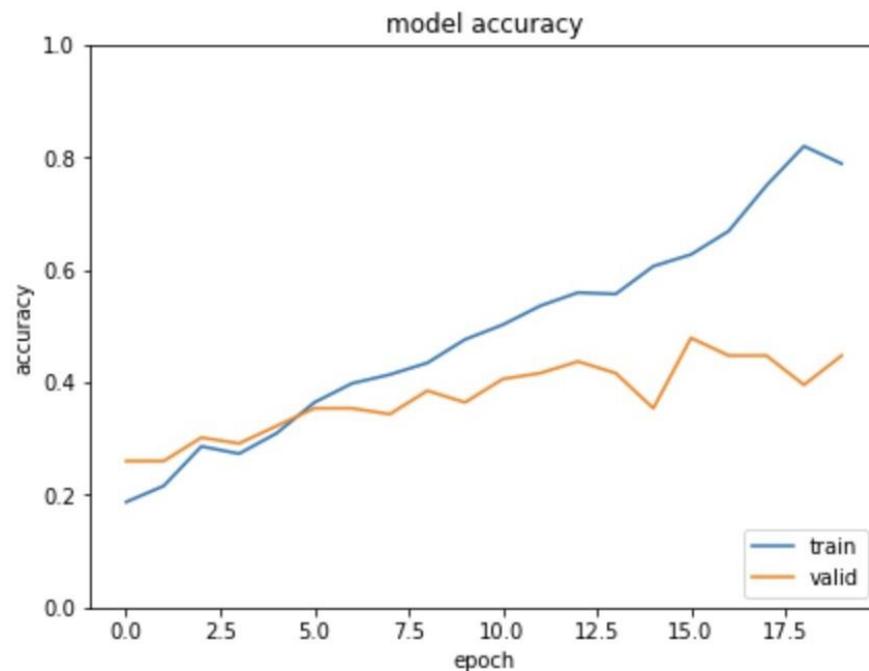
```
In [53]: history = vgg_based_model.fit(train_images_p, train_labels, epochs=20, validation_data=(test_images_p, test_labels))
```

```
uracy: 0.4167
Epoch 15/20
12/12 [=====] - 40s 3s/step - loss: 1.0874 - accuracy: 0.6068 - val_loss: 1.7943 - val_accuracy: 0.3542
Epoch 16/20
12/12 [=====] - 40s 3s/step - loss: 1.0594 - accuracy: 0.6276 - val_loss: 1.5384 - val_accuracy: 0.4792
Epoch 17/20
12/12 [=====] - 40s 3s/step - loss: 0.8631 - accuracy: 0.6693 - val_loss: 1.8091 - val_accuracy: 0.4479
Epoch 18/20
12/12 [=====] - 40s 3s/step - loss: 0.7004 - accuracy: 0.7500 - val_loss: 1.7563 - val_accuracy: 0.4479
Epoch 19/20
12/12 [=====] - 40s 3s/step - loss: 0.5487 - accuracy: 0.8203 - val_loss: 1.8636 - val_accuracy: 0.3958
Epoch 20/20
12/12 [=====] - 40s 3s/step - loss: 0.6232 - accuracy: 0.7891 - val_loss: 1.9675 - val_accuracy: 0.4479
```

```
In [55]: vgg_based_model.save('vgg16.h5')
```

TRAINING AND LOSS CURVES

```
n [54]: save_train_loss_curves(history, 'vgg_loss_acc.png')
```



PERFORMANCE

Confusion Matrix:

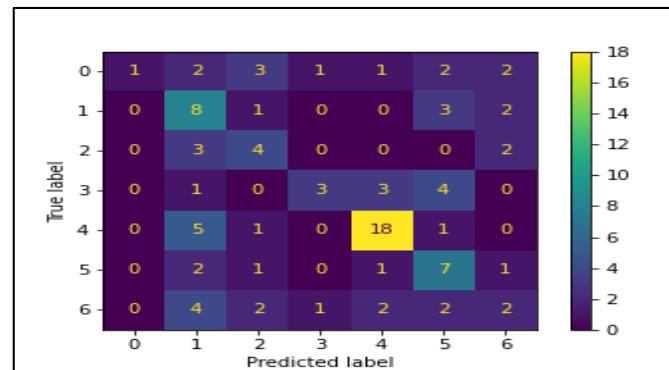
```
[[ 1  2  3  1  1  2  2]
 [ 0  8  1  0  0  3  2]
 [ 0  3  4  0  0  0  2]
 [ 0  1  0  3  3  4  0]
 [ 0  5  1  0  18 1  0]
 [ 0  2  1  0  1  7  1]
 [ 0  4  2  1  2  2  2]]
```

Performance Evaluation:

	precision	recall	f1-score	support
0	1.00	0.08	0.15	12
1	0.32	0.57	0.41	14
2	0.33	0.44	0.38	9
3	0.60	0.27	0.37	11
4	0.72	0.72	0.72	25
5	0.37	0.58	0.45	12
6	0.22	0.15	0.18	13
accuracy			0.45	96
macro avg	0.51	0.40	0.38	96
weighted avg	0.54	0.45	0.43	96

Accuracy Score:

0.4479166666666667



RES NET- 50

IMPORTS, DATA PREPROCESSING & ARCHITECTURE

RESNET-50

```
In [57]: from tensorflow.keras.applications import resnet50
```

```
In [58]: train_images_pr, test_images_pr = resnet50.preprocess_input(np.copy(train_images)), resnet50.preprocess_input(np.cop
```

```
In [17]: resnet_based_model = Sequential([
    layers.Rescaling(1./255, input_shape=(img_height, img_width, 3)),
    resnet50.ResNet50(include_top=False, weights="imagenet", input_shape=(img_height, img_width, 3), pooling="max"),
    layers.Flatten(),
    layers.Dense(512, activation='relu'),
    layers.Dropout(0.2),
    layers.Dense(128, activation='relu'),
    layers.Dropout(0.2),
    layers.Dense(num_classes, activation="softmax")
])
```

COMPILATION AND SUMMARY

```
In [18]: resnet_based_model.compile(optimizer=keras.optimizers.Adam(learning_rate=0.0005, epsilon=0.01),
                                loss=tf.keras.losses.SparseCategoricalCrossentropy(),
                                metrics=['accuracy'])
```

```
In [19]: resnet_based_model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
rescaling (Rescaling)	(None, 224, 224, 3)	0
resnet50 (Functional)	(None, 2048)	23587712
flatten (Flatten)	(None, 2048)	0
dense (Dense)	(None, 512)	1049088
dropout (Dropout)	(None, 512)	0
dense_1 (Dense)	(None, 128)	65664
dropout_1 (Dropout)	(None, 128)	0
dense_2 (Dense)	(None, 7)	903

Total params: 24,703,367

Trainable params: 24,650,247

Non-trainable params: 53,120

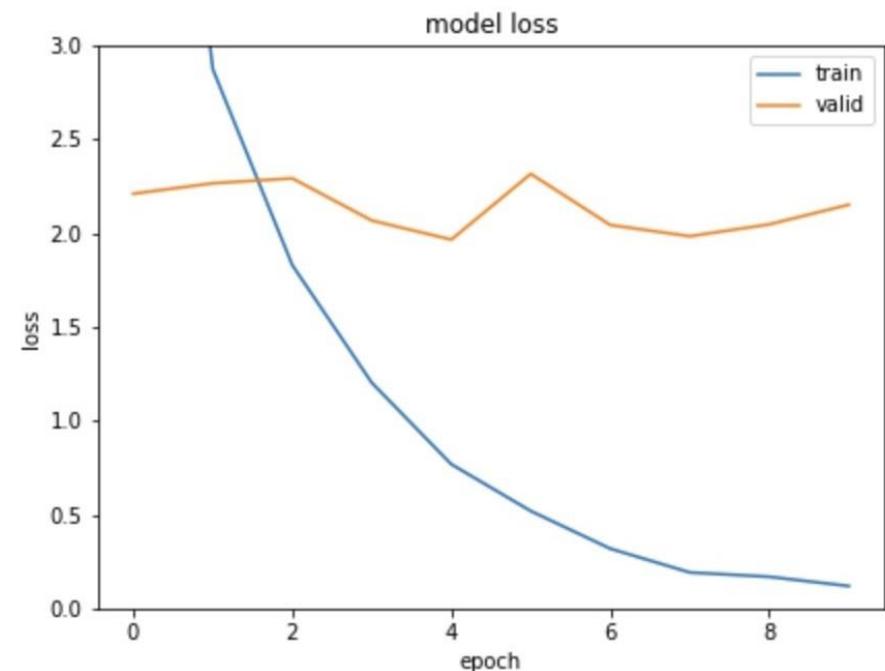
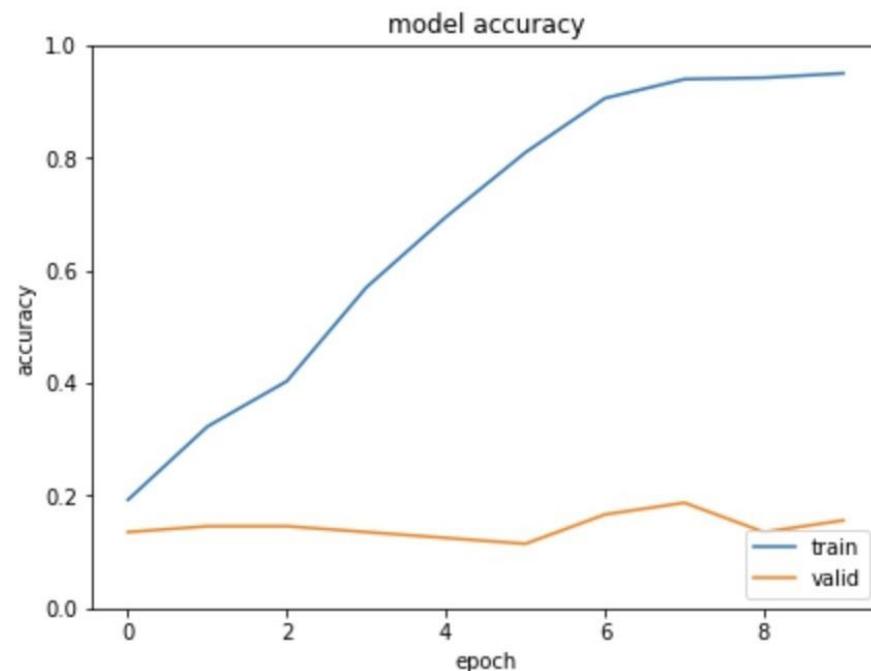
TRAINING AND SAVING

```
In [20]: history = resnet_based_model.fit(train_images_pr, train_labels, epochs=10, validation_data=(test_images_pr, test_labels))
          12/12 [=====] - 15s 1s/step - loss: 0.7679 - accuracy: 0.6953 - val_loss: 1.9651 - val_accuracy: 0.1354
          Epoch 5/10
          12/12 [=====] - 15s 1s/step - loss: 0.5185 - accuracy: 0.8099 - val_loss: 2.3155 - val_accuracy: 0.1250
          Epoch 6/10
          12/12 [=====] - 15s 1s/step - loss: 0.3179 - accuracy: 0.9062 - val_loss: 2.0430 - val_accuracy: 0.1146
          Epoch 7/10
          12/12 [=====] - 16s 1s/step - loss: 0.1912 - accuracy: 0.9401 - val_loss: 1.9831 - val_accuracy: 0.1875
          Epoch 8/10
          12/12 [=====] - 18s 1s/step - loss: 0.1681 - accuracy: 0.9427 - val_loss: 2.0465 - val_accuracy: 0.1354
          Epoch 9/10
          12/12 [=====] - 20s 2s/step - loss: 0.1185 - accuracy: 0.9505 - val_loss: 2.1512 - val_accuracy: 0.1562
```

```
In [27]: resnet_based_model.save('resnet50.h5')
```

TRAINING AND LOSS CURVES

```
In [26]: save_train_loss_curves(history, 'resnet_loss_acc.png')
```



PERFORMANCE

Confusion Matrix:

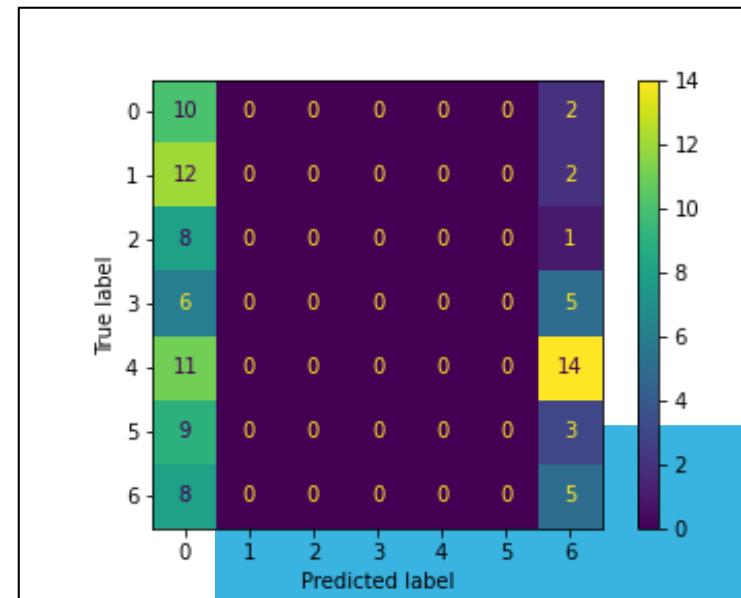
```
[[10  0  0  0  0  0  2]
 [12  0  0  0  0  0  2]
 [ 8  0  0  0  0  0  1]
 [ 6  0  0  0  0  0  5]
 [11  0  0  0  0  0  14]
 [ 9  0  0  0  0  0  3]
 [ 8  0  0  0  0  0  5]]
```

Performance Evaluation:

	precision	recall	f1-score	support
0	0.16	0.83	0.26	12
1	0.00	0.00	0.00	14
2	0.00	0.00	0.00	9
3	0.00	0.00	0.00	11
4	0.00	0.00	0.00	25
5	0.00	0.00	0.00	12
6	0.16	0.38	0.22	13
accuracy			0.16	96
macro avg	0.04	0.17	0.07	96
weighted avg	0.04	0.16	0.06	96

Accuracy Score:

0.15625



RNN (LSTM)



MODEL ARCHITECTURE AND COMPIRATION

RNN (LSTM based)

```
In [24]: lstm_based_model = Sequential([
    layers.Rescaling(1./255, input_shape=(img_height, img_width, 3)),
    layers.Resizing(64, 64),
    layers.Reshape((4096, 3)),
    layers.LSTM(units=50, return_sequences=True),
    layers.Dropout(0.2),
    layers.LSTM(units=50),
    layers.Dropout(0.2),
    layers.Dense(128, activation='relu'),
    layers.Dropout(0.2),
    layers.Dense(num_classes, activation="softmax")
])
```

```
In [25]: lstm_based_model.compile(optimizer=keras.optimizers.Adam(learning_rate=0.001),
                               loss=tf.keras.losses.SparseCategoricalCrossentropy(),
                               metrics=['accuracy'])
```

TRAINING AND SAVING

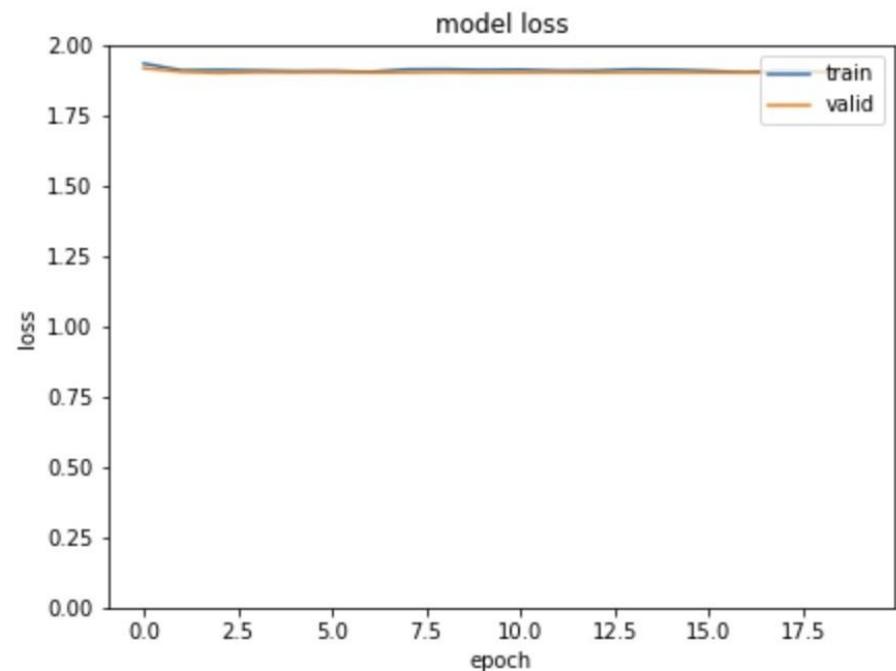
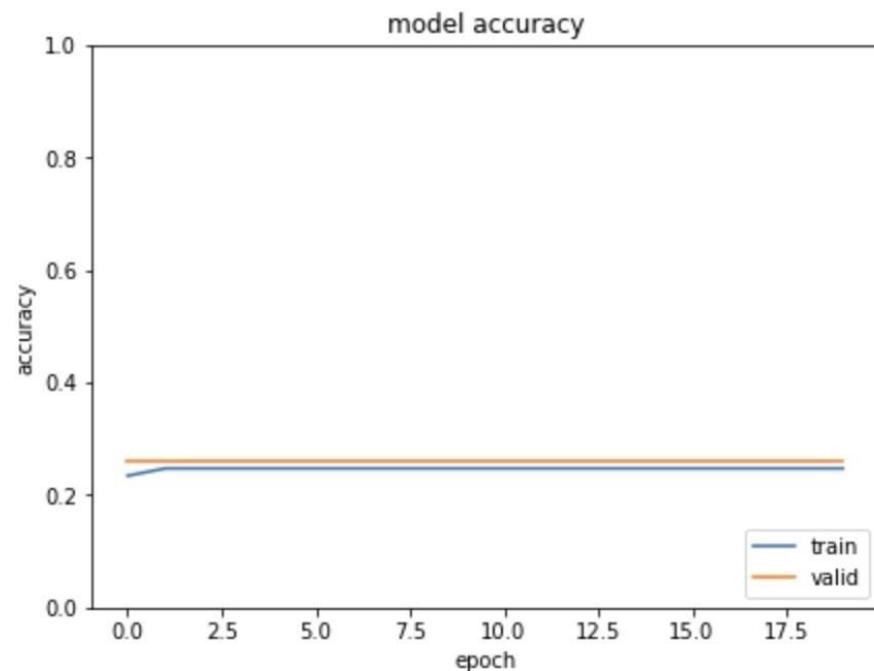
```
In [27]: history = lstm_based_model.fit(train_images, train_labels, epochs=20, validation_data=(test_images, test_labels))
          accuracy: 0.2604
Epoch 15/20
12/12 [=====] - 43s 4s/step - loss: 1.9130 - accuracy: 0.2474 - val_loss: 1.9036 - val_accuracy: 0.2604
Epoch 16/20
12/12 [=====] - 42s 4s/step - loss: 1.9091 - accuracy: 0.2474 - val_loss: 1.9030 - val_accuracy: 0.2604
Epoch 17/20
12/12 [=====] - 43s 4s/step - loss: 1.9039 - accuracy: 0.2474 - val_loss: 1.9030 - val_accuracy: 0.2604
Epoch 18/20
12/12 [=====] - 43s 4s/step - loss: 1.9095 - accuracy: 0.2474 - val_loss: 1.9016 - val_accuracy: 0.2604
Epoch 19/20
12/12 [=====] - 43s 4s/step - loss: 1.9072 - accuracy: 0.2474 - val_loss: 1.9030 - val_accuracy: 0.2604
Epoch 20/20
12/12 [=====] - 43s 4s/step - loss: 1.9059 - accuracy: 0.2474 - val_loss: 1.9033 - val_accuracy: 0.2604
```

```
In [28]: lstm_based_model.save('lstm.h5')
```

TRAINING AND LOSS CURVES

```
In [28]: lstm_based_model.save('lstm.h5')
```

```
In [29]: save_train_loss_curves(history, 'lstm_loss_acc.png')
```



PERFORMANCE

Confusion Matrix:

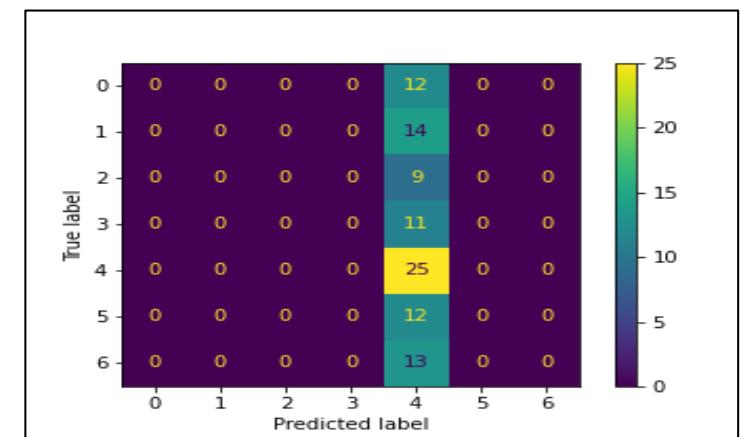
```
[[ 0  0  0  0 12  0  0]
 [ 0  0  0  0 14  0  0]
 [ 0  0  0  0  9  0  0]
 [ 0  0  0  0 11  0  0]
 [ 0  0  0  0 25  0  0]
 [ 0  0  0  0 12  0  0]
 [ 0  0  0  0 13  0  0]]
```

Performance Evaluation:

	precision	recall	f1-score	support
0	0.00	0.00	0.00	12
1	0.00	0.00	0.00	14
2	0.00	0.00	0.00	9
3	0.00	0.00	0.00	11
4	0.26	1.00	0.41	25
5	0.00	0.00	0.00	12
6	0.00	0.00	0.00	13
accuracy			0.26	96
macro avg	0.04	0.14	0.06	96
weighted avg	0.07	0.26	0.11	96

Accuracy Score:

0.2604166666666667



ALEXNET



MODEL ARCHITECTURE - I

```
In [14]: from tensorflow.keras.layers import Dense, Dropout, Activation, Flatten, Conv2D, MaxPooling2D, ZeroPadding2D, BatchNormaliz
from tensorflow.keras.regularizers import l2

In [23]: def alexnet_model(img_shape=(32, 32, 3), n_classes=10, l2_reg=0.,
weights=None):

    # Initialize model
    alexnet = Sequential()
    alexnet.add(layers.Input(shape=img_shape))
    alexnet.add(layers.Resizing(64, 64))

    # Layer 1
    alexnet.add(Conv2D(96, (11, 11), padding='same', kernel_regularizer=l2(l2_reg)))
    alexnet.add(BatchNormalization())
    alexnet.add(Activation('relu'))
    alexnet.add(MaxPooling2D(pool_size=(2, 2)))

    # Layer 2
    alexnet.add(Conv2D(256, (5, 5), padding='same'))
    alexnet.add(BatchNormalization())
    alexnet.add(Activation('relu'))
    alexnet.add(MaxPooling2D(pool_size=(2, 2)))

    # Layer 3
    alexnet.add(ZeroPadding2D((1, 1)))
    alexnet.add(Conv2D(512, (3, 3), padding='same'))
    alexnet.add(BatchNormalization())
    alexnet.add(Activation('relu'))
    alexnet.add(MaxPooling2D(pool_size=(2, 2)))

    # Layer 4
    alexnet.add(ZeroPadding2D((1, 1)))
    alexnet.add(Conv2D(1024, (3, 3), padding='same'))
    alexnet.add(BatchNormalization())
    alexnet.add(Activation('relu'))
```

MODEL ARCHITECTURE - II

```
# Layer 5
alexnet.add(ZeroPadding2D((1, 1)))
alexnet.add(Conv2D(1024, (3, 3), padding='same'))
alexnet.add(BatchNormalization())
alexnet.add(Activation('relu'))
alexnet.add(MaxPooling2D(pool_size=(2, 2)))

# Layer 6
alexnet.add(Flatten())
alexnet.add(Dense(3072))
alexnet.add(BatchNormalization())
alexnet.add(Activation('relu'))
alexnet.add(Dropout(0.5))

# Layer 7
alexnet.add(Dense(4096))
alexnet.add(BatchNormalization())
alexnet.add(Activation('relu'))
alexnet.add(Dropout(0.5))

# Layer 8
alexnet.add(Dense(n_classes))
alexnet.add(BatchNormalization())
alexnet.add(Activation('softmax'))

if weights is not None:
    alexnet.load_weights(weights)

return alexnet
```

```
In [24]: alexnet_model = alexnet_model(img_shape=(img_height, img_width, 3), n_classes=num_classes)
```

```
In [25]: alexnet_model.compile(optimizer=keras.optimizers.Adam(learning_rate=0.0002, epsilon=0.05),
                           loss=tf.keras.losses.SparseCategoricalCrossentropy(),
                           metrics=['accuracy'])
```

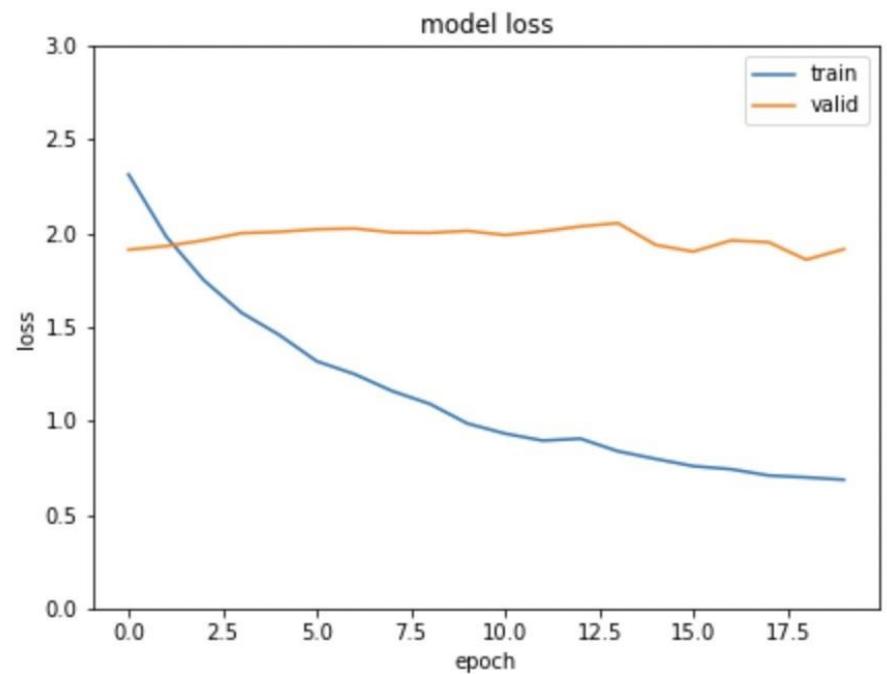
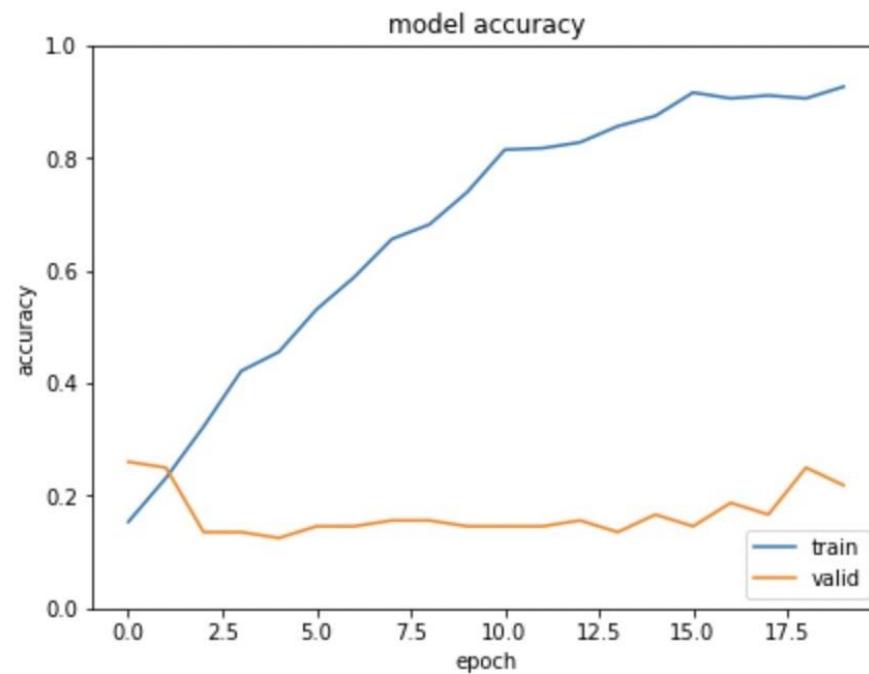
TRAINING AND SAVING

```
In [27]: history = alexnet_model.fit(train_images, train_labels, epochs=20, validation_data=(test_images, test_labels))  
Epoch 1/20  
12/12 [=====] - 9s 757ms/step - loss: 0.7957 - accuracy: 0.8555 - val_loss: 1.9355 - val_accuracy: 0.1354  
Epoch 2/20  
12/12 [=====] - 9s 757ms/step - loss: 0.7957 - accuracy: 0.8555 - val_loss: 1.9355 - val_accuracy: 0.1354  
Epoch 3/20  
12/12 [=====] - 9s 757ms/step - loss: 0.7957 - accuracy: 0.8555 - val_loss: 1.9355 - val_accuracy: 0.1354  
Epoch 4/20  
12/12 [=====] - 9s 757ms/step - loss: 0.7957 - accuracy: 0.8555 - val_loss: 1.9355 - val_accuracy: 0.1354  
Epoch 5/20  
12/12 [=====] - 9s 757ms/step - loss: 0.7957 - accuracy: 0.8555 - val_loss: 1.9355 - val_accuracy: 0.1354  
Epoch 6/20  
12/12 [=====] - 9s 757ms/step - loss: 0.7957 - accuracy: 0.8555 - val_loss: 1.9355 - val_accuracy: 0.1354  
Epoch 7/20  
12/12 [=====] - 9s 757ms/step - loss: 0.7957 - accuracy: 0.8555 - val_loss: 1.9355 - val_accuracy: 0.1354  
Epoch 8/20  
12/12 [=====] - 9s 757ms/step - loss: 0.7957 - accuracy: 0.8555 - val_loss: 1.9355 - val_accuracy: 0.1354  
Epoch 9/20  
12/12 [=====] - 9s 757ms/step - loss: 0.7957 - accuracy: 0.8555 - val_loss: 1.9355 - val_accuracy: 0.1354  
Epoch 10/20  
12/12 [=====] - 9s 757ms/step - loss: 0.7957 - accuracy: 0.8555 - val_loss: 1.9355 - val_accuracy: 0.1354  
Epoch 11/20  
12/12 [=====] - 9s 757ms/step - loss: 0.7957 - accuracy: 0.8555 - val_loss: 1.9355 - val_accuracy: 0.1354  
Epoch 12/20  
12/12 [=====] - 9s 757ms/step - loss: 0.7957 - accuracy: 0.8555 - val_loss: 1.9355 - val_accuracy: 0.1354  
Epoch 13/20  
12/12 [=====] - 9s 757ms/step - loss: 0.7957 - accuracy: 0.8555 - val_loss: 1.9355 - val_accuracy: 0.1354  
Epoch 14/20  
12/12 [=====] - 9s 757ms/step - loss: 0.7957 - accuracy: 0.8555 - val_loss: 1.9355 - val_accuracy: 0.1354  
Epoch 15/20  
12/12 [=====] - 9s 746ms/step - loss: 0.7966 - accuracy: 0.8750 - val_loss: 1.9376 - val_accuracy: 0.1667  
Epoch 16/20  
12/12 [=====] - 9s 762ms/step - loss: 0.7583 - accuracy: 0.9167 - val_loss: 1.9018 - val_accuracy: 0.1667  
Epoch 17/20  
12/12 [=====] - 9s 779ms/step - loss: 0.7418 - accuracy: 0.9062 - val_loss: 1.9614 - val_accuracy: 0.1458  
Epoch 18/20  
12/12 [=====] - 9s 781ms/step - loss: 0.7084 - accuracy: 0.9115 - val_loss: 1.9521 - val_accuracy: 0.1875  
Epoch 19/20  
12/12 [=====] - 9s 786ms/step - loss: 0.6981 - accuracy: 0.9062 - val_loss: 1.8590 - val_accuracy: 0.1667  
Epoch 20/20  
12/12 [=====] - 9s 789ms/step - loss: 0.6857 - accuracy: 0.9271 - val_loss: 1.9149 - val_accuracy: 0.2188
```

```
In [28]: alexnet_model.save('alexnet.h5')
```

TRAINING AND LOSS CURVES

```
In [31]: save_train_loss_curves(history, 'alexnet_loss_acc.png')
```



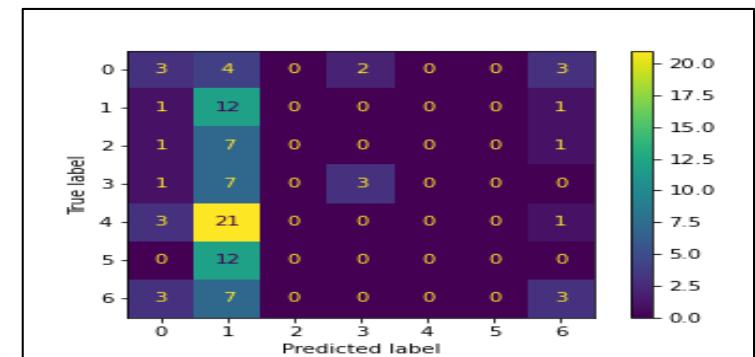
PERFORMANCE

Confusion Matrix:

```
[[ 3  4  0  2  0  0  3]
 [ 1 12  0  0  0  0  1]
 [ 1  7  0  0  0  0  1]
 [ 1  7  0  3  0  0  0]
 [ 3 21  0  0  0  0  1]
 [ 0 12  0  0  0  0  0]
 [ 3  7  0  0  0  0  3]]
```

Performance Evaluation:

	precision	recall	f1-score	support
0	0.25	0.25	0.25	12
1	0.17	0.86	0.29	14
2	0.00	0.00	0.00	9
3	0.60	0.27	0.37	11
4	0.00	0.00	0.00	25
5	0.00	0.00	0.00	12
6	0.33	0.23	0.27	13
accuracy			0.22	96
macro avg	0.19	0.23	0.17	96
weighted avg	0.17	0.22	0.15	96



Accuracy Score:

0.21875

GOOGLENET



IMPORTS, DATA PREPROCESSING AND INCEPTION MODULE GOOGLENET

```
In [34]: from tensorflow.keras.layers import Flatten, Input, Conv2D, Dense, Dropout, MaxPooling2D, AveragePooling2D, GlobalAv
from tensorflow.keras.layers import concatenate
from tensorflow.keras.models import Model
from tensorflow.keras.regularizers import l2
```

```
In [35]: train_images_normalized = train_images/255.0
test_images_normalized = test_images/255.0
```

```
In [36]: def inception(x,
                    filters_1x1,
                    filters_3x3_reduce,
                    filters_3x3,
                    filters_5x5_reduce,
                    filters_5x5,
                    filters_pool):
    path1 = layers.Conv2D(filters_1x1, (1, 1), padding='same', activation='relu')(x)
    path2 = layers.Conv2D(filters_3x3_reduce, (1, 1), padding='same', activation='relu')(x)
    path2 = layers.Conv2D(filters_3x3, (1, 1), padding='same', activation='relu')(path2)
    path3 = layers.Conv2D(filters_5x5_reduce, (1, 1), padding='same', activation='relu')(x)
    path3 = layers.Conv2D(filters_5x5, (1, 1), padding='same', activation='relu')(path3)
    path4 = layers.MaxPool2D((3, 3), strides=(1, 1), padding='same')(x)
    path4 = layers.Conv2D(filters_pool, (1, 1), padding='same', activation='relu')(path4)
    return tf.concat([path1, path2, path3, path4], axis=3)
```

MODEL ARCHITECTURE

```
In [40]: inp = layers.Input(shape=(224, 224, 3))
x = layers.Conv2D(64, 7, strides=2, padding='same', activation='relu')(inp)
x = layers.MaxPooling2D(3, strides=2)(x)
x = layers.Conv2D(64, 1, strides=1, padding='same', activation='relu')(x)
x = layers.Conv2D(192, 3, strides=1, padding='same', activation='relu')(x)
x = layers.MaxPooling2D(3, strides=2)(x)
x = inception(x, filters_1x1=64, filters_3x3_reduce=96, filters_3x3=128, filters_5x5_reduce=16, filters_5x5=32, filt
x = inception(x, filters_1x1=128, filters_3x3_reduce=128, filters_3x3=192, filters_5x5_reduce=32, filters_5x5=96, fi
x = layers.MaxPooling2D(3, strides=2)(x)
x = inception(x, filters_1x1=192, filters_3x3_reduce=96, filters_3x3=208, filters_5x5_reduce=16, filters_5x5=48, fil
aux1 = layers.AveragePooling2D((5, 5), strides=3)(x)
aux1 = layers.Conv2D(128, 1, padding='same', activation='relu')(aux1)
aux1 = layers.Flatten()(aux1)
aux1 = layers.Dense(1024, activation='relu')(aux1)
aux1 = layers.Dropout(0.7)(aux1)
aux1 = layers.Dense(10, activation='softmax')(aux1)
x = inception(x, filters_1x1=160, filters_3x3_reduce=112, filters_3x3=224, filters_5x5_reduce=24, filters_5x5=64, fi
x = inception(x, filters_1x1=128, filters_3x3_reduce=128, filters_3x3=256, filters_5x5_reduce=24, filters_5x5=64, fi
x = inception(x, filters_1x1=112, filters_3x3_reduce=144, filters_3x3=288, filters_5x5_reduce=32, filters_5x5=64, fi
aux2 = layers.AveragePooling2D((5, 5), strides=3)(x)
aux2 = layers.Conv2D(128, 1, padding='same', activation='relu')(aux2)
aux2 = layers.Flatten()(aux2)
aux2 = layers.Dense(1024, activation='relu')(aux2)
aux2 = layers.Dropout(0.7)(aux2)
aux2 = layers.Dense(10, activation='softmax')(aux2)
x = inception(x, filters_1x1=256, filters_3x3_reduce=160, filters_3x3=320, filters_5x5_reduce=32, filters_5x5=128, f
x = layers.MaxPooling2D(3, strides=2)(x)
x = inception(x, filters_1x1=256, filters_3x3_reduce=160, filters_3x3=320, filters_5x5_reduce=32, filters_5x5=128, f
x = inception(x, filters_1x1=384, filters_3x3_reduce=192, filters_3x3=384, filters_5x5_reduce=48, filters_5x5=128, f
x = layers.GlobalAveragePooling2D()(x)
x = layers.Dropout(0.4)(x)
out = layers.Dense(10, activation='softmax')(x)
```

```
In [41]: googlenet_model = Model(inputs = inp, outputs = [out, aux1, aux2])
```

COMPILATION

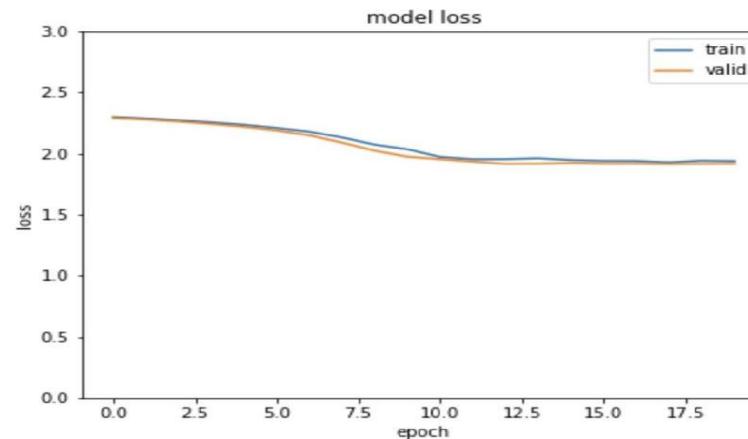
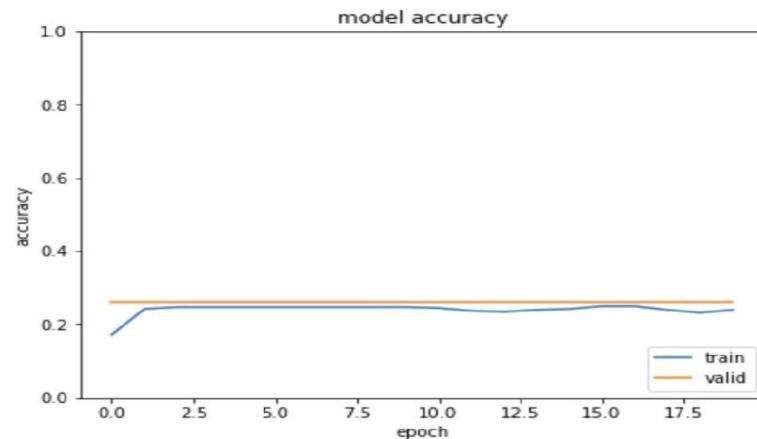
```
In [42]: googlenet_model.compile(optimizer=keras.optimizers.Adam(learning_rate=0.01, epsilon=0.05),
                               loss=[keras.losses.sparse_categorical_crossentropy,
                                     keras.losses.sparse_categorical_crossentropy,
                                     keras.losses.sparse_categorical_crossentropy],
                               loss_weights=[1, 0.3, 0.3],
                               metrics=['accuracy'])
```

TRAINING

```
In [43]: history = googlenet_model.fit(train_images_normalized, [train_labels, train_labels, train_labels], \
                                     validation_data=(test_images_normalized, [test_labels, test_labels, test_labels]), \
                                     batch_size=64, epochs=20)
0/6 [=====] - 7s 1s/step - loss: 3.1279 - dense_20_loss: 1.9389 - dense_17_loss: 1.9709 -  
dense_19_loss: 1.9865 - dense_20_accuracy: 0.2500 - dense_17_accuracy: 0.2266 - dense_19_accuracy: 0.2109 - val_los  
s: 3.0757 - val_dense_20_loss: 1.9162 - val_dense_17_loss: 1.9300 - val_dense_19_loss: 1.9349 - val_dense_20_accura  
cy: 0.2604 - val_dense_17_accuracy: 0.2604 - val_dense_19_accuracy: 0.2604  
Epoch 18/20  
6/6 [=====] - 7s 1s/step - loss: 3.0956 - dense_20_loss: 1.9276 - dense_17_loss: 1.9489 -  
dense_19_loss: 1.9446 - dense_20_accuracy: 0.2396 - dense_17_accuracy: 0.2344 - dense_19_accuracy: 0.2448 - val_los  
s: 3.0679 - val_dense_20_loss: 1.9132 - val_dense_17_loss: 1.9202 - val_dense_19_loss: 1.9288 - val_dense_20_accura  
cy: 0.2604 - val_dense_17_accuracy: 0.2604 - val_dense_19_accuracy: 0.2604  
Epoch 19/20  
6/6 [=====] - 7s 1s/step - loss: 3.1197 - dense_20_loss: 1.9413 - dense_17_loss: 1.9704 -  
dense_19_loss: 1.9575 - dense_20_accuracy: 0.2318 - dense_17_accuracy: 0.2188 - dense_19_accuracy: 0.2266 - val_los  
s: 3.0688 - val_dense_20_loss: 1.9148 - val_dense_17_loss: 1.9197 - val_dense_19_loss: 1.9271 - val_dense_20_accura  
cy: 0.2604 - val_dense_17_accuracy: 0.2604 - val_dense_19_accuracy: 0.2604  
Epoch 20/20  
6/6 [=====] - 7s 1s/step - loss: 3.1085 - dense_20_loss: 1.9374 - dense_17_loss: 1.9522 -  
dense_19_loss: 1.9514 - dense_20_accuracy: 0.2396 - dense_17_accuracy: 0.2448 - dense_19_accuracy: 0.2422 - val_los  
s: 3.0718 - val_dense_20_loss: 1.9160 - val_dense_17_loss: 1.9246 - val_dense_19_loss: 1.9281 - val_dense_20_accura  
cy: 0.2604 - val_dense_17_accuracy: 0.2604 - val_dense_19_accuracy: 0.2604
```

TRAINING AND LOSS CURVES

```
In [46]: plt.figure(figsize=(15,5))
plt.subplot(1,2,1)
plt.plot(history.history['dense_20_accuracy'])
plt.plot(history.history['val_dense_20_accuracy'])
plt.ylim(0, 1)
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'valid'], loc='lower right')
plt.subplot(1,2,2)
plt.plot(history.history['dense_20_loss'])
plt.plot(history.history['val_dense_20_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'valid'], loc='upper right')
plt.ylim([0,3])
plt.savefig("googlenet_loss_acc.png")
```



PERFORMANCE

Confusion Matrix:

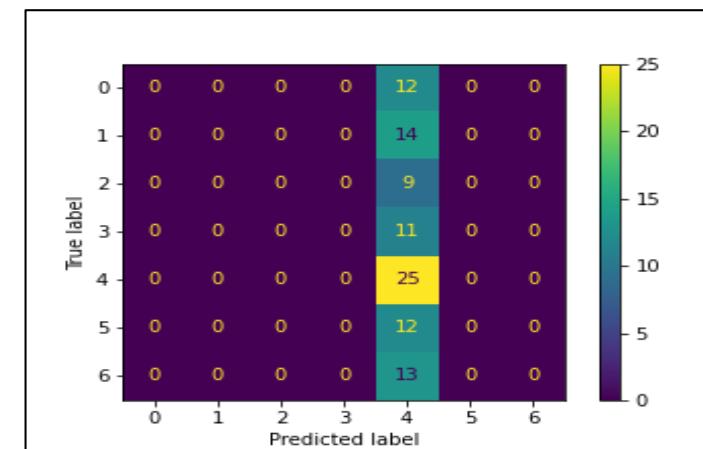
```
[[ 0  0  0  0 12  0  0]
 [ 0  0  0  0 14  0  0]
 [ 0  0  0  0  9  0  0]
 [ 0  0  0  0 11  0  0]
 [ 0  0  0  0 25  0  0]
 [ 0  0  0  0 12  0  0]
 [ 0  0  0  0 13  0  0]]
```

Performance Evaluation:

	precision	recall	f1-score	support
0	0.00	0.00	0.00	12
1	0.00	0.00	0.00	14
2	0.00	0.00	0.00	9
3	0.00	0.00	0.00	11
4	0.26	1.00	0.41	25
5	0.00	0.00	0.00	12
6	0.00	0.00	0.00	13
accuracy			0.26	96
macro avg	0.04	0.14	0.06	96
weighted avg	0.07	0.26	0.11	96

Accuracy Score:

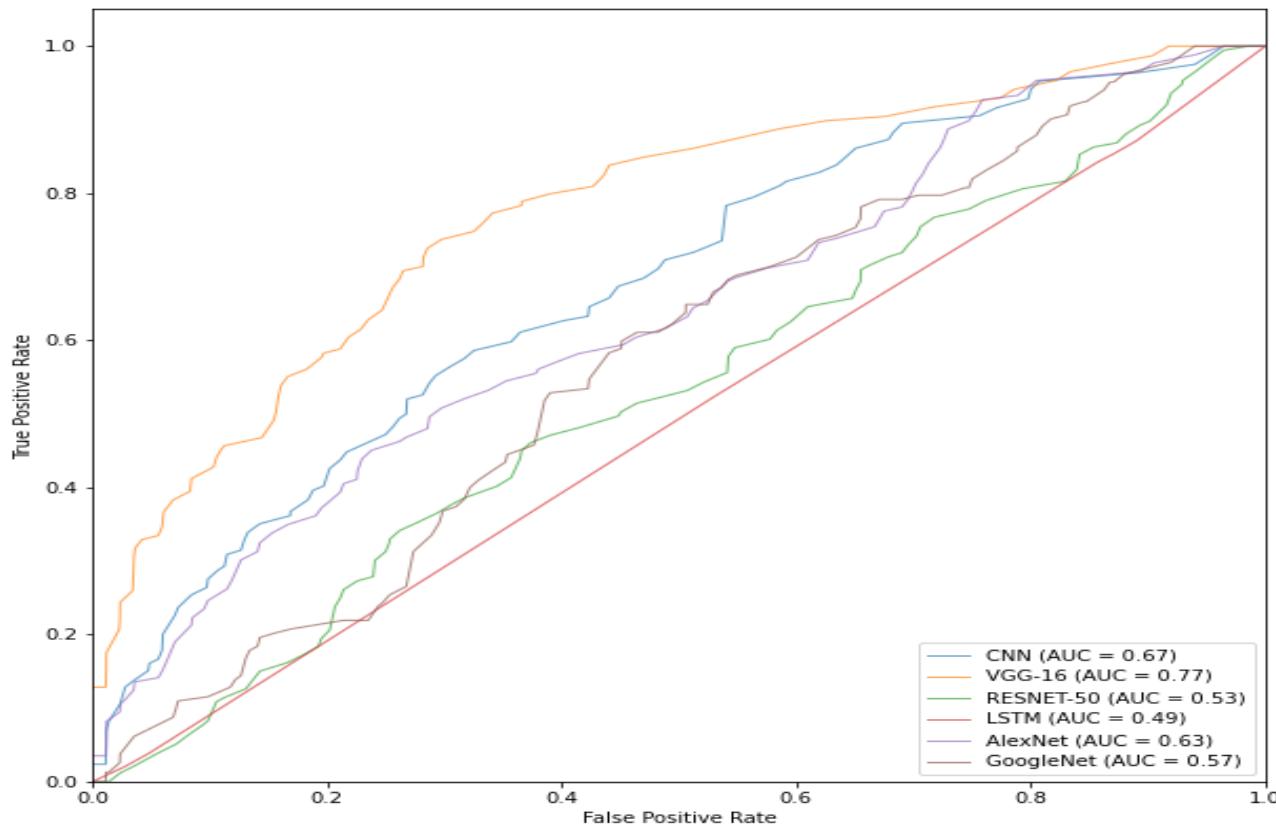
0.2604166666666667



COMPARISON BETWEEN CNN, VGG-16, RESNET-50, LSTM(RNN), ALEXNET AND GOOGLENET.



ROC CURVES AND AUC COMPARISON



ROC CURVES AND AUC COMPARISON - CODE (I)

```
In [51]: def plot_macro_avg_roc_curve(classifier, n_classes, name, X_test, y_test, single_output=True):

    y_test_binarized = label_binarize(y_test, classes=list(range(0,n_classes)))
    y_score = classifier.predict(X_test)

    #for cases where the model (like googlenet) has extra (auxillary) outputs
    if not single_output:
        y_score = y_score[0]

    # Compute ROC curve and ROC area for each class
    fpr = dict()
    tpr = dict()
    roc_auc = dict()
    for i in range(n_classes):
        fpr[i], tpr[i], _ = roc_curve(y_test_binarized[:, i], y_score[:, i])
        roc_auc[i] = auc(fpr[i], tpr[i])

    # First aggregate all false positive rates
    all_fpr = np.unique(np.concatenate([fpr[i] for i in range(n_classes)]))

    # Then interpolate all ROC curves at this points
    mean_tpr = np.zeros_like(all_fpr)
    for i in range(n_classes):
        mean_tpr += np.interp(all_fpr, fpr[i], tpr[i])

    # Finally average it and compute AUC
    mean_tpr /= n_classes

    fpr["macro"] = all_fpr
    tpr["macro"] = mean_tpr
    roc_auc["macro"] = auc(fpr["macro"], tpr["macro"])

    plt.plot(
        fpr["macro"],
        tpr["macro"],
        label="{0} (AUC = {1:0.2f})".format(name, roc_auc["macro"]),
        alpha=0.7,
        linewidth=1,
    )

    return plt
```

ROC CURVES AND AUC COMPARISON - CODE (II)

```
In [52]:
```

```
cnn = load_model('cnn.h5')
vgg_based_model = load_model('vgg16.h5')
resnet_based_model = load_model('resnet50.h5')
lstm_based_model = load_model('lstm.h5')
alexnet_model = load_model('alexnet.h5')
googlenet_model = load_model('googlenet.h5')
```

```
In [59]:
```

```
plt.figure(figsize=(10, 10))
plot_macro_avg_roc_curve(cnn, num_classes, "CNN", test_images, test_labels)
plot_macro_avg_roc_curve(vgg_based_model, num_classes, "VGG-16", test_images_p, test_labels)
plot_macro_avg_roc_curve(resnet_based_model, num_classes, "RESNET-50", test_images_pr, test_labels)
plot_macro_avg_roc_curve(lstm_based_model, num_classes, "LSTM", test_images, test_labels)
plot_macro_avg_roc_curve(alexnet_model, num_classes, "AlexNet", test_images, test_labels)
plot_macro_avg_roc_curve(googlenet_model, num_classes, "GoogleNet", test_images_normalized, test_labels, single_outp

plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.legend(loc="lower right")

plt.savefig('roc_auc.png')
```

COMPARISON OF PERFORMANCE

CLASSIFIER	Accuracy	Precision	Recall	F1-Score	AUC
CNN	0.31	0.39	0.31	0.31	0.67
VGG-16	0.45	0.54	0.45	0.43	0.77
RESNET-50	0.16	0.04	0.16	0.06	0.53
LSTM	0.26	0.07	0.26	0.11	0.49
ALEXNET	0.22	0.17	0.22	0.15	0.63
GOOGLENET	0.26	0.07	0.26	0.11	0.57

CONCLUSION:

VGG-16 performed better than the other models. However, all the models performed poorly. It appears that it is hard to generalise patterns as all humans convey emotions differently.

EMODB- CLASSIFICATION



CONVERSION TO WAVEFORM (IMAGE)



IMPORTS AND CLASS NAMES

```
In [1]: import matplotlib.pyplot as plt
import tensorflow as tf
from matplotlib.backend_bases import RendererBase
from scipy import signal
from scipy.io import wavfile
#import soundfile as sf
import os
import numpy as np
from PIL import Image
from scipy.fftpack import fft

%matplotlib inline
```

Some mappings

```
In [2]: class_names = ['anger', 'disgust', 'fear', 'happiness', 'neutral', 'sadness', 'boredom']
lsymbols = ['a', 'd', 'f', 'h', 'n', 'sa', 'su']
full = {
    'W': 'anger',
    'E': 'disgust',
    'A': 'fear',
    'L': 'boredom',
    'F': 'happiness',
    'N': 'neutral',
    'T': 'sadness',
}
```

AUDIO

PREPROCESSIN

Audio Preprocessing (Conversion to images of waveform)

```
In [3]: audio_path = './AudioData/'  
pict_Path = './ImageData/'  
samples = []  
  
In [4]: if not os.path.exists(pict_Path):  
    os.makedirs(pict_Path)  
  
In [5]: for x in class_names:  
    os.makedirs(pict_Path + '/' + x)  
  
In [6]: def wav2img_waveform(wav_path, targetdir='', figsize=(4,4)):  
    samplerate,test_sound = wavfile.read(wav_path)  
    fig = plt.figure(figsize=figsize)  
    plt.plot(test_sound)  
    plt.axis('off')  
    output_file = wav_path.split('/')[-1].split('.wav')[0]  
    output_file = targetdir + '/' + output_file  
    plt.savefig(f'{output_file}.png')  
    plt.close()  
  
In [7]: all_files = [y for y in os.listdir(audio_path) if '.wav' in y]  
for file in all_files:  
    directory = full[file[5]]  
    wav2img_waveform(audio_path + file, pict_Path + '/' + directory)
```

LOADING THE IMAGES INTO TRAIN AND TEST DATASETS

```
In [10]: class_names = ['anger', 'disgust', 'fear', 'happiness', 'neutral', 'sadness', 'boredom']
img_height = 224
img_width = 224
num_train_images = 428
num_test_images = 107
num_classes = len(class_names)
```

```
In [11]: train_ds = tf.keras.utils.image_dataset_from_directory(
    pict_Path,
    validation_split=0.2,
    subset="training",
    seed=123,
    image_size=(img_height, img_width),
    batch_size=num_train_images,
    label_mode='int')

Found 535 files belonging to 7 classes.
Using 428 files for training.
```

```
In [12]: test_ds = tf.keras.utils.image_dataset_from_directory(
    pict_Path,
    validation_split=0.2,
    subset="validation",
    seed=123,
    image_size=(img_height, img_width),
    batch_size=num_test_images,
    label_mode='int')
```

```
Found 535 files belonging to 7 classes.
Using 107 files for validation.
```

```
In [13]: train_images = train_labels = test_images = test_labels = None
```

```
In [49]: train_images, train_labels = next(train_ds.as_numpy_iterator())
```

```
In [15]: test_images, test_labels = next(test_ds.as_numpy_iterator())
```

CNN



IMPORTS AND MODEL ARCHITECTURE

Convolutional Neural Network

Importing the libraries

```
In [16]: from tensorflow import keras
from tensorflow.keras import layers, datasets
from tensorflow.keras.models import Sequential
import matplotlib.pyplot as plt
```

Building the CNN

```
In [17]: cnn = Sequential([
    layers.Rescaling(1./255, input_shape=(img_height, img_width, 3)),
    layers.Conv2D(filters=32, kernel_size=3, padding='same', activation='relu'),
    layers.Conv2D(filters=32, kernel_size=3, padding='same', activation='relu'),
    layers.MaxPooling2D(pool_size=2, strides=2),
    layers.Conv2D(32, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(pool_size=2, strides=2),
    layers.Dropout(0.2),
    layers.Flatten(),
    layers.Dense(128, activation='relu'),
    layers.Dense(num_classes, activation="softmax")
])
```

COMPILATION AND SUMMARY

Compiling the CNN

```
In [18]: cnn.compile(optimizer=keras.optimizers.Adam(learning_rate=0.0005, epsilon=0.01),
                  loss=tf.keras.losses.SparseCategoricalCrossentropy(),
                  metrics=['accuracy'])
```

```
In [19]: cnn.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
<hr/>		
rescaling (Rescaling)	(None, 224, 224, 3)	0
conv2d (Conv2D)	(None, 224, 224, 32)	896
conv2d_1 (Conv2D)	(None, 224, 224, 32)	9248
max_pooling2d (MaxPooling2D)	(None, 112, 112, 32)	0
conv2d_2 (Conv2D)	(None, 112, 112, 32)	9248
max_pooling2d_1 (MaxPooling2D)	(None, 56, 56, 32)	0
dropout (Dropout)	(None, 56, 56, 32)	0
flatten (Flatten)	(None, 100352)	0
dense (Dense)	(None, 128)	12845184
dense_1 (Dense)	(None, 7)	903
<hr/>		
Total params: 12,865,479		
Trainable params: 12,865,479		
Non-trainable params: 0		

TRAINING AND SAVING

Training the CNN on the Training set and evaluating it on the Test set

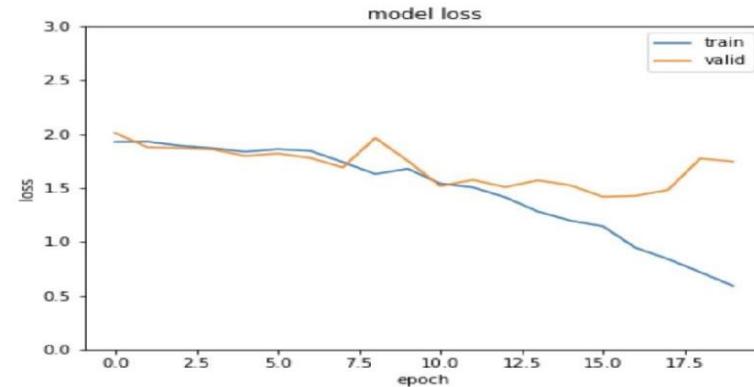
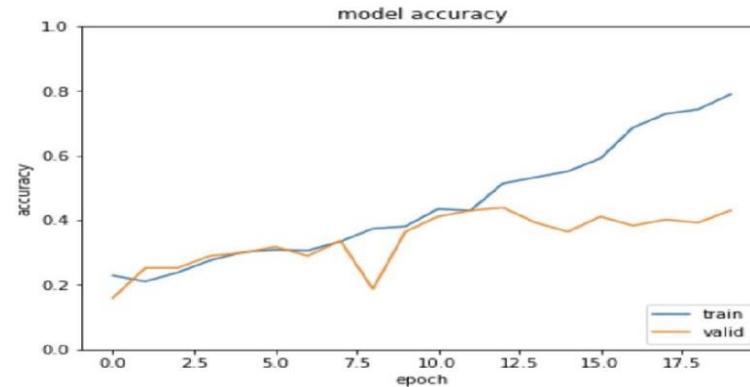
```
In [20]: history = cnn.fit(train_images, train_labels, epochs=20, validation_data=(test_images, test_labels))
Epoch 1/20
14/14 [=====] - 4s 253ms/step - loss: 1.2000 - accuracy: 0.3925 - val_loss: 1.5242 - val_accuracy: 0.3925
Epoch 15/20
14/14 [=====] - 4s 254ms/step - loss: 1.1960 - accuracy: 0.5514 - val_loss: 1.5242 - val_accuracy: 0.3645
Epoch 16/20
14/14 [=====] - 4s 253ms/step - loss: 1.1445 - accuracy: 0.5911 - val_loss: 1.4174 - val_accuracy: 0.4112
Epoch 17/20
14/14 [=====] - 4s 253ms/step - loss: 0.9447 - accuracy: 0.6869 - val_loss: 1.4263 - val_accuracy: 0.3832
Epoch 18/20
14/14 [=====] - 4s 253ms/step - loss: 0.8399 - accuracy: 0.7290 - val_loss: 1.4798 - val_accuracy: 0.4019
Epoch 19/20
14/14 [=====] - 4s 254ms/step - loss: 0.7169 - accuracy: 0.7430 - val_loss: 1.7733 - val_accuracy: 0.3925
Epoch 20/20
14/14 [=====] - 4s 252ms/step - loss: 0.5905 - accuracy: 0.7897 - val_loss: 1.7445 - val_accuracy: 0.4299
```

```
In [21]: cnn.save('cnn.h5')
```

TRAINING AND LOSS CURVES

```
In [25]: def save_train_loss_curves(history, filename):
    plt.figure(figsize=(15,5))
    plt.subplot(1,2,1)
    plt.plot(history.history['accuracy'])
    plt.plot(history.history['val_accuracy'])
    plt.ylim(0, 1)
    plt.title('model accuracy')
    plt.ylabel('accuracy')
    plt.xlabel('epoch')
    plt.legend(['train', 'valid'], loc='lower right')
    plt.subplot(1,2,2)
    plt.plot(history.history['loss'])
    plt.plot(history.history['val_loss'])
    plt.title('model loss')
    plt.ylabel('loss')
    plt.xlabel('epoch')
    plt.legend(['train', 'valid'], loc='upper right')
    plt.ylim([0,3])
    plt.savefig(filename)
```

```
In [26]: save_train_loss_curves(history, 'cnn_loss_acc.png')
```



PERFORMANCE - CODE

```
In [22]: from sklearn.metrics import classification_report, confusion_matrix, accuracy_score, ConfusionMatrixDisplay
```

```
In [23]: def get_predictions(model, inputs):
    def index_max_arr_element(arr):
        best_index = -1
        best = float('-inf')
        for i, num in enumerate(arr):
            if num > best:
                best = num
                best_index = i
        return best_index

    pred_arr = model.predict(inputs)
    pred_labels = [index_max_arr_element(row) for row in pred_arr]
    return pred_labels
```

```
In [24]: pred_labels = get_predictions(cnn, test_images)

print("Confusion Matrix:")
print(confusion_matrix(test_labels, pred_labels))
print("-----")
print("-----")
print("Performance Evaluation:")
print(classification_report(test_labels, pred_labels))
print("-----")
print("-----")
print("Accuracy Score:")
print(accuracy_score(test_labels, pred_labels))

plt.figure()
ConfusionMatrixDisplay.from_predictions(test_labels, pred_labels)
plt.savefig('./cnn_matrix.png')
```

PERFORMANCE

Confusion Matrix:

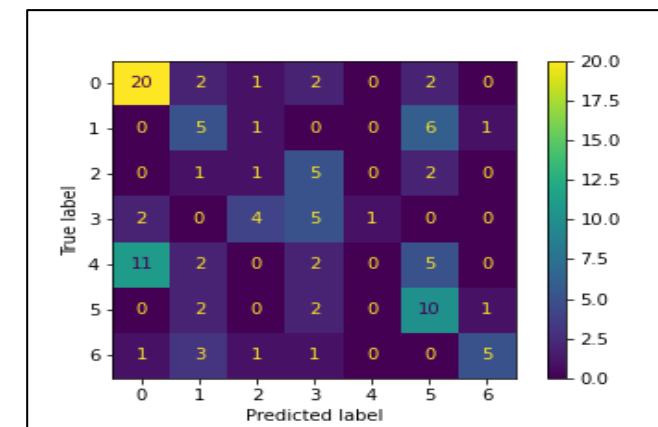
```
[[20  2  1  2  0  2  0]
 [ 0  5  1  0  0  6  1]
 [ 0  1  1  5  0  2  0]
 [ 2  0  4  5  1  0  0]
 [11  2  0  2  0  5  0]
 [ 0  2  0  2  0 10  1]
 [ 1  3  1  1  0  0  5]]
```

Performance Evaluation:

	precision	recall	f1-score	support
0	0.59	0.74	0.66	27
1	0.33	0.38	0.36	13
2	0.12	0.11	0.12	9
3	0.29	0.42	0.34	12
4	0.00	0.00	0.00	20
5	0.40	0.67	0.50	15
6	0.71	0.45	0.56	11
accuracy			0.43	107
macro avg	0.35	0.40	0.36	107
weighted avg	0.36	0.43	0.38	107

Accuracy Score:

0.42990654205607476



VGG-16

IMPORTS, DATA PREPROCESSING & ARCHITECTURE

VGG-16

```
In [27]: from tensorflow.keras.applications import vgg16
```

```
In [28]: train_images_p, test_images_p = vgg16.preprocess_input(np.copy(train_images)), vgg16.preprocess_input(np.copy(test_i
```

```
In [29]: vgg_based_model = Sequential([
    layers.Rescaling(1./255, input_shape=(img_height, img_width, 3)),
    vgg16.VGG16(include_top=False, weights="imagenet", input_shape=(img_height, img_width, 3), pooling="max"),
    layers.Flatten(),
    layers.Dense(512, activation='relu'),
    layers.Dropout(0.2),
    layers.Dense(128, activation='relu'),
    layers.Dropout(0.2),
    layers.Dense(num_classes, activation="softmax")
])
```

COMPILATION AND SUMMARY

```
In [30]: vgg_based_model.compile(optimizer=keras.optimizers.Adam(learning_rate=0.0005, epsilon=0.01),  
                           loss=tf.keras.losses.SparseCategoricalCrossentropy(),  
                           metrics=['accuracy'])
```

```
In [31]: vgg_based_model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
rescaling_1 (Rescaling)	(None, 224, 224, 3)	0
vgg16 (Functional)	(None, 512)	14714688
flatten_1 (Flatten)	(None, 512)	0
dense_2 (Dense)	(None, 512)	262656
dropout_1 (Dropout)	(None, 512)	0
dense_3 (Dense)	(None, 128)	65664
dropout_2 (Dropout)	(None, 128)	0
dense_4 (Dense)	(None, 7)	903
<hr/>		
Total params:	15,043,911	
Trainable params:	15,043,911	
Non-trainable params:	0	

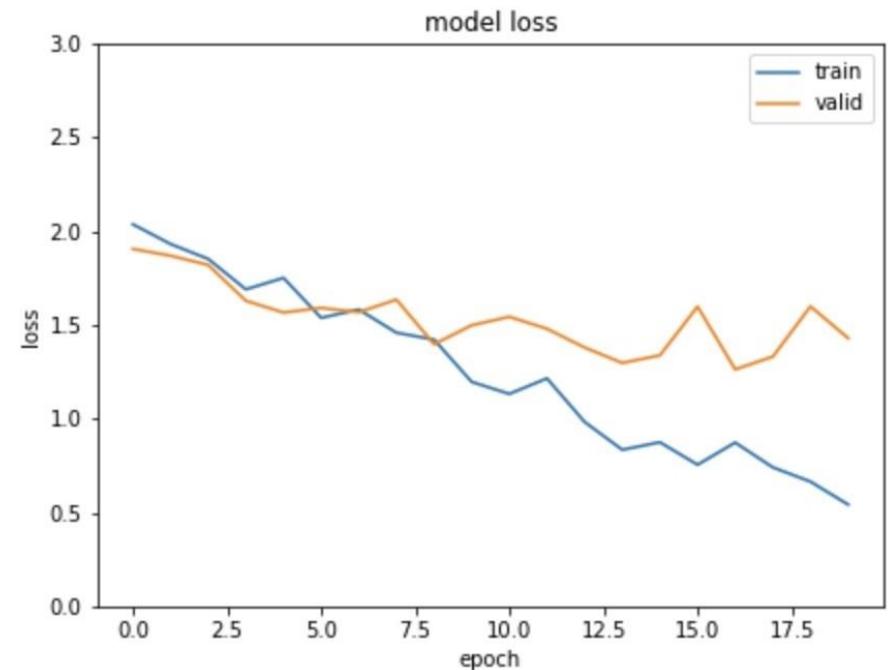
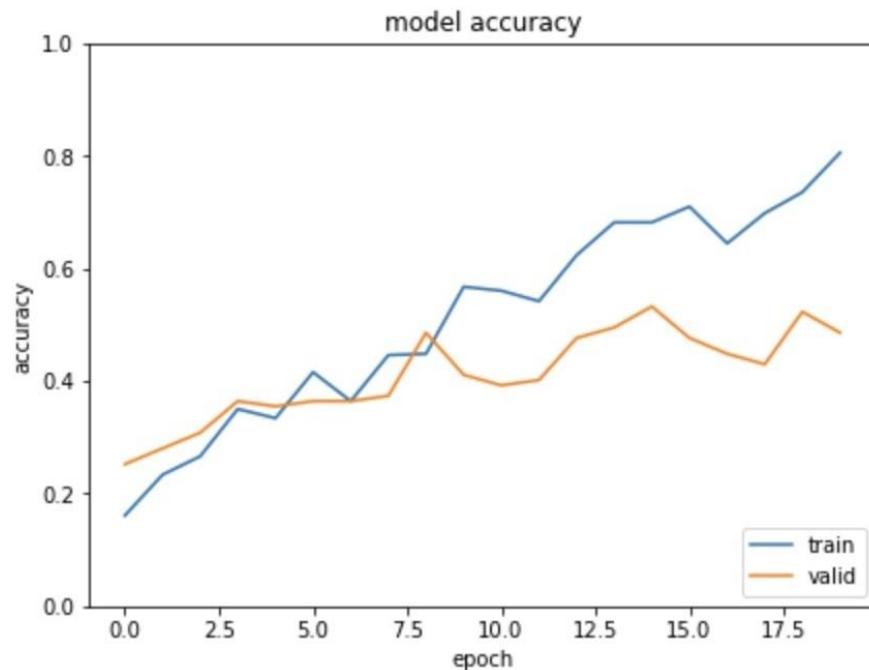
TRAINING AND SAVING

```
In [32]: history = vgg_based_model.fit(train_images_p, train_labels, epochs=20, validation_data=(test_images_p, test_labels))  
    accuracy: 0.4953  
Epoch 15/20  
14/14 [=====] - 39s 3s/step - loss: 0.8741 - accuracy: 0.6822 - val_loss: 1.3374 - val_accuracy: 0.5327  
Epoch 16/20  
14/14 [=====] - 40s 3s/step - loss: 0.7542 - accuracy: 0.7103 - val_loss: 1.5978 - val_accuracy: 0.4766  
Epoch 17/20  
14/14 [=====] - 41s 3s/step - loss: 0.8725 - accuracy: 0.6449 - val_loss: 1.2624 - val_accuracy: 0.4486  
Epoch 18/20  
14/14 [=====] - 39s 3s/step - loss: 0.7401 - accuracy: 0.6986 - val_loss: 1.3304 - val_accuracy: 0.4299  
Epoch 19/20  
14/14 [=====] - 39s 3s/step - loss: 0.6640 - accuracy: 0.7360 - val_loss: 1.5984 - val_accuracy: 0.5234  
Epoch 20/20  
14/14 [=====] - 38s 3s/step - loss: 0.5418 - accuracy: 0.8061 - val_loss: 1.4294 - val_accuracy: 0.4860
```

```
In [33]: vgg_based_model.save('vgg16.h5')
```

TRAINING AND LOSS CURVES

```
In [34]: save_train_loss_curves(history, 'vgg_loss_acc.png')
```



PERFORMANCE

Confusion Matrix:

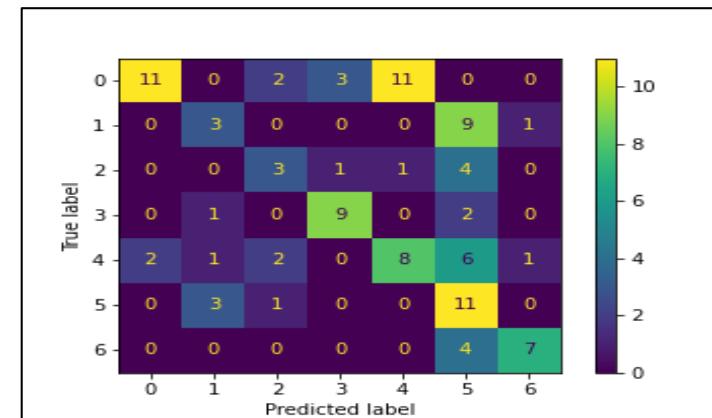
```
[[11  0  2  3 11  0  0]
 [ 0  3  0  0  0  9  1]
 [ 0  0  3  1  1  4  0]
 [ 0  1  0  9  0  2  0]
 [ 2  1  2  0  8  6  1]
 [ 0  3  1  0  0 11  0]
 [ 0  0  0  0  0  4  7]]
```

Performance Evaluation:

	precision	recall	f1-score	support
0	0.85	0.41	0.55	27
1	0.38	0.23	0.29	13
2	0.38	0.33	0.35	9
3	0.69	0.75	0.72	12
4	0.40	0.40	0.40	20
5	0.31	0.73	0.43	15
6	0.78	0.64	0.70	11
accuracy			0.49	107
macro avg	0.54	0.50	0.49	107
weighted avg	0.57	0.49	0.49	107

Accuracy Score:

0.48598130841121495



RES NET- 50



IMPORTS, DATA PREPROCESSING & ARCHITECTURE

```
In [36]: from tensorflow.keras.applications import resnet50
```

```
In [37]: train_images_pr, test_images_pr = resnet50.preprocess_input(np.copy(train_images)), resnet50.preprocess_input(np.cop
```

```
In [38]: resnet_based_model = Sequential([
    layers.Rescaling(1./255, input_shape=(img_height, img_width, 3)),
    resnet50.ResNet50(include_top=False, weights="imagenet", input_shape=(img_height, img_width, 3), pooling="max"),
    layers.Flatten(),
    layers.Dense(512, activation='relu'),
    layers.Dropout(0.2),
    layers.Dense(128, activation='relu'),
    layers.Dropout(0.2),
    layers.Dense(num_classes, activation="softmax")
])
```

COMPILATION AND SUMMARY

```
In [39]: resnet_based_model.compile(optimizer=keras.optimizers.Adam(learning_rate=0.0005, epsilon=0.01),  
                                loss=tf.keras.losses.SparseCategoricalCrossentropy(),  
                                metrics=['accuracy'])
```

```
In [40]: resnet_based_model.summary()
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
rescaling_2 (Rescaling)	(None, 224, 224, 3)	0
resnet50 (Functional)	(None, 2048)	23587712
flatten_2 (Flatten)	(None, 2048)	0
dense_5 (Dense)	(None, 512)	1049088
dropout_3 (Dropout)	(None, 512)	0
dense_6 (Dense)	(None, 128)	65664
dropout_4 (Dropout)	(None, 128)	0
dense_7 (Dense)	(None, 7)	903

Total params: 24,703,367
Trainable params: 24,650,247
Non-trainable params: 53,120

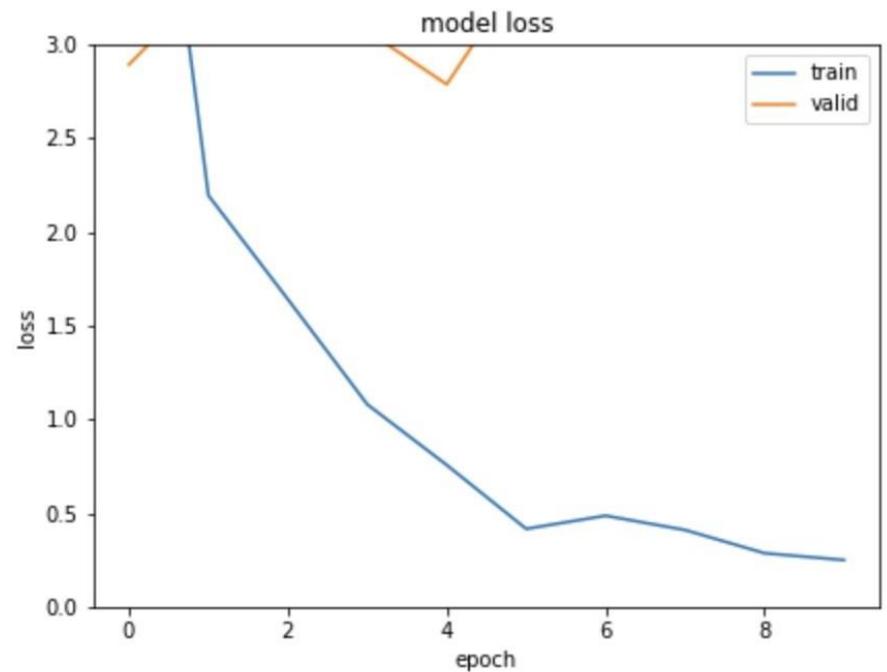
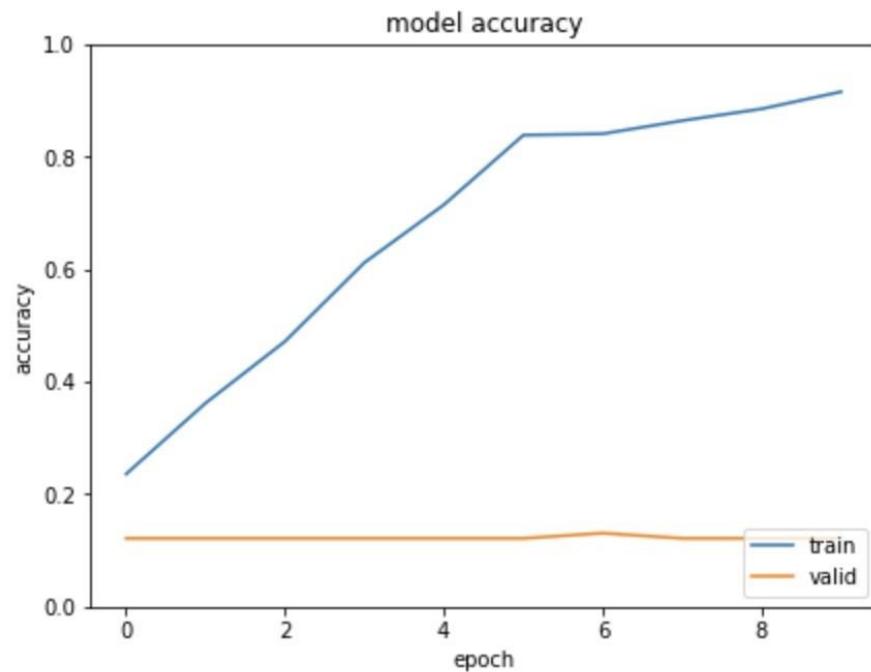
TRAINING AND SAVING

```
In [41]: history = resnet_based_model.fit(train_images_pr, train_labels, epochs=10, validation_data=(test_images_pr, test_labels))
         accuracy: 0.1215
Epoch 5/10
14/14 [=====] - 17s 1s/step - loss: 0.7559 - accuracy: 0.7150 - val_loss: 2.7873 - val_accuracy: 0.1215
Epoch 6/10
14/14 [=====] - 17s 1s/step - loss: 0.4153 - accuracy: 0.8388 - val_loss: 3.4255 - val_accuracy: 0.1215
Epoch 7/10
14/14 [=====] - 17s 1s/step - loss: 0.4854 - accuracy: 0.8411 - val_loss: 3.2850 - val_accuracy: 0.1308
Epoch 8/10
14/14 [=====] - 17s 1s/step - loss: 0.4092 - accuracy: 0.8645 - val_loss: 4.2360 - val_accuracy: 0.1215
Epoch 9/10
14/14 [=====] - 17s 1s/step - loss: 0.2854 - accuracy: 0.8855 - val_loss: 3.5137 - val_accuracy: 0.1215
Epoch 10/10
14/14 [=====] - 17s 1s/step - loss: 0.2489 - accuracy: 0.9159 - val_loss: 4.0176 - val_accuracy: 0.1215
```

```
In [42]: resnet_based_model.save('resnet50.h5')
```

TRAINING AND LOSS CURVES

```
[n 43]: save_train_loss_curves(history, 'resnet_loss_acc.png')
```



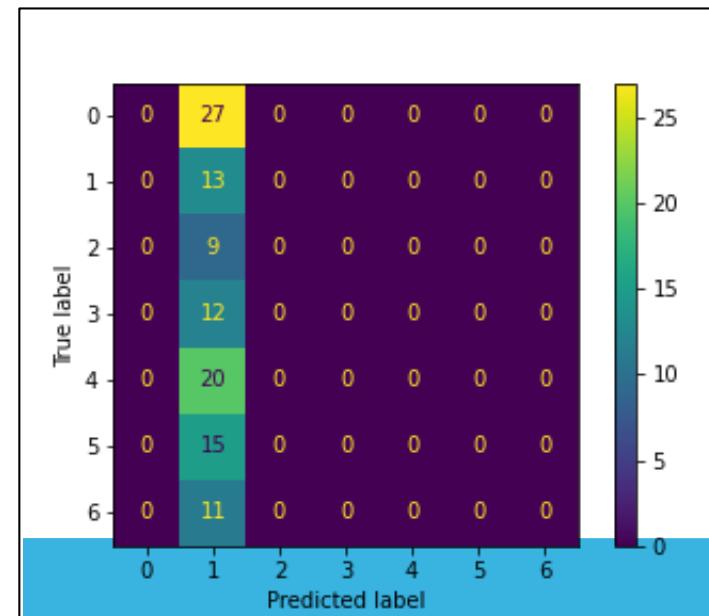
PERFORMANCE

Confusion Matrix:

```
[[ 0 27  0  0  0  0  0]
 [ 0 13  0  0  0  0  0]
 [ 0  9  0  0  0  0  0]
 [ 0 12  0  0  0  0  0]
 [ 0 20  0  0  0  0  0]
 [ 0 15  0  0  0  0  0]
 [ 0 11  0  0  0  0  0]]
```

Performance Evaluation:

	precision	recall	f1-score	support
0	0.00	0.00	0.00	27
1	0.12	1.00	0.22	13
2	0.00	0.00	0.00	9
3	0.00	0.00	0.00	12
4	0.00	0.00	0.00	20
5	0.00	0.00	0.00	15
6	0.00	0.00	0.00	11
accuracy			0.12	107
macro avg	0.02	0.14	0.03	107
weighted avg	0.01	0.12	0.03	107



Accuracy Score:

0.12149532710280374

RNN (LSTM)



MODEL ARCHITECTURE AND COMPIRATION

RNN (LSTM based)

```
In [45]: lstm_based_model = Sequential([
    layers.Rescaling(1./255, input_shape=(img_height, img_width, 3)),
    layers.Resizing(64, 64),
    layers.Reshape((4096, 3)),
    layers.LSTM(units=50, return_sequences=True),
    layers.Dropout(0.2),
    layers.LSTM(units=50),
    layers.Dropout(0.2),
    layers.Dense(128, activation='relu'),
    layers.Dropout(0.2),
    layers.Dense(num_classes, activation="softmax")
])
```

```
In [46]: lstm_based_model.compile(optimizer=keras.optimizers.Adam(learning_rate=0.001),
                               loss=tf.keras.losses.SparseCategoricalCrossentropy(),
                               metrics=['accuracy'])
```

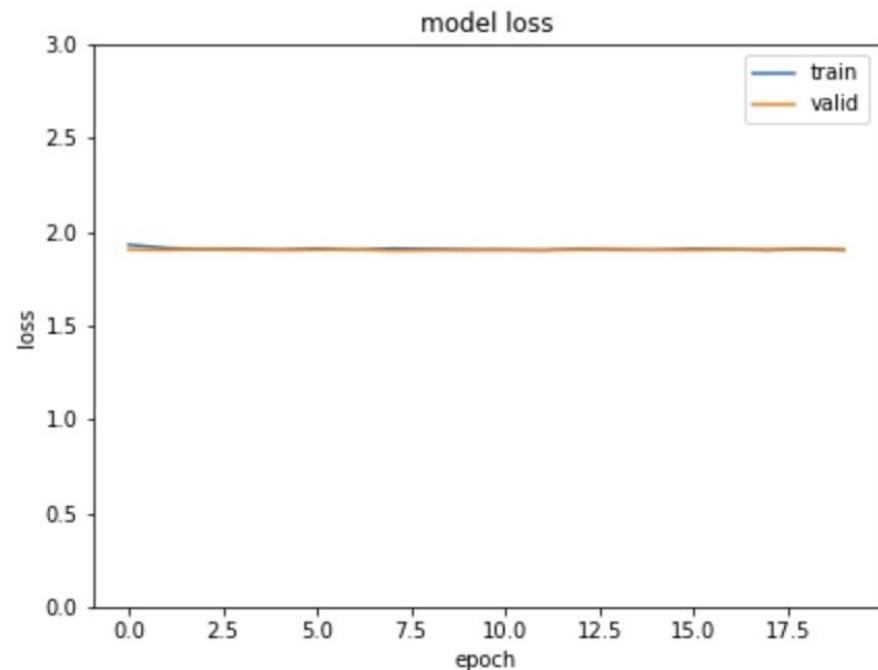
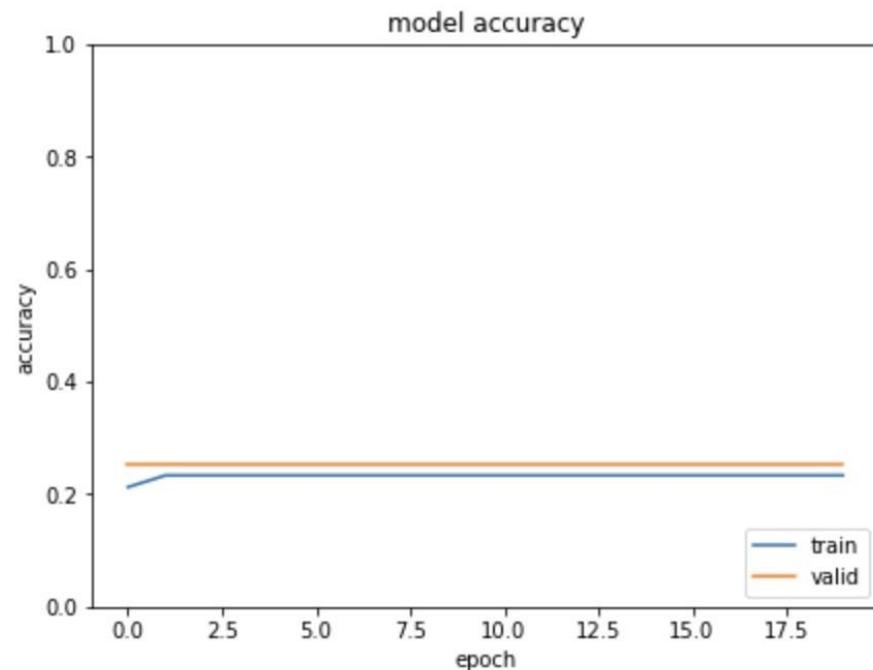
TRAINING AND SAVING

```
In [48]: history = lstm_based_model.fit(train_images, train_labels, epochs=20, validation_data=(test_images, test_labels))  
uracy: 0.2523  
Epoch 15/20  
14/14 [=====] - 52s 4s/step - loss: 1.9035 - accuracy: 0.2336 - val_loss: 1.9032 - val_accuracy: 0.2523  
uracy: 0.2523  
Epoch 16/20  
14/14 [=====] - 53s 4s/step - loss: 1.9087 - accuracy: 0.2336 - val_loss: 1.9022 - val_accuracy: 0.2523  
uracy: 0.2523  
Epoch 17/20  
14/14 [=====] - 52s 4s/step - loss: 1.9065 - accuracy: 0.2336 - val_loss: 1.9051 - val_accuracy: 0.2523  
uracy: 0.2523  
Epoch 18/20  
14/14 [=====] - 52s 4s/step - loss: 1.9017 - accuracy: 0.2336 - val_loss: 1.9059 - val_accuracy: 0.2523  
uracy: 0.2523  
Epoch 19/20  
14/14 [=====] - 52s 4s/step - loss: 1.9107 - accuracy: 0.2336 - val_loss: 1.9067 - val_accuracy: 0.2523  
uracy: 0.2523  
Epoch 20/20  
14/14 [=====] - 53s 4s/step - loss: 1.9017 - accuracy: 0.2336 - val_loss: 1.9070 - val_accuracy: 0.2523
```

```
In [50]: lstm_based_model.save('lstm.h5')
```

TRAINING AND LOSS CURVES

```
[51]: save_train_loss_curves(history, 'lstm_loss_acc.png')
```



PERFORMANCE

Confusion Matrix:

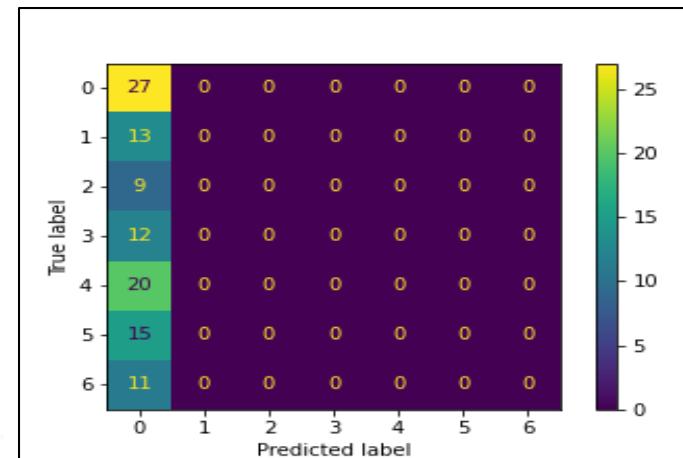
```
[[27  0  0  0  0  0  0]
 [13  0  0  0  0  0  0]
 [ 9  0  0  0  0  0  0]
 [12  0  0  0  0  0  0]
 [20  0  0  0  0  0  0]
 [15  0  0  0  0  0  0]
 [11  0  0  0  0  0  0]]
```

Performance Evaluation:

	precision	recall	f1-score	support
0	0.25	1.00	0.40	27
1	0.00	0.00	0.00	13
2	0.00	0.00	0.00	9
3	0.00	0.00	0.00	12
4	0.00	0.00	0.00	20
5	0.00	0.00	0.00	15
6	0.00	0.00	0.00	11
accuracy			0.25	107
macro avg	0.04	0.14	0.06	107
weighted avg	0.06	0.25	0.10	107

Accuracy Score:

0.2523364485981308



ALEXNET



MODEL ARCHITECTURE - I

ALEXNET

```
In [52]: from tensorflow.keras.layers import Dense, Dropout, Activation, Flatten, Conv2D, MaxPooling2D, ZeroPadding2D, BatchN  
from tensorflow.keras.regularizers import l2
```

```
In [53]: def alexnet_model(img_shape=(32, 32, 3), n_classes=10, l2_reg=0.,  
weights=None):  
  
    # Initialize model  
    alexnet = Sequential()  
    alexnet.add(layers.Input(shape=img_shape))  
    alexnet.add(layers.Resizing(64, 64))  
  
    # Layer 1  
    alexnet.add(Conv2D(96, (11, 11), padding='same', kernel_regularizer=l2(l2_reg)))  
    alexnet.add(BatchNormalization())  
    alexnet.add(Activation('relu'))  
    alexnet.add(MaxPooling2D(pool_size=(2, 2)))  
  
    # Layer 2  
    alexnet.add(Conv2D(256, (5, 5), padding='same'))  
    alexnet.add(BatchNormalization())  
    alexnet.add(Activation('relu'))  
    alexnet.add(MaxPooling2D(pool_size=(2, 2)))  
  
    # Layer 3  
    alexnet.add(ZeroPadding2D((1, 1)))  
    alexnet.add(Conv2D(512, (3, 3), padding='same'))  
    alexnet.add(BatchNormalization())  
    alexnet.add(Activation('relu'))  
    alexnet.add(MaxPooling2D(pool_size=(2, 2)))  
  
    # Layer 4  
    alexnet.add(ZeroPadding2D((1, 1)))  
    alexnet.add(Conv2D(1024, (3, 3), padding='same'))  
    alexnet.add(BatchNormalization())  
    alexnet.add(Activation('relu'))
```

MODEL ARCHITECTURE - II

```
# Layer 5
alexnet.add(ZeroPadding2D((1, 1)))
alexnet.add(Conv2D(1024, (3, 3), padding='same'))
alexnet.add(BatchNormalization())
alexnet.add(Activation('relu'))
alexnet.add(MaxPooling2D(pool_size=(2, 2)))

# Layer 6
alexnet.add(Flatten())
alexnet.add(Dense(3072))
alexnet.add(BatchNormalization())
alexnet.add(Activation('relu'))
alexnet.add(Dropout(0.5))

# Layer 7
alexnet.add(Dense(4096))
alexnet.add(BatchNormalization())
alexnet.add(Activation('relu'))
alexnet.add(Dropout(0.5))

# Layer 8
alexnet.add(Dense(n_classes))
alexnet.add(BatchNormalization())
alexnet.add(Activation('softmax'))

if weights is not None:
    alexnet.load_weights(weights)

return alexnet
```

```
[54]: alexnet_model = alexnet_model(img_shape=(img_height, img_width, 3), n_classes=num_classes)
```

```
[55]: alexnet_model.compile(optimizer=keras.optimizers.Adam(learning_rate=0.0002, epsilon=0.05),
                           loss=tf.keras.losses.SparseCategoricalCrossentropy(),
                           metrics=['accuracy'])
```

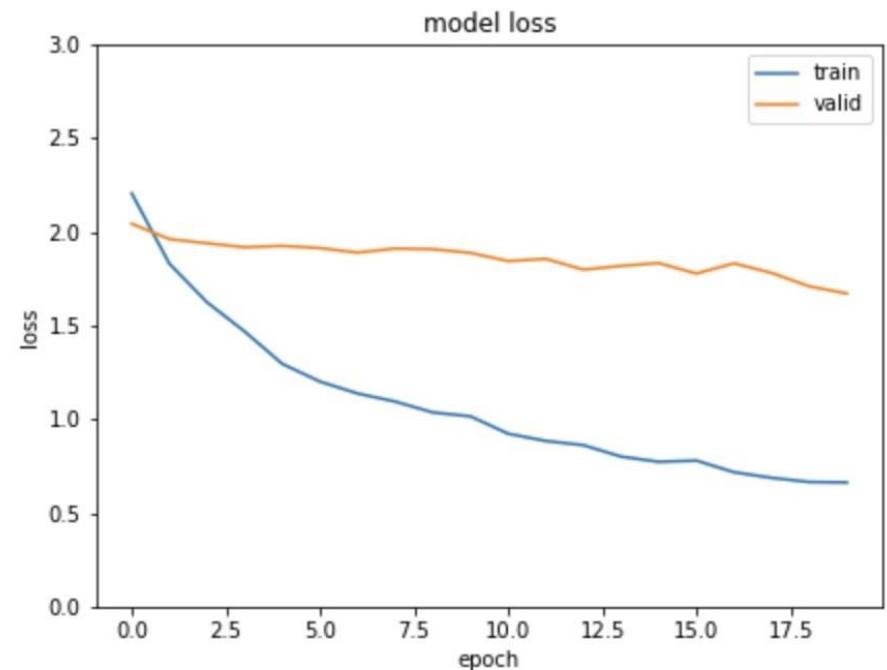
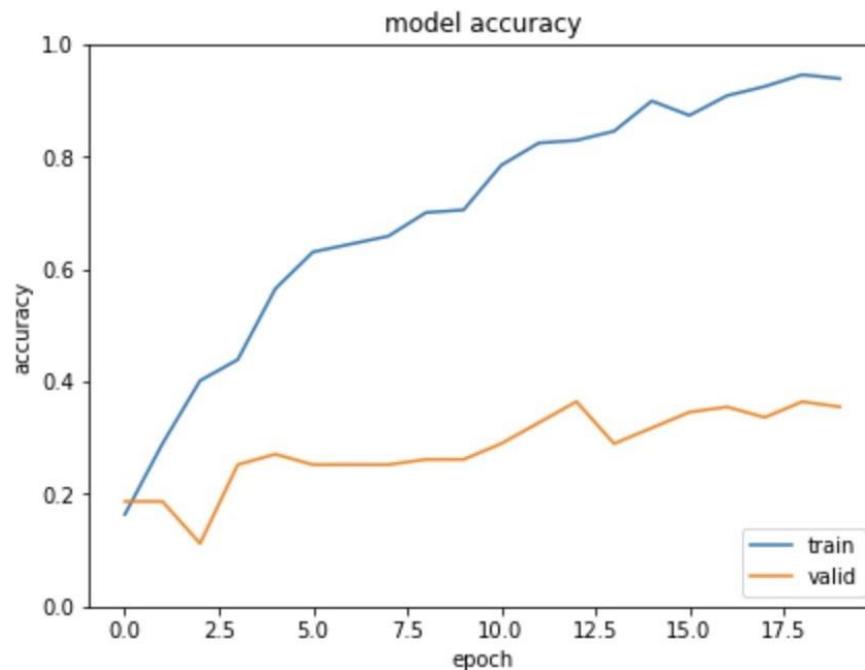
TRAINING AND SAVING

```
In [57]: history = alexnet_model.fit(train_images, train_labels, epochs=20, validation_data=(test_images, test_labels))
accuracy: 0.2897
Epoch 15/20
14/14 [=====] - 12s 840ms/step - loss: 0.7722 - accuracy: 0.8995 - val_loss: 1.8331 - val_
accuracy: 0.3178
Epoch 16/20
14/14 [=====] - 12s 871ms/step - loss: 0.7789 - accuracy: 0.8738 - val_loss: 1.7774 - val_
accuracy: 0.3458
Epoch 17/20
14/14 [=====] - 12s 865ms/step - loss: 0.7169 - accuracy: 0.9089 - val_loss: 1.8318 - val_
accuracy: 0.3551
Epoch 18/20
14/14 [=====] - 14s 1s/step - loss: 0.6867 - accuracy: 0.9252 - val_loss: 1.7806 - val_
accuracy: 0.3364
Epoch 19/20
14/14 [=====] - 14s 1s/step - loss: 0.6647 - accuracy: 0.9463 - val_loss: 1.7089 - val_
accuracy: 0.3645
Epoch 20/20
14/14 [=====] - 13s 959ms/step - loss: 0.6623 - accuracy: 0.9393 - val_loss: 1.6708 - val_
accuracy: 0.3551
```

```
In [59]: alexnet_model.save('alexnet.h5')
```

TRAINING AND LOSS CURVES

```
In [60]: save_train_loss_curves(history, 'alexnet_loss_acc.png')
```



PERFORMANCE

Confusion Matrix:

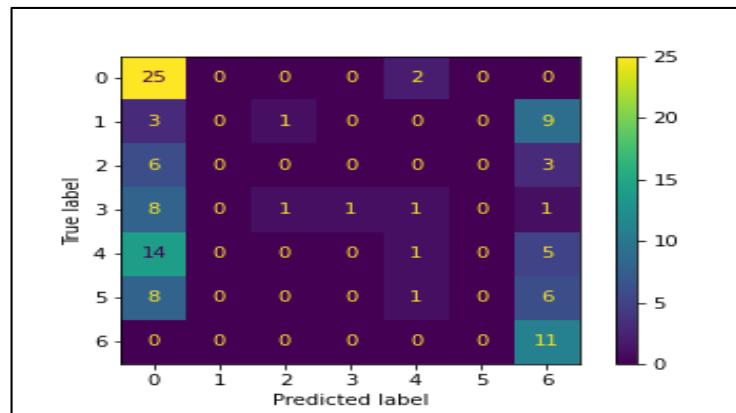
```
[[25  0  0  0  2  0  0]
 [ 3  0  1  0  0  0  9]
 [ 6  0  0  0  0  0  3]
 [ 8  0  1  1  1  0  1]
 [14  0  0  0  1  0  5]
 [ 8  0  0  0  1  0  6]
 [ 0  0  0  0  0  0 11]]
```

Performance Evaluation:

	precision	recall	f1-score	support
0	0.39	0.93	0.55	27
1	0.00	0.00	0.00	13
2	0.00	0.00	0.00	9
3	1.00	0.08	0.15	12
4	0.20	0.05	0.08	20
5	0.00	0.00	0.00	15
6	0.31	1.00	0.48	11
accuracy			0.36	107
macro avg	0.27	0.29	0.18	107
weighted avg	0.28	0.36	0.22	107

Accuracy Score:

0.35514018691588783



GOOGLENET



IMPORTS, DATA PREPROCESSING AND INCEPTION MODULE

GOOGLENET

```
In [63]: from tensorflow.keras.layers import Flatten, Input, Conv2D, Dense, Dropout, MaxPooling2D, AveragePooling2D, GlobalAvgPool1D
from tensorflow.keras.layers import concatenate
from tensorflow.keras.models import Model
from tensorflow.keras.regularizers import l2
```

```
In [64]: train_images_normalized = train_images/255.0
test_images_normalized = test_images/255.0
```

```
In [65]: def inception(x,
                  filters_1x1,
                  filters_3x3_reduce,
                  filters_3x3,
                  filters_5x5_reduce,
                  filters_5x5,
                  filters_pool):
    path1 = layers.Conv2D(filters_1x1, (1, 1), padding='same', activation='relu')(x)
    path2 = layers.Conv2D(filters_3x3_reduce, (1, 1), padding='same', activation='relu')(x)
    path2 = layers.Conv2D(filters_3x3, (1, 1), padding='same', activation='relu')(path2)
    path3 = layers.Conv2D(filters_5x5_reduce, (1, 1), padding='same', activation='relu')(x)
    path3 = layers.Conv2D(filters_5x5, (1, 1), padding='same', activation='relu')(path3)
    path4 = layers.MaxPool2D((3, 3), strides=(1, 1), padding='same')(x)
    path4 = layers.Conv2D(filters_pool, (1, 1), padding='same', activation='relu')(path4)
    return tf.concat([path1, path2, path3, path4], axis=3)
```

MODEL ARCHITECTURE

```
In [66]: inp = layers.Input(shape=(224, 224, 3))
x = layers.Conv2D(64, 7, strides=2, padding='same', activation='relu')(inp)
x = layers.MaxPooling2D(3, strides=2)(x)
x = layers.Conv2D(64, 1, strides=1, padding='same', activation='relu')(x)
x = layers.Conv2D(192, 3, strides=1, padding='same', activation='relu')(x)
x = layers.MaxPooling2D(3, strides=2)(x)
x = inception(x, filters_1x1=64, filters_3x3_reduce=96, filters_3x3=128, filters_5x5_reduce=16, filters_5x5=32, filters_5x5=32)
x = inception(x, filters_1x1=128, filters_3x3_reduce=128, filters_3x3=192, filters_5x5_reduce=32, filters_5x5=96, filters_5x5=96)
x = layers.MaxPooling2D(3, strides=2)(x)
x = inception(x, filters_1x1=192, filters_3x3_reduce=96, filters_3x3=208, filters_5x5_reduce=16, filters_5x5=48, filters_5x5=48)
aux1 = layers.AveragePooling2D((5, 5), strides=3)(x)
aux1 = layers.Conv2D(128, 1, padding='same', activation='relu')(aux1)
aux1 = layers.Flatten()(aux1)
aux1 = layers.Dense(1024, activation='relu')(aux1)
aux1 = layers.Dropout(0.7)(aux1)
aux1 = layers.Dense(10, activation='softmax')(aux1)
x = inception(x, filters_1x1=160, filters_3x3_reduce=112, filters_3x3=224, filters_5x5_reduce=24, filters_5x5=64, filters_5x5=64)
x = inception(x, filters_1x1=128, filters_3x3_reduce=128, filters_3x3=256, filters_5x5_reduce=24, filters_5x5=64, filters_5x5=64)
x = inception(x, filters_1x1=112, filters_3x3_reduce=144, filters_3x3=288, filters_5x5_reduce=32, filters_5x5=64, filters_5x5=64)
aux2 = layers.AveragePooling2D((5, 5), strides=3)(x)
aux2 = layers.Conv2D(128, 1, padding='same', activation='relu')(aux2)
aux2 = layers.Flatten()(aux2)
aux2 = layers.Dense(1024, activation='relu')(aux2)
aux2 = layers.Dropout(0.7)(aux2)
aux2 = layers.Dense(10, activation='softmax')(aux2)
x = inception(x, filters_1x1=256, filters_3x3_reduce=160, filters_3x3=320, filters_5x5_reduce=32, filters_5x5=128, filters_5x5=128)
x = layers.MaxPooling2D(3, strides=2)(x)
x = inception(x, filters_1x1=256, filters_3x3_reduce=160, filters_3x3=320, filters_5x5_reduce=32, filters_5x5=128, filters_5x5=128)
x = inception(x, filters_1x1=384, filters_3x3_reduce=192, filters_3x3=384, filters_5x5_reduce=48, filters_5x5=128, filters_5x5=128)
x = layers.GlobalAveragePooling2D()(x)
x = layers.Dropout(0.4)(x)
out = layers.Dense(10, activation='softmax')(x)
```

```
In [67]: googlenet_model = Model(inputs = inp, outputs = [out, aux1, aux2])
```

COMPILATION

```
In [67]: googlenet_model = Model(inputs = inp, outputs = [out, aux1, aux2])
```

```
In [68]: googlenet_model.compile(optimizer=keras.optimizers.Adam(learning_rate=0.01, epsilon=0.05),
                               loss=[keras.losses.sparse_categorical_crossentropy,
                                     keras.losses.sparse_categorical_crossentropy,
                                     keras.losses.sparse_categorical_crossentropy],
                               loss_weights=[1, 0.3, 0.3],
                               metrics=['accuracy'])
```

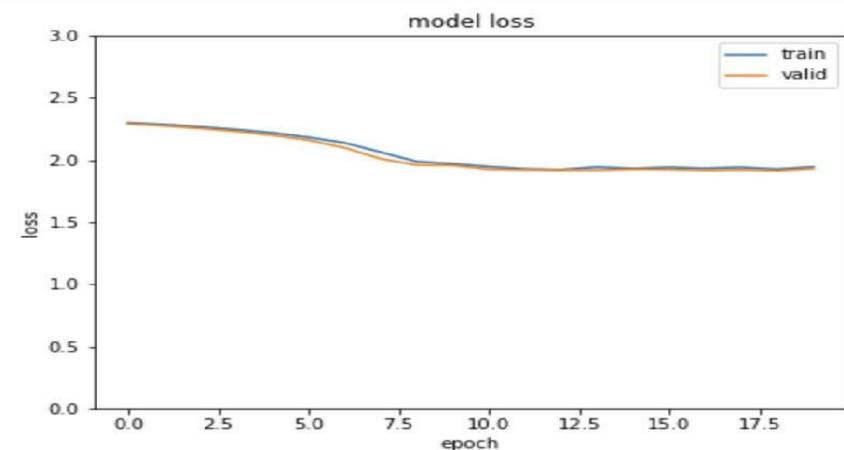
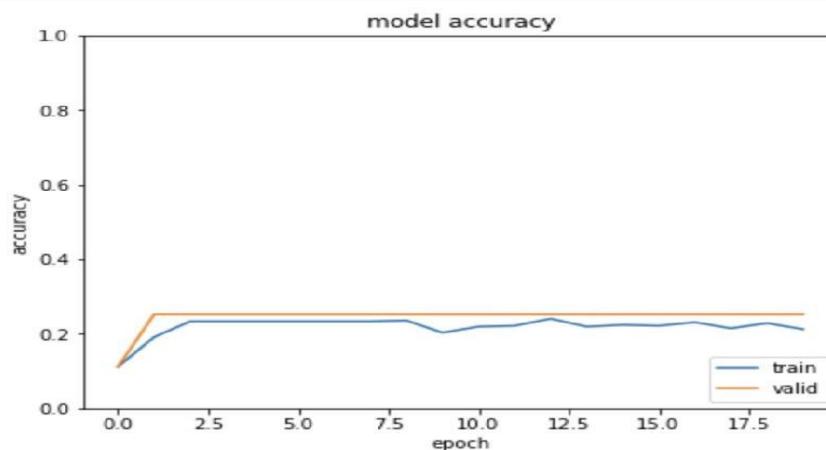
TRAINING

```
In [69]: history = googlenet_model.fit(train_images_normalized, [train_labels, train_labels, train_labels], \
                                     validation_data=(test_images_normalized, [test_labels, test_labels, test_labels]), \
                                     batch_size=64, epochs=20)

    7/7 [=====] - 10s 1s/step - loss: 3.1050 - dense_17_loss: 1.9521 - dense_14_loss: 1.9409
dense_16_loss: 1.9627 - dense_17_accuracy: 0.2313 - dense_14_accuracy: 0.2430 - dense_16_accuracy: 0.2360 - val_lo
s: 3.0746 - val_dense_17_loss: 1.9167 - val_dense_14_loss: 1.9268 - val_dense_16_loss: 1.9331 - val_dense_17_accur
cy: 0.2523 - val_dense_14_accuracy: 0.2523 - val_dense_16_accuracy: 0.2523
Epoch 18/20
    7/7 [=====] - 10s 1s/step - loss: 3.1114 - dense_17_loss: 1.9405 - dense_14_loss: 1.9557
dense_16_loss: 1.9471 - dense_17_accuracy: 0.2150 - dense_14_accuracy: 0.2430 - dense_16_accuracy: 0.2407 - val_lo
s: 3.0775 - val_dense_17_loss: 1.9196 - val_dense_14_loss: 1.9239 - val_dense_16_loss: 1.9359 - val_dense_17_accur
cy: 0.2523 - val_dense_14_accuracy: 0.2523 - val_dense_16_accuracy: 0.2523
Epoch 19/20
    7/7 [=====] - 10s 1s/step - loss: 3.1000 - dense_17_loss: 1.9256 - dense_14_loss: 1.9555
dense_16_loss: 1.9588 - dense_17_accuracy: 0.2290 - dense_14_accuracy: 0.2103 - dense_16_accuracy: 0.2033 - val_lo
s: 3.0635 - val_dense_17_loss: 1.9141 - val_dense_14_loss: 1.9109 - val_dense_16_loss: 1.9202 - val_dense_17_accur
cy: 0.2523 - val_dense_14_accuracy: 0.2523 - val_dense_16_accuracy: 0.2523
Epoch 20/20
    7/7 [=====] - 11s 2s/step - loss: 3.1215 - dense_17_loss: 1.9437 - dense_14_loss: 1.9741
dense_16_loss: 1.9520 - dense_17_accuracy: 0.2126 - dense_14_accuracy: 0.2196 - dense_16_accuracy: 0.2103 - val_lo
s: 3.0879 - val_dense_17_loss: 1.9301 - val_dense_14_loss: 1.9298 - val_dense_16_loss: 1.9295 - val_dense_17_accur
cy: 0.2523 - val_dense_14_accuracy: 0.2523 - val_dense_16_accuracy: 0.2523
```

TRAINING AND LOSS CURVES

```
In [72]: plt.figure(figsize=(15,5))
plt.subplot(1,2,1)
plt.plot(history.history['dense_17_accuracy'])
plt.plot(history.history['val_dense_17_accuracy'])
plt.ylim(0, 1)
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'valid'], loc='lower right')
plt.subplot(1,2,2)
plt.plot(history.history['dense_17_loss'])
plt.plot(history.history['val_dense_17_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'valid'], loc='upper right')
plt.ylim([0,3])
plt.savefig("googlenet_loss_acc.png")
```



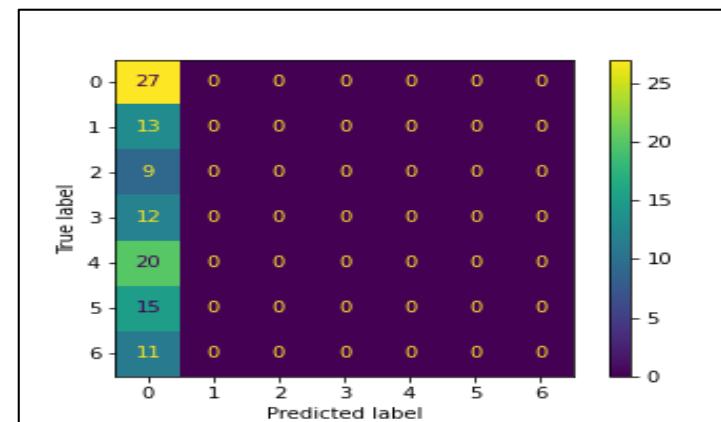
PERFORMANCE

Confusion Matrix:

```
[[27  0  0  0  0  0  0]
 [13  0  0  0  0  0  0]
 [ 9  0  0  0  0  0  0]
 [12  0  0  0  0  0  0]
 [20  0  0  0  0  0  0]
 [15  0  0  0  0  0  0]
 [11  0  0  0  0  0  0]]
```

Performance Evaluation:

	precision	recall	f1-score	support
0	0.25	1.00	0.40	27
1	0.00	0.00	0.00	13
2	0.00	0.00	0.00	9
3	0.00	0.00	0.00	12
4	0.00	0.00	0.00	20
5	0.00	0.00	0.00	15
6	0.00	0.00	0.00	11
accuracy			0.25	107
macro avg	0.04	0.14	0.06	107
weighted avg	0.06	0.25	0.10	107



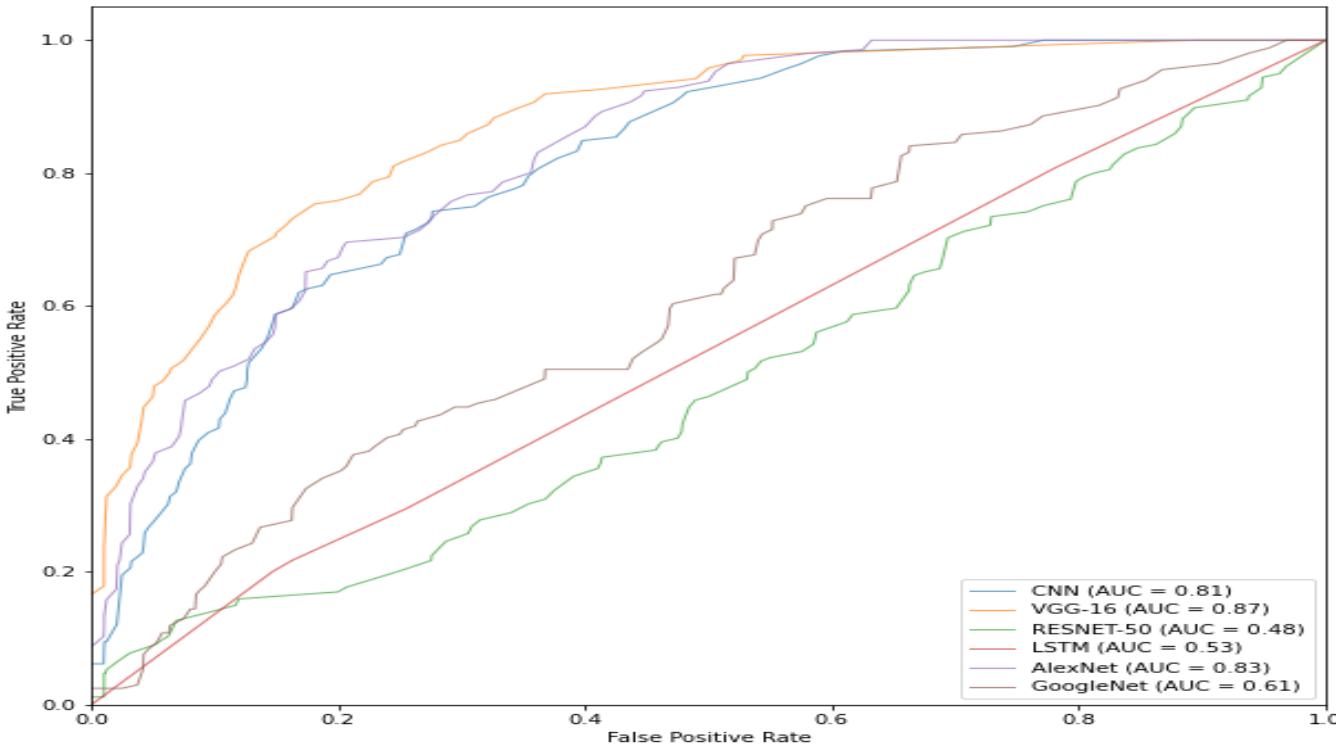
Accuracy Score:

0.2523364485981308

COMPARISON BETWEEN CNN, VGG-16, RESNET-50, LSTM(RNN), ALEXNET AND GOOGLENET.



ROC CURVES AND AUC COMPARISON



ROC CURVES AND AUC COMPARISON - CODE (I)

```
In [77]: def plot_macro_avg_roc_curve(classifier, n_classes, name, X_test, y_test, single_output=True):
    y_test_binarized = label_binarize(y_test, classes=list(range(0,n_classes)))
    y_score = classifier.predict(X_test)

    #for cases where the model (like googlenet) has extra (auxillary) outputs
    if not single_output:
        y_score = y_score[0]

    # Compute ROC curve and ROC area for each class
    fpr = dict()
    tpr = dict()
    roc_auc = dict()
    for i in range(n_classes):
        fpr[i], tpr[i], _ = roc_curve(y_test_binarized[:, i], y_score[:, i])
        roc_auc[i] = auc(fpr[i], tpr[i])

    # First aggregate all false positive rates
    all_fpr = np.unique(np.concatenate([fpr[i] for i in range(n_classes)]))

    # Then interpolate all ROC curves at this points
    mean_tpr = np.zeros_like(all_fpr)
    for i in range(n_classes):
        mean_tpr += np.interp(all_fpr, fpr[i], tpr[i])

    # Finally average it and compute AUC
    mean_tpr /= n_classes

    fpr["macro"] = all_fpr
    tpr["macro"] = mean_tpr
    roc_auc["macro"] = auc(fpr["macro"], tpr["macro"])

    plt.plot(
        fpr["macro"],
        tpr["macro"],
        label="{0} (AUC = {1:0.2f})".format(name, roc_auc["macro"]),
        alpha=0.7,
        linewidth=1,
    )

    return plt
```

ROC CURVES AND AUC COMPARISON - CODE (II)

```
In [78]:
```

```
cnn = load_model('cnn.h5')
vgg_based_model = load_model('vgg16.h5')
resnet_based_model = load_model('resnet50.h5')
lstm_based_model = load_model('lstm.h5')
alexnet_model = load_model('alexnet.h5')
googlenet_model = load_model('googlenet.h5')
```

```
In [79]:
```

```
plt.figure(figsize=(10, 10))
plot_macro_avg_roc_curve(cnn, num_classes, "CNN", test_images, test_labels)
plot_macro_avg_roc_curve(vgg_based_model, num_classes, "VGG-16", test_images_p, test_labels)
plot_macro_avg_roc_curve(resnet_based_model, num_classes, "RESNET-50", test_images_pr, test_labels)
plot_macro_avg_roc_curve(lstm_based_model, num_classes, "LSTM", test_images, test_labels)
plot_macro_avg_roc_curve(alexnet_model, num_classes, "AlexNet", test_images, test_labels)
plot_macro_avg_roc_curve(googlenet_model, num_classes, "GoogleNet", test_images_normalized, test_labels, single_outp

plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.legend(loc="lower right")

plt.savefig('roc_auc.png')
```

COMPARISON OF PERFORMANCE

CLASSIFIER	Accuracy	Precision	Recall	F1-Score	AUC
CNN	0.43	0.68	0.43	0.38	0.81
VGG-16	0.49	0.57	0.49	0.49	0.87
RESNET-50	0.12	0.01	0.12	0.03	0.48
LSTM	0.25	0.06	0.25	0.10	0.53
ALEXNET	0.36	0.28	0.36	0.22	0.83
GOOGLENET	0.25	0.06	0.25	0.10	0.61

CONCLUSION:

VGG-16 performed better than the other models. AlexNet and the simple CNN did respectable jobs. However, all the models performed poorly. It appears that it is hard to generalise patterns as all humans convey emotions differently.