


# Asynchronous Shared-Memory Systems

# Asynchronous Shared-Memory systems

- We've covered basics of non-fault-tolerant asynchronous network algorithms:
  - How to model them.
  - Basic asynchronous network protocols---broadcast, spanning trees, leader election,...
  - Synchronizers (running synchronous algorithms in asynch networks)
  - Logical time
  - Global snapshots
- Now consider asynchronous shared-memory systems:

- Processes, interacting via shared objects, possibly subject to some access constraints.
- Shared objects are typed, e.g.:
  - Read/write (weak)
  - Read-modify-write, compare-and-swap (strong)
  - Queues, stacks, others (in between)



## Asynch Shared-Memory systems

- Theory of ASM systems has much in common with theory of asynchronous networks:
  - Similar algorithms and impossibility results.
  - Even with failures.
  - Transformations from ASM model to asynch network model allow ASM algorithms to run in asynchronous networks. • “Distributed Shared Memory”.
- Historically, theory for ASM started first.

- Arose in study of early operating systems, in which several processes can run on a single processor, sharing memory, with possibly-arbitrary interleavings of steps.
- Currently, ASM models apply to multiprocessor shared-memory systems, in which processes can run on separate processors and share memory.


## Topics

- Define the **basic system model, without failures**.
- Use it to study basic problems:
  - Mutual exclusion.
  - Other resource-allocation problems.
- **Introduce process failures** into the model.
- Use model with failures to study basic problems:
  - Distributed consensus
  - Implementing atomic objects:



- Atomic snapshot objects
- Atomic read/write registers
- Wait-free and fault-tolerant computability theory
- Modern shared-memory multiprocessors:
  - Practical issues
  - Algorithms
  - Transactional memory

## Basic ASM Model, Version 1

- Processes + objects, modeled as automata.
- Arrows:
  - Represent invocations and responses for operations on the objects.
  - Modeled as input and output actions.
- Fine-granularity model, can describe:
  - Delay between invocation and response. `invoke(read)`
  - Concurrent (overlapping) operations:  $p_1$




- Object could reorder.  $x_1$
- Could allow them to run concurrently, interfering with  $\text{respond}(v)$  each other.
- We'll begin with a simpler, coarser model:  
 $\text{invoke}(\text{write}, v)$ 
  - Object runs ops in invocation order, one at a time.—  
 In fact, collapse each operation into a single step.
- Return to the finer model later.



# Basic ASM Model, Version 2

- One big shared memory system automaton A.
- External actions at process “ports”.
- Each process  $i$  has:
  - A set **states<sub>i</sub>** of states.
  - A subset **start<sub>i</sub>** of start states.
- Each variable  $x$  has:
  - A set **values<sub>x</sub>** of values it can take on.
  - A subset **initial<sub>x</sub>** of initial values.
- Automaton A:
  - **States**: State for each process, a value for each variable.
  - **Start**: Start states, initial values.
  - **Actions**: Each action associated with one process, and some also with a single shared variable.
  - **Input/output actions**: At the external boundary.
  - **Transitions**: Correspond to local process steps and variable accesses.
- Action enabling, which variable is accessed, depend only on process state.
- Changes to variable and process state depend also on variable value.
- Must respect the type of the variable.




- **Tasks:** One or more per process (threads).

# Basic ASM Model

- **Execution of A:**

- By IOA fairness definition, each task gets infinitely many chances to take steps.
- Model environment as a separate automaton, to express restrictions on environment behavior.



- **Commonly-used variable types:**

- Read/write registers: Most basic primitive.
- Allows access using separate read and write operations.
  - Read-modify-write: More powerful primitive:
- Atomically, read variable, do local computation, write to variable. – Compare-and-swap, fetch-and-add, queues, stacks, etc.




- Different computability and complexity results hold for different variable types.


# The Mutual Exclusion Problem

- Share one resource among  $n$  user processes,  $U_1, U_2, \dots, U_n$ .
  - E.g., printer, portion of a database.
- $U_i$  has four “regions”.
  - Subsets of its states, described by portions of its code.
  - C critical; R remainder; T trying; E exit

Protocols for obtaining and relinquishing the resource




- Cycle:
- Architecture:
  - $U_i$ s and  $A$  are IOAs, compose.




# The Mutual Exclusion Problem

- Actions at user interface:
  - $\text{try}_i, \text{crit}_i, \text{exit}_i, \text{rem}_i$      $i$
  - $U_i$  interacts with  $p_i$
- Correctness conditions:
  - **Well-formedness** (Safety property):
- System obeys cyclic discipline.  <sup>$i$</sup>
- E.g., doesn't grant resource when it wasn't requested.
  - **Mutual exclusion** (Safety):




# The Mutual Exclusion Problem

- System never grants to  $> 1$  user simultaneously.
- Trace safety property.
- Or, there's no reachable system state in which  $> 1$  user is in C at once. – **Progress** (Liveness):
- From any point in a fair execution:
  - If some user is in T and no user is in C then at some later point, some user enters C.
  - If some user is in E then at some later point, some user enters R.



# The Mutual Exclusion Problem

- Well-formedness (Safety):
  - System obeys cyclic discipline.
- Mutual exclusion (Safety):– System never grants to  $> 1$  user.
- Progress (Liveness):
  - From any point in a fair execution:
- If some user is in T and no user is in C then at some later point, some user enters C.
- If some user is in E then at some later point, some user enters R. •




Conditions all constrain the system automaton  $A$ , not users.

- System determines if/when users enter C and R.
- Users determine if/when users enter T and E.

# The Mutual Exclusion Problem

- We don't state any requirements on the users, except that they preserve well-formedness.
- Well-formedness (Safety):
- Mutual exclusion (Safety):
- Progress (Liveness):
  - From any point in a fair execution:
- If some user is in T and no user is in C then at some later point, some user enters C.
- If some user is in E then at some later point, some user enters R.
- **Fairness assumption:**
  - Progress condition requires fairness assumption (all process tasks continue to get turns to take steps).
  - Needed to guarantee that some process enters C or R.



# The Mutual Exclusion Problem

- In general, in the asynchronous model, liveness properties require fairness assumptions.
- Contrast: Well-formedness and mutual exclusion are safety properties, don't depend on fairness.

# One more assumption...

- No permanently active processes.
  - Locally-controlled actions can be enabled only when user is in T or E.
  - No always-awake, dedicated processes.
  - Motivation:
- Multiprocessor settings, where users can run processes at any time, but are otherwise not involved in the protocol.
- Avoid “wasting processors”.

## Mutual Exclusion algorithm



# [Dijkstra 65]

- Based on Dekker's 2-process solution.
- Pseudocode, p. 265-266
  - Written in traditional sequential style, must somehow translate into more detailed state/transition description.
  - Shared variables: Read/write registers.
    - **turn**, in  $\{1,2,\dots,n\}$ , multi-writer multi-reader (mWmR), init anything.
    - for each process  $i$ :
      - **flag(i)**, in  $\{0,1,2\}$ , single-writer multi-reader (1WmR), init 0
      - Written by  $i$ , read by everyone.
  - Process  $i$ 's Stage 1:
    - Set **flag** := 1, test to see if **turn** =  $i$ .
    - If not, and **turn**'s current owner is seen to be inactive, then set **turn** :=  $i$ .

- Otherwise go back to to testing...
- When you see `turn = i`, move to Stage 2.

## Dijkstra's algorithm

- Stage 2:
  - Set `flag(i) := 2`.
  - Check (one at a time, any order) that no other process has `flag = 2`.
  - If check completes successfully, go to C. – If not, go back to beginning of Stage 1.
- Exit protocol:– Set `flag(i) := 0`.
- Problem with the sequential code style:
  - Unclear what constitutes an atomic step.
- E.g., need three separate steps to test `turn`, test `flag(turn)`, and set `turn`.
  - Must rewrite to make this clear:
- E.g., precondition/effect code (p. 268-269)

- E.g., sequential-style code with explicit reads and writes, one per line.

## Dijkstra's algorithm, pre/eff code

- One transition definition for each kind of atomic step.
- Explicit program counter, `pc`.
- E.g.: When `pc` is:
  - `set-flag-1i`: Sets `flag` to 1 and prepares to test `turn`.
  - `test-turni`: Tests `turn`, and either moves to Stage 2 or prepares to test the current owner's `flag`.
  - `test-flag(j)i`: Tests `j`'s `flag`, and either goes on to set `turn` or goes back to test `turn` again.
  - ...
  - `set-flag-2i`: Sets `flag` to 2 and initializes set `S`, preparing to check all other processes' `flags`.
  - `check(j)i`: If `flag(j) = 2`, go back to beginning.
  - ...

- **S** keeps track of which processes have been successfully checked in Stage 2.

## Precondition/effect code

### Shared variables:

**turn**  $\in \{1, \dots, n\}$ , initially arbitrary

for every  $i$ : **flag**( $i$ )  $\in \{0, 1, 2\}$ ,

initially 0

### Actions of process $i$ :

Input: **try** <sub>$i$</sub> , **exit** <sub>$i$</sub>

Output: **crit** <sub>$i$</sub> , **rem** <sub>$i$</sub>

Internal:  $\text{set-flag-1}_i, \text{test-turn}_i, \text{test-flag}(j)_i, \text{set-turn}_i, \text{set-flag-2}_i, \text{check}(j)_i, \text{reset}_i$

## Precondition/effect code, Dijkstra process $i$

$\text{try}_i$ : Eff:  $\text{pc} := \text{set-flag-1}$

$\text{set-flag-1}_i$ :  
Pre:  $\text{pc} = \text{set-flag-1}$   
Eff:  $\text{flag}(i) := 1$   
 $\text{pc} := \text{test-turn}$

$\text{test-turn}_i$ :

Pre:  $\text{pc} = \text{test-turn}$

Eff: if  $\text{turn} = i$  then  $\text{pc} := \text{set-flag-2}$   
else  $\text{pc} := \text{test-flag}(\text{turn})$

$\text{test-flag}(j)_i$

Pre:  $\text{pc} = \text{test-flag}(j)$

Eff: if  $\text{flag}(j) = 0$  then  $\text{pc} := \text{set-turn}$   
else  $\text{pc} := \text{test-turn}$

$\text{set-turn}_i$ :

Pre:  $\text{pc} = \text{set-turn}$

Eff:  $\text{turn} := i$   $\text{pc}$   
       $:= \text{set-flag-2}$

Pre:  $\text{pc} = \text{set-flag-2}$   
Eff:  $\text{flag}(i) :=$   
      2  $S := \{i\}$   $\text{pc}$   
       $:= \text{check}$

$\text{set-flag-2}_i :$

## More precondition/effect code, Dijkstra process i

$\text{check}(j)_i :$

Pre:  $\text{pc} = \text{check}$     Eff:  $\text{pc} := \text{reset}$   $j \notin S$

Eff: if  $\text{flag}(j) = 2$  then

$S := \emptyset$

$\text{pc} := \text{set-flag-1}$

$\text{exit}_i$

$\text{reset}_i :$

Pre:  $\text{pc} = \text{reset}$

Eff:  $\text{flag}(i) := 0$

$S := \emptyset$  else  $pc$

$:= \text{leave-exit } S := S \cup \{j\}$   
if  $|S| = n$  then  $pc := \text{leave-try}$

$\text{rem}_i :$

$\text{crit}_i :$

Eff: Pre:  $pcpc := \text{rem leave-exit}$

Pre:  $pc = \text{leave-try}$

Eff:  $pc := \text{crit}$

## Note on code style

- Explicit  $pc$  makes atomicity clear, but looks somewhat verbose/awkward.
- $pc$  is often needed in invariants.

- Alternatively: Use sequential style, with explicit reads or writes (or other operations), one per line.
- Need line numbers:
  - Play same role as **pc**.
  - Used in invariants: “If process  $i$  is at line 7 then...”

## Correctness

- **Well-formedness:** Obvious.
- **Mutual exclusion:**
  - Based on event order in executions, rather than invariants.
  - By contradiction: Assume  $U_i, U_j$  are ever in  $u_i, u_j$  in  $C$  at the same time.

Initial  
state

in  $C$

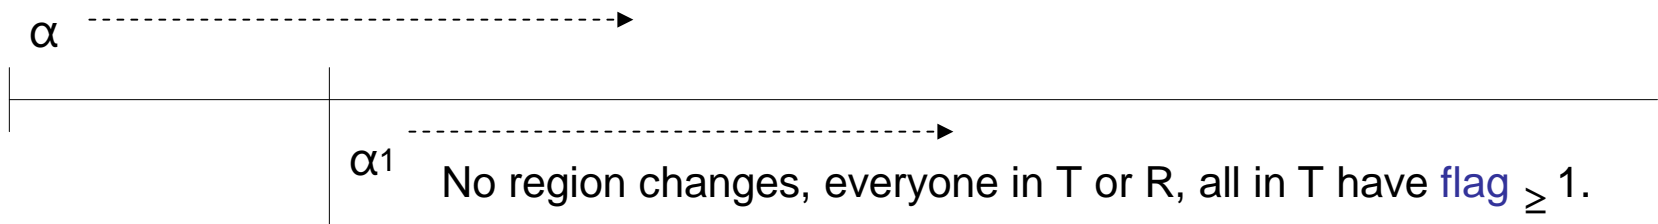


- Both must set-flag-2 before entering C; Initial consider the last time they do this.<sub>state</sub>
- WLOG, suppose set-flag-2<sub>i</sub> comes first. set-flag-2<sub>i</sub> — Then  
 $\text{flag}(i) = 2$  from that point onward (until set-flag-2 they are both in C).<sub>j</sub>
- However, j must see  $\text{flag}(i) \neq 2$ , in order to j sees  $\text{flag}(i) \neq 2$   
enter C.<sub>U<sub>i</sub>, U<sub>j</sub></sub>
- Impossible. in C

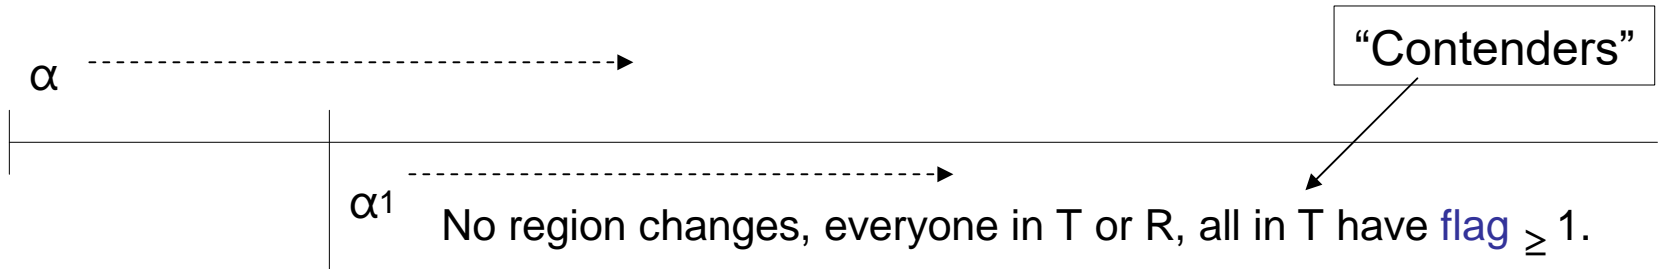
## Progress

- Interesting case: Trying region.
- Proof by contradiction:

- Suppose  $\alpha$  is a fair execution, reaches a point where some process is in T, no process is in C, and thereafter, no process ever enters C.
- Now start removing complications...
- Eventually, all regions changes stop and all in T keep their **flags**  $\geq 1$ .
- Then it must be that everyone is in T and R, and all in T have **flag**  $\geq 1$ .



# Progress, cont'd

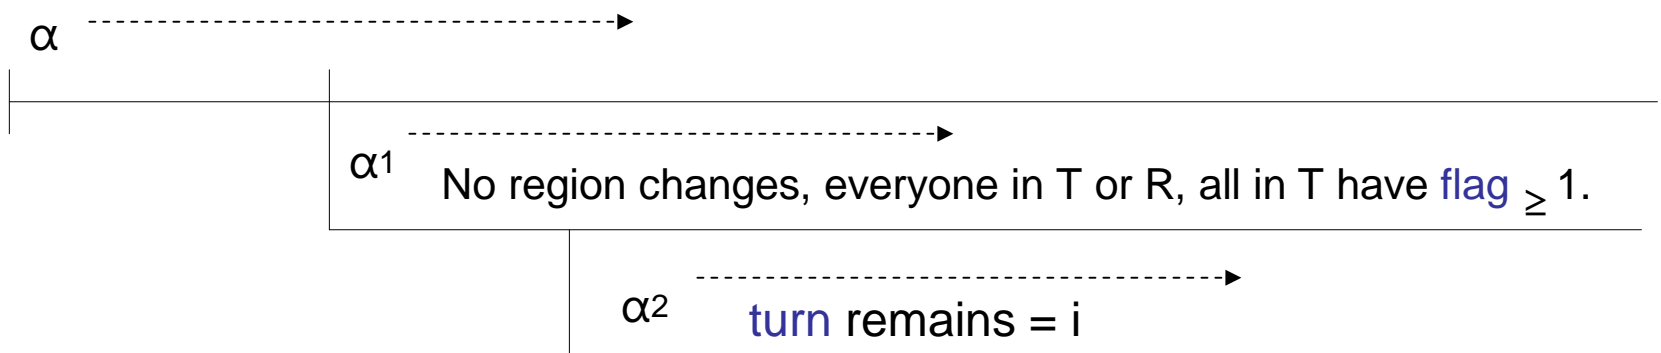


- Then whenever **turn** is reset in  $\alpha_1$ , it must be set to a contender's index.
- **Claim:** In  $\alpha_1$ , **turn** eventually acquires a contender's index.
- **Proof:**
  - Suppose not---stays non-contender forever.
  - Consider any contender  $i$ .
  - If it ever reaches test-turn, then it will set  $\text{turn} := i$ , since it sees an inactive process.
  - Why must process  $i$  reach test-turn?
- It's either that, or it succeeds in reaching C.

- But we have assumed no one reaches C.
  - Contradiction.

## Progress, cont'd

- In  $\alpha_1$ , once **turn** = contender's index, it is thereafter always = some contender's index. – Because contenders are the only processes that can change **turn**.
- May change several times.
- Eventually, **turn** stops changing (because tests come out negative), stabilizes to some value, say  $i$ .



- Thereafter, all contenders  $\neq i$  wind up looping in Stage 1.
  - If  $j$  reaches Stage 2, it returns to Stage 1, since it doesn't go to C.
  - But then  $j$ 's tests always fail, so  $j$  stays in Stage 1.
- But then nothing stops process  $i$  from entering C.

## Mutual exclusion, Proof 2

- Use invariants.
- Must show they hold after any number of steps.
- Main goal invariant:  $|\{i : pc_i = \text{crit}\}| \leq 1$ . • To prove by

induction, need more:

1. If  $pc_i = \text{crit}$  (or leave-try or reset) then  $|S_i| = n$ .
2. There do not exist  $i, j$ ,  $i \neq j$ , with  $i$  in  $S_j$  and  $j$  in  $S_i$ .

- 1 and 2 easily imply mutual exclusion.
- Proof of 1: Easy induction
- Proof of 2:
  - Needs some easy auxiliary invariants saying what **S**-values go with what **flag** values and what **pc** values.
  - Key step: When  $j$  gets added to  $S_i$ , by  $\text{check}(j)_i$  event.
- Then must have  $\text{flag}(j) \neq 2$ .
- But then  $S_j = \emptyset$  (by auxiliary invariant), so  $i \notin S_j$ , can't break invariant.

## Running Time

- Upper bound on time from when some process is in  $T$  until some process is in  $C$ .

- Assume upper bound of  $l$  on successive turns for each process task (here, all steps of each process are in one task).
- Time upper bound for [Dijkstra]:  $O(l n)$ .
- Proof: LTTR

## Adding fairness guarantees [Peterson]

- Dijkstra algorithm does not guarantee fairness in granting the resource to different users.
- Might not be important in practice, if contention is rare.
- Other algorithms add fairness guarantees.

- E.g., [Peterson]: a collection of algorithms guaranteeing lockout-freedom.
- Lockout-freedom: In any (low-level) fair execution:
  - If all users always return the resource then any user that enters T eventually enters C.
  - Any user that enters E eventually enters R.

## Peterson 2-process algorithm

- Shared variables:
  - `turn`, in  $\{0,1\}$ , 2W2R read/write register, initially arbitrary.
  - for each process  $i = 0,1$ :
    - `flag(i)`, in  $\{0,1\}$ , 1W1R register, initially 0
    - Written by  $i$ , read by  $1-i$ .
- Process  $i$ 's trying protocol:
  - Sets `flag(i) := 1`, sets `turn := i`.



- Waits for either  $\text{flag}(1-i) = 0$  or  $\text{turn} \neq i$ .

$\underbrace{\text{flag}(1-i) = 0}_{\text{Other process not active.}} \quad \underbrace{\text{turn} \neq i}_{\text{Other process has the turn variable.}}$

- Toggles between the two tests.
- Exit protocol:
  - Sets  $\text{flag}(i) := 0$

## Precondition/effect code

### Shared variables:

$\text{turn} \in \{0,1\}$ , initially arbitrary

for every  $i \in \{0,1\}$ :  $\text{flag}(i) \in \{0,1\}$ , initially 0

## Actions of process i:

Input:  $try_i$ ,  $exit_i$

Output:  $crit_i$ ,  $rem_i$

Internal:  $set-flag_i$ ,  $set-turn_i$ ,  $check-flag_i$ ,  $check-turn_i$ ,  $reset_i$

## Precondition/effect code, Peterson 2P, process i

$try_i$ :

Eff:  $pc := set-flag$

$set-flag_i$ :

Pre:  $pc = set-flag$

Eff:  $flag(i) := 1$

$pc := set-turn$

$set-turn_i$ :

Pre:  $pc = set-turn$

Eff:  $turn := i$   $pc := check-flag$

$check-flag_i$ :

Pre:  $pc = check-flag$

Eff: if  $flag(1-i) = 0$  then  $pc := leave-$   
try else  $pc := check-turn$

**check-turn<sub>i</sub> :**

Pre: **pc** = check-  
turn

Eff: if **turn** ≠ i  
then **pc** :=  
leave-try else  
**pc** := check-  
flag

Pre: **pc** = leave-try Eff: **pc** := crit

**exit<sub>i</sub>** Eff: **pc** := reset

**reset<sub>i</sub> :**

Pre: **pc** = reset  
Eff: **flag(i)** := 0 **pc** :=  
leave-exit

More precondition/effect co

Peterson 2P, process i

**rem<sub>i</sub> :**

Pre: **pc**  
= leave-

**crit<sub>i</sub> :**

exit

Eff: **pc** := rem

# Correctness: Mutual exclusion

- **Key invariant:**
  - If  $pc_i \in \{\text{leave-try, crit, reset}\}$  (essentially in C), and
  - $pc_{1-i} \in \{\text{check-flag, check-turn, leave-try, crit, reset}\}$  (engaged in the competition or in C), – then  $turn \neq i$ .
- **That is:**
  - If i has won and 1-i is currently competing then  $turn$  is set favorably for i---which means it is set to 1-i.
- **Implies mutual exclusion:** If both are in C then  $turn$  must be set both ways, contradiction.
- **Proof of invariant:** All cases of inductive step are easy.
  - E.g.: a successful  $check\text{-}turn_i$ , causing i to advance to leave-try.

- This explicitly checks that `turn`  $\neq i$ , as needed.

## Correctness: Progress

- By contradiction:
  - Suppose someone is in T, and no one is ever thereafter in C.
  - Then the execution eventually stabilizes so no new region changes occur.
  - After stabilization:
- If exactly one process is in T, then it sees the other's `flag` = 0 and enters C.

- If both processes are in T, then **turn** is set favorably to one of them, and it enters C.

## Correctness: Lockout-freedom

- Argue that neither process can enter C three times while the other stays in T, after setting its **flag** := 1.
- **Bounded bypass.**
- **Proof:** By contradiction.
  - Suppose process i is in T and has set **flag** := 1, and subsequently process (1-i) enters C three times.
  - In each of the second and third times through T, process (1-i) sets **turn** := 1-i but later sees **turn** = i.
  - That means process i must set **turn** := i at least twice during that time.

- But process  $i$  sets  $\text{turn} := i$  only once during its one execution of  $T$ .
  - Contradiction.
- Bounded bypass + progress imply lockout-freedom.

## Time complexity

- Time from when any particular process  $i$  enters  $T$  until it enters  $C$ :  $c + O(l)$ , where:
  - $c$  is an upper bound on the time any user remains in the critical section, and
  - $l$  is an upper bound on local process step time.
- Detailed proof: **See book.**
- Rough idea:

- Either process  $i$  can enter immediately, or else it has to wait for  $(1-i)$ .
- But in that case, it only has to wait for one critical-section time, since if  $(1-i)$  reenters, it will set turn favorably for  $i$ .

## Peterson n-process algorithms

- Extend 2-process algorithm for lockout-free mutual exclusion to n-process algorithm, in two ways:
  - Using linear sequence of competitions, or –
  - Using binary tree of competitions.



# Sequence of competitions

- Competitions  $1, 2, \dots, n-1$ .
  - Competition  $k$  has one loser, up to  $n-k$  winners.
  - Thus, only one can win in competition  $n-1$ , implying mutual exclusion.
- 
- Shared vars:
    - For each competition  $k$  in  $\{1, 2, \dots, n-1\}$ :
      - $\text{turn}(k)$  in  $\{1, 2, \dots, n\}$ , mWmR register, written and read by all, initially arbitrary. – For  $i$  in  $\{1, 2, \dots, n\}$ :
        - $\text{flag}(i)$  in  $\{0, 1, 2, \dots, n-1\}$ , 1WmR register, written by  $i$  and read by all, initially 0.
  - Process  $i$  trying protocol:
    - For each level  $k$ :
      - Set  $\text{flag}(i) := k$ , indicating  $i$  is competing at level  $k$ .
      - Set  $\text{turn}(k) := i$ .
      - Wait for either  $\text{turn}(k) \neq i$ , or everyone else's  $\text{flag} < k$  (check flags one at a time).

- Exit protocol:
  - Set  $\text{flag}(i) := 0$

## Correctness: Mutual exclusion

- **Definition:** Process  $i$  is a winner at level  $k$  if either:
  - $\text{level}_i > k$ , or
  - $\text{level}_i = k$  and  $\text{pc}_i \in \{\text{leave-try}, \text{crit}, \text{reset}\}$ .
- **Definition:** Process  $i$  is a competitor at level  $k$  if either:
  - Process  $i$  is a winner at level  $k$ , or
  - $\text{level}_i = k$  and  $\text{pc}_i \in \{\text{check-flag}, \text{check-turn}\}$ .
- **Invariant 1:** If process  $i$  is a winner at level  $k$ , and process  $j \neq i$  is a competitor at level  $k$ , then  $\text{turn}(k) \neq i$ .
- **Proof:** By induction, similar to 2-process case.
  - Complication: More steps to consider.

- Now have many flags, checked in many steps.
- Need auxiliary invariants saying something about what is true in the middle of checking a set of flags.

## Correctness: Mutual exclusion

- **Invariant 2:** For any  $k$ ,  $1 \leq k \leq n-1$ , there are at most  $n-k$  winners at level  $k$ .
- **Proof:** By induction, on level number, for a particular reachable state (not induction on number of steps).
  - **Basis:**  $k = 1$ :
    - Suppose false, for contradiction.
    - Then all  $n$  processes are winners at level 1.
    - Then Invariant 1 implies that  $\text{turn}(1)$  is unequal to all indices, contradiction.

– Inductive step: ...

## Correctness: Mutual exclusion

- **Invariant 2:** For any  $k$ ,  $1 \leq k \leq n-1$ , there are at most  $n - k$  winners at level  $k$ .
- **Inductive step:** Assume for  $k$ ,  $1 \leq k \leq n-2$ , show for  $k+1$ .
  - Suppose false, for contradiction.
  - Then more than  $n - (k + 1)$  processes, that is, at least  $n - k$  processes, are winners at level  $k + 1$ :  $|Win_{k+1}| \geq n - k$ .
  - Every level  $k+1$  winner is also a level  $k$  winner:  $Win_{k+1} \subseteq Win_k$ .
  - By inductive hypothesis,  $|Win_k| \leq n - k$ .
  - So  $Win_{k+1} = Win_k$ , and  $|Win_{k+1}| = |Win_k| = n - k$ .
  - **Q:** What is the value of  $turn(k+1)$  ?
- Can't be the index of any process in  $Win_{k+1}$ , by Invariant 1.
- Must be the index of some competitor at level  $k+1$  (Invariant, LTTR).

- But every competitor at level  $k+1$  is a winner at level  $k$ , so is in  $\text{Win}_k$ .
- Contradiction, since  $\text{Win}_{k+1} = \text{Win}_k$ .

## Progress, Lockout-freedom

- Lockout-freedom proof idea:
  - Let  $k$  be the highest level at which some process  $i$  gets stuck.
  - Then  $\text{turn}(k)$  must remain  $= i$ .
  - That means no one else ever reenters the competition at level  $k$ .
  - Eventually, winners from level  $k$  will finish, since  $k$  is the highest level at which anyone gets stuck.
  - Then all other flags will be  $< k$ , so  $i$  advances.
- Alternatively, prove lockout-freedom by showing a time bound for each process, from  $\rightarrow T$  until  $\rightarrow C$ . (See book)
  - Define  $T(0)$  = maximum time from when a process  $\rightarrow T$  until  $\rightarrow C$ .
  - Define  $T(k)$ ,  $1 \leq k \leq n-1$  = max time from when a process wins at level  $k$  until  $\rightarrow C$ .


- $T(n-1) \leq l$ .
- $T(k) \leq 2 T(k+1) + c + (3n+2) l$ , by detailed analysis.
- Solve recurrences, get exponential bound, good enough for showing lockout-freedom.

## Peterson Tournament Algorithm

- Assume  $n = 2^h$ .
- Processes = leaves of binary tree of height  $h$ .
- **Competitions** = internal nodes, labeled by binary strings.
- Each process engages in  $\log n$  competitions, following path up to root.
- Each process  $i$  has:  $\lambda$

- A unique competition  $x$  at each level  $k$ .
- A unique role in  $x$  ( $0 = \text{left}$ ,  $1 = \text{right}$ ).
- A set of potential opponents in  $x$ .

# Peterson Tournament Algorithm




- Shared variables:
  - For each process  $i$ ,  $\text{flag}(i)$  in  $\{0, \dots, h\}$ , indicating level, initially 0
  - For each competition  $x$ ,  $\text{turn}(x)$ , a Boolean, initially arbitrary. •

Process  $i$ 's trying protocol: For each level  $k$ :

- Set  $\text{flag}(i) := k$ .
- Set  $\text{turn}(x) := b$ , where:
  - $x$  is  $i$ 's level  $k$  competition,  $\lambda$
  - $b$  is  $i$ 's “role”, 0 or 1– Wait for either:
    - $\text{turn}(x) = \text{opposite role}$ , or
    - all  $\text{flags}$  of potential opponents in  $x$  are  $< k$ .
- Exit protocol:
  - Set  $\text{flag}(i) := 0$ .

## Correctness

- **Mutual exclusion:**
  - Similar to before.
  - Key invariant: At most one process from any particular subtree rooted at level  $k$  is currently a winner at level  $k$ .





- **Time bound** (from  $\rightarrow T$  until  $\rightarrow C$ ):  $(n-1) c + O(n^2 l)$ 
  - Implies progress, lockout-freedom.
  - Define:  $T(0) = \text{max time from } \rightarrow T \text{ until } \rightarrow C.$
  - $T(k), 1 \leq k \leq \log n = \text{max time from winning at level } k \text{ until } \rightarrow C.$
  - $T(\log n) \leq l.$
  - $T(k) \leq 2 T(k+1) + c + (2^{k+1} + 2^k + 7) l$  (see book).
- Roughly: Might need to wait for a competitor to reach C, then finish C, then for yourself to reach C.
  - Solve recurrences.

## Bounded Bypass?

- Peterson's Tournament algorithm has a low time bound from  $\rightarrow T$  until  $\rightarrow C$ :
 
$$(n-1) c + O(n^2 l)$$
- Implies lockout-freedom, progress.

- **Q:** Does it satisfy bounded bypass?
- **No!** There's no upper bound on the number of times one process could bypass another in the trying region. E.g.:
  - Process 0 enters, starts competing at level 1, then pauses.
  - Process 7 enters, quickly works its way to the top, enters C, leaves C.
  - Process 7 enters again...repeats any number of times.
  - All while process 0 is paused.
- No contradiction between small time bound and unbounded bypass.
  - Because of the way we're modeling timing of asynchronous executions, using upper bound assumptions.
  - When processes go at very different speeds, we say that the slow processes are going at normal speed, faster processes are going very fast.

## Lamport's Bakery Algorithm

- Like taking tickets in a bakery.
- Nice features:

- Uses only single-writer, multi-reader registers.
- Extends to even weaker registers, in which operations have durations, and a read that overlaps a write receives an arbitrary response.
- Guarantees lockout-freedom, in fact, almost-FIFO behavior.
- But:
  - Registers are unbounded size.
  - Algorithm can be simulated using bounded registers, but not easily (uses bounded concurrent timestamps).
- Shared variables:
  - For each process  $i$ :
    - $\text{choosing}(i)$ , a Boolean, written by  $i$ , read by all, initially 0
    - $\text{number}(i)$ , a natural number, written by  $i$ , read by all, initially 0

## Bakery Algorithm

- First part, up to  $\text{choosing}(i) := 0$  (the “Doorway”, D):

- Process  $i$  chooses a number greater than all the numbers it reads for the other processes; writes this in  $\text{number}(i)$ .
- While doing this, keeps choosing  $\text{choosing}(i) = 1$ .
- Two processes could choose the same number (unlike real bakery).
- Break ties with process ids.

- **Second part:**

- Wait to see that no others are choosing, and no one else has a smaller number.
- That is, wait to see that your ticket is the smallest.
- Never go back to the beginning of this part---just proceed step by step, waiting when necessary.

# Code

**Shared variables:** for every  $i \in \{1, \dots, n\}$ : **choosing(i)**  $\in \{0, 1\}$ , initially 0, writable by  $i$ , readable by all  $j \neq i$  **number(i)**, a natural number, initially 0, writable by  $i$ , readable by  $j \neq i$ .

**try<sub>i</sub>**

**choosing(i) := 1** **number(i) := 1 +**

**max<sub>j ≠ i</sub> number(j)** **choosing(i) := 0**

**for j ≠ i do waitfor**

**choosing(j) = 0**

**waitfor number(j) = 0 or (number(i), i) < (number(j), j)**

**crit<sub>i</sub>**

**exit<sub>i</sub>**

number(i) := 0

rem<sub>i</sub>

## Correctness: Mutual exclusion

- **Key invariant:** If process  $i$  is in  $C$ , and process  $j \neq i$  is in  $(T - D) \cup C$ ,

Trying region after doorway, or critical region


then  $(\text{number}(i), i) < (\text{number}(j), j)$ .

- **Proof:**
  - Could prove by induction.
  - Instead, give argument based on events in executions.

- This argument extends to weaker registers, with concurrent accesses.

## Correctness: Mutual exclusion

- **Invariant:** If  $i$  is in  $C$ , and  $j \neq i$  is in  $(T - D) \cup C$ , then  $(\text{number}(i), i) < (\text{number}(j), j)$ .
- **Proof:**
  - Consider a point where  $i$  is in  $C$  and  $j \neq i$  is in  $(T - D) \cup C$ .
  - Then before  $i$  entered  $C$ , it must have read  $\text{choosing}(j) = 0$ , event  $\pi$ .



- **Case 1:**  $j$  sets  $\text{choosing}(j) := 1$  (starts choosing) after  $\pi$ .
- Then  $\text{number}(i)$  is set before  $j$  starts choosing.

- So  $j$  sees the “correct”  $\text{number}(i)$  and chooses something bigger.
- That suffices.
- **Case 2:**  $j$  sets  $\text{choosing}(j) := 0$  (finishes choosing) before  $\pi$ .
  - Then when  $i$  reads  $\text{number}(j)$  in its second waitfor loop, it gets the “correct”  $\text{number}(j)$ .
  - Since  $i$  decides to enter  $C$ , it must see  $(\text{number}(i), i) < (\text{number}(j), j)$ .

## Correctness: Mutual exclusion

- **Invariant:** If  $i$  is in  $C$ , and  $j \neq i$  is in  $(T - D) \cup C$ , then  $(\text{number}(i), i) < (\text{number}(j), j)$ .
- **Proof of mutual exclusion:**
  - Apply invariant both ways.
  - Contradictory requirements.



# Liveness Conditions

- **Progress:**
  - By contradiction.
  - If not, eventually region changes stop, leaving everyone in T or R, and at least one process in T.
  - Everyone in T eventually finishes choosing.
  - Then nothing blocks the smallest (number, index) process from entering C.
- **Lockout-freedom:**
  - Consider any  $i$  that enters T
  - Eventually it finishes the doorway.
  - Thereafter, any newly-entering process picks a bigger number.
  - Progress implies that processes continue to enter C, as long as  $i$  is still in T.

- In fact, this must happen infinitely many times!
- But those with bigger numbers can't get past  $i$ , contradiction.

## FIFO Condition

- Not really FIFO ( $\rightarrow T$  vs.  $\rightarrow C$ ), but almost:
  - **FIFO after the doorway**: if  $j$  leaves  $D$  before  $i \rightarrow T$ , then  $j \rightarrow C$  before  $i \rightarrow C$ .
- But the “doorway” is an artifact of this algorithm, so this isn't a meaningful way to evaluate the algorithm!
- Maybe say “there exists a doorway such that”...
- But then we could take  $D$  to be the entire trying region, making the property trivial.
- To make the property nontrivial:

- Require D to be “wait-free”: a process is guaranteed to complete D it if it keeps taking steps, regardless of what any other processes do.
- D in the Bakery Algorithm is wait-free.
- The algorithm is **FIFO after a wait-free doorway**.

## Impact of Bakery Algorithm

- Originated important ideas:
  - Wait-freedom
- Fundamental notion for theory of fault-tolerant asynchronous distributed algorithms.
  - Weakly coherent memories

- Beginning of formal study: definitions, and some algorithmic strategies for coping with them.

## Next time...

- More mutual exclusion algorithms:
  - Lamport's Bakery Algorithm, cont'd
  - Burns' algorithm
- Number of registers needed for mutual exclusion.
- Reading: Sections 10.6-10.8

MIT OpenCourseWare  
<http://ocw.mit.edu>

For information about citing these materials or our Terms of Use, visit: <http://ocw.mit.edu/terms>.