IT/PC/B/T/411

Machine Learning

Deep Learning Basics

Lecture 01: Feedforward

Dr. Pawan Kumar Singh

Department of Information Technology

Jadavpur University

pawankrsingh.cse@gmail.com

+91-6291555693

Motivation I: representation learning

Machine learning

- Collect data and extract features
- Build model: choose hypothesis class $\mathcal{H}$ and loss function l
- Optimization: minimize the empirical loss

Features

x

Extract features

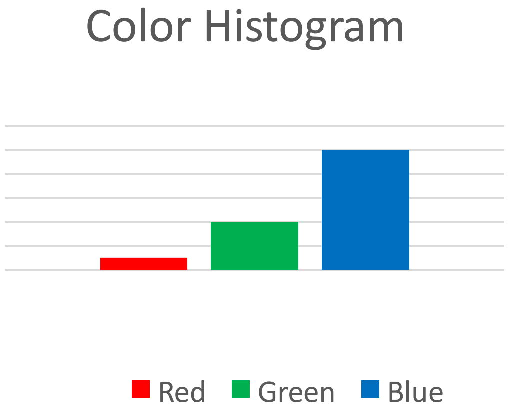

build hypothesis

$$y = w^T \phi(x)$$

Features: part of the model

Nonlinear model

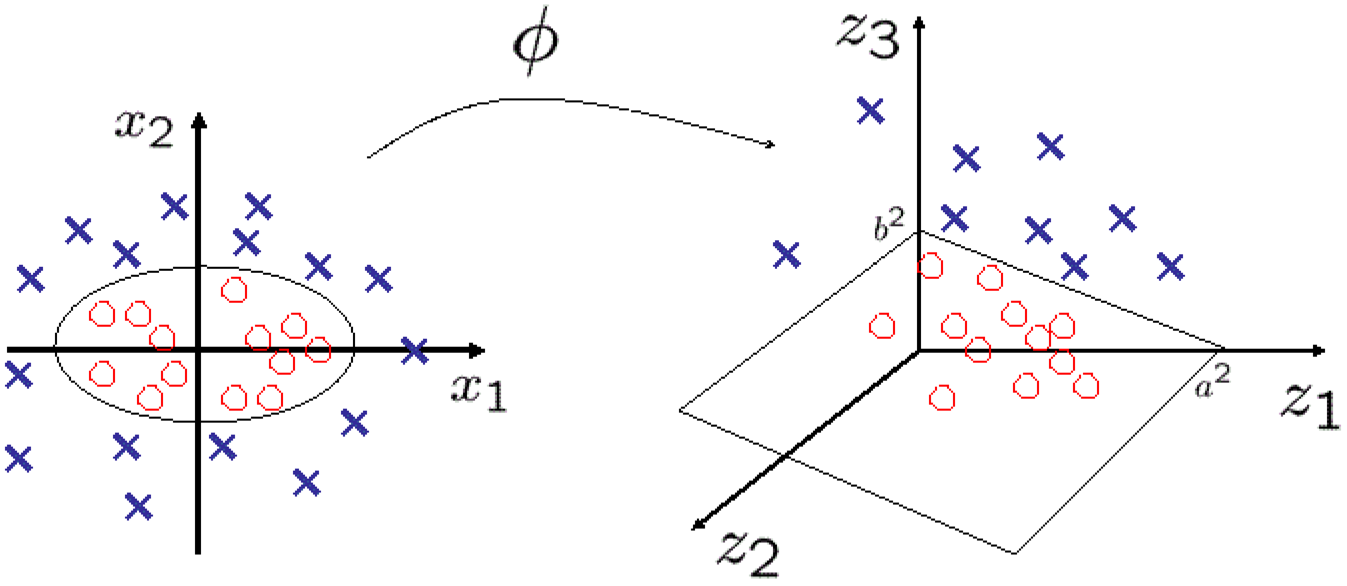

build hypothesis

$$y = w^T \phi(x)$$

Linear model

Example: Polynomial kernel SVM

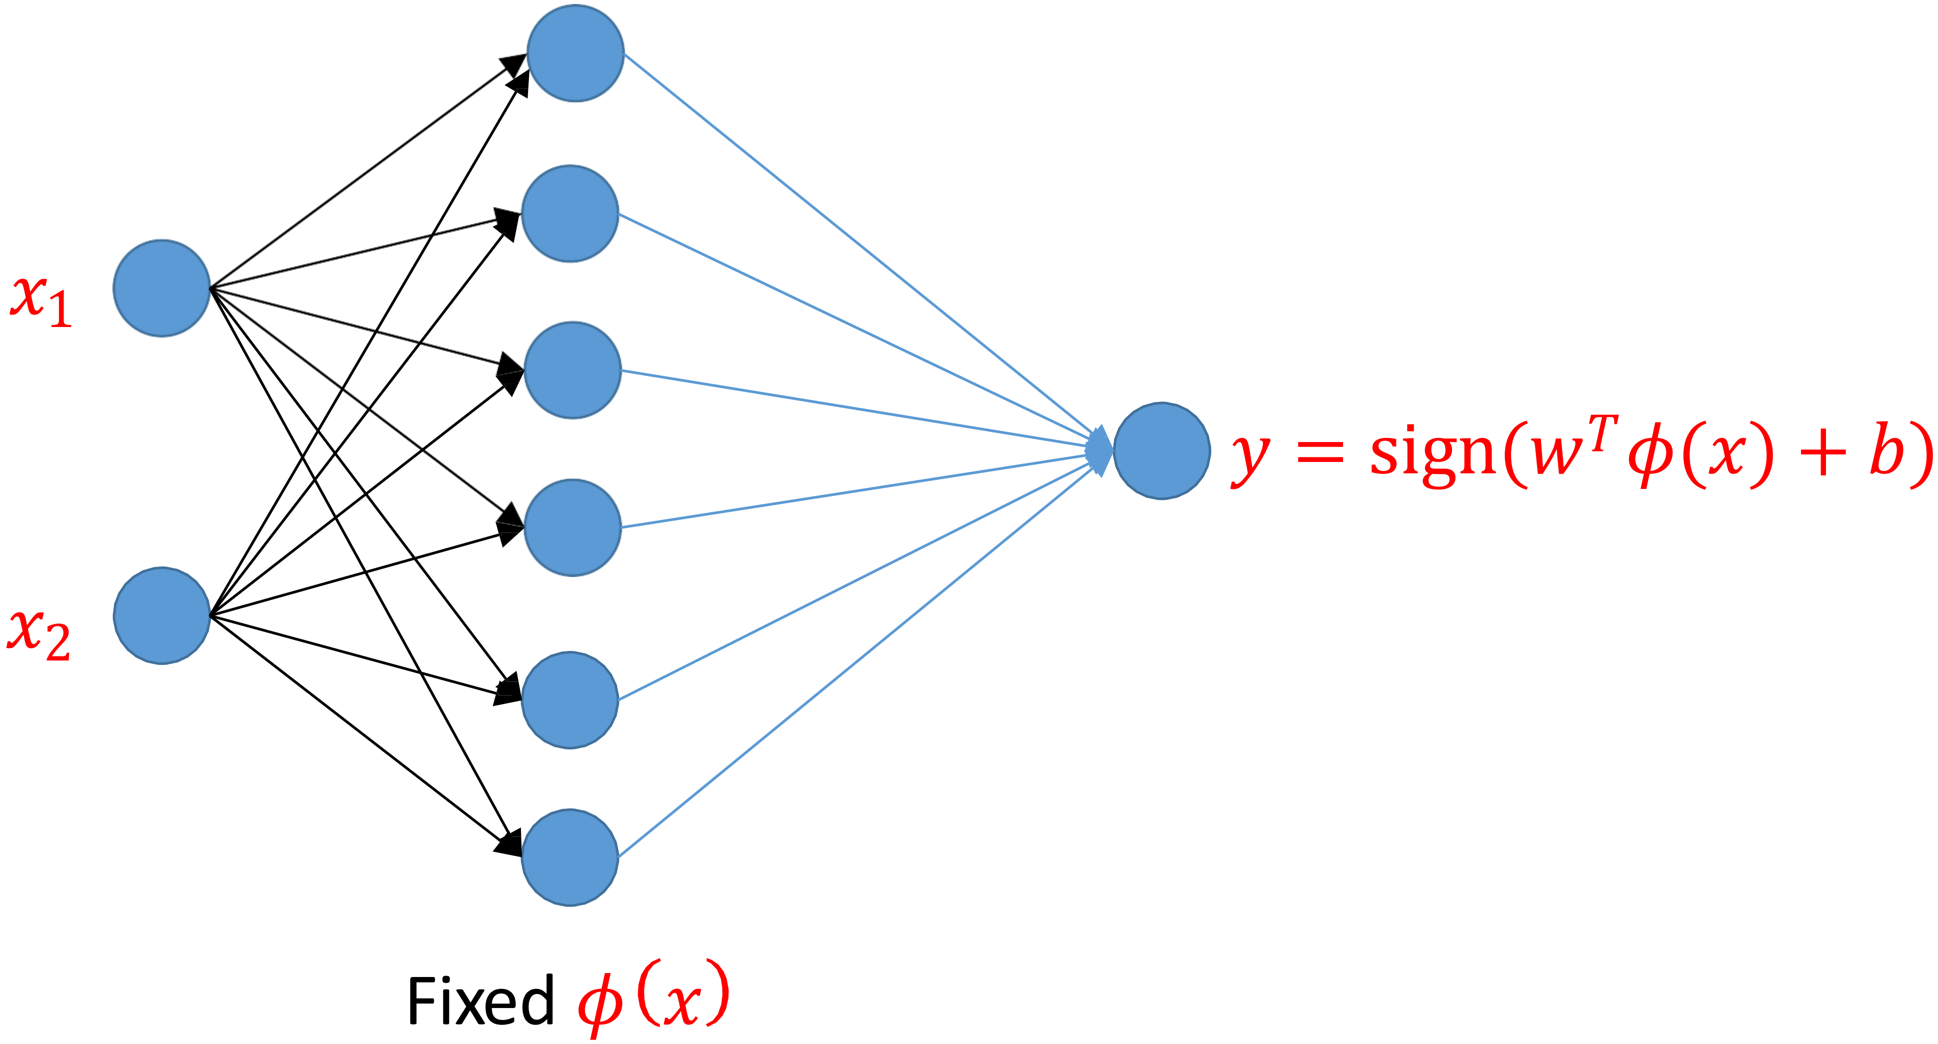

Motivation: representation learning

- Why don't we also learn $\phi(x)$?

x

Learn $\phi(x)$ $\longrightarrow$ $\phi(x)$ $\xrightarrow{\text{Learn } w}$ $y = w^T \phi(x)$

Feedforward networks

- View each dimension of $\phi(x)$ as something to be learned

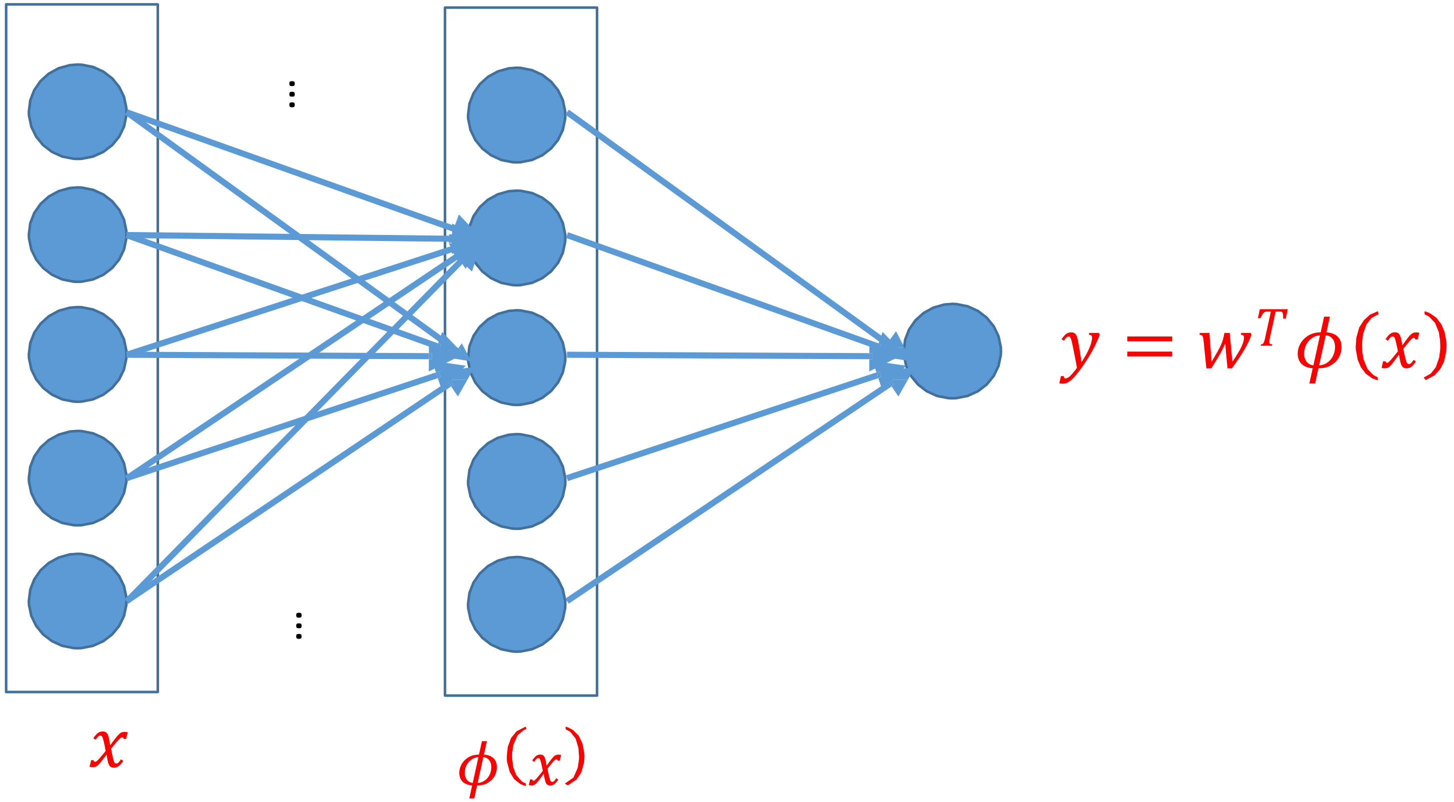

Feedforward networks

- Linear functions $\phi_i(x) = \theta_i^T x$ don't work: need some nonlinearity

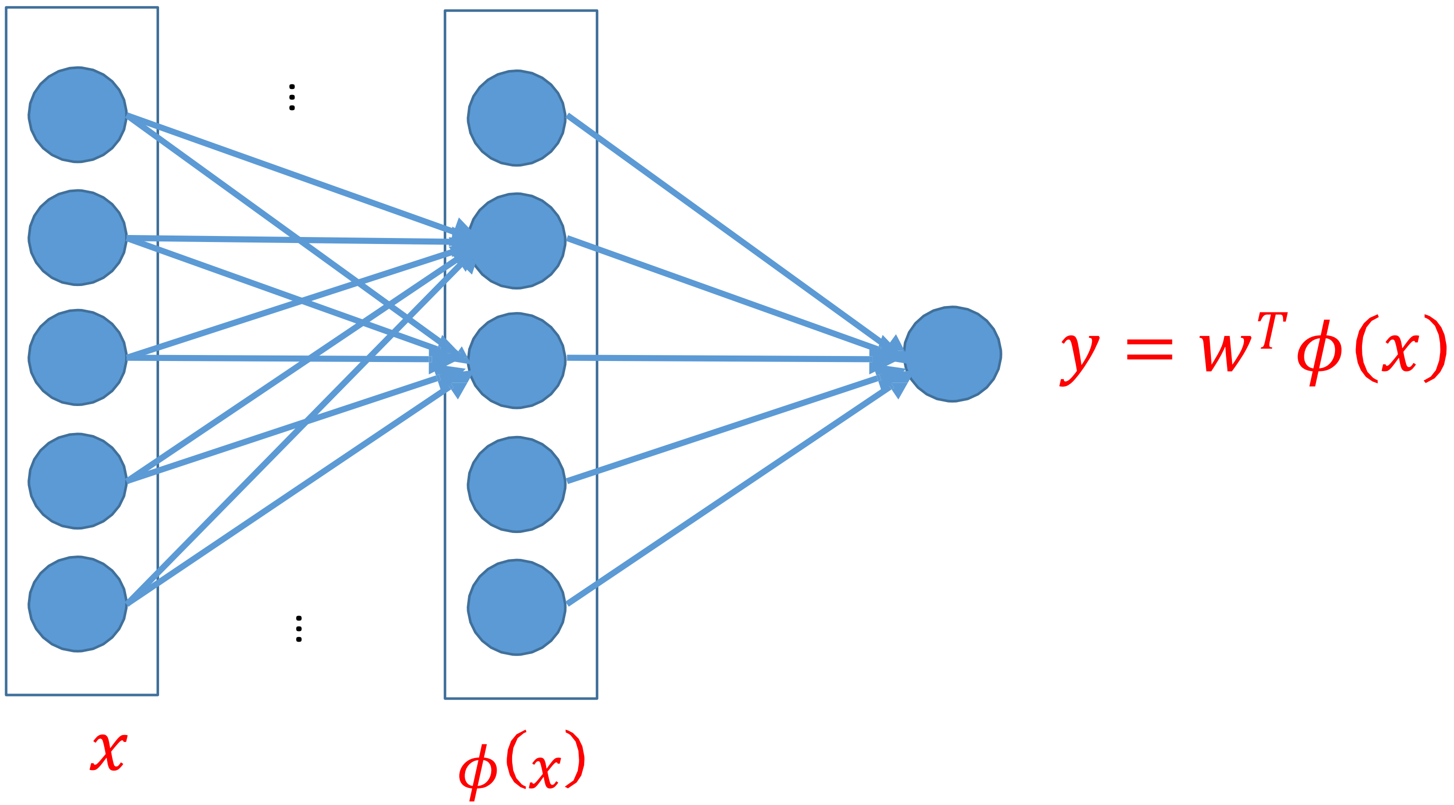

Feedforward networks

- Typically, set $\phi_i(x) = r(\theta_i^T x)$ where $r(\cdot)$ is some nonlinear function

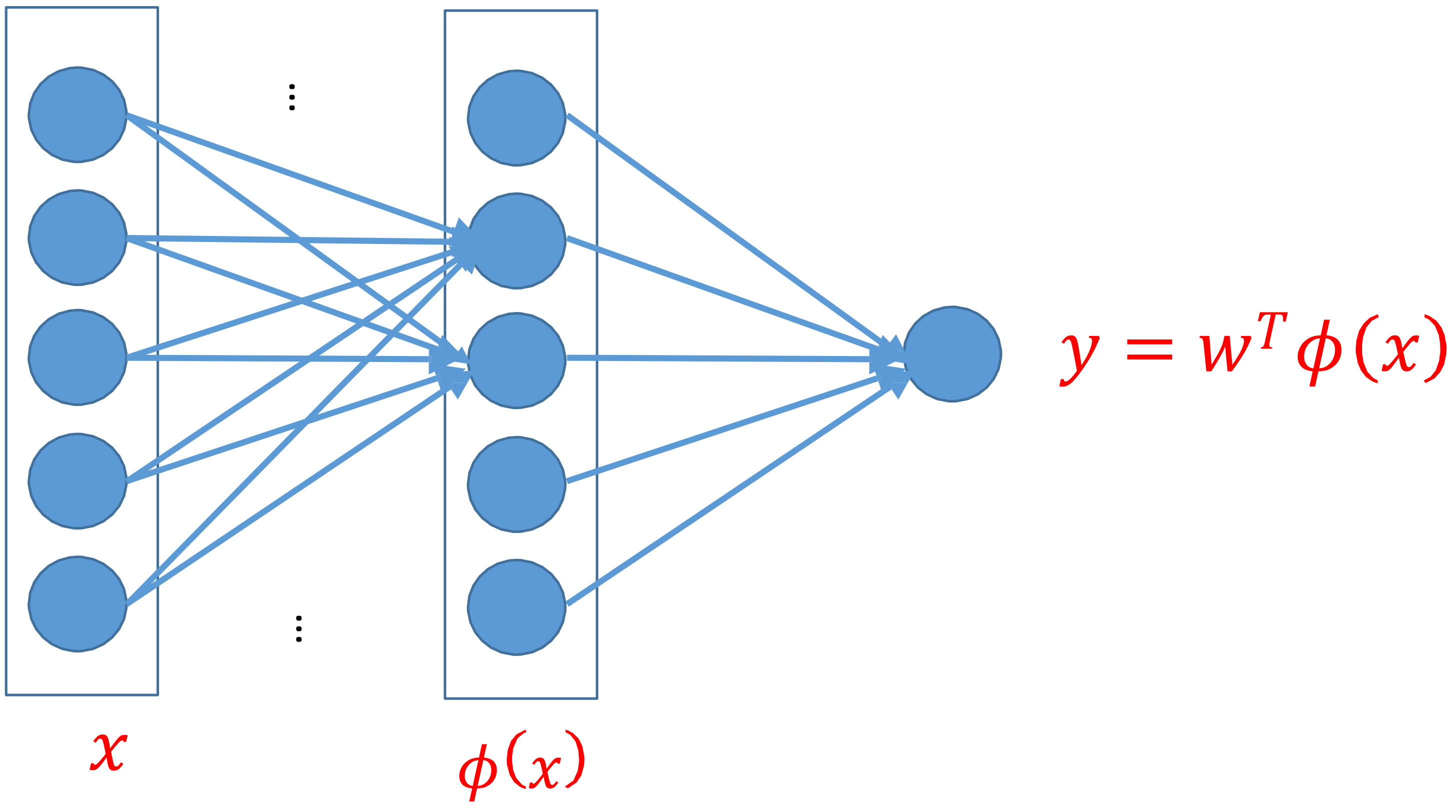

Feedforward deep networks

- What if we go deeper?

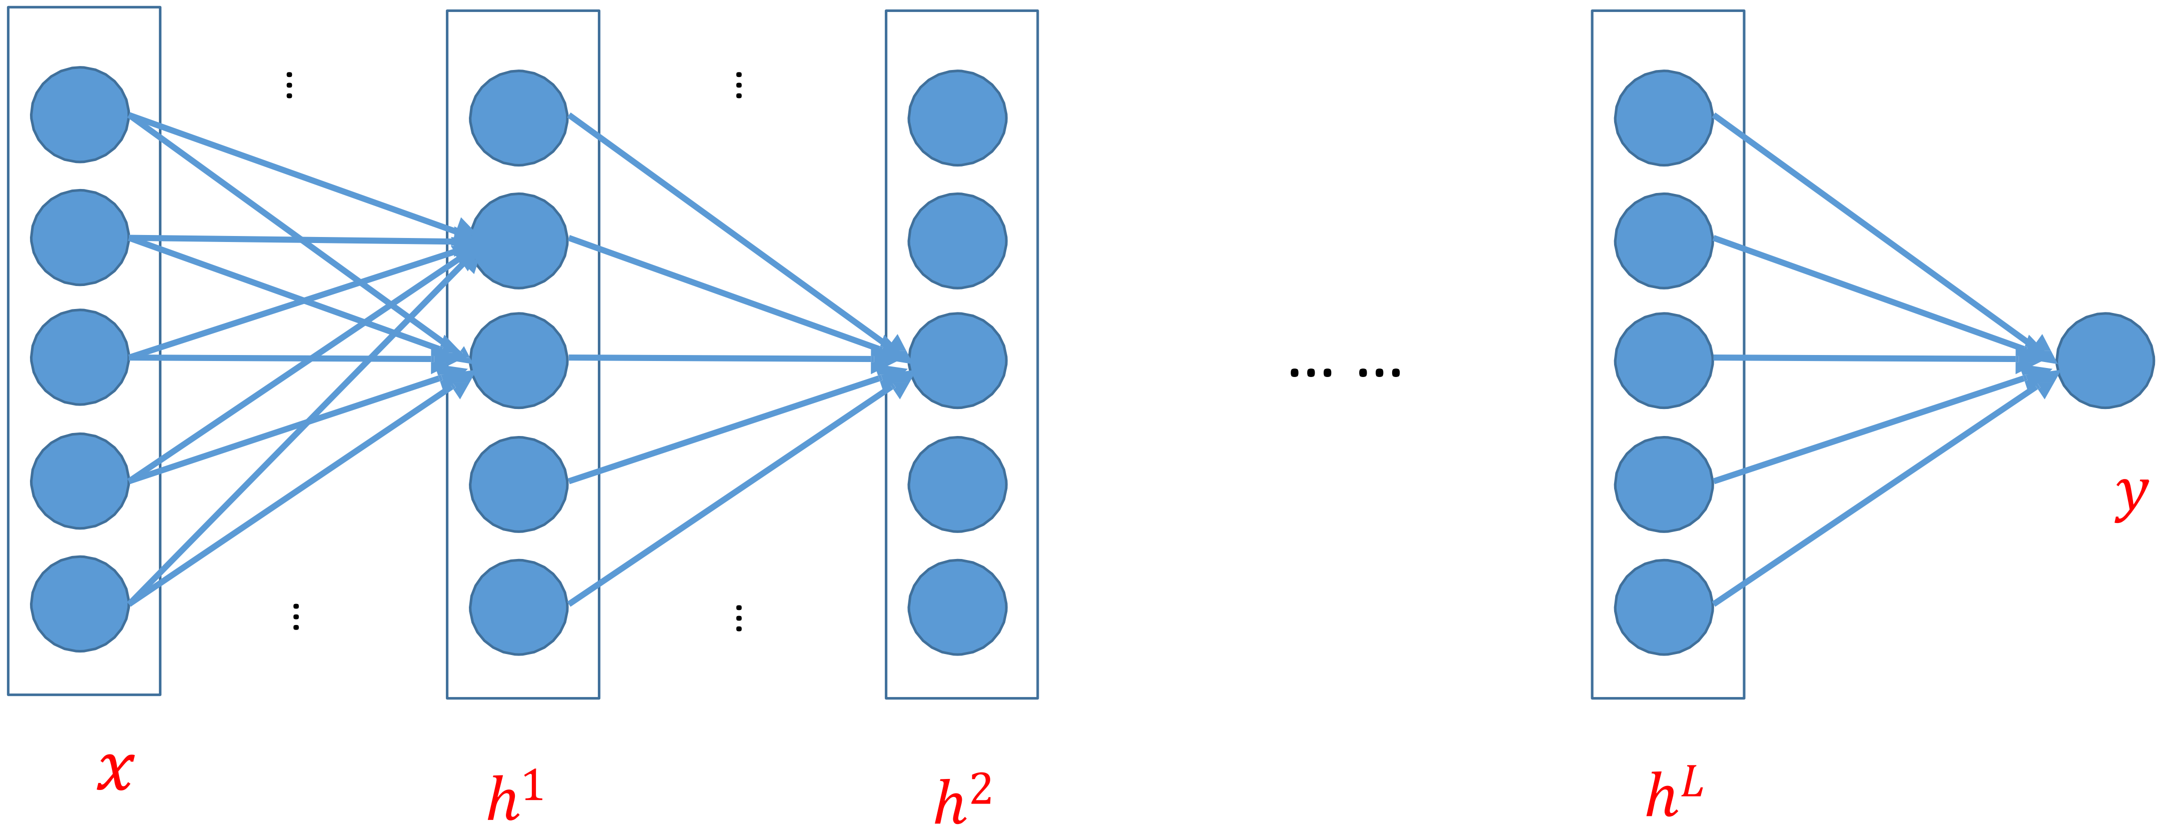

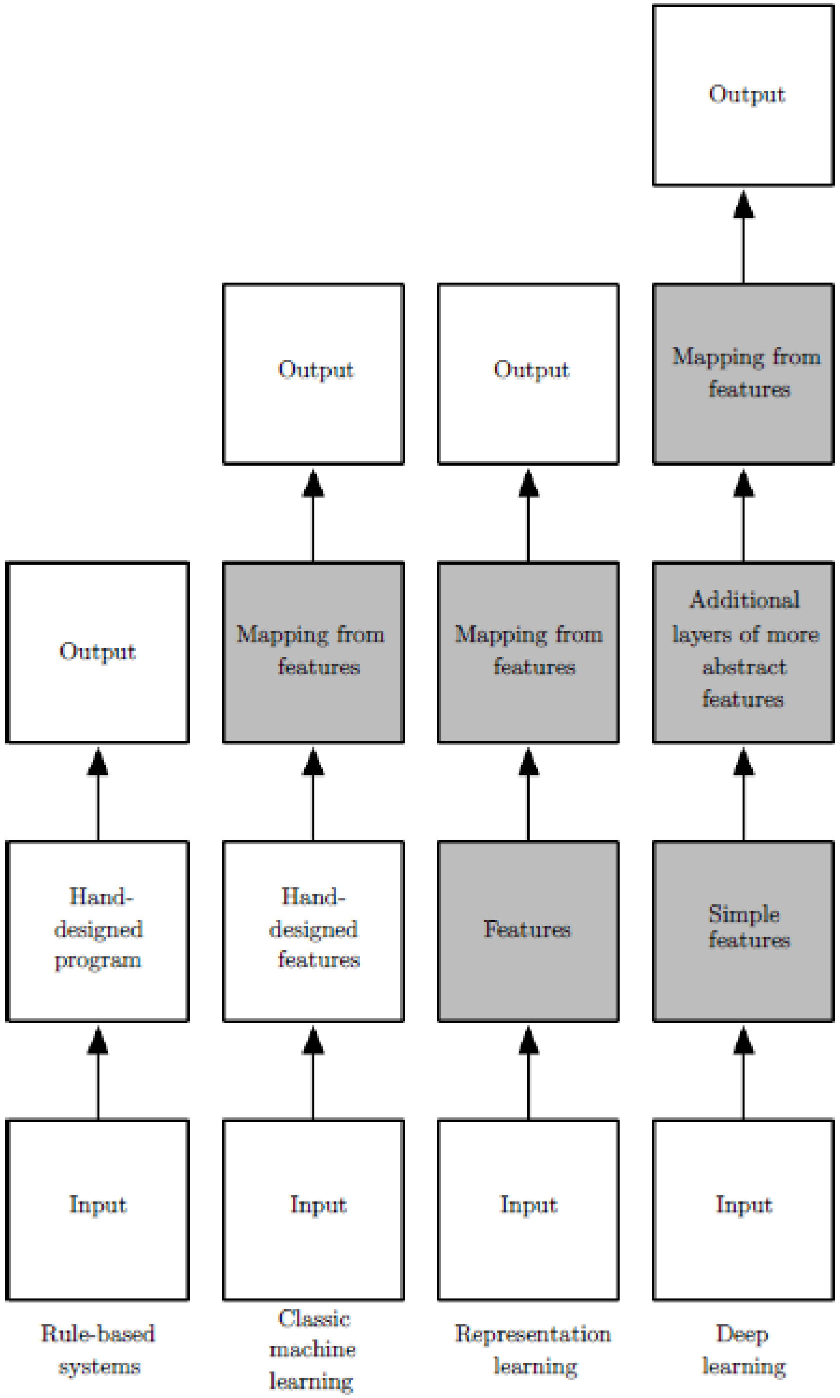

Figure from
Deep learning, by
Goodfellow, Bengio, Courville.
Dark boxes are things to be learned.

Motivation II: neurons

Motivation: neurons

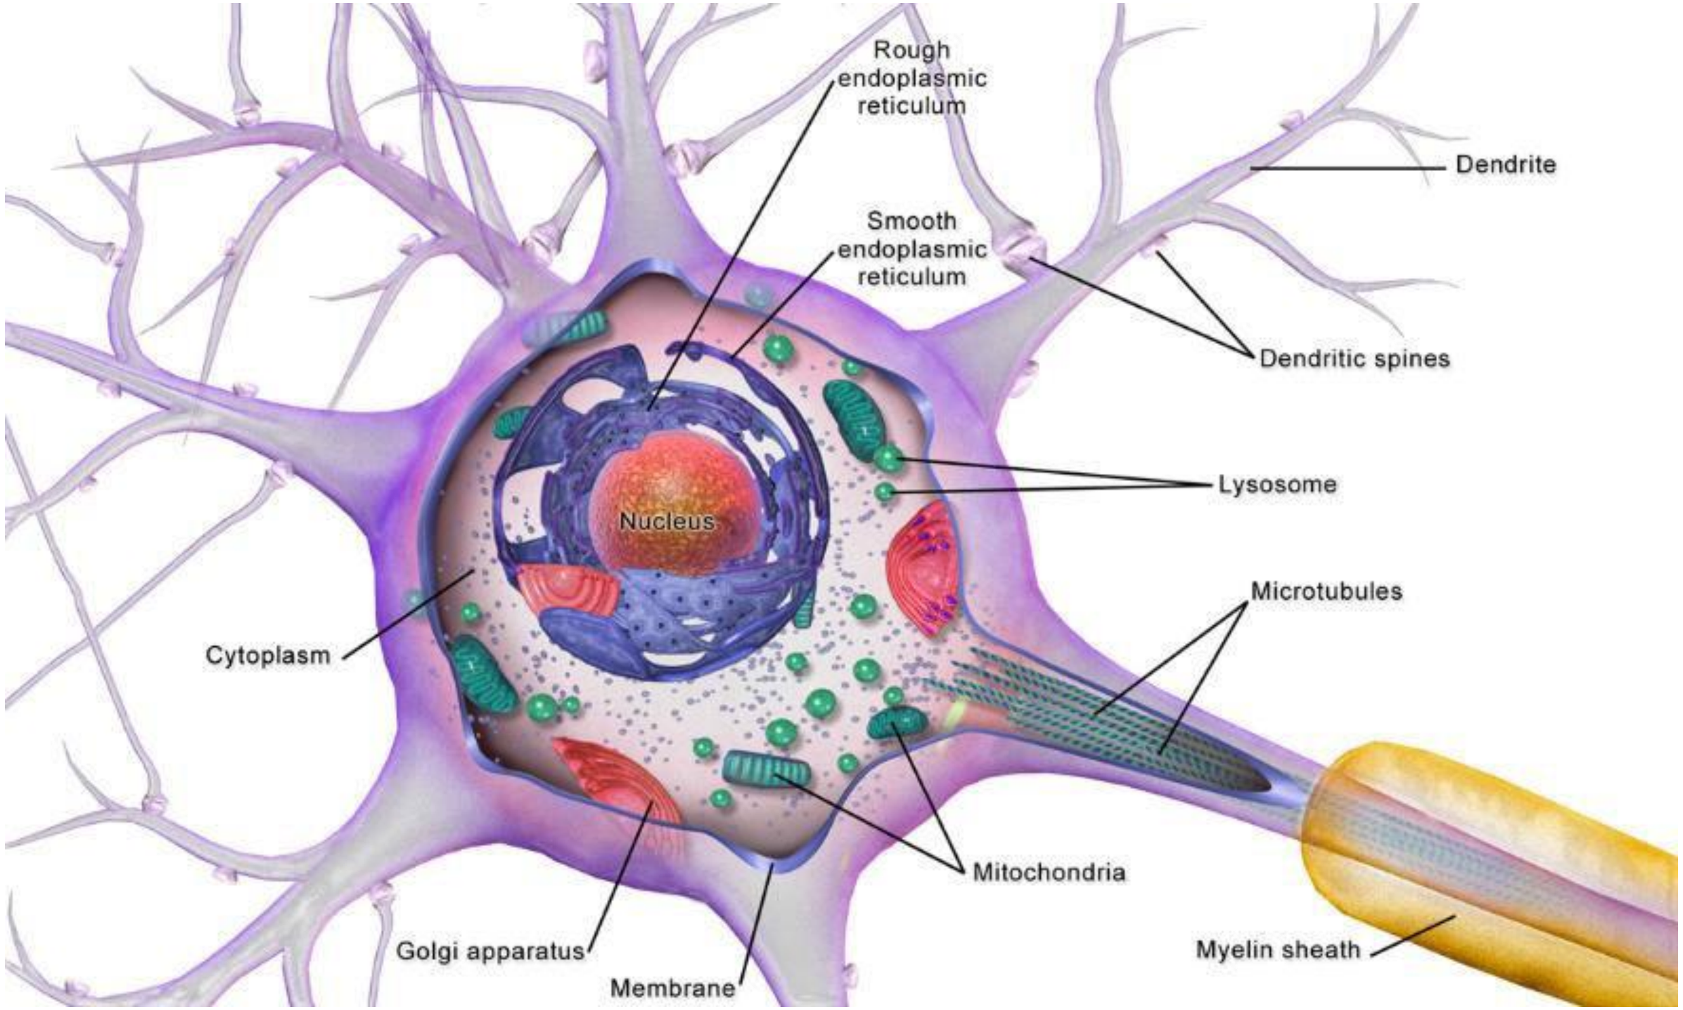

Figure from Wikipedia

Motivation: abstract neuron model

- Neuron activated when the correlation between the input and a pattern θ exceeds some threshold b
- $y = \text{threshold}(\theta^T x - b)$
  or $y = r(\theta^T x - b)$
- $r(\cdot)$ called activation function

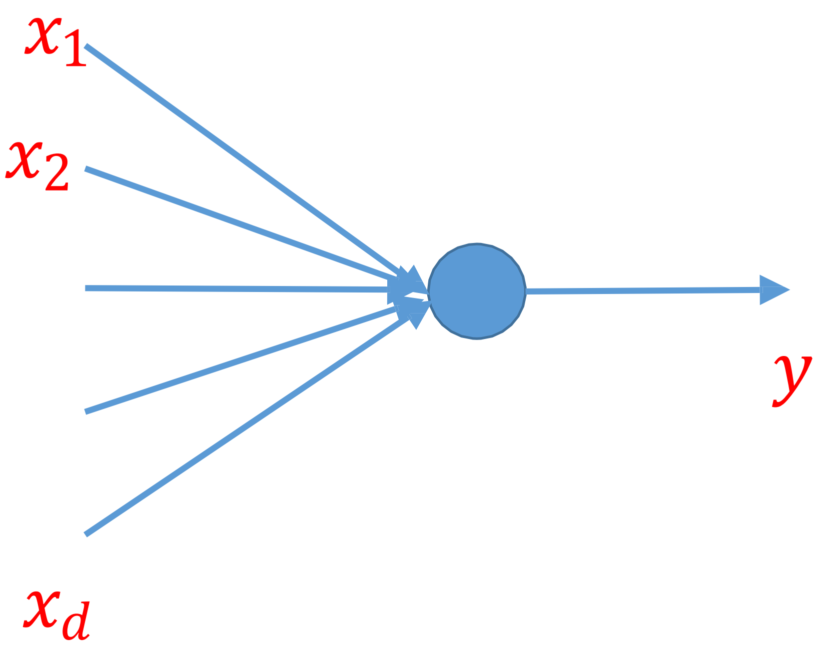

Motivation: artificial neural networks

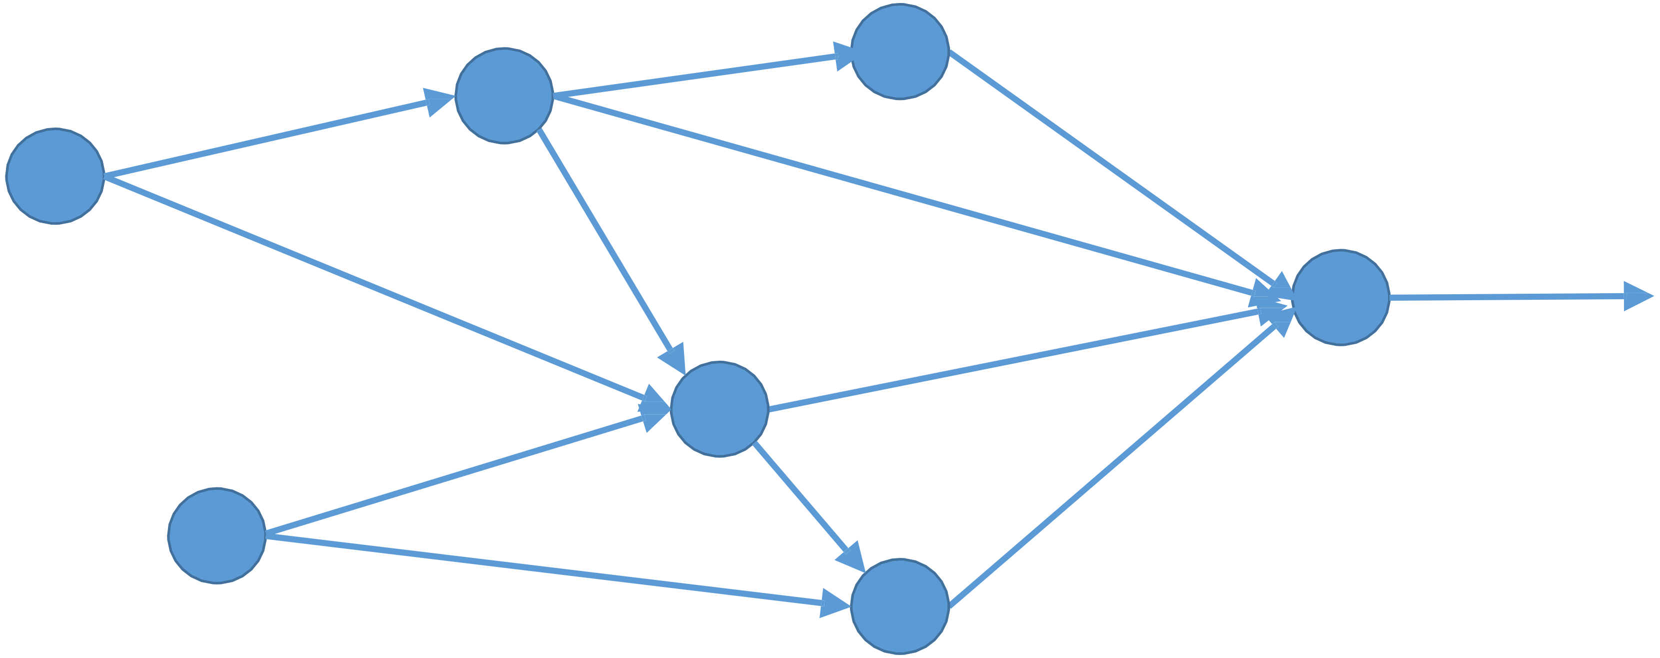

Motivation: artificial neural networks

- Put into layers: feedforward deep networks

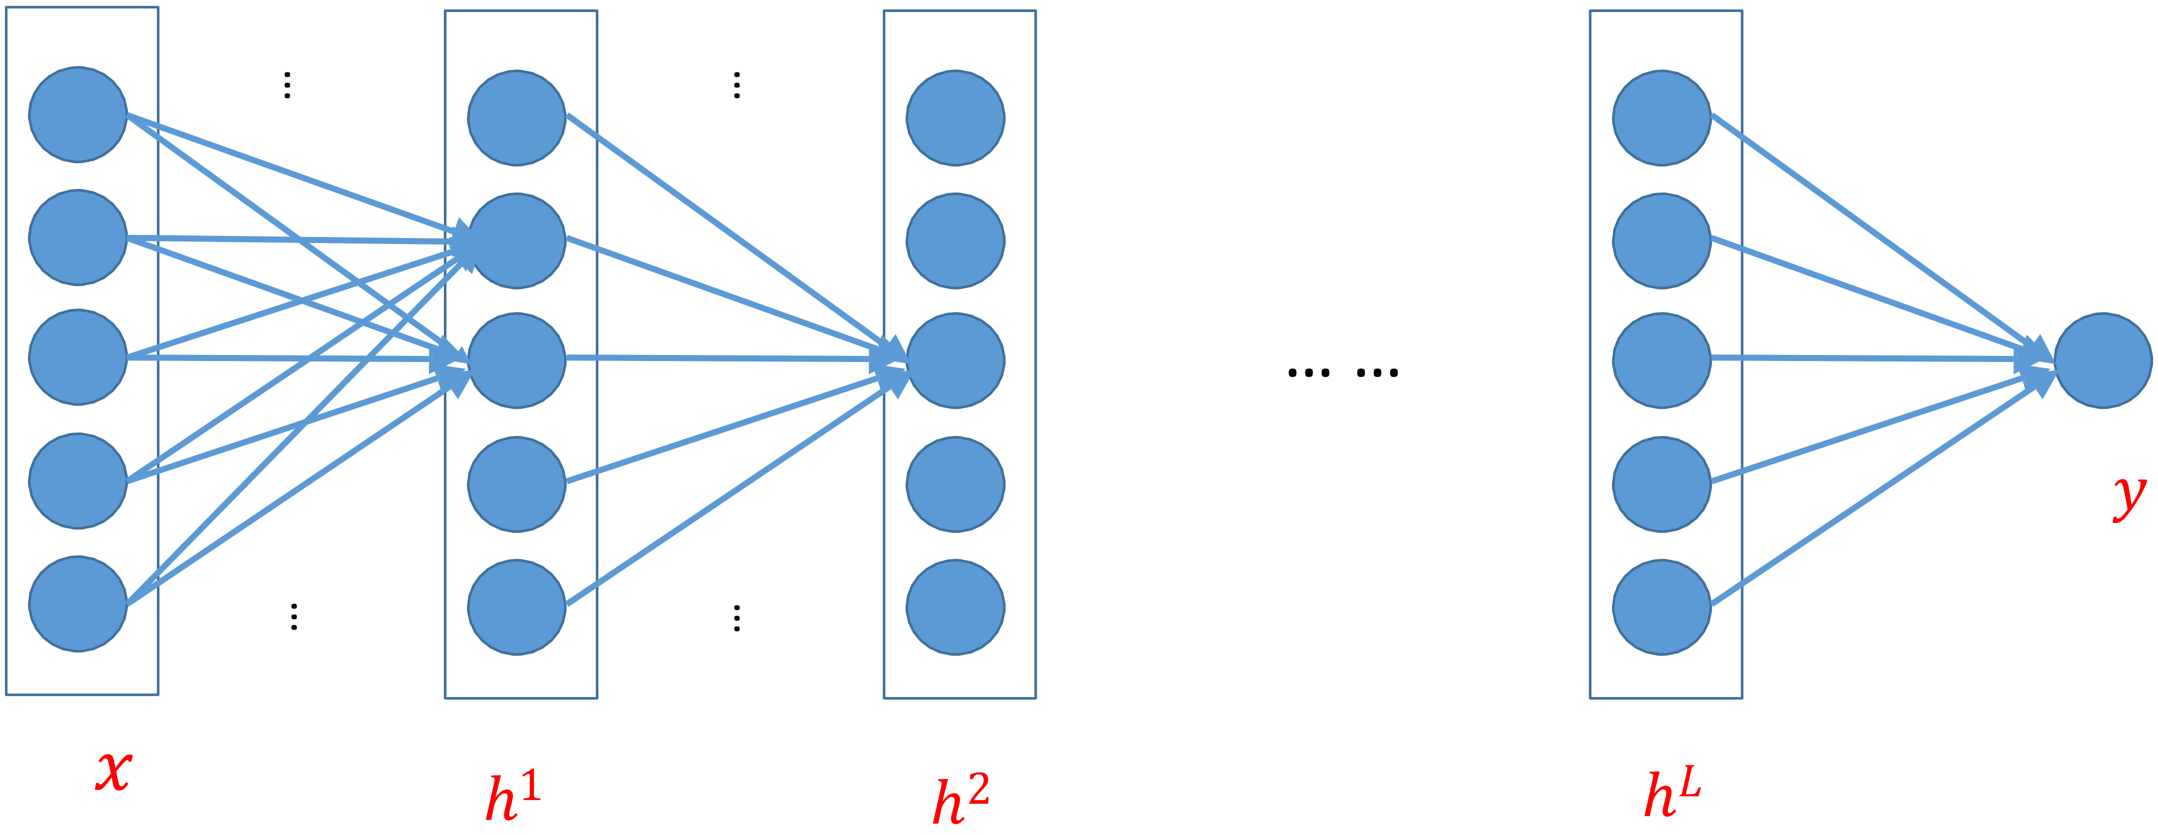

Components in Feedforward networks

Components

- Representations:
  - Input
  - Hidden variables
- Layers/weights:
  - Hidden layers
  - Output layer

Components

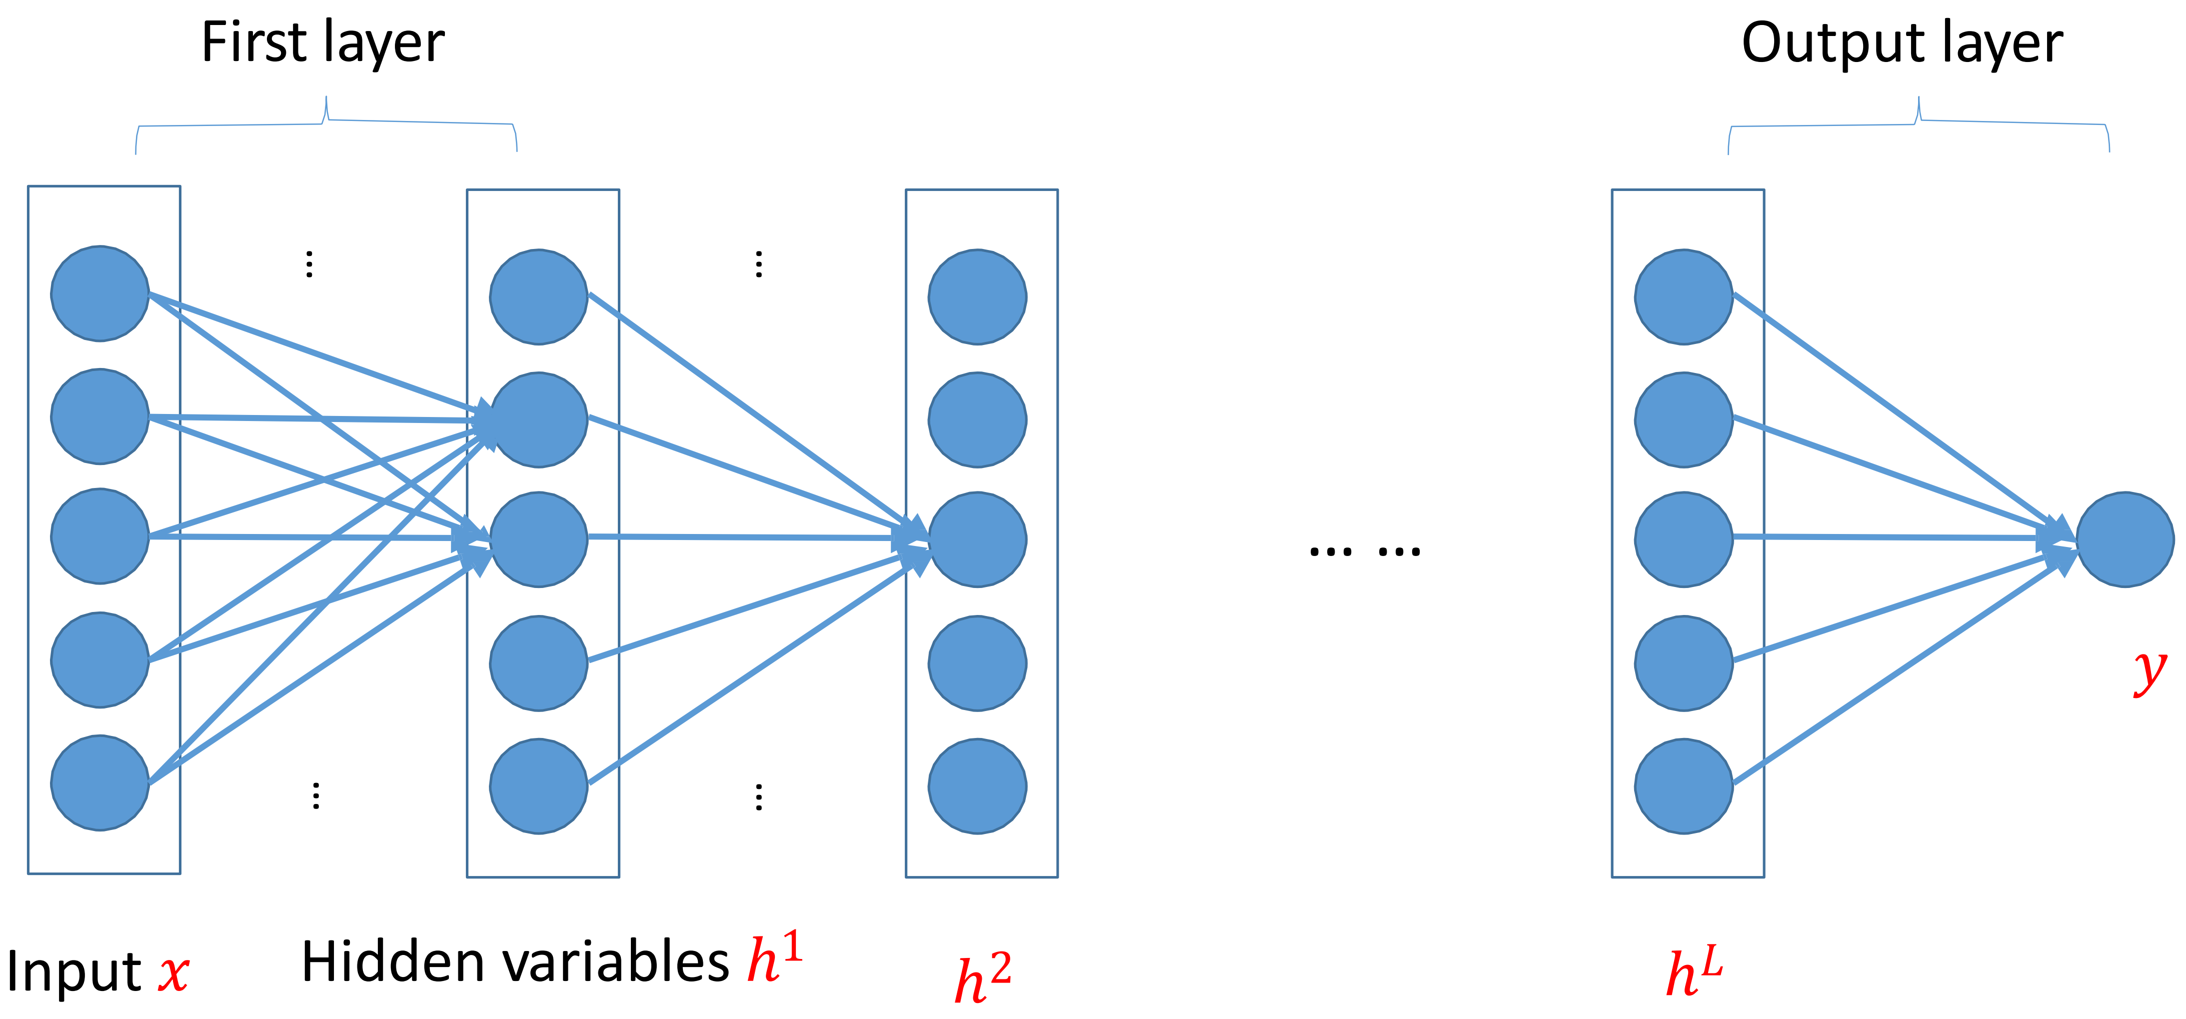

Input

- Represented as a vector
- Sometimes require some preprocessing, e.g.,
  - Subtract mean
  - Normalize to [-1,1]

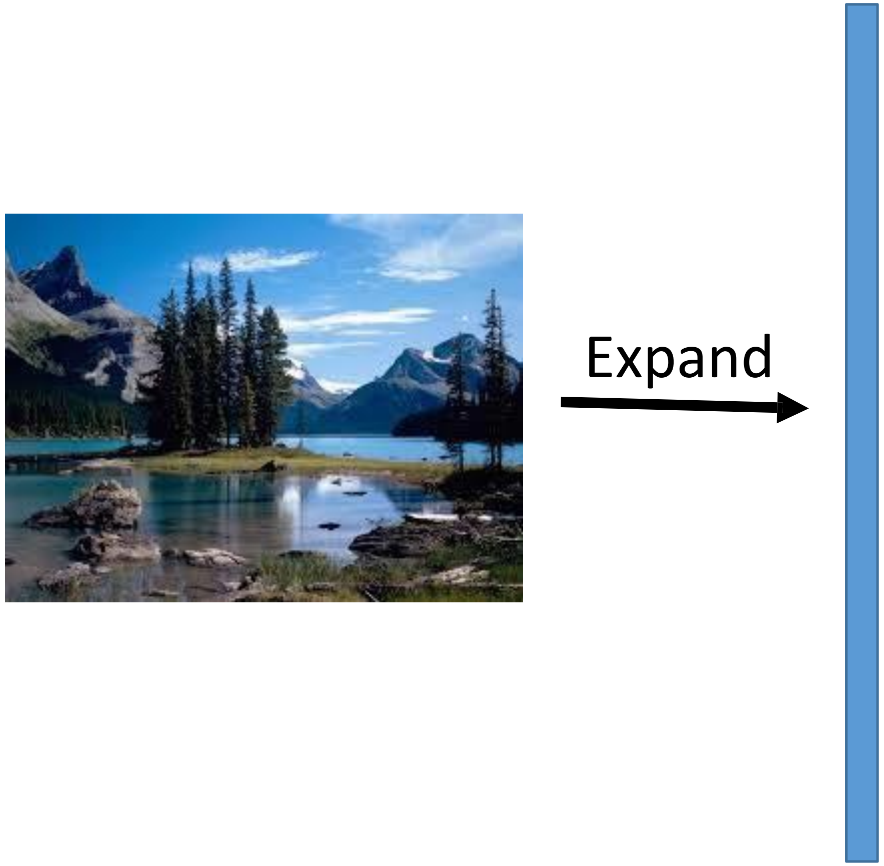

Output layers

- Regression: $y = w^T h + b$
- Linear units: no nonlinearity

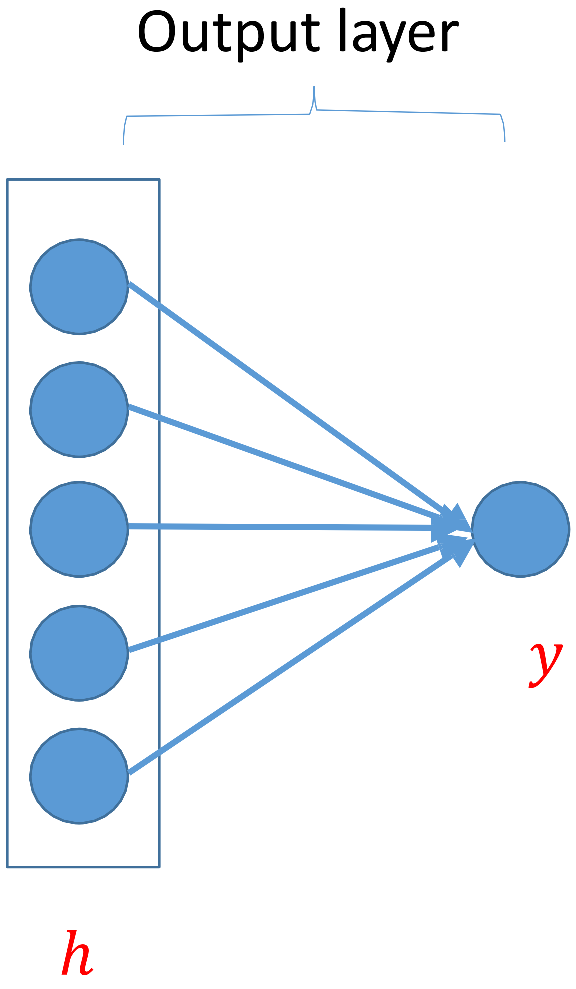

Output layers

- Multi-dimensional regression: $y = W^T h + b$
- Linear units: no nonlinearity

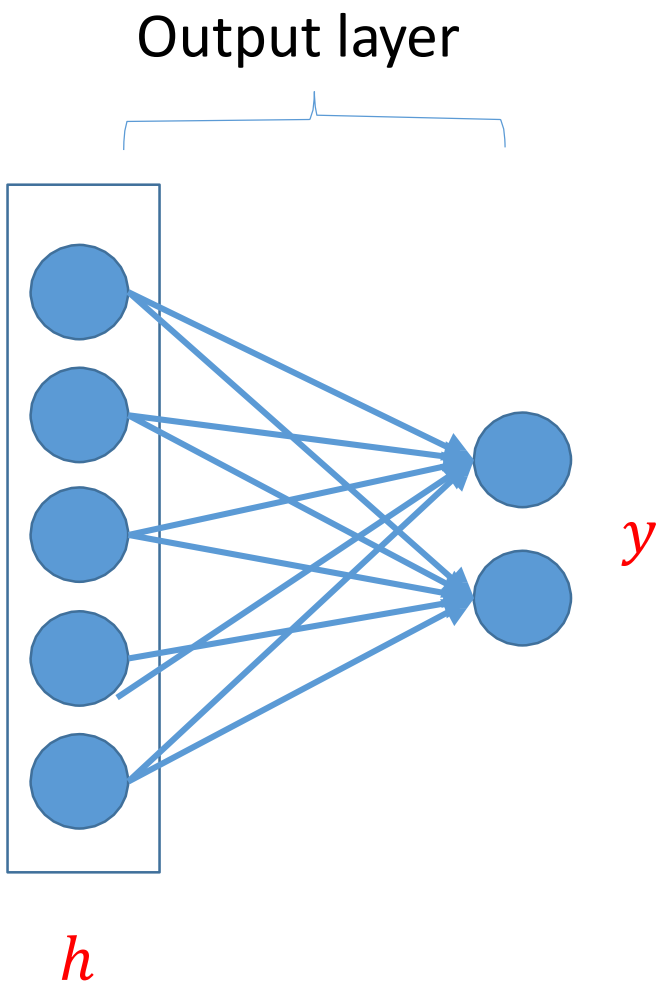

Output layers

- Binary classification: $y = \sigma(w^T h + b)$
- Corresponds to using logistic regression on h

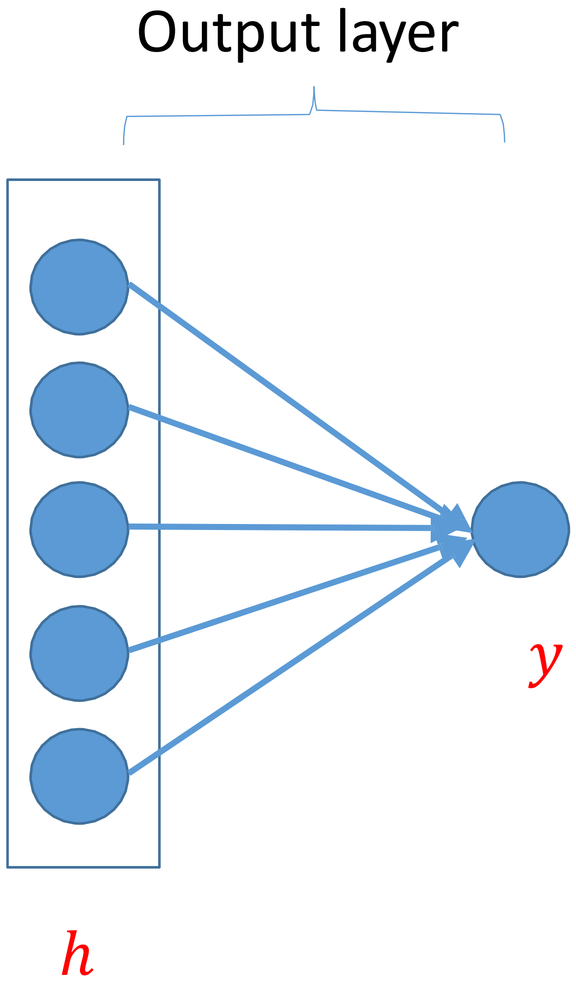

Output layers

- Multi-class classification:
- $y = \text{softmax}(z)$ where $z = W^T h + b$
- Corresponds to using multi-class logistic regression on h

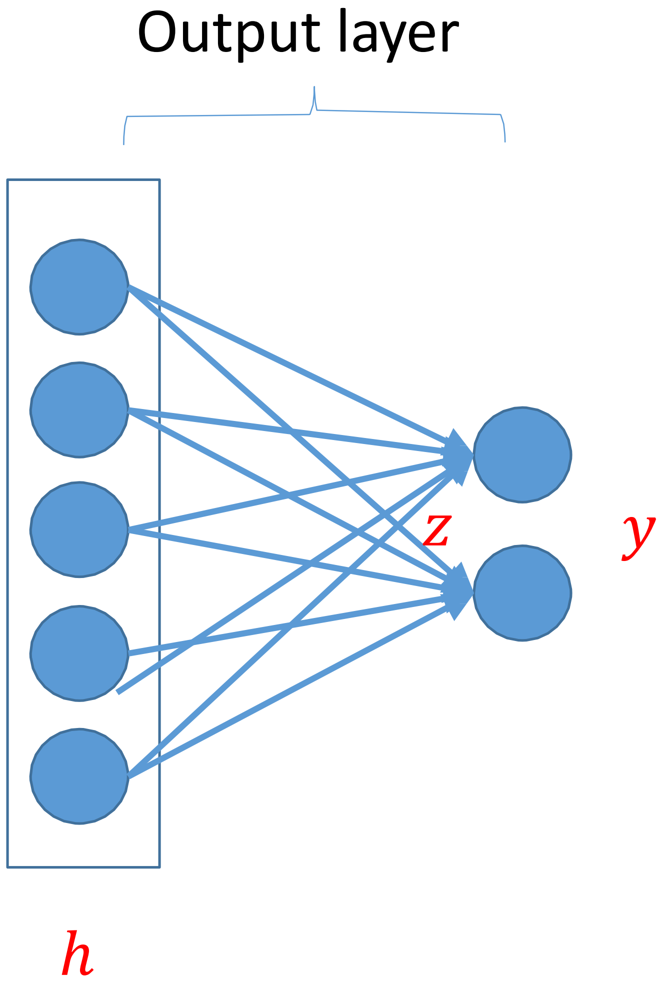

Hidden layers

- Neuron take weighted linear combination of the previous layer
- So can think of outputting one value for the next layer

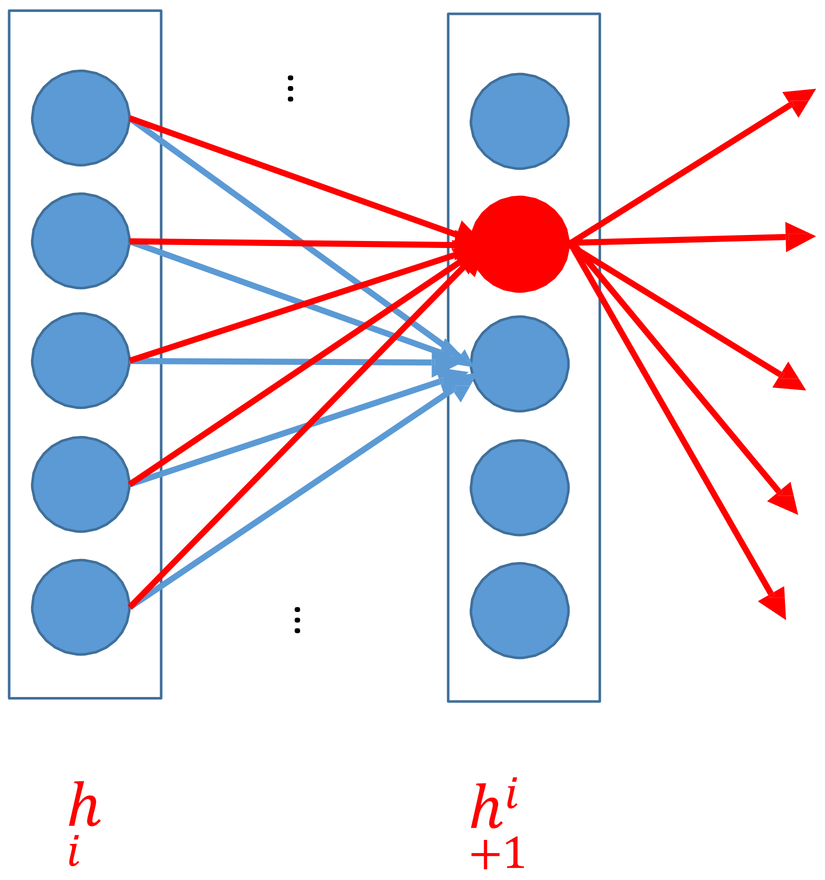

Hidden layers

- $y = r(w^T x + b)$

- Typical activation function r
  - Threshold $\mathrm{t}(z) = \mathbb{I}[z \geq 0]$
  - Sigmoid $\sigma(z) = 1/(1 + \exp(-z))$
  - Tanh $\tanh(z) = 2\sigma(2z) - 1$

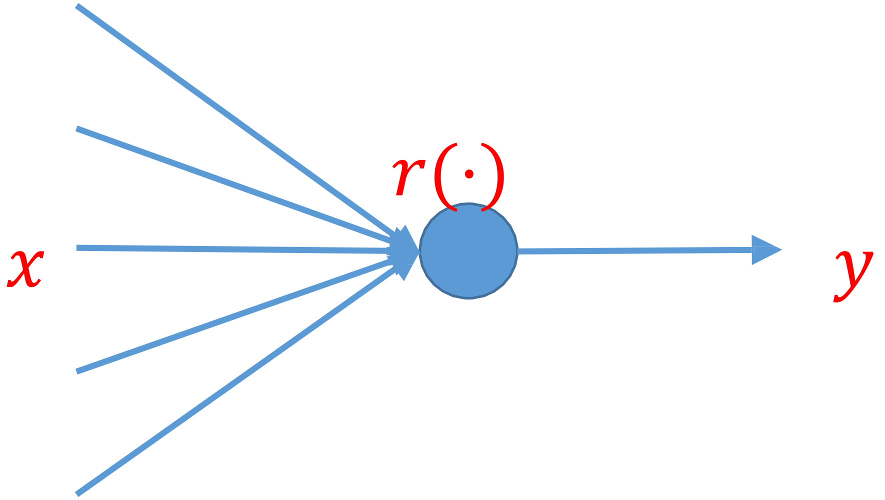

Hidden layers

- Problem: saturation

Too small gradient

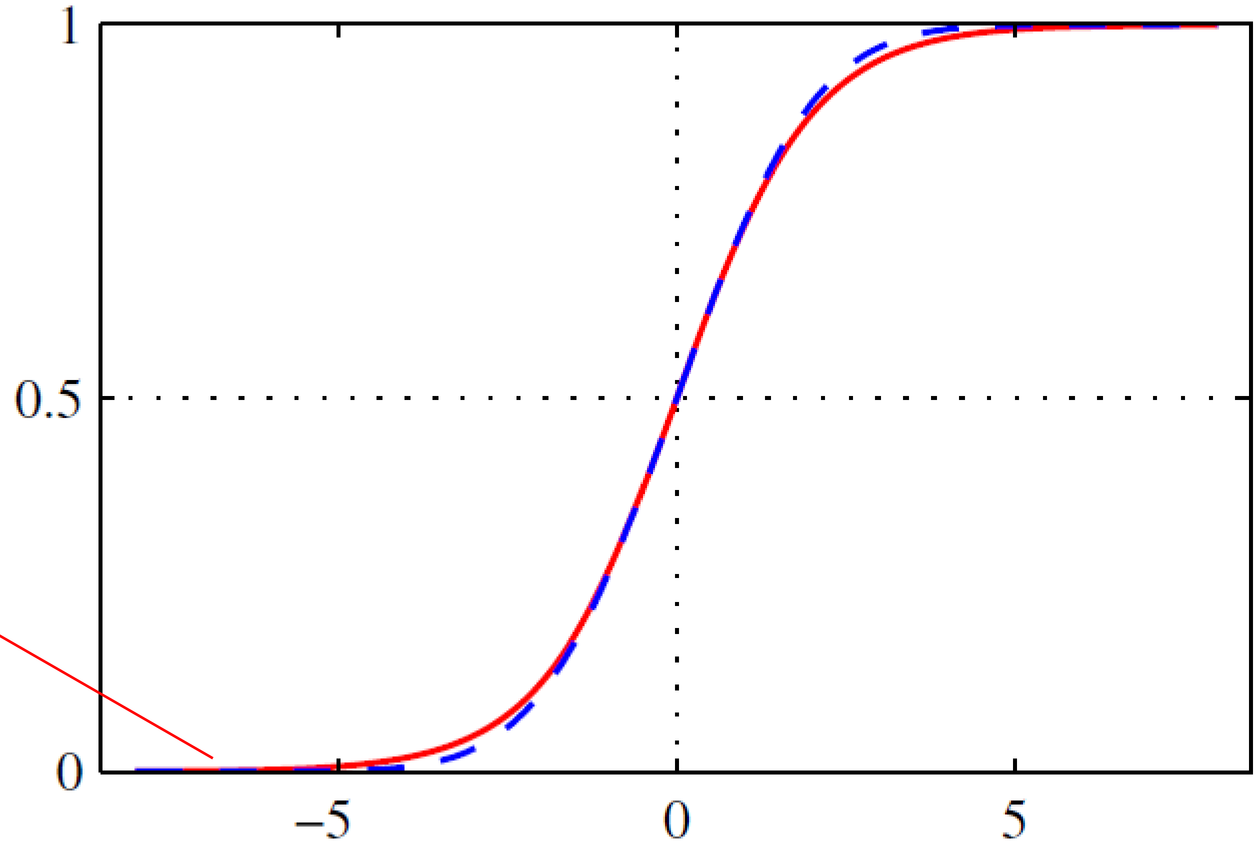

Figure borrowed from *Pattern Recognition and Machine Learning*, Bishop

Hidden layers

- Activation function ReLU (rectified linear unit)
  - $\mathrm{ReLU}(z) = \max\{z, 0\}$

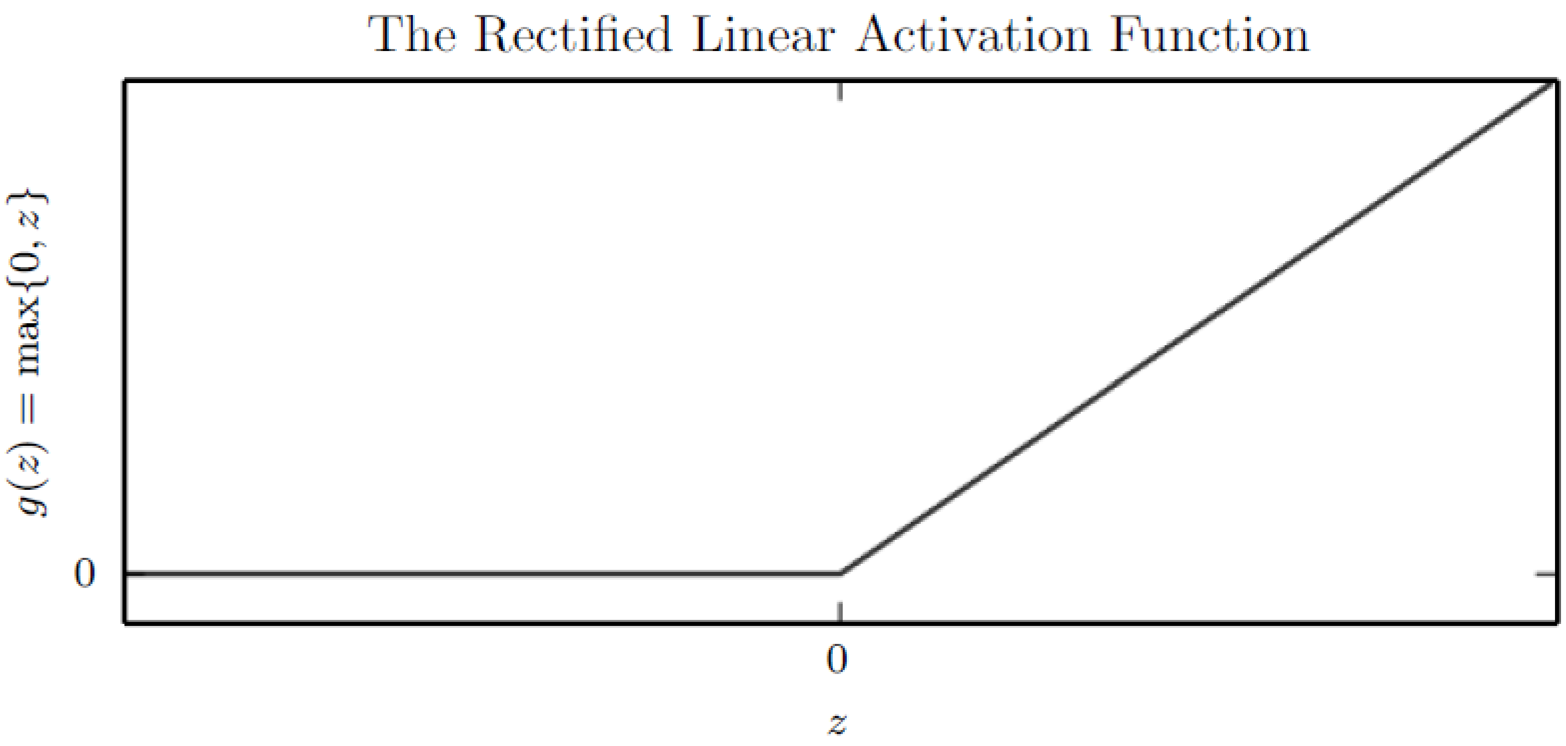

Figure from *Deep learning*, by Goodfellow, Bengio, Courville.

Hidden layers

- Activation function ReLU (rectified linear unit)
  - $\mathrm{ReLU}(z) = \max\{z, 0\}$

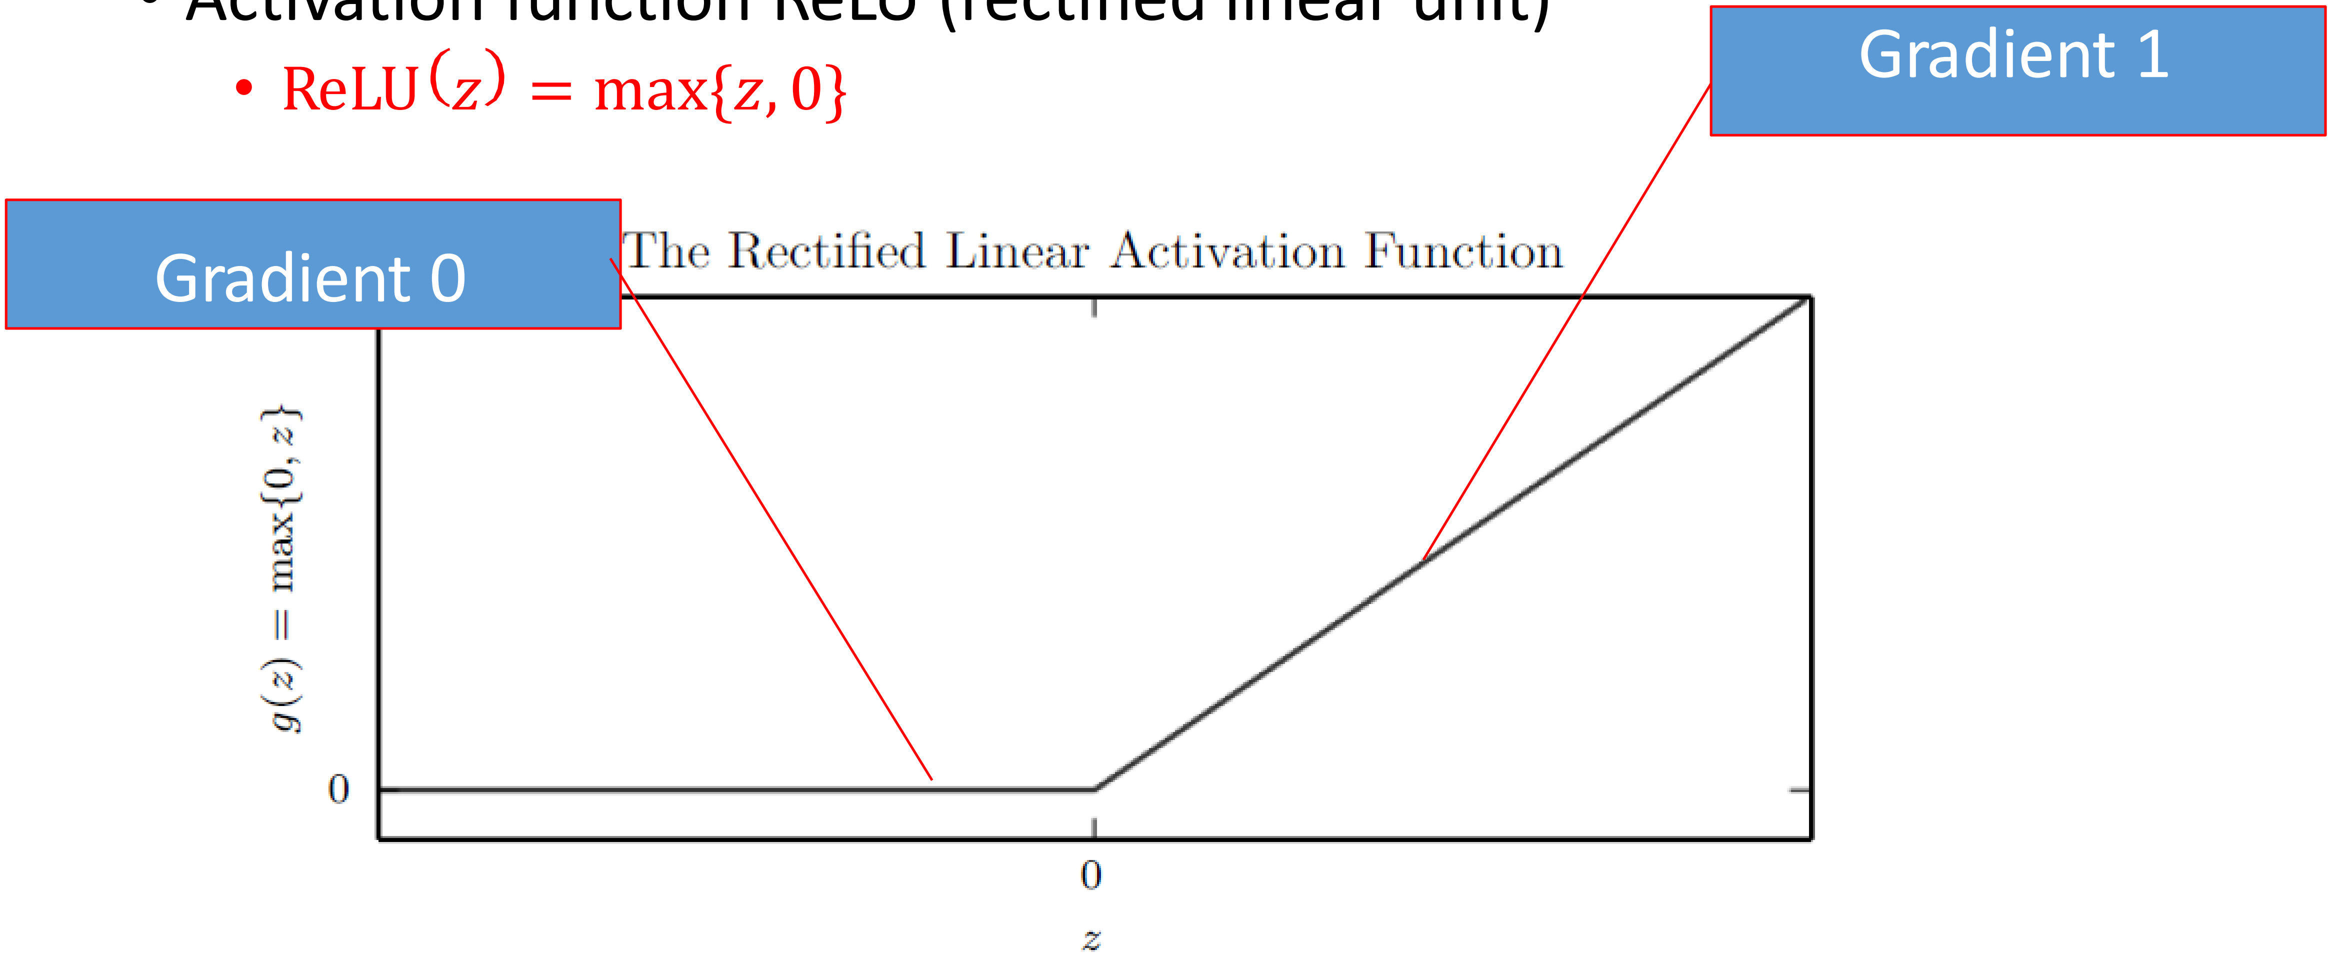

Hidden layers

- Generalizations of ReLU $\mathrm{gReLU}(z) = \max\{z, 0\} + \alpha \min\{z, 0\}$
  - Leaky-ReLU $(z) = \max\{z, 0\} + 0.01 \min\{z, 0\}$
  - Parametric-ReLU (z): α learnable